

Debugging with GDB

The GNU Source-Level Debugger

Ninth Edition, for GDB version 6.2.50`20041017

Richard Stallman, Roland Pesch, Stan Shebs, et al.

(Send bugs and comments on GDB to bug-gdb@gnu.org.)

Debugging with GDB

TeXinfo 2004-02-19.09

Copyright © 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1998, 1999, 2000, 2001,
2002, 2003, 2004 Free Software Foundation, Inc.

Published by the Free Software Foundation
59 Temple Place - Suite 330,
Boston, MA 02111-1307 USA
ISBN 1-882114-77-9

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.1 or any later version published by the Free Software Foundation; with the Invariant Sections being “Free Software” and “Free Software Needs Free Documentation”, with the Front-Cover Texts being “A GNU Manual,” and with the Back-Cover Texts as in (a) below.

(a) The Free Software Foundation’s Back-Cover Text is: “You have freedom to copy and modify this GNU Manual, like GNU software. Copies published by the Free Software Foundation raise funds for GNU development.”

Table of Contents

Summary of GDB	1
Free software	1
Free Software Needs Free Documentation	1
Contributors to GDB.....	3
1 A Sample GDB Session.....	7
2 Getting In and Out of GDB	11
2.1 Invoking GDB	11
2.1.1 Choosing files	12
2.1.2 Choosing modes	13
2.2 Quitting GDB	15
2.3 Shell commands	15
2.4 Logging output	15
3 GDB Commands.....	17
3.1 Command syntax	17
3.2 Command completion	17
3.3 Getting help	19
4 Running Programs Under GDB	23
4.1 Compiling for debugging.....	23
4.2 Starting your program	24
4.3 Your program's arguments	25
4.4 Your program's environment	26
4.5 Your program's working directory	27
4.6 Your program's input and output.....	27
4.7 Debugging an already-running process	28
4.8 Killing the child process	28
4.9 Debugging programs with multiple threads	29
4.10 Debugging programs with multiple processes	31
5 Stopping and Continuing.....	33
5.1 Breakpoints, watchpoints, and catchpoints	33
5.1.1 Setting breakpoints	34
5.1.2 Setting watchpoints	37
5.1.3 Setting catchpoints	39
5.1.4 Deleting breakpoints	41
5.1.5 Disabling breakpoints	41
5.1.6 Break conditions.....	42
5.1.7 Breakpoint command lists	44

5.1.8 Breakpoint menus	45
5.1.9 “Cannot insert breakpoints”	45
5.1.10 “Breakpoint address adjusted...”	46
5.2 Continuing and stepping	46
5.3 Signals	50
5.4 Stopping and starting multi-thread programs	51
6 Examining the Stack	53
6.1 Stack frames	53
6.2 Backtraces	54
6.3 Selecting a frame	55
6.4 Information about a frame	56
7 Examining Source Files	59
7.1 Printing source lines	59
7.2 Editing source files	60
7.2.1 Choosing your editor	61
7.3 Searching source files	61
7.4 Specifying source directories	61
7.5 Source and machine code	62
8 Examining Data	65
8.1 Expressions	65
8.2 Program variables	66
8.3 Artificial arrays	67
8.4 Output formats	68
8.5 Examining memory	69
8.6 Automatic display	70
8.7 Print settings	71
8.8 Value history	76
8.9 Convenience variables	77
8.10 Registers	78
8.11 Floating point hardware	79
8.12 Vector Unit	79
8.13 Operating system auxiliary vector	80
8.14 Memory region attributes	80
8.14.1 Attributes	81
8.14.1.1 Memory Access Mode	81
8.14.1.2 Memory Access Size	81
8.14.1.3 Data Cache	81
8.15 Copy between memory and a file	82
8.16 Character Sets	82
9 C Preprocessor Macros	87

10 Tracepoints	91
10.1 Commands to Set Tracepoints.....	91
10.1.1 Create and Delete Tracepoints.....	91
10.1.2 Enable and Disable Tracepoints.....	92
10.1.3 Tracepoint Passcounts	92
10.1.4 Tracepoint Action Lists	93
10.1.5 Listing Tracepoints	94
10.1.6 Starting and Stopping Trace Experiment	94
10.2 Using the collected data	95
10.2.1 <code>tfind n</code>	95
10.2.2 <code>tdump</code>	97
10.2.3 <code>save-tracepoints filename</code>	98
10.3 Convenience Variables for Tracepoints	98
11 Debugging Programs That Use Overlays ..	99
11.1 How Overlays Work	99
11.2 Overlay Commands	100
11.3 Automatic Overlay Debugging	102
11.4 Overlay Sample Program	103
12 Using GDB with Different Languages	105
12.1 Switching between source languages	105
12.1.1 List of filename extensions and languages	105
12.1.2 Setting the working language	106
12.1.3 Having GDB infer the source language	106
12.2 Displaying the language	106
12.3 Type and range checking	107
12.3.1 An overview of type checking	107
12.3.2 An overview of range checking	108
12.4 Supported languages	109
12.4.1 C and C++	109
12.4.1.1 C and C++ operators	110
12.4.1.2 C and C++ constants	111
12.4.1.3 C++ expressions.....	112
12.4.1.4 C and C++ defaults	113
12.4.1.5 C and C++ type and range checks	113
12.4.1.6 GDB and C	113
12.4.1.7 GDB features for C++	114
12.4.2 Objective-C	115
12.4.2.1 Method Names in Commands	115
12.4.2.2 The Print Command With Objective-C	116
12.4.3 Modula-2	116
12.4.3.1 Operators	116
12.4.3.2 Built-in functions and procedures.....	117
12.4.3.3 Constants	118
12.4.3.4 Modula-2 defaults	119
12.4.3.5 Deviations from standard Modula-2	119

12.4.3.6	Modula-2 type and range checks	119
12.4.3.7	The scope operators :: and	120
12.4.3.8	GDB and Modula-2	120
12.4.4	Ada	120
12.4.4.1	Introduction	120
12.4.4.2	Omissions from Ada	121
12.4.4.3	Additions to Ada	122
12.4.4.4	Stopping at the Very Beginning	123
12.4.4.5	Known Peculiarities of Ada Mode	123
12.5	Unsupported languages	124
13	Examining the Symbol Table	125
14	Altering Execution	131
14.1	Assignment to variables	131
14.2	Continuing at a different address	132
14.3	Giving your program a signal	133
14.4	Returning from a function	133
14.5	Calling program functions	134
14.6	Patching programs	134
15	GDB Files	135
15.1	Commands to specify files	135
15.2	Debugging Information in Separate Files	141
15.3	Errors reading symbol files	143
16	Specifying a Debugging Target	147
16.1	Active targets	147
16.2	Commands for managing targets	147
16.3	Choosing target byte order	149
16.4	Remote debugging	149
16.5	Kernel Object Display	150
17	Debugging remote programs	151
17.1	Connecting to a remote target	151
17.2	Using the <code>gdbserver</code> program	152
17.3	Using the <code>gdbserve.nlm</code> program	153
17.4	Remote configuration	153
17.5	Implementing a remote stub	154
17.5.1	What the stub can do for you	155
17.5.2	What you must do for the stub	155
17.5.3	Putting it all together	157

18 Configuration-Specific Information 159

18.1 Native	159
18.1.1 HP-UX	159
18.1.2 BSD libkvm Interface	159
18.1.3 SVR4 process information	159
18.1.4 Features for Debugging DJGPP Programs	160
18.1.5 Features for Debugging MS Windows PE executables	161
18.1.5.1 Support for DLLs without debugging symbols	162
18.1.5.2 DLL name prefixes	163
18.1.5.3 Working with minimal symbols	163
18.2 Embedded Operating Systems	164
18.2.1 Using GDB with VxWorks	164
18.2.1.1 Connecting to VxWorks	165
18.2.1.2 VxWorks download	165
18.2.1.3 Running tasks	166
18.3 Embedded Processors	166
18.3.1 ARM	166
18.3.2 Renesas H8/300	166
18.3.2.1 Connecting to Renesas boards	167
18.3.2.2 Using the E7000 in-circuit emulator	168
18.3.2.3 Special GDB commands for Renesas micros	168
18.3.3 H8/500	169
18.3.4 Renesas M32R/D	169
18.3.5 M68k	169
18.3.6 MIPS Embedded	169
18.3.7 OpenRISC 1000	171
18.3.8 PowerPC	173
18.3.9 HP PA Embedded	173
18.3.10 Renesas SH	173
18.3.11 Tsquare Sparclet	173
18.3.11.1 Setting file to debug	174
18.3.11.2 Connecting to Sparclet	174
18.3.11.3 Sparclet download	174
18.3.11.4 Running and debugging	175
18.3.12 Fujitsu Sparclite	175
18.3.13 Tandem ST2000	175
18.3.14 Zilog Z8000	175
18.4 Architectures	176
18.4.1 A29K	176
18.4.2 Alpha	176
18.4.3 MIPS	176

19 Controlling GDB	179
19.1 Prompt.....	179
19.2 Command editing.....	179
19.3 Command history.....	179
19.4 Screen size.....	181
19.5 Numbers.....	181
19.6 Configuring the current ABI.....	182
19.7 Optional warnings and messages.....	183
19.8 Optional messages about internal happenings.....	184
20 Canned Sequences of Commands	187
20.1 User-defined commands.....	187
20.2 User-defined command hooks.....	188
20.3 Command files.....	189
20.4 Commands for controlled output.....	190
21 Command Interpreters	193
22 GDB Text User Interface.....	195
22.1 TUI overview	195
22.2 TUI Key Bindings.....	196
22.3 TUI Single Key Mode	197
22.4 TUI specific commands	198
22.5 TUI configuration variables	199
23 Using GDB under GNU Emacs	201
24 The GDB/MI Interface	203
Function and Purpose	203
Notation and Terminology	203
24.3 GDB/MI Command Syntax	203
24.3.1 GDB/MI Input Syntax	203
24.3.2 GDB/MI Output Syntax	204
24.3.3 Simple Examples of GDB/MI Interaction	206
24.4 GDB/MI Compatibility with CLI	206
24.5 GDB/MI Output Records	207
24.5.1 GDB/MI Result Records	207
24.5.2 GDB/MI Stream Records	207
24.5.3 GDB/MI Out-of-band Records	207
24.6 GDB/MI Command Description Format	207
24.7 GDB/MI Breakpoint table commands	208
24.8 GDB/MI Data Manipulation	215
24.9 GDB/MI Program control	225
24.10 Miscellaneous GDB commands in GDB/MI	234
24.11 GDB/MI Stack Manipulation Commands	235
24.12 GDB/MI Symbol Query Commands	239

24.13 GDB/MI Target Manipulation Commands	242
24.14 GDB/MI Thread Commands	247
24.15 GDB/MI Tracepoint Commands	248
24.16 GDB/MI Variable Objects	248
25 GDB Annotations	253
25.1 What is an Annotation?	253
25.2 The Server Prefix	254
25.3 Annotation for GDB Input	254
25.4 Errors	254
25.5 Invalidation Notices	255
25.6 Running the Program	255
25.7 Displaying Source	256
26 Reporting Bugs in GDB	257
26.1 Have you found a bug?	257
26.2 How to report bugs	257
27 Command Line Editing	261
27.1 Introduction to Line Editing	261
27.2 Readline Interaction	261
27.2.1 Readline Bare Essentials	261
27.2.2 Readline Movement Commands	262
27.2.3 Readline Killing Commands	262
27.2.4 Readline Arguments	263
27.2.5 Searching for Commands in the History	263
27.3 Readline Init File	264
27.3.1 Readline Init File Syntax	264
27.3.2 Conditional Init Constructs	269
27.3.3 Sample Init File	269
27.4 Bindable Readline Commands	272
27.4.1 Commands For Moving	272
27.4.2 Commands For Manipulating The History	272
27.4.3 Commands For Changing Text	273
27.4.4 Killing And Yanking	275
27.4.5 Specifying Numeric Arguments	276
27.4.6 Letting Readline Type For You	276
27.4.7 Keyboard Macros	276
27.4.8 Some Miscellaneous Commands	277
27.5 Readline vi Mode	278
28 Using History Interactively	279
28.1 History Expansion	279
28.1.1 Event Designators	279
28.1.2 Word Designators	279
28.1.3 Modifiers	280

Appendix A Formatting Documentation 283**Appendix B Installing GDB 285**

B.1	Compiling GDB in another directory	286
B.2	Specifying names for hosts and targets	287
B.3	configure options	287

Appendix C Maintenance Commands 289**Appendix D GDB Remote Serial Protocol 293**

D.1	Overview	293
D.2	Packets	294
D.3	Stop Reply Packets	301
D.4	General Query Packets	302
D.5	Register Packet Format	305
D.6	Examples	305
D.7	File-I/O remote protocol extension	306
D.7.1	File-I/O Overview	306
D.7.2	Protocol basics	306
D.7.3	The F request packet	307
D.7.4	The F reply packet	307
D.7.5	Memory transfer	308
D.7.6	The Ctrl-C message	308
D.7.7	Console I/O	308
D.7.8	The isatty(3) call	309
D.7.9	The system(3) call	309
D.7.10	List of supported calls	309
	open	309
	close	310
	read	311
	write	311
	lseek	311
	rename	312
	unlink	313
	stat/fstat	313
	gettimeofday	314
	isatty	314
	system	314
D.7.11	Protocol specific representation of datatypes	314
	Integral datatypes	314
	Pointer values	315
	struct stat	315
	struct timeval	316
D.7.12	Constants	316
	Open flags	316
	mode_t values	316
	Errno values	316

Lseek flags	317
Limits	317
D.7.13 File-I/O Examples	317
Appendix E The GDB Agent Expression Mechanism	319
E.1 General Bytecode Design	319
E.2 Bytecode Descriptions	321
E.3 Using Agent Expressions	325
E.4 Varying Target Capabilities	326
E.5 Tracing on Symmetrix	326
E.6 Rationale	328
Appendix F GNU GENERAL PUBLIC LICENSE	331
Preamble	331
TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION	332
How to Apply These Terms to Your New Programs	336
Appendix G GNU Free Documentation License	337
G.1 ADDENDUM: How to use this License for your documents ..	343
Index	345

Summary of GDB

The purpose of a debugger such as GDB is to allow you to see what is going on “inside” another program while it executes—or what another program was doing at the moment it crashed.

GDB can do four main kinds of things (plus other things in support of these) to help you catch bugs in the act:

- Start your program, specifying anything that might affect its behavior.
- Make your program stop on specified conditions.
- Examine what has happened, when your program has stopped.
- Change things in your program, so you can experiment with correcting the effects of one bug and go on to learn about another.

You can use GDB to debug programs written in C and C++. For more information, see [Section 12.4 \[Supported languages\], page 109](#). For more information, see [Section 12.4.1 \[C and C++\], page 109](#).

Support for Modula-2 is partial. For information on Modula-2, see [Section 12.4.3 \[Modula-2\], page 116](#).

Debugging Pascal programs which use sets, subranges, file variables, or nested functions does not currently work. GDB does not support entering expressions, printing values, or similar features using Pascal syntax.

GDB can be used to debug programs written in Fortran, although it may be necessary to refer to some variables with a trailing underscore.

GDB can be used to debug programs written in Objective-C, using either the Apple/NeXT or the GNU Objective-C runtime.

Free software

GDB is *free software*, protected by the GNU General Public License (GPL). The GPL gives you the freedom to copy or adapt a licensed program—but every person getting a copy also gets with it the freedom to modify that copy (which means that they must get access to the source code), and the freedom to distribute further copies. Typical software companies use copyrights to limit your freedoms; the Free Software Foundation uses the GPL to preserve these freedoms.

Fundamentally, the General Public License is a license which says that you have these freedoms and that you cannot take these freedoms away from anyone else.

Free Software Needs Free Documentation

The biggest deficiency in the free software community today is not in the software—it is the lack of good free documentation that we can include with the free software. Many of our most important programs do not come with free reference manuals and free introductory texts. Documentation is an essential part of any software package; when an important free software package does not come with a free manual and a free tutorial, that is a major gap. We have many such gaps today.

Consider Perl, for instance. The tutorial manuals that people normally use are non-free. How did this come about? Because the authors of those manuals published them with restrictive terms—no copying, no modification, source files not available—which exclude them from the free software world.

That wasn't the first time this sort of thing happened, and it was far from the last. Many times we have heard a GNU user eagerly describe a manual that he is writing, his intended contribution to the community, only to learn that he had ruined everything by signing a publication contract to make it non-free.

Free documentation, like free software, is a matter of freedom, not price. The problem with the non-free manual is not that publishers charge a price for printed copies—that in itself is fine. (The Free Software Foundation sells printed copies of manuals, too.) The problem is the restrictions on the use of the manual. Free manuals are available in source code form, and give you permission to copy and modify. Non-free manuals do not allow this.

The criteria of freedom for a free manual are roughly the same as for free software. Redistribution (including the normal kinds of commercial redistribution) must be permitted, so that the manual can accompany every copy of the program, both on-line and on paper.

Permission for modification of the technical content is crucial too. When people modify the software, adding or changing features, if they are conscientious they will change the manual too—so they can provide accurate and clear documentation for the modified program. A manual that leaves you no choice but to write a new manual to document a changed version of the program is not really available to our community.

Some kinds of limits on the way modification is handled are acceptable. For example, requirements to preserve the original author's copyright notice, the distribution terms, or the list of authors, are ok. It is also no problem to require modified versions to include notice that they were modified. Even entire sections that may not be deleted or changed are acceptable, as long as they deal with nontechnical topics (like this one). These kinds of restrictions are acceptable because they don't obstruct the community's normal use of the manual.

However, it must be possible to modify all the *technical* content of the manual, and then distribute the result in all the usual media, through all the usual channels. Otherwise, the restrictions obstruct the use of the manual, it is not free, and we need another manual to replace it.

Please spread the word about this issue. Our community continues to lose manuals to proprietary publishing. If we spread the word that free software needs free reference manuals and free tutorials, perhaps the next person who wants to contribute by writing documentation will realize, before it is too late, that only free manuals contribute to the free software community.

If you are writing documentation, please insist on publishing it under the GNU Free Documentation License or another free documentation license. Remember that this decision requires your approval—you don't have to let the publisher decide. Some commercial publishers will use a free license if you insist, but they will not propose the option; it is up to you to raise the issue and say firmly that this is what you want. If the publisher you are dealing with refuses, please try other publishers. If you're not sure whether a proposed license is free, write to licensing@gnu.org.

You can encourage commercial publishers to sell more free, copylefted manuals and tutorials by buying them, and particularly by buying copies from the publishers that paid for their writing or for major improvements. Meanwhile, try to avoid buying non-free documentation at all. Check the distribution terms of a manual before you buy it, and insist that whoever seeks your business must respect your freedom. Check the history of the book, and try to reward the publishers that have paid or pay the authors to work on it.

The Free Software Foundation maintains a list of free documentation published by other publishers, at <http://www.fsf.org/doc/other-free-books.html>.

Contributors to GDB

Richard Stallman was the original author of GDB, and of many other GNU programs. Many others have contributed to its development. This section attempts to credit major contributors. One of the virtues of free software is that everyone is free to contribute to it; with regret, we cannot actually acknowledge everyone here. The file ‘**ChangeLog**’ in the GDB distribution approximates a blow-by-blow account.

Changes much prior to version 2.0 are lost in the mists of time.

Plea: Additions to this section are particularly welcome. If you or your friends (or enemies, to be evenhanded) have been unfairly omitted from this list, we would like to add your names!

So that they may not regard their many labors as thankless, we particularly thank those who shepherded GDB through major releases: Andrew Cagney (releases 6.1, 6.0, 5.3, 5.2, 5.1 and 5.0); Jim Blandy (release 4.18); Jason Molenda (release 4.17); Stan Shebs (release 4.14); Fred Fish (releases 4.16, 4.15, 4.13, 4.12, 4.11, 4.10, and 4.9); Stu Grossman and John Gilmore (releases 4.8, 4.7, 4.6, 4.5, and 4.4); John Gilmore (releases 4.3, 4.2, 4.1, 4.0, and 3.9); Jim Kingdon (releases 3.5, 3.4, and 3.3); and Randy Smith (releases 3.2, 3.1, and 3.0).

Richard Stallman, assisted at various times by Peter TerMaat, Chris Hanson, and Richard Mlynarik, handled releases through 2.8.

Michael Tiemann is the author of most of the GNU C++ support in GDB, with significant additional contributions from Per Bothner and Daniel Berlin. James Clark wrote the GNU C++ demangler. Early work on C++ was by Peter TerMaat (who also did much general update work leading to release 3.0).

GDB uses the BFD subroutine library to examine multiple object-file formats; BFD was a joint project of David V. Henkel-Wallace, Rich Pixley, Steve Chamberlain, and John Gilmore.

David Johnson wrote the original COFF support; Pace Willison did the original support for encapsulated COFF.

Brent Benson of Harris Computer Systems contributed DWARF 2 support.

Adam de Boor and Bradley Davis contributed the ISI Optimum V support. Per Bothner, Noboyuki Hikichi, and Alessandro Forin contributed MIPS support. Jean-Daniel Fekete contributed Sun 386i support. Chris Hanson improved the HP9000 support. Noboyuki Hikichi and Tomoyuki Hasei contributed Sony/News OS 3 support. David Johnson contributed Encore Umax support. Jyrki Kuoppala contributed Altos 3068 support. Jeff Law contributed HP PA and SOM support. Keith Packard contributed NS32K support. Doug Rabson contributed Acorn Risc Machine support. Bob Rusk contributed Harris

Nighthawk CX-UX support. Chris Smith contributed Convex support (and Fortran debugging). Jonathan Stone contributed Pyramid support. Michael Tiemann contributed SPARC support. Tim Tucker contributed support for the Gould NP1 and Gould Powernode. Pace Willison contributed Intel 386 support. Jay Vosburgh contributed Symmetry support. Marko Mlinar contributed OpenRISC 1000 support.

Andreas Schwab contributed M68K GNU/Linux support.

Rich Schaefer and Peter Schauer helped with support of SunOS shared libraries.

Jay Fenlason and Roland McGrath ensured that GDB and GAS agree about several machine instruction sets.

Patrick Duval, Ted Goldstein, Vikram Koka and Glenn Engel helped develop remote debugging. Intel Corporation, Wind River Systems, AMD, and ARM contributed remote debugging modules for the i960, VxWorks, A29K UDI, and RDI targets, respectively.

Brian Fox is the author of the readline libraries providing command-line editing and command history.

Andrew Beers of SUNY Buffalo wrote the language-switching code, the Modula-2 support, and contributed the Languages chapter of this manual.

Fred Fish wrote most of the support for Unix System Vr4. He also enhanced the command-completion support to cover C++ overloaded symbols.

Hitachi America (now Renesas America), Ltd. sponsored the support for H8/300, H8/500, and Super-H processors.

NEC sponsored the support for the v850, Vr4xxx, and Vr5xxx processors.

Mitsubishi (now Renesas) sponsored the support for D10V, D30V, and M32R/D processors.

Toshiba sponsored the support for the TX39 Mips processor.

Matsushita sponsored the support for the MN10200 and MN10300 processors.

Fujitsu sponsored the support for SPARClite and FR30 processors.

Kung Hsu, Jeff Law, and Rick Sladkey added support for hardware watchpoints.

Michael Snyder added support for tracepoints.

Stu Grossman wrote gdbserver.

Jim Kingdon, Peter Schauer, Ian Taylor, and Stu Grossman made nearly innumerable bug fixes and cleanups throughout GDB.

The following people at the Hewlett-Packard Company contributed support for the PA-RISC 2.0 architecture, HP-UX 10.20, 10.30, and 11.0 (narrow mode), HP's implementation of kernel threads, HP's aC++ compiler, and the Text User Interface (nee Terminal User Interface): Ben Krepp, Richard Title, John Bishop, Susan Macchia, Kathy Mann, Satish Pai, India Paul, Steve Rehrauer, and Elena Zannoni. Kim Haase provided HP-specific information in this manual.

DJ Delorie ported GDB to MS-DOS, for the DJGPP project. Robert Hoehne made significant contributions to the DJGPP port.

Cygnus Solutions has sponsored GDB maintenance and much of its development since 1991. Cygnus engineers who have worked on GDB fulltime include Mark Alexander, Jim Blandy, Per Bothner, Kevin Buettner, Edith Epstein, Chris Taylor, Fred Fish, Martin

Hunt, Jim Ingham, John Gilmore, Stu Grossman, Kung Hsu, Jim Kingdon, John Metzler, Fernando Nasser, Geoffrey Noer, Dawn Perchik, Rich Pixley, Zdenek Radouch, Keith Seitz, Stan Shebs, David Taylor, and Elena Zannoni. In addition, Dave Brolley, Ian Carmichael, Steve Chamberlain, Nick Clifton, JT Conklin, Stan Cox, DJ Delorie, Ulrich Drepper, Frank Eigler, Doug Evans, Sean Fagan, David Henkel-Wallace, Richard Henderson, Jeff Holcomb, Jeff Law, Jim Lemke, Tom Lord, Bob Manson, Michael Meissner, Jason Merrill, Catherine Moore, Drew Moseley, Ken Raeburn, Gavin Romig-Koch, Rob Savoye, Jamie Smith, Mike Stump, Ian Taylor, Angela Thomas, Michael Tiemann, Tom Tromey, Ron Unrau, Jim Wilson, and David Zuhn have made contributions both large and small.

Jim Blandy added support for preprocessor macros, while working for Red Hat.

1 A Sample GDB Session

You can use this manual at your leisure to read all about GDB. However, a handful of commands are enough to get started using the debugger. This chapter illustrates those commands.

In this sample session, we emphasize user input like this: **input**, to make it easier to pick out from the surrounding output.

One of the preliminary versions of GNU **m4** (a generic macro processor) exhibits the following bug: sometimes, when we change its quote strings from the default, the commands used to capture one macro definition within another stop working. In the following short **m4** session, we define a macro **foo** which expands to **0000**; we then use the **m4** built-in **defn** to define **bar** as the same thing. However, when we change the open quote string to **<QUOTE>** and the close quote string to **<UNQUOTE>**, the same procedure fails to define a new synonym **baz**:

```
$ cd gnu/m4
$ ./m4
define(foo,0000)

foo
0000
define(bar,defn('foo'))

bar
0000
changequote(<QUOTE>,<UNQUOTE>)

define(baz,defn(<QUOTE>foo<UNQUOTE>))
baz
C-d
m4: End of input: 0: fatal error: EOF in string
```

Let us use GDB to try to see what is going on.

```
$ gdb m4
GDB is free software and you are welcome to distribute copies
of it under certain conditions; type "show copying" to see
the conditions.
There is absolutely no warranty for GDB; type "show warranty"
for details.

GDB 6.2.50_20041017, Copyright 1999 Free Software Foundation, Inc...
(gdb)
```

GDB reads only enough symbol data to know where to find the rest when needed; as a result, the first prompt comes up very quickly. We now tell GDB to use a narrower display width than usual, so that examples fit in this manual.

```
(gdb) set width 70
```

We need to see how the **m4** built-in **changequote** works. Having looked at the source, we know the relevant subroutine is **m4_changequote**, so we set a breakpoint there with the GDB **break** command.

```
(gdb) break m4_changequote
Breakpoint 1 at 0x62f4: file builtin.c, line 879.
```

Using the **run** command, we start **m4** running under GDB control; as long as control does not reach the **m4_changequote** subroutine, the program runs as usual:

```
(gdb) run
Starting program: /work/Editorial/gdb/gnu/m4/m4
define(foo,0000)

foo
0000
```

To trigger the breakpoint, we call `changequote`. GDB suspends execution of `m4`, displaying information about the context where it stops.

```
changequote(<QUOTE>,<UNQUOTE>)

Breakpoint 1, m4_changequote (argc=3, argv=0x33c70)
  at builtin.c:879
  879      if (bad_argc(TOKEN_DATA_TEXT(argv[0]),argc,1,3))
```

Now we use the command `n` (`next`) to advance execution to the next line of the current function.

```
(gdb) n
882      set_quotes((argc >= 2) ? TOKEN_DATA_TEXT(argv[1])\
: nil,
```

`set_quotes` looks like a promising subroutine. We can go into it by using the command `s` (`step`) instead of `next`. `step` goes to the next line to be executed in *any* subroutine, so it steps into `set_quotes`.

```
(gdb) s
set_quotes (lq=0x34c78 "<QUOTE>", rq=0x34c88 "<UNQUOTE>")
  at input.c:530
  530      if (lquote != def_lquote)
```

The display that shows the subroutine where `m4` is now suspended (and its arguments) is called a stack frame display. It shows a summary of the stack. We can use the `backtrace` command (which can also be spelled `bt`), to see where we are in the stack as a whole: the `backtrace` command displays a stack frame for each active subroutine.

```
(gdb) bt
#0  set_quotes (lq=0x34c78 "<QUOTE>", rq=0x34c88 "<UNQUOTE>")
  at input.c:530
#1  0x6344 in m4_changequote (argc=3, argv=0x33c70)
  at builtin.c:882
#2  0x8174 in expand_macro (sym=0x33320) at macro.c:242
#3  0x7a88 in expand_token (obs=0x0, t=209696, td=0xf7ffffa30)
  at macro.c:71
#4  0x79dc in expand_input () at macro.c:40
#5  0x2930 in main (argc=0, argv=0xf7ffffb20) at m4.c:195
```

We step through a few more lines to see what happens. The first two times, we can use '`s`'; the next two times we use `n` to avoid falling into the `xstrdup` subroutine.

```
(gdb) s
0x3b5c  532      if (rquote != def_rquote)
(gdb) s
0x3b80  535      lquote = (lq == nil || *lq == '\0') ? \
def_lquote : xstrdup(lq);
(gdb) n
536      rquote = (rq == nil || *rq == '\0') ? def_rquote\
: xstrdup(rq);
(gdb) n
538      len_lquote = strlen(rquote);
```

The last line displayed looks a little odd; we can examine the variables `lquote` and `rquote` to see if they are in fact the new left and right quotes we specified. We use the command `p` (print) to see their values.

```
(gdb) p lquote
$1 = 0x35d40 "<QUOTE>"
(gdb) p rquote
$2 = 0x35d50 "<UNQUOTE>"
```

`lquote` and `rquote` are indeed the new left and right quotes. To look at some context, we can display ten lines of source surrounding the current line with the `l` (list) command.

```
(gdb) l
533           xfree(rquote);
534
535           lquote = (lq == nil || *lq == '\0') ? def_lquote\
: xstrdup (lq);
536           rquote = (rq == nil || *rq == '\0') ? def_rquote\
: xstrdup (rq);
537
538           len_lquote = strlen(rquote);
539           len_rquote = strlen(lquote);
540       }
541
542     void
```

Let us step past the two lines that set `len_lquote` and `len_rquote`, and then examine the values of those variables.

```
(gdb) n
539           len_rquote = strlen(lquote);
(gdb) n
540       }
(gdb) p len_lquote
$3 = 9
(gdb) p len_rquote
$4 = 7
```

That certainly looks wrong, assuming `len_lquote` and `len_rquote` are meant to be the lengths of `lquote` and `rquote` respectively. We can set them to better values using the `p` command, since it can print the value of any expression—and that expression can include subroutine calls and assignments.

```
(gdb) p len_lquote=strlen(lquote)
$5 = 7
(gdb) p len_rquote=strlen(rquote)
$6 = 9
```

Is that enough to fix the problem of using the new quotes with the `m4` built-in `defn`? We can allow `m4` to continue executing with the `c` (`continue`) command, and then try the example that caused trouble initially:

```
(gdb) c
Continuing.

define(baz,defn(<QUOTE>foo<UNQUOTE>))

baz
0000
```

Success! The new quotes now work just as well as the default ones. The problem seems to have been just the two typos defining the wrong lengths. We allow `m4` exit by giving it an EOF as input:

```
C-d  
Program exited normally.
```

The message ‘`Program exited normally.`’ is from GDB; it indicates `m4` has finished executing. We can end our GDB session with the GDB `quit` command.

```
(gdb) quit
```

2 Getting In and Out of GDB

This chapter discusses how to start GDB, and how to get out of it. The essentials are:

- type ‘`gdb`’ to start GDB.
- type `quit` or `C-d` to exit.

2.1 Invoking GDB

Invoke GDB by running the program `gdb`. Once started, GDB reads commands from the terminal until you tell it to exit.

You can also run `gdb` with a variety of arguments and options, to specify more of your debugging environment at the outset.

The command-line options described here are designed to cover a variety of situations; in some environments, some of these options may effectively be unavailable.

The most usual way to start GDB is with one argument, specifying an executable program:

```
gdb program
```

You can also start with both an executable program and a core file specified:

```
gdb program core
```

You can, instead, specify a process ID as a second argument, if you want to debug a running process:

```
gdb program 1234
```

would attach GDB to process 1234 (unless you also have a file named ‘1234’; GDB does check for a core file first).

Taking advantage of the second command-line argument requires a fairly complete operating system; when you use GDB as a remote debugger attached to a bare board, there may not be any notion of “process”, and there is often no way to get a core dump. GDB will warn you if it is unable to attach or to read core dumps.

You can optionally have `gdb` pass any arguments after the executable file to the inferior using `--args`. This option stops option processing.

```
gdb --args gcc -O2 -c foo.c
```

This will cause `gdb` to debug `gcc`, and to set `gcc`’s command-line arguments (see [Section 4.3 \[Arguments\], page 25](#)) to ‘`-O2 -c foo.c`’.

You can run `gdb` without printing the front material, which describes GDB’s non-warranty, by specifying `-silent`:

```
gdb -silent
```

You can further control how GDB starts up by using command-line options. GDB itself can remind you of the options available.

Type

```
gdb -help
```

to display all available options and briefly describe their use (‘`gdb -h`’ is a shorter equivalent).

All options and command line arguments you give are processed in sequential order. The order makes a difference when the ‘`-x`’ option is used.

2.1.1 Choosing files

When GDB starts, it reads any arguments other than options as specifying an executable file and core file (or process ID). This is the same as if the arguments were specified by the ‘`-se`’ and ‘`-c`’ (or ‘`-p`’) options respectively. (GDB reads the first argument that does not have an associated option flag as equivalent to the ‘`-se`’ option followed by that argument; and the second argument that does not have an associated option flag, if any, as equivalent to the ‘`-c`’/‘`-p`’ option followed by that argument.) If the second argument begins with a decimal digit, GDB will first attempt to attach to it as a process, and if that fails, attempt to open it as a corefile. If you have a corefile whose name begins with a digit, you can prevent GDB from treating it as a pid by prefixing it with ‘`./`’, eg. ‘`./12345`’.

If GDB has not been configured to included core file support, such as for most embedded targets, then it will complain about a second argument and ignore it.

Many options have both long and short forms; both are shown in the following list. GDB also recognizes the long forms if you truncate them, so long as enough of the option is present to be unambiguous. (If you prefer, you can flag option arguments with ‘`--`’ rather than ‘`-`’, though we illustrate the more usual convention.)

```
-symbols file
-s file      Read symbol table from file file.
-exec file
-e file      Use file file as the executable file to execute when appropriate, and for examining
                pure data in conjunction with a core dump.
-se file     Read symbol table from file file and use it as the executable file.
-core file
-c file      Use file file as a core dump to examine.
-c number
-pid number
-p number
                Connect to process ID number, as with the attach command. If there is no
                such process, GDB will attempt to open a core file named number.
-command file
-x file      Execute GDB commands from file file. See Section 20.3 \[Command files\], page 189.
-directory directory
-d directory
                Add directory to the path to search for source files.
-m
-mapped     Warning: this option depends on operating system facilities that are not supported on all systems.
                If memory-mapped files are available on your system through the mmap system
                call, you can use this option to have GDB write the symbols from your program
                into a reusable file in the current directory. If the program you are debugging
                is called ‘/tmp/fred’, the mapped symbol file is ‘/tmp/fred.syms’. Future
                GDB debugging sessions notice the presence of this file, and can quickly map
```

in symbol information from it, rather than reading the symbol table from the executable program.

The ‘.syms’ file is specific to the host machine where GDB is run. It holds an exact image of the internal GDB symbol table. It cannot be shared across multiple host platforms.

-r

-readnow Read each symbol file’s entire symbol table immediately, rather than the default, which is to read it incrementally as it is needed. This makes startup slower, but makes future operations faster.

You typically combine the **-mapped** and **-readnow** options in order to build a ‘.syms’ file that contains complete symbol information. (See [Section 15.1 \[Commands to specify files\]](#), page 135, for information on ‘.syms’ files.) A simple GDB invocation to do nothing but build a ‘.syms’ file for future use is:

```
gdb -batch -nx -mapped -readnow programname
```

2.1.2 Choosing modes

You can run GDB in various alternative modes—for example, in batch mode or quiet mode.

-nx

-n Do not execute commands found in any initialization files. Normally, GDB executes the commands in these files after all the command options and arguments have been processed. See [Section 20.3 \[Command files\]](#), page 189.

-quiet

-silent

-q “Quiet”. Do not print the introductory and copyright messages. These messages are also suppressed in batch mode.

-batch Run in batch mode. Exit with status 0 after processing all the command files specified with ‘-x’ (and all commands from initialization files, if not inhibited with ‘-n’). Exit with nonzero status if an error occurs in executing the GDB commands in the command files.

Batch mode may be useful for running GDB as a filter, for example to download and run a program on another computer; in order to make this more useful, the message

Program exited normally.

(which is ordinarily issued whenever a program running under GDB control terminates) is not issued when running in batch mode.

-nowindows

-nw “No windows”. If GDB comes with a graphical user interface (GUI) built in, then this option tells GDB to only use the command-line interface. If no GUI is available, this option has no effect.

-windows

-w If GDB includes a GUI, then this option requires it to be used if possible.

-cd directory

Run GDB using *directory* as its working directory, instead of the current directory.

-fullname

-f GNU Emacs sets this option when it runs GDB as a subprocess. It tells GDB to output the full file name and line number in a standard, recognizable fashion each time a stack frame is displayed (which includes each time your program stops). This recognizable format looks like two ‘\032’ characters, followed by the file name, line number and character position separated by colons, and a newline. The Emacs-to-GDB interface program uses the two ‘\032’ characters as a signal to display the source code for the frame.

-epoch The Epoch Emacs-GDB interface sets this option when it runs GDB as a subprocess. It tells GDB to modify its print routines so as to allow Epoch to display values of expressions in a separate window.

-annotate level

This option sets the *annotation level* inside GDB. Its effect is identical to using ‘`set annotate level`’ (see [Chapter 25 \[Annotations\], page 253](#)). The annotation *level* controls how much information GDB prints together with its prompt, values of expressions, source lines, and other types of output. Level 0 is the normal, level 1 is for use when GDB is run as a subprocess of GNU Emacs, level 3 is the maximum annotation suitable for programs that control GDB, and level 2 has been deprecated.

The annotation mechanism has largely been superseeded by GDB/MI (see [Chapter 24 \[GDB/MI\], page 203](#)).

--args Change interpretation of command line so that arguments following the executable file are passed as command line arguments to the inferior. This option stops option processing.

-baud bps

-b bps Set the line speed (baud rate or bits per second) of any serial interface used by GDB for remote debugging.

-tty device**-t device**

Run using *device* for your program’s standard input and output.

-tui Activate the *Text User Interface* when starting. The Text User Interface manages several text windows on the terminal, showing source, assembly, registers and GDB command outputs (see [Chapter 22 \[GDB Text User Interface\], page 195](#)). Alternatively, the Text User Interface can be enabled by invoking the program ‘`gdbtui`’. Do not use this option if you run GDB from Emacs (see [Chapter 23 \[Using GDB under GNU Emacs\], page 201](#)).

-interpreter interp

Use the interpreter *interp* for interface with the controlling program or device. This option is meant to be set by programs which communicate with GDB using it as a back end. See [Chapter 21 \[Command Interpreters\], page 193](#).

‘`--interpreter=mi`’ (or ‘`--interpreter=mi2`’) causes GDB to use the GDB/MI interface (see [Chapter 24 \[The GDB/MI Interface\], page 203](#)) included since GDBN version 6.0. The previous GDB/MI interface, included in GDB version

5.3 and selected with ‘`--interpreter=mi1`’, is deprecated. Earlier GDB/MI interfaces are no longer supported.

-write Open the executable and core files for both reading and writing. This is equivalent to the ‘`set write on`’ command inside GDB (see [Section 14.6 \[Patching\], page 134](#)).

-statistics

This option causes GDB to print statistics about time and memory usage after it completes each command and returns to the prompt.

-version This option causes GDB to print its version number and no-warranty blurb, and exit.

2.2 Quitting GDB

quit [*expression*]

q To exit GDB, use the `quit` command (abbreviated `q`), or type an end-of-file character (usually `C-d`). If you do not supply *expression*, GDB will terminate normally; otherwise it will terminate using the result of *expression* as the error code.

An interrupt (often `C-c`) does not exit from GDB, but rather terminates the action of any GDB command that is in progress and returns to GDB command level. It is safe to type the interrupt character at any time because GDB does not allow it to take effect until a time when it is safe.

If you have been using GDB to control an attached process or device, you can release it with the `detach` command (see [Section 4.7 \[Debugging an already-running process\], page 28](#)).

2.3 Shell commands

If you need to execute occasional shell commands during your debugging session, there is no need to leave or suspend GDB; you can just use the `shell` command.

shell *command string*

Invoke a standard shell to execute *command string*. If it exists, the environment variable `SHELL` determines which shell to run. Otherwise GDB uses the default shell (‘`/bin/sh`’ on Unix systems, ‘`COMMAND.COM`’ on MS-DOS, etc.).

The utility `make` is often needed in development environments. You do not have to use the `shell` command for this purpose in GDB:

make *make-args*

Execute the `make` program with the specified arguments. This is equivalent to ‘`shell make make-args`’.

2.4 Logging output

You may want to save the output of GDB commands to a file. There are several commands to control GDB’s logging.

```
set logging on
    Enable logging.

set logging off
    Disable logging.

set logging file file
    Change the name of the current logfile. The default logfile is ‘gdb.txt’.

set logging overwrite [on|off]
    By default, GDB will append to the logfile. Set overwrite if you want set
logging on to overwrite the logfile instead.

set logging redirect [on|off]
    By default, GDB output will go to both the terminal and the logfile. Set
redirect if you want output to go only to the log file.

show logging
    Show the current values of the logging settings.
```

3 GDB Commands

You can abbreviate a GDB command to the first few letters of the command name, if that abbreviation is unambiguous; and you can repeat certain GDB commands by typing just `\RET`. You can also use the `\TAB` key to get GDB to fill out the rest of a word in a command (or to show you the alternatives available, if there is more than one possibility).

3.1 Command syntax

A GDB command is a single line of input. There is no limit on how long it can be. It starts with a command name, which is followed by arguments whose meaning depends on the command name. For example, the command `step` accepts an argument which is the number of times to step, as in ‘`step 5`’. You can also use the `step` command with no arguments. Some commands do not allow any arguments.

GDB command names may always be truncated if that abbreviation is unambiguous. Other possible command abbreviations are listed in the documentation for individual commands. In some cases, even ambiguous abbreviations are allowed; for example, `s` is specially defined as equivalent to `step` even though there are other commands whose names start with `s`. You can test abbreviations by using them as arguments to the `help` command.

A blank line as input to GDB (typing just `\RET`) means to repeat the previous command. Certain commands (for example, `run`) will not repeat this way; these are commands whose unintentional repetition might cause trouble and which you are unlikely to want to repeat.

The `list` and `x` commands, when you repeat them with `\RET`, construct new arguments rather than repeating exactly as typed. This permits easy scanning of source or memory.

GDB can also use `\RET` in another way: to partition lengthy output, in a way similar to the common utility `more` (see [Section 19.4 \[Screen size\], page 181](#)). Since it is easy to press one `\RET` too many in this situation, GDB disables command repetition after any command that generates this sort of display.

Any text from a `#` to the end of the line is a comment; it does nothing. This is useful mainly in command files (see [Section 20.3 \[Command files\], page 189](#)).

The `C-o` binding is useful for repeating a complex sequence of commands. This command accepts the current line, like `RET`, and then fetches the next line relative to the current line from the history for editing.

3.2 Command completion

GDB can fill in the rest of a word in a command for you, if there is only one possibility; it can also show you what the valid possibilities are for the next word in a command, at any time. This works for GDB commands, GDB subcommands, and the names of symbols in your program.

Press the `\TAB` key whenever you want GDB to fill out the rest of a word. If there is only one possibility, GDB fills in the word, and waits for you to finish the command (or press `\RET` to enter it). For example, if you type

```
(gdb) info bre \TAB
```

GDB fills in the rest of the word ‘`breakpoints`’, since that is the only `info` subcommand beginning with ‘`bre`’:

```
(gdb) info breakpoints
```

You can either press **RET** at this point, to run the `info breakpoints` command, or backspace and enter something else, if ‘`breakpoints`’ does not look like the command you expected. (If you were sure you wanted `info breakpoints` in the first place, you might as well just type **RET** immediately after ‘`info bre`’, to exploit command abbreviations rather than command completion).

If there is more than one possibility for the next word when you press **TAB**, GDB sounds a bell. You can either supply more characters and try again, or just press **TAB** a second time; GDB displays all the possible completions for that word. For example, you might want to set a breakpoint on a subroutine whose name begins with ‘`make_`’, but when you type `b make_`**TAB** GDB just sounds the bell. Typing **TAB** again displays all the function names in your program that begin with those characters, for example:

```
(gdb) b make_ TAB
GDB sounds bell; press TAB again, to see:
make_a_section_from_file      make_environ
make_abs_section               make_function_type
make_blockvector                make_pointer_type
make_cleanup                   make_reference_type
make_command                   make_symbol_completion_list
(gdb) b make_
```

After displaying the available possibilities, GDB copies your partial input (‘`b make_`’ in the example) so you can finish the command.

If you just want to see the list of alternatives in the first place, you can press **M-?** rather than pressing **TAB** twice. **M-?** means **META** **?**. You can type this either by holding down a key designated as the **META** shift on your keyboard (if there is one) while typing **?**, or as **ESC** followed by **?**.

Sometimes the string you need, while logically a “word”, may contain parentheses or other characters that GDB normally excludes from its notion of a word. To permit word completion to work in this situation, you may enclose words in ‘ (single quote marks) in GDB commands.

The most likely situation where you might need this is in typing the name of a C++ function. This is because C++ allows function overloading (multiple definitions of the same function, distinguished by argument type). For example, when you want to set a breakpoint you may need to distinguish whether you mean the version of `name` that takes an `int` parameter, `name(int)`, or the version that takes a `float` parameter, `name(float)`. To use the word-completion facilities in this situation, type a single quote ‘ at the beginning of the function name. This alerts GDB that it may need to consider more information than usual when you press **TAB** or **M-?** to request word completion:

```
(gdb) b 'bubble( M-?
          bubble(double,double)    bubble(int,int)
(gdb) b 'bubble(
```

In some cases, GDB can tell that completing a name requires using quotes. When this happens, GDB inserts the quote for you (while completing as much as it can) if you do not type the quote in the first place:

```
(gdb) b bub TAB
GDB alters your input line to the following, and rings a bell:
(gdb) b 'bubble(
```

In general, GDB can tell that a quote is needed (and inserts it) if you have not yet started typing the argument list when you ask for completion on an overloaded symbol.

For more information about overloaded functions, see [Section 12.4.1.3 \[C++ expressions\]](#), [page 112](#). You can use the command `set overload-resolution off` to disable overload resolution; see [Section 12.4.1.7 \[GDB features for C++\]](#), [page 114](#).

3.3 Getting help

You can always ask GDB itself for information on its commands, using the command `help`.

```
help
h      You can use help (abbreviated h) with no arguments to display a short list of
          named classes of commands:
(gdb) help
List of classes of commands:

aliases -- Aliases of other commands
breakpoints -- Making program stop at certain points
data -- Examining data
files -- Specifying and examining files
internals -- Maintenance commands
obscure -- Obscure features
running -- Running the program
stack -- Examining the stack
status -- Status inquiries
support -- Support facilities
tracepoints -- Tracing of program execution without

                      stopping the program
user-defined -- User-defined commands

Type "help" followed by a class name for a list of
commands in that class.
Type "help" followed by command name for full
documentation.
Command name abbreviations are allowed if unambiguous.
(gdb)
```

```
help class
```

Using one of the general help classes as an argument, you can get a list of the individual commands in that class. For example, here is the help display for the class `status`:

```
(gdb) help status
Status inquiries.

List of commands:

info -- Generic command for showing things
about the program being debugged
show -- Generic command for showing things
about the debugger

Type "help" followed by command name for full
documentation.
Command name abbreviations are allowed if unambiguous.
(gdb)
```

help *command*

With a command name as **help** argument, GDB displays a short paragraph on how to use that command.

apropos *args*

The **apropos** *args* command searches through all of the GDB commands, and their documentation, for the regular expression specified in *args*. It prints out all matches found. For example:

```
apropos reload
```

results in:

```
set symbol-reloading -- Set dynamic symbol table reloading
                     multiple times in one run
show symbol-reloading -- Show dynamic symbol table reloading
                     multiple times in one run
```

complete *args*

The **complete** *args* command lists all the possible completions for the beginning of a command. Use *args* to specify the beginning of the command you want completed. For example:

```
complete i
```

results in:

```
if
ignore
info
inspect
```

This is intended for use by GNU Emacs.

In addition to **help**, you can use the GDB commands **info** and **show** to inquire about the state of your program, or the state of GDB itself. Each command supports many topics of inquiry; this manual introduces each of them in the appropriate context. The listings under **info** and under **show** in the Index point to all the sub-commands. See [Index], page 345.

info This command (abbreviated **i**) is for describing the state of your program. For example, you can list the arguments given to your program with **info args**, list the registers currently in use with **info registers**, or list the breakpoints you have set with **info breakpoints**. You can get a complete list of the **info** sub-commands with **help info**.

set You can assign the result of an expression to an environment variable with **set**. For example, you can set the GDB prompt to a \$-sign with **set prompt \$**.

show In contrast to **info**, **show** is for describing the state of GDB itself. You can change most of the things you can **show**, by using the related command **set**; for example, you can control what number system is used for displays with **set radix**, or simply inquire which is currently in use with **show radix**.

To display all the settable parameters and their current values, you can use **show** with no arguments; you may also use **info set**. Both commands produce the same display.

Here are three miscellaneous **show** subcommands, all of which are exceptional in lacking corresponding **set** commands:

show version

Show what version of GDB is running. You should include this information in GDB bug-reports. If multiple versions of GDB are in use at your site, you may need to determine which version of GDB you are running; as GDB evolves, new commands are introduced, and old ones may wither away. Also, many system vendors ship variant versions of GDB, and there are variant versions of GDB in GNU/Linux distributions as well. The version number is the same as the one announced when you start GDB.

show copying

Display information about permission for copying GDB.

show warranty

Display the GNU “NO WARRANTY” statement, or a warranty, if your version of GDB comes with one.

4 Running Programs Under GDB

When you run a program under GDB, you must first generate debugging information when you compile it.

You may start GDB with its arguments, if any, in an environment of your choice. If you are doing native debugging, you may redirect your program’s input and output, debug an already running process, or kill a child process.

4.1 Compiling for debugging

In order to debug a program effectively, you need to generate debugging information when you compile it. This debugging information is stored in the object file; it describes the data type of each variable or function and the correspondence between source line numbers and addresses in the executable code.

To request debugging information, specify the ‘`-g`’ option when you run the compiler.

Most compilers do not include information about preprocessor macros in the debugging information if you specify the ‘`-g`’ flag alone, because this information is rather large. Version 3.1 of GCC, the GNU C compiler, provides macro information if you specify the options ‘`-gdwarf-2`’ and ‘`-g3`’; the former option requests debugging information in the Dwarf 2 format, and the latter requests “extra information”. In the future, we hope to find more compact ways to represent macro information, so that it can be included with ‘`-g`’ alone.

Many C compilers are unable to handle the ‘`-g`’ and ‘`-O`’ options together. Using those compilers, you cannot generate optimized executables containing debugging information.

GCC, the GNU C compiler, supports ‘`-g`’ with or without ‘`-O`’, making it possible to debug optimized code. We recommend that you *always* use ‘`-g`’ whenever you compile a program. You may think your program is correct, but there is no sense in pushing your luck.

When you debug a program compiled with ‘`-g -O`’, remember that the optimizer is rearranging your code; the debugger shows you what is really there. Do not be too surprised when the execution path does not exactly match your source file! An extreme example: if you define a variable, but never use it, GDB never sees that variable—because the compiler optimizes it out of existence.

Some things do not work as well with ‘`-g -O`’ as with just ‘`-g`’, particularly on machines with instruction scheduling. If in doubt, recompile with ‘`-g`’ alone, and if this fixes the problem, please report it to us as a bug (including a test case!). See [Section 8.2 \[Variables\], page 66](#), for more information about debugging optimized code.

Older versions of the GNU C compiler permitted a variant option ‘`-gg`’ for debugging information. GDB no longer supports this format; if your GNU C compiler has this option, do not use it.

4.2 Starting your program

run

r

Use the `run` command to start your program under GDB. You must first specify the program name (except on VxWorks) with an argument to GDB (see [Chapter 2 \[Getting In and Out of GDB\], page 11](#)), or by using the `file` or `exec-file` command (see [Section 15.1 \[Commands to specify files\], page 135](#)).

If you are running your program in an execution environment that supports processes, `run` creates an inferior process and makes that process run your program. (In environments without processes, `run` jumps to the start of your program.)

The execution of a program is affected by certain information it receives from its superior. GDB provides ways to specify this information, which you must do *before* starting your program. (You can change it after starting your program, but such changes only affect your program the next time you start it.) This information may be divided into four categories:

The *arguments*.

Specify the arguments to give your program as the arguments of the `run` command. If a shell is available on your target, the shell is used to pass the arguments, so that you may use normal conventions (such as wildcard expansion or variable substitution) in describing the arguments. In Unix systems, you can control which shell is used with the `SHELL` environment variable. See [Section 4.3 \[Your program's arguments\], page 25](#).

The *environment*.

Your program normally inherits its environment from GDB, but you can use the GDB commands `set environment` and `unset environment` to change parts of the environment that affect your program. See [Section 4.4 \[Your program's environment\], page 26](#).

The *working directory*.

Your program inherits its working directory from GDB. You can set the GDB working directory with the `cd` command in GDB. See [Section 4.5 \[Your program's working directory\], page 27](#).

The *standard input and output*.

Your program normally uses the same device for standard input and standard output as GDB is using. You can redirect input and output in the `run` command line, or you can use the `tty` command to set a different device for your program. See [Section 4.6 \[Your program's input and output\], page 27](#).

Warning: While input and output redirection work, you cannot use pipes to pass the output of the program you are debugging to another program; if you attempt this, GDB is likely to wind up debugging the wrong program.

When you issue the `run` command, your program begins to execute immediately. See [Chapter 5 \[Stopping and continuing\], page 33](#), for discussion of how to arrange for your program to stop. Once your program has stopped, you may call functions in your program, using the `print` or `call` commands. See [Chapter 8 \[Examining Data\], page 65](#).

If the modification time of your symbol file has changed since the last time GDB read its symbols, GDB discards its symbol table, and reads it again. When it does this, GDB tries to retain your current breakpoints.

start The name of the main procedure can vary from language to language. With C or C++, the main procedure name is always `main`, but other languages such as Ada do not require a specific name for their main procedure. The debugger provides a convenient way to start the execution of the program and to stop at the beginning of the main procedure, depending on the language used.

The ‘`start`’ command does the equivalent of setting a temporary breakpoint at the beginning of the main procedure and then invoking the ‘`run`’ command.

Some programs contain an elaboration phase where some startup code is executed before the main program is called. This depends on the languages used to write your program. In C++ for instance, constructors for static and global objects are executed before `main` is called. It is therefore possible that the debugger stops before reaching the main procedure. However, the temporary breakpoint will remain to halt execution.

Specify the arguments to give to your program as arguments to the ‘`start`’ command. These arguments will be given verbatim to the underlying ‘`run`’ command. Note that the same arguments will be reused if no argument is provided during subsequent calls to ‘`start`’ or ‘`run`’.

It is sometimes necessary to debug the program during elaboration. In these cases, using the `start` command would stop the execution of your program too late, as the program would have already completed the elaboration phase. Under these circumstances, insert breakpoints in your elaboration code before running your program.

4.3 Your program’s arguments

The arguments to your program can be specified by the arguments of the `run` command. They are passed to a shell, which expands wildcard characters and performs redirection of I/O, and thence to your program. Your `SHELL` environment variable (if it exists) specifies what shell GDB uses. If you do not define `SHELL`, GDB uses the default shell (‘`/bin/sh`’ on Unix).

On non-Unix systems, the program is usually invoked directly by GDB, which emulates I/O redirection via the appropriate system calls, and the wildcard characters are expanded by the startup code of the program, not by the shell.

`run` with no arguments uses the same arguments used by the previous `run`, or those set by the `set args` command.

set args Specify the arguments to be used the next time your program is run. If `set args` has no arguments, `run` executes your program with no arguments. Once you have run your program with arguments, using `set args` before the next `run` is the only way to run it again without arguments.

show args Show the arguments to give your program when it is started.

4.4 Your program's environment

The *environment* consists of a set of environment variables and their values. Environment variables conventionally record such things as your user name, your home directory, your terminal type, and your search path for programs to run. Usually you set up environment variables with the shell and they are inherited by all the other programs you run. When debugging, it can be useful to try running your program with a modified environment without having to start GDB over again.

path *directory*

Add *directory* to the front of the PATH environment variable (the search path for executables) that will be passed to your program. The value of PATH used by GDB does not change. You may specify several directory names, separated by whitespace or by a system-dependent separator character (‘:’ on Unix, ‘;’ on MS-DOS and MS-Windows). If *directory* is already in the path, it is moved to the front, so it is searched sooner.

You can use the string ‘\$cwd’ to refer to whatever is the current working directory at the time GDB searches the path. If you use ‘.’ instead, it refers to the directory where you executed the **path** command. GDB replaces ‘.’ in the *directory* argument (with the current path) before adding *directory* to the search path.

show paths

Display the list of search paths for executables (the PATH environment variable).

show environment [*varname*]

Print the value of environment variable *varname* to be given to your program when it starts. If you do not supply *varname*, print the names and values of all environment variables to be given to your program. You can abbreviate **environment** as **env**.

set environment *varname* [=*value*]

Set environment variable *varname* to *value*. The value changes for your program only, not for GDB itself. *value* may be any string; the values of environment variables are just strings, and any interpretation is supplied by your program itself. The *value* parameter is optional; if it is eliminated, the variable is set to a null value.

For example, this command:

```
set env USER = foo
```

tells the debugged program, when subsequently run, that its user is named ‘foo’. (The spaces around ‘=’ are used for clarity here; they are not actually required.)

unset environment *varname*

Remove variable *varname* from the environment to be passed to your program. This is different from ‘**set env varname =’**; **unset environment** removes the variable from the environment, rather than assigning it an empty value.

Warning: On Unix systems, GDB runs your program using the shell indicated by your SHELL environment variable if it exists (or /bin/sh if not). If your SHELL variable names a

shell that runs an initialization file—such as ‘`.cshrc`’ for C-shell, or ‘`.bashrc`’ for BASH—any variables you set in that file affect your program. You may wish to move setting of environment variables to files that are only run when you sign on, such as ‘`.login`’ or ‘`.profile`’.

4.5 Your program’s working directory

Each time you start your program with `run`, it inherits its working directory from the current working directory of GDB. The GDB working directory is initially whatever it inherited from its parent process (typically the shell), but you can specify a new working directory in GDB with the `cd` command.

The GDB working directory also serves as a default for the commands that specify files for GDB to operate on. See [Section 15.1 \[Commands to specify files\]](#), page 135.

`cd directory`
Set the GDB working directory to *directory*.

`pwd` Print the GDB working directory.

4.6 Your program’s input and output

By default, the program you run under GDB does input and output to the same terminal that GDB uses. GDB switches the terminal to its own terminal modes to interact with you, but it records the terminal modes your program was using and switches back to them when you continue running your program.

`info terminal`
Displays information recorded by GDB about the terminal modes your program is using.

You can redirect your program’s input and/or output using shell redirection with the `run` command. For example,

`run > outfile`

starts your program, diverting its output to the file ‘`outfile`’.

Another way to specify where your program should do input and output is with the `tty` command. This command accepts a file name as argument, and causes this file to be the default for future `run` commands. It also resets the controlling terminal for the child process, for future `run` commands. For example,

`tty /dev/ttys`

directs that processes started with subsequent `run` commands default to do input and output on the terminal ‘`/dev/ttys`’ and have that as their controlling terminal.

An explicit redirection in `run` overrides the `tty` command’s effect on the input/output device, but not its effect on the controlling terminal.

When you use the `tty` command or redirect input in the `run` command, only the input for your program is affected. The input for GDB still comes from your terminal.

4.7 Debugging an already-running process

`attach process-id`

This command attaches to a running process—one that was started outside GDB. (`info files` shows your active targets.) The command takes as argument a process ID. The usual way to find out the process-ID of a Unix process is with the `ps` utility, or with the ‘`jobs -l`’ shell command.

`attach` does not repeat if you press `\RET` a second time after executing the command.

To use `attach`, your program must be running in an environment which supports processes; for example, `attach` does not work for programs on bare-board targets that lack an operating system. You must also have permission to send the process a signal.

When you use `attach`, the debugger finds the program running in the process first by looking in the current working directory, then (if the program is not found) by using the source file search path (see [Section 7.4 \[Specifying source directories\]](#), page 61). You can also use the `file` command to load the program. See [Section 15.1 \[Commands to Specify Files\]](#), page 135.

The first thing GDB does after arranging to debug the specified process is to stop it. You can examine and modify an attached process with all the GDB commands that are ordinarily available when you start processes with `run`. You can insert breakpoints; you can step and continue; you can modify storage. If you would rather the process continue running, you may use the `continue` command after attaching GDB to the process.

`detach` When you have finished debugging the attached process, you can use the `detach` command to release it from GDB control. Detaching the process continues its execution. After the `detach` command, that process and GDB become completely independent once more, and you are ready to `attach` another process or start one with `run`. `detach` does not repeat if you press `\RET` again after executing the command.

If you exit GDB or use the `run` command while you have an attached process, you kill that process. By default, GDB asks for confirmation if you try to do either of these things; you can control whether or not you need to confirm by using the `set confirm` command (see [Section 19.7 \[Optional warnings and messages\]](#), page 183).

4.8 Killing the child process

`kill` Kill the child process in which your program is running under GDB.

This command is useful if you wish to debug a core dump instead of a running process. GDB ignores any core dump file while your program is running.

On some operating systems, a program cannot be executed outside GDB while you have breakpoints set on it inside GDB. You can use the `kill` command in this situation to permit running your program outside the debugger.

The `kill` command is also useful if you wish to recompile and relink your program, since on many systems it is impossible to modify an executable file while it is running in a process. In this case, when you next type `run`, GDB notices that the file has changed, and reads the symbol table again (while trying to preserve your current breakpoint settings).

4.9 Debugging programs with multiple threads

In some operating systems, such as HP-UX and Solaris, a single program may have more than one *thread* of execution. The precise semantics of threads differ from one operating system to another, but in general the threads of a single program are akin to multiple processes—except that they share one address space (that is, they can all examine and modify the same variables). On the other hand, each thread has its own registers and execution stack, and perhaps private memory.

GDB provides these facilities for debugging multi-thread programs:

- automatic notification of new threads
- ‘`thread threadno`’, a command to switch among threads
- ‘`info threads`’, a command to inquire about existing threads
- ‘`thread apply [threadno] [all] args`’, a command to apply a command to a list of threads
- thread-specific breakpoints

Warning: These facilities are not yet available on every GDB configuration where the operating system supports threads. If your GDB does not support threads, these commands have no effect. For example, a system without thread support shows no output from ‘`info threads`’, and always rejects the `thread` command, like this:

```
(gdb) info threads
(gdb) thread 1
Thread ID 1 not known. Use the "info threads" command to
see the IDs of currently known threads.
```

The GDB thread debugging facility allows you to observe all threads while your program runs—but whenever GDB takes control, one thread in particular is always the focus of debugging. This thread is called the *current thread*. Debugging commands show program information from the perspective of the current thread.

Whenever GDB detects a new thread in your program, it displays the target system’s identification for the thread with a message in the form ‘[New *systag*]’. *systag* is a thread identifier whose form varies depending on the particular system. For example, on LynxOS, you might see

```
[New process 35 thread 27]
```

when GDB notices a new thread. In contrast, on an SGI system, the *systag* is simply something like ‘process 368’, with no further qualifier.

For debugging purposes, GDB associates its own thread number—always a single integer—with each thread in your program.

`info threads`

Display a summary of all threads currently in your program. GDB displays for each thread (in this order):

1. the thread number assigned by GDB
2. the target system’s thread identifier (*systag*)
3. the current stack frame summary for that thread

An asterisk '*' to the left of the GDB thread number indicates the current thread.

For example,

```
(gdb) info threads
 3 process 35 thread 27 0x34e5 in sigpause ()
 2 process 35 thread 23 0x34e5 in sigpause ()
* 1 process 35 thread 13 main (argc=1, argv=0x7fffffff8)
  at threadtest.c:68
```

On HP-UX systems:

For debugging purposes, GDB associates its own thread number—a small integer assigned in thread-creation order—with each thread in your program.

Whenever GDB detects a new thread in your program, it displays both GDB's thread number and the target system's identification for the thread with a message in the form '[New *systag*]'. *systag* is a thread identifier whose form varies depending on the particular system. For example, on HP-UX, you see

```
[New thread 2 (system thread 26594)]
```

when GDB notices a new thread.

info threads

Display a summary of all threads currently in your program. GDB displays for each thread (in this order):

1. the thread number assigned by GDB
2. the target system's thread identifier (*systag*)
3. the current stack frame summary for that thread

An asterisk '*' to the left of the GDB thread number indicates the current thread.

For example,

```
(gdb) info threads
* 3 system thread 26607 worker (wptr=0x7b09c318 "@") \
  at quicksort.c:137
  2 system thread 26606 0x7b0030d8 in __ksleep () \
  from /usr/lib/libc.2
  1 system thread 27905 0x7b003498 in _brk () \
  from /usr/lib/libc.2
```

thread *threadno*

Make thread number *threadno* the current thread. The command argument *threadno* is the internal GDB thread number, as shown in the first field of the 'info threads' display. GDB responds by displaying the system identifier of the thread you selected, and its current stack frame summary:

```
(gdb) thread 2
[Switching to process 35 thread 23]
0x34e5 in sigpause ()
```

As with the '[New ...]' message, the form of the text after 'Switching to' depends on your system's conventions for identifying threads.

thread apply [*threadno*] [*all*] *args*

The **thread apply** command allows you to apply a command to one or more threads. Specify the numbers of the threads that you want affected with the

command argument *threadno*. *threadno* is the internal GDB thread number, as shown in the first field of the ‘`info threads`’ display. To apply a command to all threads, use `thread apply all args`.

Whenever GDB stops your program, due to a breakpoint or a signal, it automatically selects the thread where that breakpoint or signal happened. GDB alerts you to the context switch with a message of the form ‘[Switching to *systag*]’ to identify the thread.

See [Section 5.4 \[Stopping and starting multi-thread programs\], page 51](#), for more information about how GDB behaves when you stop and start programs with multiple threads.

See [Section 5.1.2 \[Setting watchpoints\], page 37](#), for information about watchpoints in programs with multiple threads.

4.10 Debugging programs with multiple processes

On most systems, GDB has no special support for debugging programs which create additional processes using the `fork` function. When a program forks, GDB will continue to debug the parent process and the child process will run unimpeded. If you have set a breakpoint in any code which the child then executes, the child will get a `SIGTRAP` signal which (unless it catches the signal) will cause it to terminate.

However, if you want to debug the child process there is a workaround which isn’t too painful. Put a call to `sleep` in the code which the child process executes after the fork. It may be useful to sleep only if a certain environment variable is set, or a certain file exists, so that the delay need not occur when you don’t want to run GDB on the child. While the child is sleeping, use the `ps` program to get its process ID. Then tell GDB (a new invocation of GDB if you are also debugging the parent process) to attach to the child process (see [Section 4.7 \[Attach\], page 28](#)). From that point on you can debug the child process just like any other process which you attached to.

On some systems, GDB provides support for debugging programs that create additional processes using the `fork` or `vfork` functions. Currently, the only platforms with this feature are HP-UX (11.x and later only?) and GNU/Linux (kernel version 2.5.60 and later).

By default, when a program forks, GDB will continue to debug the parent process and the child process will run unimpeded.

If you want to follow the child process instead of the parent process, use the command `set follow-fork-mode`.

`set follow-fork-mode mode`

Set the debugger response to a program call of `fork` or `vfork`. A call to `fork` or `vfork` creates a new process. The *mode* can be:

parent The original process is debugged after a fork. The child process runs unimpeded. This is the default.

child The new process is debugged after a fork. The parent process runs unimpeded.

`show follow-fork-mode`

Display the current debugger response to a `fork` or `vfork` call.

If you ask to debug a child process and a `vfork` is followed by an `exec`, GDB executes the new target up to the first breakpoint in the new target. If you have a breakpoint set on `main` in your original program, the breakpoint will also be set on the child process's `main`.

When a child process is spawned by `vfork`, you cannot debug the child or parent until an `exec` call completes.

If you issue a `run` command to GDB after an `exec` call executes, the new target restarts. To restart the parent process, use the `file` command with the parent executable name as its argument.

You can use the `catch` command to make GDB stop whenever a `fork`, `vfork`, or `exec` call is made. See [Section 5.1.3 \[Setting catchpoints\]](#), page 39.

5 Stopping and Continuing

The principal purposes of using a debugger are so that you can stop your program before it terminates; or so that, if your program runs into trouble, you can investigate and find out why.

Inside GDB, your program may stop for any of several reasons, such as a signal, a breakpoint, or reaching a new line after a GDB command such as `step`. You may then examine and change variables, set new breakpoints or remove old ones, and then continue execution. Usually, the messages shown by GDB provide ample explanation of the status of your program—but you can also explicitly request this information at any time.

`info program`

Display information about the status of your program: whether it is running or not, what process it is, and why it stopped.

5.1 Breakpoints, watchpoints, and catchpoints

A *breakpoint* makes your program stop whenever a certain point in the program is reached. For each breakpoint, you can add conditions to control in finer detail whether your program stops. You can set breakpoints with the `break` command and its variants (see [Section 5.1.1 \[Setting breakpoints\], page 34](#)), to specify the place where your program should stop by line number, function name or exact address in the program.

In HP-UX, SunOS 4.x, SVR4, and Alpha OSF/1 configurations, you can set breakpoints in shared libraries before the executable is run. There is a minor limitation on HP-UX systems: you must wait until the executable is run in order to set breakpoints in shared library routines that are not called directly by the program (for example, routines that are arguments in a `pthread_create` call).

A *watchpoint* is a special breakpoint that stops your program when the value of an expression changes. You must use a different command to set watchpoints (see [Section 5.1.2 \[Setting watchpoints\], page 37](#)), but aside from that, you can manage a watchpoint like any other breakpoint: you enable, disable, and delete both breakpoints and watchpoints using the same commands.

You can arrange to have values from your program displayed automatically whenever GDB stops at a breakpoint. See [Section 8.6 \[Automatic display\], page 70](#).

A *catchpoint* is another special breakpoint that stops your program when a certain kind of event occurs, such as the throwing of a C++ exception or the loading of a library. As with watchpoints, you use a different command to set a catchpoint (see [Section 5.1.3 \[Setting catchpoints\], page 39](#)), but aside from that, you can manage a catchpoint like any other breakpoint. (To stop when your program receives a signal, use the `handle` command; see [Section 5.3 \[Signals\], page 50](#).)

GDB assigns a number to each breakpoint, watchpoint, or catchpoint when you create it; these numbers are successive integers starting with one. In many of the commands for controlling various features of breakpoints you use the breakpoint number to say which breakpoint you want to change. Each breakpoint may be *enabled* or *disabled*; if disabled, it has no effect on your program until you enable it again.

Some GDB commands accept a range of breakpoints on which to operate. A breakpoint range is either a single breakpoint number, like ‘5’, or two such numbers, in increasing

order, separated by a hyphen, like ‘5–7’. When a breakpoint range is given to a command, all breakpoints in that range are operated on.

5.1.1 Setting breakpoints

Breakpoints are set with the `break` command (abbreviated `b`). The debugger convenience variable ‘\$bpnum’ records the number of the breakpoint you’ve set most recently; see [Section 8.9 \[Convenience variables\], page 77](#), for a discussion of what you can do with convenience variables.

You have several ways to say where the breakpoint should go.

`break function`

Set a breakpoint at entry to function *function*. When using source languages that permit overloading of symbols, such as C++, *function* may refer to more than one possible place to break. See [Section 5.1.8 \[Breakpoint menus\], page 45](#), for a discussion of that situation.

`break +offset`

`break -offset`

Set a breakpoint some number of lines forward or back from the position at which execution stopped in the currently selected *stack frame*. (See [Section 6.1 \[Frames\], page 53](#), for a description of stack frames.)

`break linenum`

Set a breakpoint at line *linenum* in the current source file. The current source file is the last file whose source text was printed. The breakpoint will stop your program just before it executes any of the code on that line.

`break filename:linenum`

Set a breakpoint at line *linenum* in source file *filename*.

`break filename:function`

Set a breakpoint at entry to function *function* found in file *filename*. Specifying a file name as well as a function name is superfluous except when multiple files contain similarly named functions.

`break *address`

Set a breakpoint at address *address*. You can use this to set breakpoints in parts of your program which do not have debugging information or source files.

`break`

When called without any arguments, `break` sets a breakpoint at the next instruction to be executed in the selected stack frame (see [Chapter 6 \[Examining the Stack\], page 53](#)). In any selected frame but the innermost, this makes your program stop as soon as control returns to that frame. This is similar to the effect of a `finish` command in the frame inside the selected frame—except that `finish` does not leave an active breakpoint. If you use `break` without an argument in the innermost frame, GDB stops the next time it reaches the current location; this may be useful inside loops.

GDB normally ignores breakpoints when it resumes execution, until at least one instruction has been executed. If it did not do this, you would be unable to proceed past a breakpoint without first disabling the breakpoint. This rule applies whether or not the breakpoint already existed when your program stopped.

break ... if cond

Set a breakpoint with condition *cond*; evaluate the expression *cond* each time the breakpoint is reached, and stop only if the value is nonzero—that is, if *cond* evaluates as true. ‘...’ stands for one of the possible arguments described above (or no argument) specifying where to break. See [Section 5.1.6 \[Break conditions\], page 42](#), for more information on breakpoint conditions.

tbreak args

Set a breakpoint enabled only for one stop. *args* are the same as for the **break** command, and the breakpoint is set in the same way, but the breakpoint is automatically deleted after the first time your program stops there. See [Section 5.1.5 \[Disabling breakpoints\], page 41](#).

hbreak args

Set a hardware-assisted breakpoint. *args* are the same as for the **break** command and the breakpoint is set in the same way, but the breakpoint requires hardware support and some target hardware may not have this support. The main purpose of this is EPROM/ROM code debugging, so you can set a breakpoint at an instruction without changing the instruction. This can be used with the new trap-generation provided by SPARClite DSU and some x86-based targets. These targets will generate traps when a program accesses some data or instruction address that is assigned to the debug registers. However the hardware breakpoint registers can take a limited number of breakpoints. For example, on the DSU, only two data breakpoints can be set at a time, and GDB will reject this command if more than two are used. Delete or disable unused hardware breakpoints before setting new ones (see [Section 5.1.5 \[Disabling\], page 41](#)). See [Section 5.1.6 \[Break conditions\], page 42](#). See [\[set remote hardware-breakpoint-limit\]](#), page 153.

thbreak args

Set a hardware-assisted breakpoint enabled only for one stop. *args* are the same as for the **hbreak** command and the breakpoint is set in the same way. However, like the **tbreak** command, the breakpoint is automatically deleted after the first time your program stops there. Also, like the **hbreak** command, the breakpoint requires hardware support and some target hardware may not have this support. See [Section 5.1.5 \[Disabling breakpoints\], page 41](#). See also [Section 5.1.6 \[Break conditions\], page 42](#).

rbreak regex

Set breakpoints on all functions matching the regular expression *regex*. This command sets an unconditional breakpoint on all matches, printing a list of all breakpoints it set. Once these breakpoints are set, they are treated just like the breakpoints set with the **break** command. You can delete them, disable them, or make them conditional the same way as any other breakpoint.

The syntax of the regular expression is the standard one used with tools like ‘**grep**’. Note that this is different from the syntax used by shells, so for instance **foo*** matches all functions that include an **fo** followed by zero or more **os**. There is an implicit **.*** leading and trailing the regular expression you supply, so to match only functions that begin with **foo**, use **^foo**.

When debugging C++ programs, `rbreak` is useful for setting breakpoints on overloaded functions that are not members of any special classes.

The `rbreak` command can be used to set breakpoints in **all** the functions in a program, like this:

```
(gdb) rbreak .
```

```
info breakpoints [n]
info break [n]
info watchpoints [n]
```

Print a table of all breakpoints, watchpoints, and catchpoints set and not deleted, with the following columns for each breakpoint:

Breakpoint Numbers

Type Breakpoint, watchpoint, or catchpoint.

Disposition

Whether the breakpoint is marked to be disabled or deleted when hit.

Enabled or Disabled

Enabled breakpoints are marked with ‘y’. ‘n’ marks breakpoints that are not enabled.

Address Where the breakpoint is in your program, as a memory address. If the breakpoint is pending (see below for details) on a future load of a shared library, the address will be listed as ‘<PENDING>’.

What Where the breakpoint is in the source for your program, as a file and line number. For a pending breakpoint, the original string passed to the breakpoint command will be listed as it cannot be resolved until the appropriate shared library is loaded in the future.

If a breakpoint is conditional, `info break` shows the condition on the line following the affected breakpoint; breakpoint commands, if any, are listed after that. A pending breakpoint is allowed to have a condition specified for it. The condition is not parsed for validity until a shared library is loaded that allows the pending breakpoint to resolve to a valid location.

`info break` with a breakpoint number *n* as argument lists only that breakpoint. The convenience variable `$_` and the default examining-address for the `x` command are set to the address of the last breakpoint listed (see [Section 8.5 \[Examining memory\], page 69](#)).

`info break` displays a count of the number of times the breakpoint has been hit. This is especially useful in conjunction with the `ignore` command. You can ignore a large number of breakpoint hits, look at the breakpoint info to see how many times the breakpoint was hit, and then run again, ignoring one less than that number. This will get you quickly to the last hit of that breakpoint.

GDB allows you to set any number of breakpoints at the same place in your program. There is nothing silly or meaningless about this. When the breakpoints are conditional, this is even useful (see [Section 5.1.6 \[Break conditions\], page 42](#)).

If a specified breakpoint location cannot be found, it may be due to the fact that the location is in a shared library that is yet to be loaded. In such a case, you may want GDB to create a special breakpoint (known as a *pending breakpoint*) that attempts to resolve itself in the future when an appropriate shared library gets loaded.

Pending breakpoints are useful to set at the start of your GDB session for locations that you know will be dynamically loaded later by the program being debugged. When shared libraries are loaded, a check is made to see if the load resolves any pending breakpoint locations. If a pending breakpoint location gets resolved, a regular breakpoint is created and the original pending breakpoint is removed.

GDB provides some additional commands for controlling pending breakpoint support:

set breakpoint pending auto

This is the default behavior. When GDB cannot find the breakpoint location, it queries you whether a pending breakpoint should be created.

set breakpoint pending on

This indicates that an unrecognized breakpoint location should automatically result in a pending breakpoint being created.

set breakpoint pending off

This indicates that pending breakpoints are not to be created. Any unrecognized breakpoint location results in an error. This setting does not affect any pending breakpoints previously created.

show breakpoint pending

Show the current behavior setting for creating pending breakpoints.

Normal breakpoint operations apply to pending breakpoints as well. You may specify a condition for a pending breakpoint and/or commands to run when the breakpoint is reached. You can also enable or disable the pending breakpoint. When you specify a condition for a pending breakpoint, the parsing of the condition will be deferred until the point where the pending breakpoint location is resolved. Disabling a pending breakpoint tells GDB to not attempt to resolve the breakpoint on any subsequent shared library load. When a pending breakpoint is re-enabled, GDB checks to see if the location is already resolved. This is done because any number of shared library loads could have occurred since the time the breakpoint was disabled and one or more of these loads could resolve the location.

GDB itself sometimes sets breakpoints in your program for special purposes, such as proper handling of `longjmp` (in C programs). These internal breakpoints are assigned negative numbers, starting with `-1`; ‘`info breakpoints`’ does not display them. You can see these breakpoints with the GDB maintenance command ‘`maint info breakpoints`’ (see [maint info breakpoints], page 289).

5.1.2 Setting watchpoints

You can use a watchpoint to stop execution whenever the value of an expression changes, without having to predict a particular place where this may happen.

Depending on your system, watchpoints may be implemented in software or hardware. GDB does software watchpointing by single-stepping your program and testing the variable’s value each time, which is hundreds of times slower than normal execution. (But this may

still be worth it, to catch errors where you have no clue what part of your program is the culprit.)

On some systems, such as HP-UX, GNU/Linux and some other x86-based targets, GDB includes support for hardware watchpoints, which do not slow down the running of your program.

watch *expr*

Set a watchpoint for an expression. GDB will break when *expr* is written into by the program and its value changes.

rwatch *expr*

Set a watchpoint that will break when *watch expr* is read by the program.

awatch *expr*

Set a watchpoint that will break when *expr* is either read or written into by the program.

info watchpoints

This command prints a list of watchpoints, breakpoints, and catchpoints; it is the same as **info break**.

GDB sets a *hardware watchpoint* if possible. Hardware watchpoints execute very quickly, and the debugger reports a change in value at the exact instruction where the change occurs. If GDB cannot set a hardware watchpoint, it sets a software watchpoint, which executes more slowly and reports the change in value at the next statement, not the instruction, after the change occurs.

When you issue the **watch** command, GDB reports

Hardware watchpoint num: *expr*

if it was able to set a hardware watchpoint.

Currently, the **awatch** and **rwatch** commands can only set hardware watchpoints, because accesses to data that don't change the value of the watched expression cannot be detected without examining every instruction as it is being executed, and GDB does not do that currently. If GDB finds that it is unable to set a hardware breakpoint with the **awatch** or **rwatch** command, it will print a message like this:

Expression cannot be implemented with read/access watchpoint.

Sometimes, GDB cannot set a hardware watchpoint because the data type of the watched expression is wider than what a hardware watchpoint on the target machine can handle. For example, some systems can only watch regions that are up to 4 bytes wide; on such systems you cannot set hardware watchpoints for an expression that yields a double-precision floating-point number (which is typically 8 bytes wide). As a work-around, it might be possible to break the large region into a series of smaller ones and watch them with separate watchpoints.

If you set too many hardware watchpoints, GDB might be unable to insert all of them when you resume the execution of your program. Since the precise number of active watchpoints is unknown until such time as the program is about to be resumed, GDB might not be able to warn you about this when you set the watchpoints, and the warning will be printed only when the program is resumed:

```
Hardware watchpoint num: Could not insert watchpoint
```

If this happens, delete or disable some of the watchpoints.

The SPARClite DSU will generate traps when a program accesses some data or instruction address that is assigned to the debug registers. For the data addresses, DSU facilitates the `watch` command. However the hardware breakpoint registers can only take two data watchpoints, and both watchpoints must be the same kind. For example, you can set two watchpoints with `watch` commands, two with `rwatch` commands, or two with `awatch` commands, but you cannot set one watchpoint with one command and the other with a different command. GDB will reject the command if you try to mix watchpoints. Delete or disable unused watchpoint commands before setting new ones.

If you call a function interactively using `print` or `call`, any watchpoints you have set will be inactive until GDB reaches another kind of breakpoint or the call completes.

GDB automatically deletes watchpoints that watch local (automatic) variables, or expressions that involve such variables, when they go out of scope, that is, when the execution leaves the block in which these variables were defined. In particular, when the program being debugged terminates, *all* local variables go out of scope, and so only watchpoints that watch global variables remain set. If you rerun the program, you will need to set all such watchpoints again. One way of doing that would be to set a code breakpoint at the entry to the `main` function and when it breaks, set all the watchpoints.

Warning: In multi-thread programs, watchpoints have only limited usefulness.

With the current watchpoint implementation, GDB can only watch the value of an expression *in a single thread*. If you are confident that the expression can only change due to the current thread's activity (and if you are also confident that no other thread can become current), then you can use watchpoints as usual. However, GDB may not notice when a non-current thread's activity changes the expression.

HP-UX Warning: In multi-thread programs, software watchpoints have only limited usefulness. If GDB creates a software watchpoint, it can only watch the value of an expression *in a single thread*. If you are confident that the expression can only change due to the current thread's activity (and if you are also confident that no other thread can become current), then you can use software watchpoints as usual. However, GDB may not notice when a non-current thread's activity changes the expression. (Hardware watchpoints, in contrast, watch an expression in all threads.)

See [set remote hardware-watchpoint-limit], page 153.

5.1.3 Setting catchpoints

You can use *catchpoints* to cause the debugger to stop for certain kinds of program events, such as C++ exceptions or the loading of a shared library. Use the `catch` command to set a catchpoint.

`catch event`

Stop when event occurs. event can be any of the following:

`throw` The throwing of a C++ exception.

`catch` The catching of a C++ exception.

```

exec      A call to exec. This is currently only available for HP-UX.
fork       A call to fork. This is currently only available for HP-UX.
vfork     A call to vfork. This is currently only available for HP-UX.
load
load libname
          The dynamic loading of any shared library, or the loading of the
          library libname. This is currently only available for HP-UX.

unload
unload libname
          The unloading of any dynamically loaded shared library, or the
          unloading of the library libname. This is currently only available
          for HP-UX.

tcatch event
          Set a catchpoint that is enabled only for one stop. The catchpoint is automatically deleted after the first time the event is caught.

Use the info break command to list the current catchpoints.

There are currently some limitations to C++ exception handling (catch throw and catch catch) in GDB:


- If you call a function interactively, GDB normally returns control to you when the function has finished executing. If the call raises an exception, however, the call may bypass the mechanism that returns control to you and cause your program either to abort or to simply continue running until it hits a breakpoint, catches a signal that GDB is listening for, or exits. This is the case even if you set a catchpoint for the exception; catchpoints on exceptions are disabled within interactive calls.
- You cannot raise an exception interactively.
- You cannot install an exception handler interactively.

```

Sometimes **catch** is not the best way to debug exception handling: if you need to know exactly where an exception is raised, it is better to stop *before* the exception handler is called, since that way you can see the stack before any unwinding takes place. If you set a breakpoint in an exception handler instead, it may not be easy to find out where the exception was raised.

To stop just before an exception handler is called, you need some knowledge of the implementation. In the case of GNU C++, exceptions are raised by calling a library function named **__raise_exception** which has the following ANSI C interface:

```

/* addr is where the exception identifier is stored.
   id is the exception identifier. */
void __raise_exception (void **addr, void *id);

```

To make the debugger catch all exceptions before any stack unwinding takes place, set a breakpoint on **__raise_exception** (see [Section 5.1 \[Breakpoints; watchpoints; and exceptions\], page 33](#)).

With a conditional breakpoint (see [Section 5.1.6 \[Break conditions\], page 42](#)) that depends on the value of *id*, you can stop your program when a specific exception is raised. You can use multiple conditional breakpoints to stop your program when any of a number of exceptions are raised.

5.1.4 Deleting breakpoints

It is often necessary to eliminate a breakpoint, watchpoint, or catchpoint once it has done its job and you no longer want your program to stop there. This is called *deleting* the breakpoint. A breakpoint that has been deleted no longer exists; it is forgotten.

With the `clear` command you can delete breakpoints according to where they are in your program. With the `delete` command you can delete individual breakpoints, watchpoints, or catchpoints by specifying their breakpoint numbers.

It is not necessary to delete a breakpoint to proceed past it. GDB automatically ignores breakpoints on the first instruction to be executed when you continue execution without changing the execution address.

`clear` Delete any breakpoints at the next instruction to be executed in the selected stack frame (see [Section 6.3 \[Selecting a frame\], page 55](#)). When the innermost frame is selected, this is a good way to delete a breakpoint where your program just stopped.

`clear function`

`clear filename:function`

Delete any breakpoints set at entry to the function *function*.

`clear linenum`

`clear filename:linenum`

Delete any breakpoints set at or within the code of the specified line.

`delete [breakpoints] [range...]`

Delete the breakpoints, watchpoints, or catchpoints of the breakpoint ranges specified as arguments. If no argument is specified, delete all breakpoints (GDB asks confirmation, unless you have `set confirm off`). You can abbreviate this command as `d`.

5.1.5 Disabling breakpoints

Rather than deleting a breakpoint, watchpoint, or catchpoint, you might prefer to *disable* it. This makes the breakpoint inoperative as if it had been deleted, but remembers the information on the breakpoint so that you can *enable* it again later.

You disable and enable breakpoints, watchpoints, and catchpoints with the `enable` and `disable` commands, optionally specifying one or more breakpoint numbers as arguments. Use `info break` or `info watch` to print a list of breakpoints, watchpoints, and catchpoints if you do not know which numbers to use.

A breakpoint, watchpoint, or catchpoint can have any of four different states of enablement:

- Enabled. The breakpoint stops your program. A breakpoint set with the `break` command starts out in this state.
- Disabled. The breakpoint has no effect on your program.
- Enabled once. The breakpoint stops your program, but then becomes disabled.
- Enabled for deletion. The breakpoint stops your program, but immediately after it does so it is deleted permanently. A breakpoint set with the `tbreak` command starts out in this state.

You can use the following commands to enable or disable breakpoints, watchpoints, and catchpoints:

`disable [breakpoints] [range...]`

Disable the specified breakpoints—or all breakpoints, if none are listed. A disabled breakpoint has no effect but is not forgotten. All options such as ignore-counts, conditions and commands are remembered in case the breakpoint is enabled again later. You may abbreviate `disable` as `dis`.

`enable [breakpoints] [range...]`

Enable the specified breakpoints (or all defined breakpoints). They become effective once again in stopping your program.

`enable [breakpoints] once range...`

Enable the specified breakpoints temporarily. GDB disables any of these breakpoints immediately after stopping your program.

`enable [breakpoints] delete range...`

Enable the specified breakpoints to work once, then die. GDB deletes any of these breakpoints as soon as your program stops there.

Except for a breakpoint set with `tbreak` (see [Section 5.1.1 \[Setting breakpoints\], page 34](#)), breakpoints that you set are initially enabled; subsequently, they become disabled or enabled only when you use one of the commands above. (The command `until` can set and delete a breakpoint of its own, but it does not change the state of your other breakpoints; see [Section 5.2 \[Continuing and stepping\], page 46](#).)

5.1.6 Break conditions

The simplest sort of breakpoint breaks every time your program reaches a specified place. You can also specify a *condition* for a breakpoint. A condition is just a Boolean expression in your programming language (see [Section 8.1 \[Expressions\], page 65](#)). A breakpoint with a condition evaluates the expression each time your program reaches it, and your program stops only if the condition is *true*.

This is the converse of using assertions for program validation; in that situation, you want to stop when the assertion is violated—that is, when the condition is false. In C, if you want to test an assertion expressed by the condition `assert`, you should set the condition ‘`! assert`’ on the appropriate breakpoint.

Conditions are also accepted for watchpoints; you may not need them, since a watchpoint is inspecting the value of an expression anyhow—but it might be simpler, say, to just set a watchpoint on a variable name, and specify a condition that tests whether the new value is an interesting one.

Break conditions can have side effects, and may even call functions in your program. This can be useful, for example, to activate functions that log program progress, or to use your own print functions to format special data structures. The effects are completely predictable unless there is another enabled breakpoint at the same address. (In that case, GDB might see the other breakpoint first and stop your program without checking the condition of this one.) Note that breakpoint commands are usually more convenient and flexible than break conditions for the purpose of performing side effects when a breakpoint is reached (see [Section 5.1.7 \[Breakpoint command lists\], page 44](#)).

Break conditions can be specified when a breakpoint is set, by using ‘`if`’ in the arguments to the `break` command. See [Section 5.1.1 \[Setting breakpoints\]](#), page 34. They can also be changed at any time with the `condition` command.

You can also use the `if` keyword with the `watch` command. The `catch` command does not recognize the `if` keyword; `condition` is the only way to impose a further condition on a catchpoint.

`condition bnum expression`

Specify `expression` as the break condition for breakpoint, watchpoint, or catchpoint number `bnum`. After you set a condition, breakpoint `bnum` stops your program only if the value of `expression` is true (nonzero, in C). When you use `condition`, GDB checks `expression` immediately for syntactic correctness, and to determine whether symbols in it have referents in the context of your breakpoint. If `expression` uses symbols not referenced in the context of the breakpoint, GDB prints an error message:

```
No symbol "foo" in current context.
```

GDB does not actually evaluate `expression` at the time the `condition` command (or a command that sets a breakpoint with a condition, like `break if ...`) is given, however. See [Section 8.1 \[Expressions\]](#), page 65.

`condition bnum`

Remove the condition from breakpoint number `bnum`. It becomes an ordinary unconditional breakpoint.

A special case of a breakpoint condition is to stop only when the breakpoint has been reached a certain number of times. This is so useful that there is a special way to do it, using the `ignore count` of the breakpoint. Every breakpoint has an ignore count, which is an integer. Most of the time, the ignore count is zero, and therefore has no effect. But if your program reaches a breakpoint whose ignore count is positive, then instead of stopping, it just decrements the ignore count by one and continues. As a result, if the ignore count value is `n`, the breakpoint does not stop the next `n` times your program reaches it.

`ignore bnum count`

Set the ignore count of breakpoint number `bnum` to `count`. The next `count` times the breakpoint is reached, your program’s execution does not stop; other than to decrement the ignore count, GDB takes no action.

To make the breakpoint stop the next time it is reached, specify a count of zero.

When you use `continue` to resume execution of your program from a breakpoint, you can specify an ignore count directly as an argument to `continue`, rather than using `ignore`. See [Section 5.2 \[Continuing and stepping\]](#), page 46.

If a breakpoint has a positive ignore count and a condition, the condition is not checked. Once the ignore count reaches zero, GDB resumes checking the condition.

You could achieve the effect of the ignore count with a condition such as ‘`$foo-- <= 0`’ using a debugger convenience variable that is decremented each time. See [Section 8.9 \[Convenience variables\]](#), page 77.

Ignore counts apply to breakpoints, watchpoints, and catchpoints.

5.1.7 Breakpoint command lists

You can give any breakpoint (or watchpoint or catchpoint) a series of commands to execute when your program stops due to that breakpoint. For example, you might want to print the values of certain expressions, or enable other breakpoints.

```
commands [bnum]
... command-list ...
end      Specify a list of commands for breakpoint number bnum. The commands themselves appear on the following lines. Type a line containing just end to terminate the commands.
```

To remove all commands from a breakpoint, type `commands` and follow it immediately with `end`; that is, give no commands.

With no *bnum* argument, `commands` refers to the last breakpoint, watchpoint, or catchpoint set (not to the breakpoint most recently encountered).

Pressing `<RET>` as a means of repeating the last GDB command is disabled within a *command-list*.

You can use breakpoint commands to start your program up again. Simply use the `continue` command, or `step`, or any other command that resumes execution.

Any other commands in the command list, after a command that resumes execution, are ignored. This is because any time you resume execution (even with a simple `next` or `step`), you may encounter another breakpoint—which could have its own command list, leading to ambiguities about which list to execute.

If the first command you specify in a command list is `silent`, the usual message about stopping at a breakpoint is not printed. This may be desirable for breakpoints that are to print a specific message and then continue. If none of the remaining commands print anything, you see no sign that the breakpoint was reached. `silent` is meaningful only at the beginning of a breakpoint command list.

The commands `echo`, `output`, and `printf` allow you to print precisely controlled output, and are often useful in silent breakpoints. See [Section 20.4 \[Commands for controlled output\], page 190](#).

For example, here is how you could use breakpoint commands to print the value of `x` at entry to `foo` whenever `x` is positive.

```
break foo if x>0
commands
silent
printf "x is %d\n",x
cont
end
```

One application for breakpoint commands is to compensate for one bug so you can test for another. Put a breakpoint just after the erroneous line of code, give it a condition to detect the case in which something erroneous has been done, and give it commands to assign correct values to any variables that need them. End with the `continue` command so that your program does not stop, and start with the `silent` command so that no output is produced. Here is an example:

```
break 403
commands
```

```

silent
set x = y + 4
cont
end

```

5.1.8 Breakpoint menus

Some programming languages (notably C++ and Objective-C) permit a single function name to be defined several times, for application in different contexts. This is called *overloading*. When a function name is overloaded, ‘**break function**’ is not enough to tell GDB where you want a breakpoint. If you realize this is a problem, you can use something like ‘**break function(types)**’ to specify which particular version of the function you want. Otherwise, GDB offers you a menu of numbered choices for different possible breakpoints, and waits for your selection with the prompt ‘>’. The first two options are always ‘[0] cancel’ and ‘[1] all’. Typing 1 sets a breakpoint at each definition of *function*, and typing 0 aborts the **break** command without setting any new breakpoints.

For example, the following session excerpt shows an attempt to set a breakpoint at the overloaded symbol `String::after`. We choose three particular definitions of that function name:

```

(gdb) b String::after
[0] cancel
[1] all
[2] file:String.cc; line number:867
[3] file:String.cc; line number:860
[4] file:String.cc; line number:875
[5] file:String.cc; line number:853
[6] file:String.cc; line number:846
[7] file:String.cc; line number:735
> 2 4 6
Breakpoint 1 at 0xb26c: file String.cc, line 867.
Breakpoint 2 at 0xb344: file String.cc, line 875.
Breakpoint 3 at 0xafcc: file String.cc, line 846.
Multiple breakpoints were set.
Use the "delete" command to delete unwanted
breakpoints.
(gdb)

```

5.1.9 “Cannot insert breakpoints”

Under some operating systems, breakpoints cannot be used in a program if any other process is running that program. In this situation, attempting to run or continue a program with a breakpoint causes GDB to print an error message:

```

Cannot insert breakpoints.
The same program may be running in another process.

```

When this happens, you have three ways to proceed:

1. Remove or disable the breakpoints, then continue.
2. Suspend GDB, and copy the file containing your program to a new name. Resume GDB and use the **exec-file** command to specify that GDB should run your program under that name. Then start your program again.
3. Relink your program so that the text segment is nonsharable, using the linker option ‘**-N**’. The operating system limitation may not apply to nonsharable executables.

A similar message can be printed if you request too many active hardware-assisted breakpoints and watchpoints:

```
Stopped; cannot insert breakpoints.  
You may have requested too many hardware breakpoints and watchpoints.
```

This message is printed when you attempt to resume the program, since only then GDB knows exactly how many hardware breakpoints and watchpoints it needs to insert.

When this message is printed, you need to disable or remove some of the hardware-assisted breakpoints and watchpoints, and then continue.

5.1.10 “Breakpoint address adjusted...”

Some processor architectures place constraints on the addresses at which breakpoints may be placed. For architectures thus constrained, GDB will attempt to adjust the breakpoint’s address to comply with the constraints dictated by the architecture.

One example of such an architecture is the Fujitsu FR-V. The FR-V is a VLIW architecture in which a number of RISC-like instructions may be bundled together for parallel execution. The FR-V architecture constrains the location of a breakpoint instruction within such a bundle to the instruction with the lowest address. GDB honors this constraint by adjusting a breakpoint’s address to the first in the bundle.

It is not uncommon for optimized code to have bundles which contain instructions from different source statements, thus it may happen that a breakpoint’s address will be adjusted from one source statement to another. Since this adjustment may significantly alter GDB’s breakpoint related behavior from what the user expects, a warning is printed when the breakpoint is first set and also when the breakpoint is hit.

A warning like the one below is printed when setting a breakpoint that’s been subject to address adjustment:

```
warning: Breakpoint address adjusted from 0x00010414 to 0x00010410.
```

Such warnings are printed both for user settable and GDB’s internal breakpoints. If you see one of these warnings, you should verify that a breakpoint set at the adjusted address will have the desired affect. If not, the breakpoint in question may be removed and other breakpoints may be set which will have the desired behavior. E.g., it may be sufficient to place the breakpoint at a later instruction. A conditional breakpoint may also be useful in some cases to prevent the breakpoint from triggering too often.

GDB will also issue a warning when stopping at one of these adjusted breakpoints:

```
warning: Breakpoint 1 address previously adjusted from 0x00010414  
to 0x00010410.
```

When this warning is encountered, it may be too late to take remedial action except in cases where the breakpoint is hit earlier or more frequently than expected.

5.2 Continuing and stepping

Continuing means resuming program execution until your program completes normally. In contrast, *stepping* means executing just one more “step” of your program, where “step” may mean either one line of source code, or one machine instruction (depending on what particular command you use). Either when continuing or when stepping, your program may stop even sooner, due to a breakpoint or a signal. (If it stops due to a signal, you may want to use `handle`, or use ‘`signal 0`’ to resume execution. See [Section 5.3 \[Signals\]](#), page 50.)

```
continue [ignore-count]
c [ignore-count]
fg [ignore-count]
```

Resume program execution, at the address where your program last stopped; any breakpoints set at that address are bypassed. The optional argument *ignore-count* allows you to specify a further number of times to ignore a breakpoint at this location; its effect is like that of `ignore` (see [Section 5.1.6 \[Break conditions\], page 42](#)).

The argument *ignore-count* is meaningful only when your program stopped due to a breakpoint. At other times, the argument to `continue` is ignored.

The synonyms `c` and `fg` (for *foreground*, as the debugged program is deemed to be the foreground program) are provided purely for convenience, and have exactly the same behavior as `continue`.

To resume execution at a different place, you can use `return` (see [Section 14.4 \[Returning from a function\], page 133](#)) to go back to the calling function; or `jump` (see [Section 14.2 \[Continuing at a different address\], page 132](#)) to go to an arbitrary location in your program.

A typical technique for using stepping is to set a breakpoint (see [Section 5.1 \[Breakpoints; watchpoints; and catchpoints\], page 33](#)) at the beginning of the function or the section of your program where a problem is believed to lie, run your program until it stops at that breakpoint, and then step through the suspect area, examining the variables that are interesting, until you see the problem happen.

step Continue running your program until control reaches a different source line, then stop it and return control to GDB. This command is abbreviated `s`.

Warning: If you use the `step` command while control is within a function that was compiled without debugging information, execution proceeds until control reaches a function that does have debugging information. Likewise, it will not step into a function which is compiled without debugging information. To step through functions without debugging information, use the `stepli` command, described below.

The `step` command only stops at the first instruction of a source line. This prevents the multiple stops that could otherwise occur in `switch` statements, `for` loops, etc. `step` continues to stop if a function that has debugging information is called within the line. In other words, `step` *steps inside* any functions called within the line.

Also, the `step` command only enters a function if there is line number information for the function. Otherwise it acts like the `next` command. This avoids problems when using `cc -g1` on MIPS machines. Previously, `step` entered subroutines if there was any debugging information about the routine.

step count

Continue running as in `step`, but do so *count* times. If a breakpoint is reached, or a signal not related to stepping occurs before *count* steps, stepping stops right away.

next [count]

Continue to the next source line in the current (innermost) stack frame. This is similar to **step**, but function calls that appear within the line of code are executed without stopping. Execution stops when control reaches a different line of code at the original stack level that was executing when you gave the **next** command. This command is abbreviated **n**.

An argument *count* is a repeat count, as for **step**.

The **next** command only stops at the first instruction of a source line. This prevents multiple stops that could otherwise occur in **switch** statements, **for** loops, etc.

set step-mode**set step-mode on**

The **set step-mode on** command causes the **step** command to stop at the first instruction of a function which contains no debug line information rather than stepping over it.

This is useful in cases where you may be interested in inspecting the machine instructions of a function which has no symbolic info and do not want GDB to automatically skip over this function.

set step-mode off

Causes the **step** command to step over any functions which contains no debug information. This is the default.

finish

Continue running until just after function in the selected stack frame returns. Print the returned value (if any).

Contrast this with the **return** command (see [Section 14.4 \[Returning from a function\], page 133](#)).

until**u**

Continue running until a source line past the current line, in the current stack frame, is reached. This command is used to avoid single stepping through a loop more than once. It is like the **next** command, except that when **until** encounters a jump, it automatically continues execution until the program counter is greater than the address of the jump.

This means that when you reach the end of a loop after single stepping though it, **until** makes your program continue execution until it exits the loop. In contrast, a **next** command at the end of a loop simply steps back to the beginning of the loop, which forces you to step through the next iteration.

until always stops your program if it attempts to exit the current stack frame. **until** may produce somewhat counterintuitive results if the order of machine code does not match the order of the source lines. For example, in the following excerpt from a debugging session, the **f (frame)** command shows that execution is stopped at line 206; yet when we use **until**, we get to line 195:

```
(gdb) f
#0  main (argc=4, argv=0xf7ffffae8) at m4.c:206
206          expand_input();
(gdb) until
```

```
195      for ( ; argc > 0; NEXTARG) {
```

This happened because, for execution efficiency, the compiler had generated code for the loop closure test at the end, rather than the start, of the loop—even though the test in a C `for`-loop is written before the body of the loop. The `until` command appeared to step back to the beginning of the loop when it advanced to this expression; however, it has not really gone to an earlier statement—not in terms of the actual machine code.

`until` with no argument works by means of single instruction stepping, and hence is slower than `until` with an argument.

```
until location
u location
```

Continue running your program until either the specified location is reached, or the current stack frame returns. *location* is any of the forms of argument acceptable to `break` (see [Section 5.1.1 \[Setting breakpoints\]](#), page 34). This form of the command uses breakpoints, and hence is quicker than `until` without an argument. The specified location is actually reached only if it is in the current frame. This implies that `until` can be used to skip over recursive function invocations. For instance in the code below, if the current location is line 96, issuing `until 99` will execute the program up to line 99 in the same invocation of factorial, i.e. after the inner invocations have returned.

```
94 int factorial (int value)
95 {
96     if (value > 1) {
97         value *= factorial (value - 1);
98     }
99     return (value);
100 }
```

```
advance location
```

Continue running the program up to the given location. An argument is required, anything of the same form as arguments for the `break` command. Execution will also stop upon exit from the current stack frame. This command is similar to `until`, but `advance` will not skip over recursive function calls, and the target location doesn't have to be in the same frame as the current one.

```
stepi
stepi arg
```

`si` Execute one machine instruction, then stop and return to the debugger.

It is often useful to do ‘`display/i $pc`’ when stepping by machine instructions. This makes GDB automatically display the next instruction to be executed, each time your program stops. See [Section 8.6 \[Automatic display\]](#), page 70.

An argument is a repeat count, as in `step`.

```
nexti
nexti arg
```

`ni` Execute one machine instruction, but if it is a function call, proceed until the function returns.

An argument is a repeat count, as in `next`.

5.3 Signals

A signal is an asynchronous event that can happen in a program. The operating system defines the possible kinds of signals, and gives each kind a name and a number. For example, in Unix **SIGINT** is the signal a program gets when you type an interrupt character (often *C-c*); **SIGSEGV** is the signal a program gets from referencing a place in memory far away from all the areas in use; **SIGALRM** occurs when the alarm clock timer goes off (which happens only if your program has requested an alarm).

Some signals, including **SIGALRM**, are a normal part of the functioning of your program. Others, such as **SIGSEGV**, indicate errors; these signals are *fatal* (they kill your program immediately) if the program has not specified in advance some other way to handle the signal. **SIGINT** does not indicate an error in your program, but it is normally fatal so it can carry out the purpose of the interrupt: to kill the program.

GDB has the ability to detect any occurrence of a signal in your program. You can tell GDB in advance what to do for each kind of signal.

Normally, GDB is set up to let the non-erroneous signals like **SIGALRM** be silently passed to your program (so as not to interfere with their role in the program's functioning) but to stop your program immediately whenever an error signal happens. You can change these settings with the `handle` command.

`info signals`

`info handle`

Print a table of all the kinds of signals and how GDB has been told to handle each one. You can use this to see the signal numbers of all the defined types of signals.

`info handle` is an alias for `info signals`.

`handle signal keywords...`

Change the way GDB handles signal *signal*. *signal* can be the number of a signal or its name (with or without the ‘SIG’ at the beginning); a list of signal numbers of the form ‘*low-high*’; or the word ‘*all*’, meaning all the known signals. The *keywords* say what change to make.

The keywords allowed by the `handle` command can be abbreviated. Their full names are:

`nostop` GDB should not stop your program when this signal happens. It may still print a message telling you that the signal has come in.

`stop` GDB should stop your program when this signal happens. This implies the `print` keyword as well.

`print` GDB should print a message when this signal happens.

`noprint` GDB should not mention the occurrence of the signal at all. This implies the `nostop` keyword as well.

`pass`

`noignore` GDB should allow your program to see this signal; your program can handle the signal, or else it may terminate if the signal is fatal and not handled. `pass` and `noignore` are synonyms.

`nopass`

`ignore` GDB should not allow your program to see this signal. `nopass` and `ignore` are synonyms.

When a signal stops your program, the signal is not visible to the program until you continue. Your program sees the signal then, if `pass` is in effect for the signal in question *at that time*. In other words, after GDB reports a signal, you can use the `handle` command with `pass` or `nopass` to control whether your program sees that signal when you continue.

The default is set to `nostop`, `noprint`, `pass` for non-erroneous signals such as `SIGALRM`, `SIGWINCH` and `SIGCHLD`, and to `stop`, `print`, `pass` for the erroneous signals.

You can also use the `signal` command to prevent your program from seeing a signal, or cause it to see a signal it normally would not see, or to give it any signal at any time. For example, if your program stopped due to some sort of memory reference error, you might store correct values into the erroneous variables and continue, hoping to see more execution; but your program would probably terminate immediately as a result of the fatal signal once it saw the signal. To prevent this, you can continue with ‘signal 0’. See [Section 14.3 \[Giving your program a signal\]](#), page 133.

5.4 Stopping and starting multi-thread programs

When your program has multiple threads (see [Section 4.9 \[Debugging programs with multiple threads\]](#), page 29), you can choose whether to set breakpoints on all threads, or on a particular thread.

```
break linespec thread threadno
break linespec thread threadno if ...
```

linespec specifies source lines; there are several ways of writing them, but the effect is always to specify some source line.

Use the qualifier ‘`thread threadno`’ with a breakpoint command to specify that you only want GDB to stop the program when a particular thread reaches this breakpoint. `threadno` is one of the numeric thread identifiers assigned by GDB, shown in the first column of the ‘`info threads`’ display.

If you do not specify ‘`thread threadno`’ when you set a breakpoint, the breakpoint applies to *all* threads of your program.

You can use the `thread` qualifier on conditional breakpoints as well; in this case, place ‘`thread threadno`’ before the breakpoint condition, like this:

```
(gdb) break frik.c:13 thread 28 if bartab > lim
```

Whenever your program stops under GDB for any reason, *all* threads of execution stop, not just the current thread. This allows you to examine the overall state of the program, including switching between threads, without worrying that things may change underfoot.

There is an unfortunate side effect. If one thread stops for a breakpoint, or for some other reason, and another thread is blocked in a system call, then the system call may return prematurely. This is a consequence of the interaction between multiple threads and the signals that GDB uses to implement breakpoints and other events that stop execution.

To handle this problem, your program should check the return value of each system call and react appropriately. This is good programming style anyways.

For example, do not write code like this:

```
sleep (10);
```

The call to `sleep` will return early if a different thread stops at a breakpoint or for some other reason.

Instead, write this:

```
int unslept = 10;
while (unslept > 0)
    unslept = sleep (unslept);
```

A system call is allowed to return early, so the system is still conforming to its specification. But GDB does cause your multi-threaded program to behave differently than it would without GDB.

Also, GDB uses internal breakpoints in the thread library to monitor certain events such as thread creation and thread destruction. When such an event happens, a system call in another thread may return prematurely, even though your program does not appear to stop.

Conversely, whenever you restart the program, *all* threads start executing. *This is true even when single-stepping* with commands like `step` or `next`.

In particular, GDB cannot single-step all threads in lockstep. Since thread scheduling is up to your debugging target's operating system (not controlled by GDB), other threads may execute more than one statement while the current thread completes a single step. Moreover, in general other threads stop in the middle of a statement, rather than at a clean statement boundary, when the program stops.

You might even find your program stopped in another thread after continuing or even single-stepping. This happens whenever some other thread runs into a breakpoint, a signal, or an exception before the first thread completes whatever you requested.

On some OSes, you can lock the OS scheduler and thus allow only a single thread to run.

`set scheduler-locking mode`

Set the scheduler locking mode. If it is `off`, then there is no locking and any thread may run at any time. If `on`, then only the current thread may run when the inferior is resumed. The `step` mode optimizes for single-stepping. It stops other threads from “seizing the prompt” by preempting the current thread while you are stepping. Other threads will only rarely (or never) get a chance to run when you step. They are more likely to run when you ‘`next`’ over a function call, and they are completely free to run when you use commands like ‘`continue`’, ‘`until`’, or ‘`finish`’. However, unless another thread hits a breakpoint during its timeslice, they will never steal the GDB prompt away from the thread that you are debugging.

`show scheduler-locking`

Display the current scheduler locking mode.

6 Examining the Stack

When your program has stopped, the first thing you need to know is where it stopped and how it got there.

Each time your program performs a function call, information about the call is generated. That information includes the location of the call in your program, the arguments of the call, and the local variables of the function being called. The information is saved in a block of data called a *stack frame*. The stack frames are allocated in a region of memory called the *call stack*.

When your program stops, the GDB commands for examining the stack allow you to see all of this information.

One of the stack frames is *selected* by GDB and many GDB commands refer implicitly to the selected frame. In particular, whenever you ask GDB for the value of a variable in your program, the value is found in the selected frame. There are special GDB commands to select whichever frame you are interested in. See [Section 6.3 \[Selecting a frame\], page 55](#).

When your program stops, GDB automatically selects the currently executing frame and describes it briefly, similar to the `frame` command (see [Section 6.4 \[Information about a frame\], page 56](#)).

6.1 Stack frames

The call stack is divided up into contiguous pieces called *stack frames*, or *frames* for short; each frame is the data associated with one call to one function. The frame contains the arguments given to the function, the function's local variables, and the address at which the function is executing.

When your program is started, the stack has only one frame, that of the function `main`. This is called the *initial* frame or the *outermost* frame. Each time a function is called, a new frame is made. Each time a function returns, the frame for that function invocation is eliminated. If a function is recursive, there can be many frames for the same function. The frame for the function in which execution is actually occurring is called the *innermost* frame. This is the most recently created of all the stack frames that still exist.

Inside your program, stack frames are identified by their addresses. A stack frame consists of many bytes, each of which has its own address; each kind of computer has a convention for choosing one byte whose address serves as the address of the frame. Usually this address is kept in a register called the *frame pointer register* while execution is going on in that frame.

GDB assigns numbers to all existing stack frames, starting with zero for the innermost frame, one for the frame that called it, and so on upward. These numbers do not really exist in your program; they are assigned by GDB to give you a way of designating stack frames in GDB commands.

Some compilers provide a way to compile functions so that they operate without stack frames. (For example, the `gcc` option

```
'-fomit-frame-pointer'
```

generates functions without a frame.) This is occasionally done with heavily used library functions to save the frame setup time. GDB has limited facilities for dealing with these function invocations. If the innermost function invocation has no stack frame, GDB

nevertheless regards it as though it had a separate frame, which is numbered zero as usual, allowing correct tracing of the function call chain. However, GDB has no provision for frameless functions elsewhere in the stack.

frame args

The **frame** command allows you to move from one stack frame to another, and to print the stack frame you select. **args** may be either the address of the frame or the stack frame number. Without an argument, **frame** prints the current stack frame.

select-frame

The **select-frame** command allows you to move from one stack frame to another without printing the frame. This is the silent version of **frame**.

6.2 Backtraces

A backtrace is a summary of how your program got where it is. It shows one line per frame, for many frames, starting with the currently executing frame (frame zero), followed by its caller (frame one), and on up the stack.

backtrace

bt Print a backtrace of the entire stack: one line per frame for all frames in the stack.

You can stop the backtrace at any time by typing the system interrupt character, normally *C-c*.

backtrace n

bt n Similar, but print only the innermost *n* frames.

backtrace -n

bt -n Similar, but print only the outermost *n* frames.

The names **where** and **info stack** (abbreviated **info s**) are additional aliases for **backtrace**.

Each line in the backtrace shows the frame number and the function name. The program counter value is also shown—unless you use **set print address off**. The backtrace also shows the source file name and line number, as well as the arguments to the function. The program counter value is omitted if it is at the beginning of the code for that line number.

Here is an example of a backtrace. It was made with the command ‘**bt 3**’, so it shows the innermost three frames.

```
#0  m4_traceon (obs=0x24eb0, argc=1, argv=0x2b8c8)
    at builtin.c:993
#1  0x6e38 in expand_macro (sym=0x2b600) at macro.c:242
#2  0x6840 in expand_token (obs=0x0, t=177664, td=0xf7fffb08)
    at macro.c:71
(More stack frames follow...)
```

The display for frame zero does not begin with a program counter value, indicating that your program has stopped at the beginning of the code for line 993 of **builtin.c**.

Most programs have a standard user entry point—a place where system libraries and startup code transition into user code. For C this is **main**. When GDB finds the entry

function in a backtrace it will terminate the backtrace, to avoid tracing into highly system-specific (and generally uninteresting) code.

If you need to examine the startup code, or limit the number of levels in a backtrace, you can change this behavior:

```
set backtrace past-main
set backtrace past-main on
```

Backtraces will continue past the user entry point.

```
set backtrace past-main off
```

Backtraces will stop when they encounter the user entry point. This is the default.

```
show backtrace past-main
```

Display the current user entry point backtrace policy.

```
set backtrace limit n
```

```
set backtrace limit 0
```

Limit the backtrace to *n* levels. A value of zero means unlimited.

```
show backtrace limit
```

Display the current limit on backtrace levels.

6.3 Selecting a frame

Most commands for examining the stack and other data in your program work on whichever stack frame is selected at the moment. Here are the commands for selecting a stack frame; all of them finish by printing a brief description of the stack frame just selected.

```
frame n
```

```
f n
```

Select frame number *n*. Recall that frame zero is the innermost (currently executing) frame, frame one is the frame that called the innermost one, and so on. The highest-numbered frame is the one for `main`.

```
frame addr
```

```
f addr
```

Select the frame at address *addr*. This is useful mainly if the chaining of stack frames has been damaged by a bug, making it impossible for GDB to assign numbers properly to all frames. In addition, this can be useful when your program has multiple stacks and switches between them.

On the SPARC architecture, `frame` needs two addresses to select an arbitrary frame: a frame pointer and a stack pointer.

On the MIPS and Alpha architecture, it needs two addresses: a stack pointer and a program counter.

On the 29k architecture, it needs three addresses: a register stack pointer, a program counter, and a memory stack pointer.

```
up n
```

Move *n* frames up the stack. For positive numbers *n*, this advances toward the outermost frame, to higher frame numbers, to frames that have existed longer. *n* defaults to one.

```
down n
```

Move *n* frames down the stack. For positive numbers *n*, this advances toward the innermost frame, to lower frame numbers, to frames that were created more recently. *n* defaults to one. You may abbreviate `down` as `do`.

All of these commands end by printing two lines of output describing the frame. The first line shows the frame number, the function name, the arguments, and the source file and line number of execution in that frame. The second line shows the text of that source line.

For example:

```
(gdb) up
#1 0x22f0 in main (argc=1, argv=0xf7ffffbf4, env=0xf7ffffbf4)
      at env.c:10
10          read_input_file (argv[i]);
```

After such a printout, the `list` command with no arguments prints ten lines centered on the point of execution in the frame. You can also edit the program at the point of execution with your favorite editing program by typing `edit`. See [Section 7.1 \[Printing source lines\], page 59](#), for details.

```
up-silently n
down-silently n
```

These two commands are variants of `up` and `down`, respectively; they differ in that they do their work silently, without causing display of the new frame. They are intended primarily for use in GDB command scripts, where the output might be unnecessary and distracting.

6.4 Information about a frame

There are several other commands to print information about the selected stack frame.

`frame`

`f`

When used without any argument, this command does not change which frame is selected, but prints a brief description of the currently selected stack frame. It can be abbreviated `f`. With an argument, this command is used to select a stack frame. See [Section 6.3 \[Selecting a frame\], page 55](#).

`info frame`

`info f`

This command prints a verbose description of the selected stack frame, including:

- the address of the frame
- the address of the next frame down (called by this frame)
- the address of the next frame up (caller of this frame)
- the language in which the source code corresponding to this frame is written
- the address of the frame's arguments
- the address of the frame's local variables
- the program counter saved in it (the address of execution in the caller frame)
- which registers were saved in the frame

The verbose description is useful when something has gone wrong that has made the stack format fail to fit the usual conventions.

info frame *addr*
info f *addr*

Print a verbose description of the frame at address *addr*, without selecting that frame. The selected frame remains unchanged by this command. This requires the same kind of address (more than one for some architectures) that you specify in the **frame** command. See [Section 6.3 \[Selecting a frame\]](#), page 55.

info args Print the arguments of the selected frame, each on a separate line.

info locals

Print the local variables of the selected frame, each on a separate line. These are all variables (declared either static or automatic) accessible at the point of execution of the selected frame.

info catch

Print a list of all the exception handlers that are active in the current stack frame at the current point of execution. To see other exception handlers, visit the associated frame (using the **up**, **down**, or **frame** commands); then type **info catch**. See [Section 5.1.3 \[Setting breakpoints\]](#), page 39.

7 Examining Source Files

GDB can print parts of your program’s source, since the debugging information recorded in the program tells GDB what source files were used to build it. When your program stops, GDB spontaneously prints the line where it stopped. Likewise, when you select a stack frame (see [Section 6.3 \[Selecting a frame\], page 55](#)), GDB prints the line where execution in that frame has stopped. You can print other portions of source files by explicit command.

If you use GDB through its GNU Emacs interface, you may prefer to use Emacs facilities to view source; see [Chapter 23 \[Using GDB under GNU Emacs\], page 201](#).

7.1 Printing source lines

To print lines from a source file, use the `list` command (abbreviated `l`). By default, ten lines are printed. There are several ways to specify what part of the file you want to print.

Here are the forms of the `list` command most commonly used:

`list linenum`

Print lines centered around line number *linenum* in the current source file.

`list function`

Print lines centered around the beginning of function *function*.

`list`

Print more lines. If the last lines printed were printed with a `list` command, this prints lines following the last lines printed; however, if the last line printed was a solitary line printed as part of displaying a stack frame (see [Chapter 6 \[Examining the Stack\], page 53](#)), this prints lines centered around that line.

`list -`

Print lines just before the lines last printed.

By default, GDB prints ten source lines with any of these forms of the `list` command. You can change this using `set listsize`:

`set listsize count`

Make the `list` command display *count* source lines (unless the `list` argument explicitly specifies some other number).

`show listsize`

Display the number of lines that `list` prints.

Repeating a `list` command with `(RET)` discards the argument, so it is equivalent to typing just `list`. This is more useful than listing the same lines again. An exception is made for an argument of ‘-’; that argument is preserved in repetition so that each repetition moves up in the source file.

In general, the `list` command expects you to supply zero, one or two *linespecs*. Linespecs specify source lines; there are several ways of writing them, but the effect is always to specify some source line. Here is a complete description of the possible arguments for `list`:

`list linespec`

Print lines centered around the line specified by *linespec*.

`list first,last`

Print lines from *first* to *last*. Both arguments are *linespecs*.

list ,last	Print lines ending with <i>last</i> .
list <i>first</i>,	Print lines starting with <i>first</i> .
list +	Print lines just after the lines last printed.
list -	Print lines just before the lines last printed.
list	As described in the preceding table.

Here are the ways of specifying a single source line—all the kinds of linespec.

number	Specifies line <i>number</i> of the current source file. When a list command has two linespecs, this refers to the same source file as the first linespec.
+offset	Specifies the line <i>offset</i> lines after the last line printed. When used as the second linespec in a list command that has two, this specifies the line <i>offset</i> lines down from the first linespec.
-offset	Specifies the line <i>offset</i> lines before the last line printed.
filename:number	Specifies line <i>number</i> in the source file <i>filename</i> .
function	Specifies the line that begins the body of the function <i>function</i> . For example: in C, this is the line with the open brace.
filename:function	Specifies the line of the open-brace that begins the body of the function <i>function</i> in the file <i>filename</i> . You only need the file name with a function name to avoid ambiguity when there are identically named functions in different source files.
*address	Specifies the line containing the program address <i>address</i> . <i>address</i> may be any expression.

7.2 Editing source files

To edit the lines in a source file, use the **edit** command. The editing program of your choice is invoked with the current line set to the active line in the program. Alternatively, there are several ways to specify what part of the file you want to print if you want to see other parts of the program.

Here are the forms of the **edit** command most commonly used:

edit	Edit the current source file at the active line number in the program.
edit <i>number</i>	Edit the current source file with <i>number</i> as the active line number.
edit <i>function</i>	Edit the file containing <i>function</i> at the beginning of its definition.
edit <i>filename:number</i>	Specifies line <i>number</i> in the source file <i>filename</i> .

edit filename:function

Specifies the line that begins the body of the function *function* in the file *filename*. You only need the file name with a function name to avoid ambiguity when there are identically named functions in different source files.

edit *address

Specifies the line containing the program address *address*. *address* may be any expression.

7.2.1 Choosing your editor

You can customize GDB to use any editor you want¹. By default, it is ‘/bin/ex’, but you can change this by setting the environment variable `EDITOR` before using GDB. For example, to configure GDB to use the `vi` editor, you could use these commands with the `sh` shell:

```
EDITOR=/usr/bin/vi
export EDITOR
gdb ...
```

or in the `csh` shell,

```
setenv EDITOR /usr/bin/vi
gdb ...
```

7.3 Searching source files

There are two commands for searching through the current source file for a regular expression.

forward-search regexp**search regexp**

The command ‘`forward-search regexp`’ checks each line, starting with the one following the last line listed, for a match for *regexp*. It lists the line that is found. You can use the synonym ‘`search regexp`’ or abbreviate the command name as `fo`.

reverse-search regexp

The command ‘`reverse-search regexp`’ checks each line, starting with the one before the last line listed and going backward, for a match for *regexp*. It lists the line that is found. You can abbreviate this command as `rev`.

7.4 Specifying source directories

Executable programs sometimes do not record the directories of the source files from which they were compiled, just the names. Even when they do, the directories could be moved between the compilation and your debugging session. GDB has a list of directories to search for source files; this is called the *source path*. Each time GDB wants a source file, it tries all the directories in the list, in the order they are present in the list, until it finds a file with the desired name.

For example, suppose an executable references the file ‘`/usr/src/foo-1.0/lib/foo.c`’, and our source path is ‘`/mnt/cross`’. The file is first looked up literally; if this fails,

¹ The only restriction is that your editor (say `ex`), recognizes the following command-line syntax:

```
ex +number file
```

The optional numeric value `+number` specifies the number of the line in the file where to start editing.

`'/mnt/cross/usr/src/foo-1.0/lib/foo.c'` is tried; if this fails, `'/mnt/cross/foo.c'` is opened; if this fails, an error message is printed. GDB does not look up the parts of the source file name, such as `'/mnt/cross/src/foo-1.0/lib/foo.c'`. Likewise, the subdirectories of the source path are not searched: if the source path is `'/mnt/cross'`, and the binary refers to `'foo.c'`, GDB would not find it under `'/mnt/cross/usr/src/foo-1.0/lib'`.

Plain file names, relative file names with leading directories, file names containing dots, etc. are all treated as described above; for instance, if the source path is `'/mnt/cross'`, and the source file is recorded as `'../lib/foo.c'`, GDB would first try `'../lib/foo.c'`, then `'/mnt/cross/../lib/foo.c'`, and after that—`'/mnt/cross/foo.c'`.

Note that the executable search path is *not* used to locate the source files. Neither is the current working directory, unless it happens to be in the source path.

Whenever you reset or rearrange the source path, GDB clears out any information it has cached about where source files are found and where each line is in the file.

When you start GDB, its source path includes only ‘`cdir`’ and ‘`cwd`’, in that order. To add other directories, use the `directory` command.

```
directory dirname ...
dir dirname ...
```

Add directory `dirname` to the front of the source path. Several directory names may be given to this command, separated by ‘`:`’ (‘`;`’ on MS-DOS and MS-Windows, where ‘`:`’ usually appears as part of absolute file names) or white-space. You may specify a directory that is already in the source path; this moves it forward, so GDB searches it sooner.

You can use the string ‘`$cdir`’ to refer to the compilation directory (if one is recorded), and ‘`$cwd`’ to refer to the current working directory. ‘`$cwd`’ is not the same as ‘`.`’—the former tracks the current working directory as it changes during your GDB session, while the latter is immediately expanded to the current directory at the time you add an entry to the source path.

`directory`

Reset the source path to empty again. This requires confirmation.

`show directories`

Print the source path: show which directories it contains.

If your source path is cluttered with directories that are no longer of interest, GDB may sometimes cause confusion by finding the wrong versions of source. You can correct the situation as follows:

1. Use `directory` with no argument to reset the source path to empty.
2. Use `directory` with suitable arguments to reinstall the directories you want in the source path. You can add all the directories in one command.

7.5 Source and machine code

You can use the command `info line` to map source lines to program addresses (and vice versa), and the command `disassemble` to display a range of addresses as machine instructions. When run under GNU Emacs mode, the `info line` command causes the arrow to point to the line specified. Also, `info line` prints addresses in symbolic form as well as hex.

info line *linespec*

Print the starting and ending addresses of the compiled code for source line *linespec*. You can specify source lines in any of the ways understood by the **list** command (see [Section 7.1 \[Printing source lines\], page 59](#)).

For example, we can use **info line** to discover the location of the object code for the first line of function `m4_changequote`:

```
(gdb) info line m4_changequote
Line 895 of "builtin.c" starts at pc 0x634c and ends at 0x6350.
```

We can also inquire (using `*addr` as the form for *linespec*) what source line covers a particular address:

```
(gdb) info line *0x63ff
Line 926 of "builtin.c" starts at pc 0x63e4 and ends at 0x6404.
```

After **info line**, the default address for the **x** command is changed to the starting address of the line, so that ‘`x/i`’ is sufficient to begin examining the machine code (see [Section 8.5 \[Examining memory\], page 69](#)). Also, this address is saved as the value of the convenience variable `$_` (see [Section 8.9 \[Convenience variables\], page 77](#)).

disassemble

This specialized command dumps a range of memory as machine instructions. The default memory range is the function surrounding the program counter of the selected frame. A single argument to this command is a program counter value; GDB dumps the function surrounding this value. Two arguments specify a range of addresses (first inclusive, second exclusive) to dump.

The following example shows the disassembly of a range of addresses of HP PA-RISC 2.0 code:

```
(gdb) disas 0x32c4 0x32e4
Dump of assembler code from 0x32c4 to 0x32e4:
0x32c4 <main+204>:    addil 0,dp
0x32c8 <main+208>:    ldw 0x22c(sr0,r1),r26
0x32cc <main+212>:    ldil 0x3000,r31
0x32d0 <main+216>:    ble 0x3f8(sr4,r31)
0x32d4 <main+220>:    ldo 0(r31),rp
0x32d8 <main+224>:    addil -0x800,dp
0x32dc <main+228>:    ldo 0x588(r1),r26
0x32e0 <main+232>:    ldil 0x3000,r31
End of assembler dump.
```

Some architectures have more than one commonly-used set of instruction mnemonics or other syntax.

set disassembly-flavor *instruction-set*

Select the instruction set to use when disassembling the program via the **disassemble** or **x/i** commands.

Currently this command is only defined for the Intel x86 family. You can set *instruction-set* to either `intel` or `att`. The default is `att`, the AT&T flavor used by default by Unix assemblers for x86-based targets.

8 Examining Data

The usual way to examine data in your program is with the `print` command (abbreviated `p`), or its synonym `inspect`. It evaluates and prints the value of an expression of the language your program is written in (see [Chapter 12 \[Using GDB with Different Languages\], page 105](#)).

```
print expr  
print /f expr
```

`expr` is an expression (in the source language). By default the value of `expr` is printed in a format appropriate to its data type; you can choose a different format by specifying ‘`/f`’, where `f` is a letter specifying the format; see [Section 8.4 \[Output formats\], page 68](#).

```
print
```

`print /f` If you omit `expr`, GDB displays the last value again (from the *value history*; see [Section 8.8 \[Value history\], page 76](#)). This allows you to conveniently inspect the same value in an alternative format.

A more low-level way of examining data is with the `x` command. It examines data in memory at a specified address and prints it in a specified format. See [Section 8.5 \[Examining memory\], page 69](#).

If you are interested in information about types, or about how the fields of a struct or a class are declared, use the `ptype exp` command rather than `print`. See [Chapter 13 \[Examining the Symbol Table\], page 125](#).

8.1 Expressions

`print` and many other GDB commands accept an expression and compute its value. Any kind of constant, variable or operator defined by the programming language you are using is valid in an expression in GDB. This includes conditional expressions, function calls, casts, and string constants. It also includes preprocessor macros, if you compiled your program to include this information; see [Section 4.1 \[Compilation\], page 23](#).

GDB supports array constants in expressions input by the user. The syntax is `{element, element...}`. For example, you can use the command `print {1, 2, 3}` to build up an array in memory that is `malloced` in the target program.

Because C is so widespread, most of the expressions shown in examples in this manual are in C. See [Chapter 12 \[Using GDB with Different Languages\], page 105](#), for information on how to use expressions in other languages.

In this section, we discuss operators that you can use in GDB expressions regardless of your programming language.

Casts are supported in all languages, not just in C, because it is so useful to cast a number into a pointer in order to examine a structure at that address in memory.

GDB supports these operators, in addition to those common to programming languages:

- © ‘`@`’ is a binary operator for treating parts of memory as arrays. See [Section 8.3 \[Artificial arrays\], page 67](#), for more information.

`::` ‘`::`’ allows you to specify a variable in terms of the file or function where it is defined. See [Section 8.2 \[Program variables\], page 66](#).

`{type} addr`

Refers to an object of type *type* stored at address *addr* in memory. *addr* may be any expression whose value is an integer or pointer (but parentheses are required around binary operators, just as in a cast). This construct is allowed regardless of what kind of data is normally supposed to reside at *addr*.

8.2 Program variables

The most common kind of expression to use is the name of a variable in your program.

Variables in expressions are understood in the selected stack frame (see [Section 6.3 \[Selecting a frame\], page 55](#)); they must be either:

- global (or file-static)

or

- visible according to the scope rules of the programming language from the point of execution in that frame

This means that in the function

```
foo (a)
      int a;
{
  bar (a);
  {
    int b = test ();
    bar (b);
  }
}
```

you can examine and use the variable **a** whenever your program is executing within the function **foo**, but you can only use or examine the variable **b** while your program is executing inside the block where **b** is declared.

There is an exception: you can refer to a variable or function whose scope is a single source file even if the current execution point is not in this file. But it is possible to have more than one such variable or function with the same name (in different source files). If that happens, referring to that name has unpredictable effects. If you wish, you can specify a static variable in a particular function or file, using the colon-colon (`::`) notation:

```
file::variable
function::variable
```

Here *file* or *function* is the name of the context for the static *variable*. In the case of file names, you can use quotes to make sure GDB parses the file name as a single word—for example, to print a global value of **x** defined in ‘**f2.c**’:

```
(gdb) p 'f2.c'::x
```

This use of ‘`::`’ is very rarely in conflict with the very similar use of the same notation in C++. GDB also supports use of the C++ scope resolution operator in GDB expressions.

Warning: Occasionally, a local variable may appear to have the wrong value at certain points in a function—just after entry to a new scope, and just before exit.

You may see this problem when you are stepping by machine instructions. This is because, on most machines, it takes more than one instruction to set up a stack frame (including local variable definitions); if you are stepping by machine instructions, variables may appear to have the wrong values until the stack frame is completely built. On exit, it usually also takes more than one machine instruction to destroy a stack frame; after you begin stepping through that group of instructions, local variable definitions may be gone.

This may also happen when the compiler does significant optimizations. To be sure of always seeing accurate values, turn off all optimization when compiling.

Another possible effect of compiler optimizations is to optimize unused variables out of existence, or assign variables to registers (as opposed to memory addresses). Depending on the support for such cases offered by the debug info format used by the compiler, GDB might not be able to display values for such local variables. If that happens, GDB will print a message like this:

```
No symbol "foo" in current context.
```

To solve such problems, either recompile without optimizations, or use a different debug info format, if the compiler supports several such formats. For example, GCC, the GNU C/C++ compiler, usually supports the ‘`-gstabs+`’ option. ‘`-gstabs+`’ produces debug info in a format that is superior to formats such as COFF. You may be able to use DWARF 2 (‘`-gdwarf-2`’), which is also an effective form for debug info. See section “Options for Debugging Your Program or GNU CC” in *Using GNU CC*. See Section 12.4.1 [Debugging C++], page 109, for more info about debug info formats that are best suited to C++ programs.

8.3 Artificial arrays

It is often useful to print out several successive objects of the same type in memory; a section of an array, or an array of dynamically determined size for which only a pointer exists in the program.

You can do this by referring to a contiguous span of memory as an *artificial array*, using the binary operator ‘`@`’. The left operand of ‘`@`’ should be the first element of the desired array and be an individual object. The right operand should be the desired length of the array. The result is an array value whose elements are all of the type of the left argument. The first element is actually the left argument; the second element comes from bytes of memory immediately following those that hold the first element, and so on. Here is an example. If a program says

```
int *array = (int *) malloc (len * sizeof (int));
```

you can print the contents of `array` with

```
p *array@len
```

The left operand of ‘`@`’ must reside in memory. Array values made with ‘`@`’ in this way behave just like other arrays in terms of subscripting, and are coerced to pointers when used in expressions. Artificial arrays most often appear in expressions via the value history (see Section 8.8 [Value history], page 76), after printing one out.

Another way to create an artificial array is to use a cast. This re-interprets a value as if it were an array. The value need not be in memory:

```
(gdb) p/x (short[2])0x12345678
$1 = {0x1234, 0x5678}
```

As a convenience, if you leave the array length out (as in ‘`(type [])value`’) GDB calculates the size to fill the value (as ‘`sizeof(value)/sizeof(type)`’):

```
(gdb) p/x (short[])0x12345678
$2 = {0x1234, 0x5678}
```

Sometimes the artificial array mechanism is not quite enough; in moderately complex data structures, the elements of interest may not actually be adjacent—for example, if you are interested in the values of pointers in an array. One useful work-around in this situation is to use a convenience variable (see [Section 8.9 \[Convenience variables\], page 77](#)) as a counter in an expression that prints the first interesting value, and then repeat that expression via [RET](#). For instance, suppose you have an array `dtab` of pointers to structures, and you are interested in the values of a field `fv` in each structure. Here is an example of what you might type:

```
set $i = 0
p dtab[$i++]->fv
RET
RET
...

```

8.4 Output formats

By default, GDB prints a value according to its data type. Sometimes this is not what you want. For example, you might want to print a number in hex, or a pointer in decimal. Or you might want to view data in memory at a certain address as a character string or as an instruction. To do these things, specify an *output format* when you print a value.

The simplest use of output formats is to say how to print a value already computed. This is done by starting the arguments of the `print` command with a slash and a format letter. The format letters supported are:

- x Regard the bits of the value as an integer, and print the integer in hexadecimal.
- d Print as integer in signed decimal.
- u Print as integer in unsigned decimal.
- o Print as integer in octal.
- t Print as integer in binary. The letter ‘t’ stands for “two”.¹
- a Print as an address, both absolute in hexadecimal and as an offset from the nearest preceding symbol. You can use this format used to discover where (in what function) an unknown address is located:

```
(gdb) p/a 0x54320
$3 = 0x54320 <_initialize_vx+396>
```

The command `info symbol 0x54320` yields similar results. See [Chapter 13 \[Symbols\], page 125](#).

- c Regard as an integer and print it as a character constant.
- f Regard the bits of the value as a floating point number and print using typical floating point syntax.

For example, to print the program counter in hex (see [Section 8.10 \[Registers\], page 78](#)), type

¹ ‘b’ cannot be used because these format letters are also used with the `x` command, where ‘b’ stands for “byte”; see [Section 8.5 \[Examining memory\], page 69](#).

`p/x $pc`

Note that no space is required before the slash; this is because command names in GDB cannot contain a slash.

To reprint the last value in the value history with a different format, you can use the `print` command with just a format and no expression. For example, ‘`p/x`’ reprints the last value in hex.

8.5 Examining memory

You can use the command `x` (for “examine”) to examine memory in any of several formats, independently of your program’s data types.

`x/nfu addr`

`x addr`

`x` Use the `x` command to examine memory.

n, *f*, and *u* are all optional parameters that specify how much memory to display and how to format it; *addr* is an expression giving the address where you want to start displaying memory. If you use defaults for *nfu*, you need not type the slash ‘/’. Several commands set convenient defaults for *addr*.

n, the repeat count

The repeat count is a decimal integer; the default is 1. It specifies how much memory (counting by units *u*) to display.

f, the display format

The display format is one of the formats used by `print`, ‘`s`’ (null-terminated string), or ‘`i`’ (machine instruction). The default is ‘`x`’ (hexadecimal) initially. The default changes each time you use either `x` or `print`.

u, the unit size

The unit size is any of

`b` Bytes.

`h` Halfwords (two bytes).

`w` Words (four bytes). This is the initial default.

`g` Giant words (eight bytes).

Each time you specify a unit size with `x`, that size becomes the default unit the next time you use `x`. (For the ‘`s`’ and ‘`i`’ formats, the unit size is ignored and is normally not written.)

addr, starting display address

addr is the address where you want GDB to begin displaying memory. The expression need not have a pointer value (though it may); it is always interpreted as an integer address of a byte of memory. See [Section 8.1 \[Expressions\]](#), [page 65](#), for more information on expressions. The default for *addr* is usually just after the last address examined—but several other commands also set the default address: `info breakpoints` (to the address of the last breakpoint listed), `info line` (to the starting address of a line), and `print` (if you use it to display a value from memory).

For example, ‘`x/3uh 0x54320`’ is a request to display three halfwords (`h`) of memory, formatted as unsigned decimal integers (`u`), starting at address `0x54320`. ‘`x/4xw $sp`’ prints the four words (`w`) of memory above the stack pointer (here, `‘$sp’`; see [Section 8.10 \[Registers\]](#), page 78) in hexadecimal (`‘x’`).

Since the letters indicating unit sizes are all distinct from the letters specifying output formats, you do not have to remember whether unit size or format comes first; either order works. The output specifications ‘`4xw`’ and ‘`4wx`’ mean exactly the same thing. (However, the count `n` must come first; ‘`wx4`’ does not work.)

Even though the unit size `u` is ignored for the formats ‘`s`’ and ‘`i`’, you might still want to use a count `n`; for example, ‘`3i`’ specifies that you want to see three machine instructions, including any operands. The command `disassemble` gives an alternative way of inspecting machine instructions; see [Section 7.5 \[Source and machine code\]](#), page 62.

All the defaults for the arguments to `x` are designed to make it easy to continue scanning memory with minimal specifications each time you use `x`. For example, after you have inspected three machine instructions with ‘`x/3i addr`’, you can inspect the next seven with just ‘`x/7`’. If you use `RET` to repeat the `x` command, the repeat count `n` is used again; the other arguments default as for successive uses of `x`.

The addresses and contents printed by the `x` command are not saved in the value history because there is often too much of them and they would get in the way. Instead, GDB makes these values available for subsequent use in expressions as values of the convenience variables `$_` and `$__`. After an `x` command, the last address examined is available for use in expressions in the convenience variable `$_`. The contents of that address, as examined, are available in the convenience variable `$__`.

If the `x` command has a repeat count, the address and contents saved are from the last memory unit printed; this is not the same as the last address printed if several units were printed on the last line of output.

8.6 Automatic display

If you find that you want to print the value of an expression frequently (to see how it changes), you might want to add it to the *automatic display list* so that GDB prints its value each time your program stops. Each expression added to the list is given a number to identify it; to remove an expression from the list, you specify that number. The automatic display looks like this:

```
2: foo = 38
3: bar[5] = (struct hack *) 0x3804
```

This display shows item numbers, expressions and their current values. As with displays you request manually using `x` or `print`, you can specify the output format you prefer; in fact, `display` decides whether to use `print` or `x` depending on how elaborate your format specification is—it uses `x` if you specify a unit size, or one of the two formats (`‘i’` and `‘s’`) that are only supported by `x`; otherwise it uses `print`.

`display expr`

Add the expression `expr` to the list of expressions to display each time your program stops. See [Section 8.1 \[Expressions\]](#), page 65.

`display` does not repeat if you press `RET` again after using it.

display/ fmt expr

For *fmt* specifying only a display format and not a size or count, add the expression *expr* to the auto-display list but arrange to display it each time in the specified format *fmt*. See [Section 8.4 \[Output formats\], page 68](#).

display/ fmt addr

For *fmt* ‘*i*’ or ‘*s*’, or including a unit-size or a number of units, add the expression *addr* as a memory address to be examined each time your program stops. Examining means in effect doing ‘*x/ fmt addr*’. See [Section 8.5 \[Examining memory\], page 69](#).

For example, ‘`display/i $pc`’ can be helpful, to see the machine instruction about to be executed each time execution stops (‘\$pc’ is a common name for the program counter; see [Section 8.10 \[Registers\], page 78](#)).

undisplay dnums . . .**delete display dnums . . .**

Remove item numbers *dnums* from the list of expressions to display.

`undisplay` does not repeat if you press `\RET` after using it. (Otherwise you would just get the error ‘No display number . . .’.)

disable display dnums . . .

Disable the display of item numbers *dnums*. A disabled display item is not printed automatically, but is not forgotten. It may be enabled again later.

enable display dnums . . .

Enable display of item numbers *dnums*. It becomes effective once again in auto display of its expression, until you specify otherwise.

display Display the current values of the expressions on the list, just as is done when your program stops.

info display

Print the list of expressions previously set up to display automatically, each one with its item number, but without showing the values. This includes disabled expressions, which are marked as such. It also includes expressions which would not be displayed right now because they refer to automatic variables not currently available.

If a display expression refers to local variables, then it does not make sense outside the lexical context for which it was set up. Such an expression is disabled when execution enters a context where one of its variables is not defined. For example, if you give the command `display last_char` while inside a function with an argument `last_char`, GDB displays this argument while your program continues to stop inside that function. When it stops elsewhere—where there is no variable `last_char`—the display is disabled automatically. The next time your program stops where `last_char` is meaningful, you can enable the display expression once again.

8.7 Print settings

GDB provides the following ways to control how arrays, structures, and symbols are printed. These settings are useful for debugging programs in any language:

```
set print address
set print address on
```

GDB prints memory addresses showing the location of stack traces, structure values, pointer values, breakpoints, and so forth, even when it also displays the contents of those addresses. The default is `on`. For example, this is what a stack frame display looks like with `set print address on`:

```
(gdb) f
#0  set_quotes (lq=0x34c78 "<<", rq=0x34c88 ">>")
    at input.c:530
530      if (lquote != def_lquote)
```

```
set print address off
```

Do not print addresses when displaying their contents. For example, this is the same stack frame displayed with `set print address off`:

```
(gdb) set print addr off
(gdb) f
#0  set_quotes (lq=<<, rq=>>) at input.c:530
530      if (lquote != def_lquote)
```

You can use ‘`set print address off`’ to eliminate all machine dependent displays from the GDB interface. For example, with `print address off`, you should get the same text for backtraces on all machines—whether or not they involve pointer arguments.

```
show print address
```

Show whether or not addresses are to be printed.

When GDB prints a symbolic address, it normally prints the closest earlier symbol plus an offset. If that symbol does not uniquely identify the address (for example, it is a name whose scope is a single source file), you may need to clarify. One way to do this is with `info line`, for example ‘`info line *0x4537`’. Alternately, you can set GDB to print the source file and line number when it prints a symbolic address:

```
set print symbol-filename on
```

Tell GDB to print the source file name and line number of a symbol in the symbolic form of an address.

```
set print symbol-filename off
```

Do not print source file name and line number of a symbol. This is the default.

```
show print symbol-filename
```

Show whether or not GDB will print the source file name and line number of a symbol in the symbolic form of an address.

Another situation where it is helpful to show symbol filenames and line numbers is when disassembling code; GDB shows you the line number and source file that corresponds to each instruction.

Also, you may wish to see the symbolic form only if the address being printed is reasonably close to the closest earlier symbol:

```
set print max-symbolic-offset max-offset
```

Tell GDB to only display the symbolic form of an address if the offset between the closest earlier symbol and the address is less than `max-offset`. The default

is 0, which tells GDB to always print the symbolic form of an address if any symbol precedes it.

show print max-symbolic-offset

Ask how large the maximum offset is that GDB prints in a symbolic address.

If you have a pointer and you are not sure where it points, try ‘**set print symbol-filename on**’. Then you can determine the name and source file location of the variable where it points, using ‘**p/a pointer**’. This interprets the address in symbolic form. For example, here GDB shows that a variable **ptt** points at another variable **t**, defined in ‘**hi2.c**’:

```
(gdb) set print symbol-filename on
(gdb) p/a ptt
$4 = 0xe008 <t in hi2.c>
```

Warning: For pointers that point to a local variable, ‘**p/a**’ does not show the symbol name and filename of the referent, even with the appropriate **set print** options turned on.

Other settings control how different kinds of objects are printed:

set print array

set print array on

Pretty print arrays. This format is more convenient to read, but uses more space. The default is off.

set print array off

Return to compressed format for arrays.

show print array

Show whether compressed or pretty format is selected for displaying arrays.

set print elements *number-of-elements*

Set a limit on how many elements of an array GDB will print. If GDB is printing a large array, it stops printing after it has printed the number of elements set by the **set print elements** command. This limit also applies to the display of strings. When GDB starts, this limit is set to 200. Setting *number-of-elements* to zero means that the printing is unlimited.

show print elements

Display the number of elements of a large array that GDB will print. If the number is 0, then the printing is unlimited.

set print null-stop

Cause GDB to stop printing the characters of an array when the first NULL is encountered. This is useful when large arrays actually contain only short strings. The default is off.

set print pretty on

Cause GDB to print structures in an indented format with one member per line, like this:

```
$1 = {
    next = 0x0,
    flags = {
        sweet = 1,
        sour = 1
    },
    meat = 0x54 "Pork"
}
```

set print pretty off

Cause GDB to print structures in a compact format, like this:

```
$1 = {next = 0x0, flags = {sweet = 1, sour = 1}, \
meat = 0x54 "Pork"}
```

This is the default format.

show print pretty

Show which format GDB is using to print structures.

set print sevenbit-strings on

Print using only seven-bit characters; if this option is set, GDB displays any eight-bit characters (in strings or character values) using the notation \nnn. This setting is best if you are working in English (ASCII) and you use the high-order bit of characters as a marker or “meta” bit.

set print sevenbit-strings off

Print full eight-bit characters. This allows the use of more international character sets, and is the default.

show print sevenbit-strings

Show whether or not GDB is printing only seven-bit characters.

set print union on

Tell GDB to print unions which are contained in structures. This is the default setting.

set print union off

Tell GDB not to print unions which are contained in structures.

show print union

Ask GDB whether or not it will print unions which are contained in structures.

For example, given the declarations

```
typedef enum {Tree, Bug} Species;
typedef enum {Big_tree, Acorn, Seedling} Tree_forms;
typedef enum {Caterpillar, Cocoon, Butterfly}
    Bug_forms;
```

```
struct thing {
    Species it;
    union {
        Tree_forms tree;
        Bug_forms bug;
    } form;
};
```

```
struct thing foo = {Tree, {Acorn}};
```

with **set print union on** in effect ‘p foo’ would print

```
$1 = {it = Tree, form = {tree = Acorn, bug = Cocoon}}
and with set print union off in effect it would print
$1 = {it = Tree, form = {...}}
```

These settings are of interest when debugging C++ programs:

set print demangle

set print demangle on

Print C++ names in their source form rather than in the encoded (“mangled”) form passed to the assembler and linker for type-safe linkage. The default is on.

show print demangle

Show whether C++ names are printed in mangled or demangled form.

set print asm-demangle

set print asm-demangle on

Print C++ names in their source form rather than their mangled form, even in assembler code printouts such as instruction disassemblies. The default is off.

show print asm-demangle

Show whether C++ names in assembly listings are printed in mangled or demangled form.

set demangle-style style

Choose among several encoding schemes used by different compilers to represent C++ names. The choices for *style* are currently:

auto Allow GDB to choose a decoding style by inspecting your program.

gnu Decode based on the GNU C++ compiler (`g++`) encoding algorithm. This is the default.

hp Decode based on the HP ANSI C++ (`aCC`) encoding algorithm.

lucid Decode based on the Lucid C++ compiler (`lcc`) encoding algorithm.

arm Decode using the algorithm in the *C++ Annotated Reference Manual*. **Warning:** this setting alone is not sufficient to allow debugging `cfront`-generated executables. GDB would require further enhancement to permit that.

If you omit *style*, you will see a list of possible formats.

show demangle-style

Display the encoding style currently in use for decoding C++ symbols.

set print object

set print object on

When displaying a pointer to an object, identify the *actual* (derived) type of the object rather than the *declared* type, using the virtual function table.

set print object off

Display only the declared type of objects, without reference to the virtual function table. This is the default setting.

```

show print object
    Show whether actual, or declared, object types are displayed.

set print static-members
set print static-members on
    Print static members when displaying a C++ object. The default is on.

set print static-members off
    Do not print static members when displaying a C++ object.

show print static-members
    Show whether C++ static members are printed, or not.

set print vtbl
set print vtbl on
    Pretty print C++ virtual function tables. The default is off. (The vtbl commands do not work on programs compiled with the HP ANSI C++ compiler (aCC).)

set print vtbl off
    Do not pretty print C++ virtual function tables.

show print vtbl
    Show whether C++ virtual function tables are pretty printed, or not.

```

8.8 Value history

Values printed by the `print` command are saved in the GDB *value history*. This allows you to refer to them in other expressions. Values are kept until the symbol table is re-read or discarded (for example with the `file` or `symbol-file` commands). When the symbol table changes, the value history is discarded, since the values may contain pointers back to the types defined in the symbol table.

The values printed are given *history numbers* by which you can refer to them. These are successive integers starting with one. `print` shows you the history number assigned to a value by printing ‘`$num =` ’ before the value; here `num` is the history number.

To refer to any previous value, use ‘\$’ followed by the value’s history number. The way `print` labels its output is designed to remind you of this. Just \$ refers to the most recent value in the history, and \$\$ refers to the value before that. \$\$n refers to the *n*th value from the end; \$\$2 is the value just prior to \$\$, \$\$1 is equivalent to \$\$, and \$\$0 is equivalent to \$.

For example, suppose you have just printed a pointer to a structure and want to see the contents of the structure. It suffices to type

```
p *$
```

If you have a chain of structures where the component `next` points to the next one, you can print the contents of the next one with this:

```
p *$.next
```

You can print successive links in the chain by repeating this command—which you can do by just typing [RET](#).

Note that the history records values, not expressions. If the value of `x` is 4 and you type these commands:

```
print x
set x=5
```

then the value recorded in the value history by the `print` command remains 4 even though the value of `x` has changed.

`show values`

Print the last ten values in the value history, with their item numbers. This is like ‘`p $$9`’ repeated ten times, except that `show values` does not change the history.

`show values n`

Print ten history values centered on history item number `n`.

`show values +`

Print ten history values just after the values last printed. If no more values are available, `show values +` produces no display.

Pressing `(RET)` to repeat `show values n` has exactly the same effect as ‘`show values +`’.

8.9 Convenience variables

GDB provides *convenience variables* that you can use within GDB to hold on to a value and refer to it later. These variables exist entirely within GDB; they are not part of your program, and setting a convenience variable has no direct effect on further execution of your program. That is why you can use them freely.

Convenience variables are prefixed with ‘\$’. Any name preceded by ‘\$’ can be used for a convenience variable, unless it is one of the predefined machine-specific register names (see [Section 8.10 \[Registers\]](#), page 78). (Value history references, in contrast, are *numbers* preceded by ‘\$’. See [Section 8.8 \[Value history\]](#), page 76.)

You can save a value in a convenience variable with an assignment expression, just as you would set a variable in your program. For example:

```
set $foo = *object_ptr
```

would save in `$foo` the value contained in the object pointed to by `object_ptr`.

Using a convenience variable for the first time creates it, but its value is `void` until you assign a new value. You can alter the value with another assignment at any time.

Convenience variables have no fixed types. You can assign a convenience variable any type of value, including structures and arrays, even if that variable already has a value of a different type. The convenience variable, when used as an expression, has the type of its current value.

`show convenience`

Print a list of convenience variables used so far, and their values. Abbreviated `show conv.`

One of the ways to use a convenience variable is as a counter to be incremented or a pointer to be advanced. For example, to print a field from successive elements of an array of structures:

```
set $i = 0
print bar[$i++]->contents
```

Repeat that command by typing `\RET`.

Some convenience variables are created automatically by GDB and given values likely to be useful.

\$_ The variable `$_` is automatically set by the `x` command to the last address examined (see [Section 8.5 \[Examining memory\], page 69](#)). Other commands which provide a default address for `x` to examine also set `$_` to that address; these commands include `info line` and `info breakpoint`. The type of `$_` is `void *` except when set by the `x` command, in which case it is a pointer to the type of `$__`.

\$__ The variable `$__` is automatically set by the `x` command to the value found in the last address examined. Its type is chosen to match the format in which the data was printed.

`$_exitcode`

The variable `$_exitcode` is automatically set to the exit code when the program being debugged terminates.

On HP-UX systems, if you refer to a function or variable name that begins with a dollar sign, GDB searches for a user or system name first, before it searches for a convenience variable.

8.10 Registers

You can refer to machine register contents, in expressions, as variables with names starting with ‘\$’. The names of registers are different for each machine; use `info registers` to see the names used on your machine.

`info registers`

Print the names and values of all registers except floating-point and vector registers (in the selected stack frame).

`info all-registers`

Print the names and values of all registers, including floating-point and vector registers (in the selected stack frame).

`info registers regname ...`

Print the relativized value of each specified register `regname`. As discussed in detail below, register values are normally relative to the selected stack frame. `regname` may be any register name valid on the machine you are using, with or without the initial ‘\$’.

GDB has four “standard” register names that are available (in expressions) on most machines—whenever they do not conflict with an architecture’s canonical mnemonics for registers. The register names `$pc` and `$sp` are used for the program counter register and the stack pointer. `$fp` is used for a register that contains a pointer to the current stack frame, and `$ps` is used for a register that contains the processor status. For example, you could print the program counter in hex with

`p/x $pc`

or print the instruction to be executed next with

```
x/i $pc
```

or add four to the stack pointer² with

```
set $sp += 4
```

Whenever possible, these four standard register names are available on your machine even though the machine has different canonical mnemonics, so long as there is no conflict. The `info registers` command shows the canonical names. For example, on the SPARC, `info registers` displays the processor status register as `$psr` but you can also refer to it as `$ps`; and on x86-based machines `$ps` is an alias for the `EFLAGS` register.

GDB always considers the contents of an ordinary register as an integer when the register is examined in this way. Some machines have special registers which can hold nothing but floating point; these registers are considered to have floating point values. There is no way to refer to the contents of an ordinary register as floating point value (although you can *print* it as a floating point value with ‘`print/f $regname`’).

Some registers have distinct “raw” and “virtual” data formats. This means that the data format in which the register contents are saved by the operating system is not the same one that your program normally sees. For example, the registers of the 68881 floating point coprocessor are always saved in “extended” (raw) format, but all C programs expect to work with “double” (virtual) format. In such cases, GDB normally works with the virtual format only (the format that makes sense for your program), but the `info registers` command prints the data in both formats.

Normally, register values are relative to the selected stack frame (see [Section 6.3 \[Selecting a frame\], page 55](#)). This means that you get the value that the register would contain if all stack frames farther in were exited and their saved registers restored. In order to see the true contents of hardware registers, you must select the innermost frame (with ‘`frame 0`’).

However, GDB must deduce where registers are saved, from the machine code generated by your compiler. If some registers are not saved, or if GDB is unable to locate the saved registers, the selected stack frame makes no difference.

8.11 Floating point hardware

Depending on the configuration, GDB may be able to give you more information about the status of the floating point hardware.

```
info float
```

Display hardware-dependent information about the floating point unit. The exact contents and layout vary depending on the floating point chip. Currently, ‘`info float`’ is supported on the ARM and x86 machines.

8.12 Vector Unit

Depending on the configuration, GDB may be able to give you more information about the status of the vector unit.

² This is a way of removing one word from the stack, on machines where stacks grow downward in memory (most machines, nowadays). This assumes that the innermost stack frame is selected; setting `$sp` is not allowed when other stack frames are selected. To pop entire frames off the stack, regardless of machine architecture, use `return`; see [Section 14.4 \[Returning from a function\], page 133](#).

info vector

Display information about the vector unit. The exact contents and layout vary depending on the hardware.

8.13 Operating system auxiliary vector

Some operating systems supply an *auxiliary* vector to programs at startup. This is akin to the arguments and environment that you specify for a program, but contains a system-dependent variety of binary values that tell system libraries important details about the hardware, operating system, and process. Each value's purpose is identified by an integer tag; the meanings are well-known but system-specific. Depending on the configuration and operating system facilities, GDB may be able to show you this information.

info auxv Display the auxiliary vector of the inferior, which can be either a live process or a core dump file. GDB prints each tag value numerically, and also shows names and text descriptions for recognized tags. Some values in the vector are numbers, some bit masks, and some pointers to strings or other data. GDB displays each value in the most appropriate form for a recognized tag, and in hexadecimal for an unrecognized tag.

8.14 Memory region attributes

Memory region attributes allow you to describe special handling required by regions of your target's memory. GDB uses attributes to determine whether to allow certain types of memory accesses; whether to use specific width accesses; and whether to cache target memory.

Defined memory regions can be individually enabled and disabled. When a memory region is disabled, GDB uses the default attributes when accessing memory in that region. Similarly, if no memory regions have been defined, GDB uses the default attributes when accessing all memory.

When a memory region is defined, it is given a number to identify it; to enable, disable, or remove a memory region, you specify that number.

mem lower upper attributes...

Define memory region bounded by *lower* and *upper* with attributes *attributes*. . . . Note that *upper == 0* is a special case: it is treated as the the target's maximum memory address. (0xffff on 16 bit targets, 0xffffffff on 32 bit targets, etc.)

delete mem *nums*...

Remove memory regions *nums*. . . .

disable mem *nums*...

Disable memory regions *nums*. . . . A disabled memory region is not forgotten. It may be enabled again later.

enable mem *nums*...

Enable memory regions *nums*. . . .

info mem Print a table of all defined memory regions, with the following columns for each region.

*Memory Region Number
Enabled or Disabled.*

Enabled memory regions are marked with ‘y’. Disabled memory regions are marked with ‘n’.

Lo Address

The address defining the inclusive lower bound of the memory region.

Hi Address

The address defining the exclusive upper bound of the memory region.

Attributes The list of attributes set for this memory region.

8.14.1 Attributes

8.14.1.1 Memory Access Mode

The access mode attributes set whether GDB may make read or write accesses to a memory region.

While these attributes prevent GDB from performing invalid memory accesses, they do nothing to prevent the target system, I/O DMA, etc. from accessing memory.

- ro** Memory is read only.
- wo** Memory is write only.
- rw** Memory is read/write. This is the default.

8.14.1.2 Memory Access Size

The access size attributes tells GDB to use specific sized accesses in the memory region. Often memory mapped device registers require specific sized accesses. If no access size attribute is specified, GDB may use accesses of any size.

- 8** Use 8 bit memory accesses.
- 16** Use 16 bit memory accesses.
- 32** Use 32 bit memory accesses.
- 64** Use 64 bit memory accesses.

8.14.1.3 Data Cache

The data cache attributes set whether GDB will cache target memory. While this generally improves performance by reducing debug protocol overhead, it can lead to incorrect results because GDB does not know about volatile variables or memory mapped device registers.

- cache** Enable GDB to cache target memory.
- nocache** Disable GDB from caching target memory. This is the default.

8.15 Copy between memory and a file

You can use the commands `dump`, `append`, and `restore` to copy data between target memory and a file. The `dump` and `append` commands write data to a file, and the `restore` command reads data from a file back into the inferior's memory. Files may be in binary, Motorola S-record, Intel hex, or Tektronix Hex format; however, GDB can only append to binary files.

```
dump [format] memory filename start_addr end_addr
dump [format] value filename expr
```

Dump the contents of memory from *start_addr* to *end_addr*, or the value of *expr*, to *filename* in the given format.

The *format* parameter may be any one of:

<code>binary</code>	Raw binary form.
<code>ihex</code>	Intel hex format.
<code>srec</code>	Motorola S-record format.
<code>tekhex</code>	Tektronix Hex format.

GDB uses the same definitions of these formats as the GNU binary utilities, like ‘`objdump`’ and ‘`objcopy`’. If *format* is omitted, GDB dumps the data in raw binary form.

```
append [binary] memory filename start_addr end_addr
append [binary] value filename expr
```

Append the contents of memory from *start_addr* to *end_addr*, or the value of *expr*, to *filename*, in raw binary form. (GDB can only append data to files in raw binary form.)

```
restore filename [binary] bias start end
```

Restore the contents of file *filename* into memory. The `restore` command can automatically recognize any known BFD file format, except for raw binary. To restore a raw binary file you must specify the optional keyword `binary` after the *filename*.

If *bias* is non-zero, its value will be added to the addresses contained in the file. Binary files always start at address zero, so they will be restored at address *bias*. Other bfd files have a built-in location; they will be restored at offset *bias* from that location.

If *start* and/or *end* are non-zero, then only data between file offset *start* and file offset *end* will be restored. These offsets are relative to the addresses in the file, before the *bias* argument is applied.

8.16 Character Sets

If the program you are debugging uses a different character set to represent characters and strings than the one GDB uses itself, GDB can automatically translate between the character sets for you. The character set GDB uses we call the *host character set*; the one the inferior program uses we call the *target character set*.

For example, if you are running GDB on a GNU/Linux system, which uses the ISO Latin 1 character set, but you are using GDB’s remote protocol (see [Section 16.4 \[Remote\]](#), page 149)

to debug a program running on an IBM mainframe, which uses the EBCDIC character set, then the host character set is Latin-1, and the target character set is EBCDIC. If you give GDB the command `set target-charset EBCDIC-US`, then GDB translates between EBCDIC and Latin 1 as you print character or string values, or use character and string literals in expressions.

GDB has no way to automatically recognize which character set the inferior program uses; you must tell it, using the `set target-charset` command, described below.

Here are the commands for controlling GDB's character set support:

set target-charset charset

Set the current target character set to *charset*. We list the character set names GDB recognizes below, but if you type `set target-charset` followed by `\TAB\TAB`, GDB will list the target character sets it supports.

set host-charset charset

Set the current host character set to *charset*.

By default, GDB uses a host character set appropriate to the system it is running on; you can override that default using the `set host-charset` command.

GDB can only use certain character sets as its host character set. We list the character set names GDB recognizes below, and indicate which can be host character sets, but if you type `set target-charset` followed by `\TAB\TAB`, GDB will list the host character sets it supports.

set charset charset

Set the current host and target character sets to *charset*. As above, if you type `set charset` followed by `\TAB\TAB`, GDB will list the name of the character sets that can be used for both host and target.

show charset

Show the names of the current host and target charsets.

show host-charset

Show the name of the current host charset.

show target-charset

Show the name of the current target charset.

GDB currently includes support for the following character sets:

ASCII Seven-bit U.S. ASCII. GDB can use this as its host character set.

ISO-8859-1

The ISO Latin 1 character set. This extends ASCII with accented characters needed for French, German, and Spanish. GDB can use this as its host character set.

EBCDIC-US

IBM1047 Variants of the EBCDIC character set, used on some of IBM's mainframe operating systems. (GNU/Linux on the S/390 uses U.S. ASCII.) GDB cannot use these as its host character set.

Note that these are all single-byte character sets. More work inside GDB is needed to support multi-byte or variable-width character encodings, like the UTF-8 and UCS-2 encodings of Unicode.

Here is an example of GDB's character set support in action. Assume that the following source code has been placed in the file ‘`charset-test.c`’:

```
#include <stdio.h>

char ascii_hello[]
= {72, 101, 108, 108, 111, 44, 32, 119,
   111, 114, 108, 100, 33, 10, 0};
char ibm1047_hello[]
= {200, 133, 147, 147, 150, 107, 64, 166,
   150, 153, 147, 132, 90, 37, 0};

main ()
{
    printf ("Hello, world!\n");
}
```

In this program, `ascii_hello` and `ibm1047_hello` are arrays containing the string ‘Hello, world!’ followed by a newline, encoded in the ASCII and IBM1047 character sets.

We compile the program, and invoke the debugger on it:

```
$ gcc -g charset-test.c -o charset-test
$ gdb -nw charset-test
GNU gdb 2001-12-19-cvs
Copyright 2001 Free Software Foundation, Inc.
...
(gdb)
```

We can use the `show charset` command to see what character sets GDB is currently using to interpret and display characters and strings:

```
(gdb) show charset
The current host and target character set is 'ISO-8859-1'.
(gdb)
```

For the sake of printing this manual, let's use ASCII as our initial character set:

```
(gdb) set charset ASCII
(gdb) show charset
The current host and target character set is 'ASCII'.
(gdb)
```

Let's assume that ASCII is indeed the correct character set for our host system — in other words, let's assume that if GDB prints characters using the ASCII character set, our terminal will display them properly. Since our current target character set is also ASCII, the contents of `ascii_hello` print legibly:

```
(gdb) print ascii_hello
$1 = 0x401698 "Hello, world!\n"
(gdb) print ascii_hello[0]
$2 = 72 'H'
(gdb)
```

GDB uses the target character set for character and string literals you use in expressions:

```
(gdb) print '+'
$3 = 43 '+'
(gdb)
```

The ASCII character set uses the number 43 to encode the ‘+’ character.

GDB relies on the user to tell it which character set the target program uses. If we print `ibm1047_hello` while our target character set is still ASCII, we get gibberish:

```
(gdb) print ibm1047_hello
$4 = 0x4016a8 "\310\205\223\223\226k@\246\226\231\223\204Z"
(gdb) print ibm1047_hello[0]
$5 = 200 '\310'
(gdb)
```

If we invoke the `set target-charset` followed by `\TAB\TAB`, GDB tells us the character sets it supports:

```
(gdb) set target-charset
ASCII      EBCDIC-US   IBM1047    ISO-8859-1
(gdb) set target-charset
```

We can select IBM1047 as our target character set, and examine the program's strings again. Now the ASCII string is wrong, but GDB translates the contents of `ibm1047_hello` from the target character set, IBM1047, to the host character set, ASCII, and they display correctly:

```
(gdb) set target-charset IBM1047
(gdb) show charset
The current host character set is 'ASCII'.
The current target character set is 'IBM1047'.
(gdb) print ascii_hello
$6 = 0x401698 "\110\145%\?054\040\167?\162%\144\041\012"
(gdb) print ascii_hello[0]
$7 = 72 '\110'
(gdb) print ibm1047_hello
$8 = 0x4016a8 "Hello, world!\n"
(gdb) print ibm1047_hello[0]
$9 = 200 'H'
(gdb)
```

As above, GDB uses the target character set for character and string literals you use in expressions:

```
(gdb) print '+'
$10 = 78 '+'
(gdb)
```

The IBM1047 character set uses the number 78 to encode the '+' character.

9 C Preprocessor Macros

Some languages, such as C and C++, provide a way to define and invoke “preprocessor macros” which expand into strings of tokens. GDB can evaluate expressions containing macro invocations, show the result of macro expansion, and show a macro’s definition, including where it was defined.

You may need to compile your program specially to provide GDB with information about preprocessor macros. Most compilers do not include macros in their debugging information, even when you compile with the ‘-g’ flag. See [Section 4.1 \[Compilation\], page 23](#).

A program may define a macro at one point, remove that definition later, and then provide a different definition after that. Thus, at different points in the program, a macro may have different definitions, or have no definition at all. If there is a current stack frame, GDB uses the macros in scope at that frame’s source code line. Otherwise, GDB uses the macros in scope at the current listing location; see [Section 7.1 \[List\], page 59](#).

At the moment, GDB does not support the ## token-splicing operator, the # stringification operator, or variable-arity macros.

Whenever GDB evaluates an expression, it always expands any macro invocations present in the expression. GDB also provides the following commands for working with macros explicitly.

`macro expand expression`

`macro exp expression`

Show the results of expanding all preprocessor macro invocations in *expression*.

Since GDB simply expands macros, but does not parse the result, *expression* need not be a valid expression; it can be any string of tokens.

`macro expand-once expression`

`macro exp1 expression`

(*This command is not yet implemented.*) Show the results of expanding those preprocessor macro invocations that appear explicitly in *expression*. Macro invocations appearing in that expansion are left unchanged. This command allows you to see the effect of a particular macro more clearly, without being confused by further expansions. Since GDB simply expands macros, but does not parse the result, *expression* need not be a valid expression; it can be any string of tokens.

`info macro macro`

Show the definition of the macro named *macro*, and describe the source location where that definition was established.

`macro define macro replacement-list`

`macro define macro (arglist) replacement-list`

(*This command is not yet implemented.*) Introduce a definition for a preprocessor macro named *macro*, invocations of which are replaced by the tokens given in *replacement-list*. The first form of this command defines an “object-like” macro, which takes no arguments; the second form defines a “function-like” macro, which takes the arguments given in *arglist*.

A definition introduced by this command is in scope in every expression evaluated in GDB, until it is removed with the `macro undef` command, described below. The definition overrides all definitions for `macro` present in the program being debugged, as well as any previous user-supplied definition.

`macro undef macro`

(*This command is not yet implemented.*) Remove any user-supplied definition for the macro named `macro`. This command only affects definitions provided with the `macro define` command, described above; it cannot remove definitions present in the program being debugged.

Here is a transcript showing the above commands in action. First, we show our source files:

```
$ cat sample.c
#include <stdio.h>
#include "sample.h"

#define M 42
#define ADD(x) (M + x)

main ()
{
#define N 28
    printf ("Hello, world!\n");
#undef N
    printf ("We're so creative.\n");
#define N 1729
    printf ("Goodbye, world!\n");
}
$ cat sample.h
#define Q <
$
```

Now, we compile the program using the GNU C compiler, GCC. We pass the ‘`-gdwarf-2`’ and ‘`-g3`’ flags to ensure the compiler includes information about preprocessor macros in the debugging information.

```
$ gcc -gdwarf-2 -g3 sample.c -o sample
$
```

Now, we start GDB on our sample program:

```
$ gdb -nw sample
GNU gdb 2002-05-06-cvs
Copyright 2002 Free Software Foundation, Inc.
GDB is free software, ...
(gdb)
```

We can expand macros and examine their definitions, even when the program is not running. GDB uses the current listing position to decide which macro definitions are in scope:

```
(gdb) list main
3
4     #define M 42
5     #define ADD(x) (M + x)
6
7     main ()
8 {
```

```

9      #define N 28
10     printf ("Hello, world!\n");
11     #undef N
12     printf ("We're so creative.\n");
(gdb) info macro ADD
Defined at /home/jimb/gdb/macros/play/sample.c:5
#define ADD(x) (M + x)
(gdb) info macro Q
Defined at /home/jimb/gdb/macros/play/sample.h:1
    included at /home/jimb/gdb/macros/play/sample.c:2
#define Q <
(gdb) macro expand ADD(1)
expands to: (42 + 1)
(gdb) macro expand-once ADD(1)
expands to: once (M + 1)
(gdb)

```

In the example above, note that `macro expand-once` expands only the macro invocation explicit in the original text — the invocation of `ADD` — but does not expand the invocation of the macro `M`, which was introduced by `ADD`.

Once the program is running, GDB uses the macro definitions in force at the source line of the current stack frame:

```

(gdb) break main
Breakpoint 1 at 0x8048370: file sample.c, line 10.
(gdb) run
Starting program: /home/jimb/gdb/macros/play/sample

Breakpoint 1, main () at sample.c:10
10     printf ("Hello, world!\n");
(gdb)

```

At line 10, the definition of the macro `N` at line 9 is in force:

```

(gdb) info macro N
Defined at /home/jimb/gdb/macros/play/sample.c:9
#define N 28
(gdb) macro expand N Q M
expands to: 28 < 42
(gdb) print N Q M
$1 = 1
(gdb)

```

As we step over directives that remove `N`'s definition, and then give it a new definition, GDB finds the definition (or lack thereof) in force at each point:

```

(gdb) next
Hello, world!
12     printf ("We're so creative.\n");
(gdb) info macro N
The symbol 'N' has no definition as a C/C++ preprocessor macro
at /home/jimb/gdb/macros/play/sample.c:12
(gdb) next
We're so creative.
14     printf ("Goodbye, world!\n");
(gdb) info macro N
Defined at /home/jimb/gdb/macros/play/sample.c:13
#define N 1729
(gdb) macro expand N Q M
expands to: 1729 < 42
(gdb) print N Q M

```

```
$2 = 0  
(gdb)
```

10 Tracepoints

In some applications, it is not feasible for the debugger to interrupt the program’s execution long enough for the developer to learn anything helpful about its behavior. If the program’s correctness depends on its real-time behavior, delays introduced by a debugger might cause the program to change its behavior drastically, or perhaps fail, even when the code itself is correct. It is useful to be able to observe the program’s behavior without interrupting it.

Using GDB’s `trace` and `collect` commands, you can specify locations in the program, called *tracepoints*, and arbitrary expressions to evaluate when those tracepoints are reached. Later, using the `tfind` command, you can examine the values those expressions had when the program hit the tracepoints. The expressions may also denote objects in memory—structures or arrays, for example—whose values GDB should record; while visiting a particular tracepoint, you may inspect those objects as if they were in memory at that moment. However, because GDB records these values without interacting with you, it can do so quickly and unobtrusively, hopefully not disturbing the program’s behavior.

The tracepoint facility is currently available only for remote targets. See [Chapter 16 \[Targets\], page 147](#). In addition, your remote target must know how to collect trace data. This functionality is implemented in the remote stub; however, none of the stubs distributed with GDB support tracepoints as of this writing.

This chapter describes the tracepoint commands and features.

10.1 Commands to Set Tracepoints

Before running such a *trace experiment*, an arbitrary number of tracepoints can be set. Like a breakpoint (see [Section 5.1.1 \[Set Breaks\], page 34](#)), a tracepoint has a number assigned to it by GDB. Like with breakpoints, tracepoint numbers are successive integers starting from one. Many of the commands associated with tracepoints take the tracepoint number as their argument, to identify which tracepoint to work on.

For each tracepoint, you can specify, in advance, some arbitrary set of data that you want the target to collect in the trace buffer when it hits that tracepoint. The collected data can include registers, local variables, or global data. Later, you can use GDB commands to examine the values these data had at the time the tracepoint was hit.

This section describes commands to set tracepoints and associated conditions and actions.

10.1.1 Create and Delete Tracepoints

trace The `trace` command is very similar to the `break` command. Its argument can be a source line, a function name, or an address in the target program. See [Section 5.1.1 \[Set Breaks\], page 34](#). The `trace` command defines a tracepoint, which is a point in the target program where the debugger will briefly stop, collect some data, and then allow the program to continue. Setting a tracepoint or changing its commands doesn’t take effect until the next `tstart` command; thus, you cannot change the tracepoint attributes once a trace experiment is running.

Here are some examples of using the `trace` command:

```
(gdb) trace foo.c:121    // a source file and line number
(gdb) trace +2           // 2 lines forward
(gdb) trace my_function // first source line of function
(gdb) trace *my_function // EXACT start address of function
(gdb) trace *0x2117c4    // an address
```

You can abbreviate `trace` as `tr`.

The convenience variable `$tpnum` records the tracepoint number of the most recently set tracepoint.

`delete tracepoint [num]`

Permanently delete one or more tracepoints. With no argument, the default is to delete all tracepoints.

Examples:

```
(gdb) delete trace 1 2 3 // remove three tracepoints
(gdb) delete trace      // remove all tracepoints
```

You can abbreviate this command as `del tr`.

10.1.2 Enable and Disable Tracepoints

`disable tracepoint [num]`

Disable tracepoint *num*, or all tracepoints if no argument *num* is given. A disabled tracepoint will have no effect during the next trace experiment, but it is not forgotten. You can re-enable a disabled tracepoint using the `enable tracepoint` command.

`enable tracepoint [num]`

Enable tracepoint *num*, or all tracepoints. The enabled tracepoints will become effective the next time a trace experiment is run.

10.1.3 Tracepoint Passcounts

`passcount [n [num]]`

Set the passcount of a tracepoint. The passcount is a way to automatically stop a trace experiment. If a tracepoint's passcount is *n*, then the trace experiment will be automatically stopped on the *n*'th time that tracepoint is hit. If the tracepoint number *num* is not specified, the `passcount` command sets the passcount of the most recently defined tracepoint. If no passcount is given, the trace experiment will run until stopped explicitly by the user.

Examples:

```
(gdb) passcount 5 2 // Stop on the 5th execution of
                     // tracepoint 2

(gdb) passcount 12 // Stop on the 12th execution of the
                  // most recently defined tracepoint.

(gdb) trace foo
(gdb) pass 3
(gdb) trace bar
```

```
(gdb) pass 2
(gdb) trace baz
(gdb) pass 1      // Stop tracing when foo has been
                  // executed 3 times OR when bar has
                  // been executed 2 times
                  // OR when baz has been executed 1 time.
```

10.1.4 Tracepoint Action Lists

actions [num]

This command will prompt for a list of actions to be taken when the tracepoint is hit. If the tracepoint number *num* is not specified, this command sets the actions for the one that was most recently defined (so that you can define a tracepoint and then say **actions** without bothering about its number). You specify the actions themselves on the following lines, one action at a time, and terminate the actions list with a line containing just **end**. So far, the only defined actions are **collect** and **while-stepping**.

To remove all actions from a tracepoint, type ‘**actions num**’ and follow it immediately with ‘**end**’.

```
(gdb) collect data // collect some data
(gdb) while-stepping 5 // single-step 5 times, collect data
(gdb) end              // signals the end of actions.
```

In the following example, the action list begins with **collect** commands indicating the things to be collected when the tracepoint is hit. Then, in order to single-step and collect additional data following the tracepoint, a **while-stepping** command is used, followed by the list of things to be collected while stepping. The **while-stepping** command is terminated by its own separate **end** command. Lastly, the action list is terminated by an **end** command.

```
(gdb) trace foo
(gdb) actions
Enter actions for tracepoint 1, one per line:
> collect bar,baz
> collect $regs
> while-stepping 12
  > collect $fp, $sp
  > end
end
```

collect expr1, expr2, ...

Collect values of the given expressions when the tracepoint is hit. This command accepts a comma-separated list of any valid expressions. In addition to global, static, or local variables, the following special arguments are supported:

\$regs	collect all registers
\$args	collect all function arguments
\$locals	collect all local variables.

You can give several consecutive **collect** commands, each one with a single argument, or one **collect** command with several arguments separated by commas: the effect is the same.

The command `info scope` (see Chapter 13 [Symbols], page 125) is particularly useful for figuring out what data to collect.

`while-stepping n`

Perform *n* single-step traces after the tracepoint, collecting new data at each step. The `while-stepping` command is followed by the list of what to collect while stepping (followed by its own `end` command):

```
> while-stepping 12
> collect $regs, myglobal
> end
>
```

You may abbreviate `while-stepping` as `ws` or `stepping`.

10.1.5 Listing Tracepoints

`info tracepoints [num]`

Display information about the tracepoint *num*. If you don't specify a tracepoint number, displays information about all the tracepoints defined so far. For each tracepoint, the following information is shown:

- its number
- whether it is enabled or disabled
- its address
- its passcount as given by the `passcount n` command
- its step count as given by the `while-stepping n` command
- where in the source files is the tracepoint set
- its action list as given by the `actions` command

```
(gdb) info trace
Num Enb Address      PassC StepC What
 1  y  0x002117c4 0    0    <gdb_asm>
 2  y  0x0020dc64 0    0    in g_test at g_test.c:1375
 3  y  0x0020b1f4 0    0    in get_data at ../foo.c:41
(gdb)
```

This command can be abbreviated `info tp`.

10.1.6 Starting and Stopping Trace Experiment

`tstart` This command takes no arguments. It starts the trace experiment, and begins collecting data. This has the side effect of discarding all the data collected in the trace buffer during the previous trace experiment.

`tstop` This command takes no arguments. It ends the trace experiment, and stops collecting data.

Note: a trace experiment and data collection may stop automatically if any tracepoint's passcount is reached (see Section 10.1.3 [Tracepoint Passcounts], page 92), or if the trace buffer becomes full.

`tstatus` This command displays the status of the current trace data collection.

Here is an example of the commands we described so far:

```
(gdb) trace gdb_c_test
(gdb) actions
Enter actions for tracepoint #1, one per line.
> collect $regs,$locals,$args
> while-stepping 11
  > collect $regs
  > end
> end
(gdb) tstart
[time passes ...]
(gdb) tstop
```

10.2 Using the collected data

After the tracepoint experiment ends, you use GDB commands for examining the trace data. The basic idea is that each tracepoint collects a trace *snapshot* every time it is hit and another snapshot every time it single-steps. All these snapshots are consecutively numbered from zero and go into a buffer, and you can examine them later. The way you examine them is to *focus* on a specific trace snapshot. When the remote stub is focused on a trace snapshot, it will respond to all GDB requests for memory and registers by reading from the buffer which belongs to that snapshot, rather than from *real* memory or registers of the program being debugged. This means that **all** GDB commands (`print`, `info registers`, `backtrace`, etc.) will behave as if we were currently debugging the program state as it was when the tracepoint occurred. Any requests for data that are not in the buffer will fail.

10.2.1 `tfind n`

The basic command for selecting a trace snapshot from the buffer is `tfind n`, which finds trace snapshot number *n*, counting from zero. If no argument *n* is given, the next snapshot is selected.

Here are the various forms of using the `tfind` command.

`tfind start`

Find the first snapshot in the buffer. This is a synonym for `tfind 0` (since 0 is the number of the first snapshot).

`tfind none`

Stop debugging trace snapshots, resume *live* debugging.

`tfind end` Same as ‘`tfind none`’.

`tfind` No argument means find the next trace snapshot.

`tfind -` Find the previous trace snapshot before the current one. This permits retracing earlier steps.

`tfind tracepoint num`

Find the next snapshot associated with tracepoint *num*. Search proceeds forward from the last examined trace snapshot. If no argument *num* is given, it means find the next snapshot collected for the same tracepoint as the current snapshot.

`tfind pc addr`

Find the next snapshot associated with the value *addr* of the program counter. Search proceeds forward from the last examined trace snapshot. If no argument

addr is given, it means find the next snapshot with the same value of PC as the current snapshot.

tfind outside *addr1*, *addr2*

Find the next snapshot whose PC is outside the given range of addresses.

tfind range *addr1*, *addr2*

Find the next snapshot whose PC is between *addr1* and *addr2*.

tfind line [*file*]:*n*

Find the next snapshot associated with the source line *n*. If the optional argument *file* is given, refer to line *n* in that source file. Search proceeds forward from the last examined trace snapshot. If no argument *n* is given, it means find the next line other than the one currently being examined; thus saying **tfind line** repeatedly can appear to have the same effect as stepping from line to line in a *live* debugging session.

The default arguments for the **tfind** commands are specifically designed to make it easy to scan through the trace buffer. For instance, **tfind** with no argument selects the next trace snapshot, and **tfind -** with no argument selects the previous trace snapshot. So, by giving one **tfind** command, and then simply hitting **<RET>** repeatedly you can examine all the trace snapshots in order. Or, by saying **tfind -** and then hitting **<RET>** repeatedly you can examine the snapshots in reverse order. The **tfind line** command with no argument selects the snapshot for the next source line executed. The **tfind pc** command with no argument selects the next snapshot with the same program counter (PC) as the current frame. The **tfind tracepoint** command with no argument selects the next trace snapshot collected by the same tracepoint as the current one.

In addition to letting you scan through the trace buffer manually, these commands make it easy to construct GDB scripts that scan through the trace buffer and print out whatever collected data you are interested in. Thus, if we want to examine the PC, FP, and SP registers from each trace frame in the buffer, we can say this:

```
(gdb) tfind start
(gdb) while ($trace_frame != -1)
> printf "Frame %d, PC = %08X, SP = %08X, FP = %08X\n", \
        $trace_frame, $pc, $sp, $fp
> tfind
> end

Frame 0, PC = 0020DC64, SP = 0030BF3C, FP = 0030BF44
Frame 1, PC = 0020DC6C, SP = 0030BF38, FP = 0030BF44
Frame 2, PC = 0020DC70, SP = 0030BF34, FP = 0030BF44
Frame 3, PC = 0020DC74, SP = 0030BF30, FP = 0030BF44
Frame 4, PC = 0020DC78, SP = 0030BF2C, FP = 0030BF44
Frame 5, PC = 0020DC7C, SP = 0030BF28, FP = 0030BF44
Frame 6, PC = 0020DC80, SP = 0030BF24, FP = 0030BF44
Frame 7, PC = 0020DC84, SP = 0030BF20, FP = 0030BF44
Frame 8, PC = 0020DC88, SP = 0030BF1C, FP = 0030BF44
Frame 9, PC = 0020DC8E, SP = 0030BF18, FP = 0030BF44
Frame 10, PC = 00203F6C, SP = 0030BE3C, FP = 0030BF14
```

Or, if we want to examine the variable X at each source line in the buffer:

```
(gdb) tfind start
(gdb) while ($trace_frame != -1)
```

```
> printf "Frame %d, X == %d\n", $trace_frame, X
> tfind line
> end

Frame 0, X = 1
Frame 7, X = 2
Frame 13, X = 255
```

10.2.2 tdump

This command takes no arguments. It prints all the data collected at the current trace snapshot.

```
(gdb) trace 444
(gdb) actions
Enter actions for tracepoint #2, one per line:
> collect $regs, $locals, $args, gdb_long_test
> end

(gdb) tstart

(gdb) tfind line 444
#0  gdb_test (p1=0x11, p2=0x22, p3=0x33, p4=0x44, p5=0x55, p6=0x66)
at gdb_test.c:444
444      printp( "%s: arguments = 0x%X 0x%X 0x%X 0x%X 0x%X 0x%X\n", )

(gdb) tdump
Data collected at tracepoint 2, trace frame 1:
d0          0xc4aa0085      -995491707
d1          0x18      24
d2          0x80      128
d3          0x33      51
d4          0x71aea3d      119204413
d5          0x22      34
d6          0xe0      224
d7          0x380035 3670069
a0          0x19e24a 1696330
a1          0x3000668     50333288
a2          0x100     256
a3          0x322000 3284992
a4          0x3000698     50333336
a5          0x1ad3cc 1758156
fp          0x30bf3c 0x30bf3c
sp          0x30bf34 0x30bf34
ps          0x0      0
pc          0x20b2c8 0x20b2c8
fpcontrol   0x0      0
fpstatus    0x0      0
fpiaddr    0x0      0
p = 0x20e5b4 "gdb-test"
p1 = (void *) 0x11
p2 = (void *) 0x22
p3 = (void *) 0x33
p4 = (void *) 0x44
p5 = (void *) 0x55
p6 = (void *) 0x66
gdb_long_test = 17 '\021'

(gdb)
```

10.2.3 save-tracepoints *filename*

This command saves all current tracepoint definitions together with their actions and passcounts, into a file ‘*filename*’ suitable for use in a later debugging session. To read the saved tracepoint definitions, use the `source` command (see [Section 20.3 \[Command Files\], page 189](#)).

10.3 Convenience Variables for Tracepoints

(int) `$trace_frame`

The current trace snapshot (a.k.a. *frame*) number, or -1 if no snapshot is selected.

(int) `$tracepoint`

The tracepoint for the current trace snapshot.

(int) `$trace_line`

The line number for the current trace snapshot.

(char []) `$trace_file`

The source file for the current trace snapshot.

(char []) `$trace_func`

The name of the function containing `$tracepoint`.

Note: `$trace_file` is not suitable for use in `printf`, use `output` instead.

Here’s a simple example of using these convenience variables for stepping through all the trace snapshots and printing some of their data.

```
(gdb) tfind start  
  
(gdb) while $trace_frame != -1  
> output $trace_file  
> printf ", line %d (tracepoint #%d)\n", $trace_line, $tracepoint  
> tfind  
> end
```

11 Debugging Programs That Use Overlays


If your program is too large to fit completely in your target system's memory, you can sometimes use *overlays* to work around this problem. GDB provides some support for debugging programs that use overlays.

11.1 How Overlays Work

Suppose you have a computer whose instruction address space is only 64 kilobytes long, but which has much more memory which can be accessed by other means: special instructions, segment registers, or memory management hardware, for example. Suppose further that you want to adapt a program which is larger than 64 kilobytes to run on this system.

One solution is to identify modules of your program which are relatively independent, and need not call each other directly; call these modules *overlays*. Separate the overlays from the main program, and place their machine code in the larger memory. Place your main program in instruction memory, but leave at least enough space there to hold the largest overlay as well.

Now, to call a function located in an overlay, you must first copy that overlay's machine code from the large memory into the space set aside for it in the instruction memory, and then jump to its entry point there.



The diagram (see [A code overlay], page 99) shows a system with separate data and instruction address spaces. To map an overlay, the program copies its code from the larger address space to the instruction address space. Since the overlays shown here all use the same mapped address, only one may be mapped at a time. For a system with a single address space for data and instructions, the diagram would be similar, except that the program variables and heap would share an address space with the main program and the overlay area.

An overlay loaded into instruction memory and ready for use is called a *mapped* overlay; its *mapped address* is its address in the instruction memory. An overlay not present (or only partially present) in instruction memory is called *unmapped*; its *load address* is its address in the larger memory. The mapped address is also called the *virtual memory address*, or *VMA*; the load address is also called the *load memory address*, or *LMA*.

Unfortunately, overlays are not a completely transparent way to adapt a program to limited instruction memory. They introduce a new set of global constraints you must keep in mind as you design your program:

- Before calling or returning to a function in an overlay, your program must make sure that overlay is actually mapped. Otherwise, the call or return will transfer control to the right address, but in the wrong overlay, and your program will probably crash.
- If the process of mapping an overlay is expensive on your system, you will need to choose your overlays carefully to minimize their effect on your program’s performance.
- The executable file you load onto your system must contain each overlay’s instructions, appearing at the overlay’s load address, not its mapped address. However, each overlay’s instructions must be relocated and its symbols defined as if the overlay were at its mapped address. You can use GNU linker scripts to specify different load and relocation addresses for pieces of your program; see [section “Overlay Description” in *Using ld: the GNU linker*](#).
- The procedure for loading executable files onto your system must be able to load their contents into the larger address space as well as the instruction and data spaces.

The overlay system described above is rather simple, and could be improved in many ways:

- If your system has suitable bank switch registers or memory management hardware, you could use those facilities to make an overlay’s load area contents simply appear at their mapped address in instruction space. This would probably be faster than copying the overlay to its mapped area in the usual way.
- If your overlays are small enough, you could set aside more than one overlay area, and have more than one overlay mapped at a time.
- You can use overlays to manage data, as well as instructions. In general, data overlays are even less transparent to your design than code overlays: whereas code overlays only require care when you call or return to functions, data overlays require care every time you access the data. Also, if you change the contents of a data overlay, you must copy its contents back out to its load address before you can copy a different data overlay into the same mapped area.

11.2 Overlay Commands

To use GDB’s overlay support, each overlay in your program must correspond to a separate section of the executable file. The section’s virtual memory address and load memory address must be the overlay’s mapped and load addresses. Identifying overlays with sections allows GDB to determine the appropriate address of a function or variable, depending on whether the overlay is mapped or not.

GDB’s overlay commands all start with the word `overlay`; you can abbreviate this as `ov` or `ovly`. The commands are:

overlay off

Disable GDB’s overlay support. When overlay support is disabled, GDB assumes that all functions and variables are always present at their mapped addresses. By default, GDB’s overlay support is disabled.

overlay manual

Enable *manual* overlay debugging. In this mode, GDB relies on you to tell it which overlays are mapped, and which are not, using the `overlay map-overlay` and `overlay unmap-overlay` commands described below.

overlay map-overlay *overlay***overlay map *overlay***

Tell GDB that *overlay* is now mapped; *overlay* must be the name of the object file section containing the overlay. When an overlay is mapped, GDB assumes it can find the overlay’s functions and variables at their mapped addresses. GDB assumes that any other overlays whose mapped ranges overlap that of *overlay* are now unmapped.

overlay unmap-overlay *overlay***overlay unmap *overlay***

Tell GDB that *overlay* is no longer mapped; *overlay* must be the name of the object file section containing the overlay. When an overlay is unmapped, GDB assumes it can find the overlay’s functions and variables at their load addresses.

overlay auto

Enable *automatic* overlay debugging. In this mode, GDB consults a data structure the overlay manager maintains in the inferior to see which overlays are mapped. For details, see [Section 11.3 \[Automatic Overlay Debugging\]](#), page 102.

overlay load-target**overlay load**

Re-read the overlay table from the inferior. Normally, GDB re-reads the table GDB automatically each time the inferior stops, so this command should only be necessary if you have changed the overlay mapping yourself using GDB. This command is only useful when using automatic overlay debugging.

overlay list-overlays**overlay list**

Display a list of the overlays currently mapped, along with their mapped addresses, load addresses, and sizes.

Normally, when GDB prints a code address, it includes the name of the function the address falls in:

```
(gdb) print main  
$3 = {int ()} 0x11a0 <main>
```

When overlay debugging is enabled, GDB recognizes code in unmapped overlays, and prints the names of unmapped functions with asterisks around them. For example, if `foo` is a function in an unmapped overlay, GDB prints it this way:

```
(gdb) overlay list  
No sections are mapped.  
(gdb) print foo
```

```
$5 = {int (int)} 0x100000 <*foo*>
```

When `foo`'s overlay is mapped, GDB prints the function's name normally:

```
(gdb) overlay list
Section .ov.foo.text, loaded at 0x100000 - 0x100034,
      mapped at 0x1016 - 0x104a
(gdb) print foo
$6 = {int (int)} 0x1016 <foo>
```

When overlay debugging is enabled, GDB can find the correct address for functions and variables in an overlay, whether or not the overlay is mapped. This allows most GDB commands, like `break` and `disassemble`, to work normally, even on unmapped code. However, GDB's breakpoint support has some limitations:

- You can set breakpoints in functions in unmapped overlays, as long as GDB can write to the overlay at its load address.
- GDB can not set hardware or simulator-based breakpoints in unmapped overlays. However, if you set a breakpoint at the end of your overlay manager (and tell GDB which overlays are now mapped, if you are using manual overlay management), GDB will re-set its breakpoints properly.

11.3 Automatic Overlay Debugging

GDB can automatically track which overlays are mapped and which are not, given some simple co-operation from the overlay manager in the inferior. If you enable automatic overlay debugging with the `overlay auto` command (see [Section 11.2 \[Overlay Commands\], page 100](#)), GDB looks in the inferior's memory for certain variables describing the current state of the overlays.

Here are the variables your overlay manager must define to support GDB's automatic overlay debugging:

`_ovly_table:`

This variable must be an array of the following structures:

```
struct
{
    /* The overlay's mapped address. */
    unsigned long vma;

    /* The size of the overlay, in bytes. */
    unsigned long size;

    /* The overlay's load address. */
    unsigned long lma;

    /* Non-zero if the overlay is currently mapped;
       zero otherwise. */
    unsigned long mapped;
}
```

`_novlys:` This variable must be a four-byte signed integer, holding the total number of elements in `_ovly_table`.

To decide whether a particular overlay is mapped or not, GDB looks for an entry in `_ovly_table` whose `vma` and `lma` members equal the VMA and LMA of the overlay's section

in the executable file. When GDB finds a matching entry, it consults the entry's `mapped` member to determine whether the overlay is currently mapped.

In addition, your overlay manager may define a function called `_ovly_debug_event`. If this function is defined, GDB will silently set a breakpoint there. If the overlay manager then calls this function whenever it has changed the overlay table, this will enable GDB to accurately keep track of which overlays are in program memory, and update any breakpoints that may be set in overlays. This will allow breakpoints to work even if the overlays are kept in ROM or other non-writable memory while they are not being executed.

11.4 Overlay Sample Program

When linking a program which uses overlays, you must place the overlays at their load addresses, while relocating them to run at their mapped addresses. To do this, you must write a linker script (see section “Overlay Description” in *Using ld: the GNU linker*). Unfortunately, since linker scripts are specific to a particular host system, target architecture, and target memory layout, this manual cannot provide portable sample code demonstrating GDB’s overlay support.

However, the GDB source distribution does contain an overlaid program, with linker scripts for a few systems, as part of its test suite. The program consists of the following files from ‘`gdb/testsuite/gdb.base`’:

‘`overlays.c`’

The main program file.

‘`ovlymgr.c`’

A simple overlay manager, used by ‘`overlays.c`’.

‘`foo.c`’

‘`bar.c`’

‘`baz.c`’

‘`grbx.c`’ Overlay modules, loaded and used by ‘`overlays.c`’.

‘`d10v.1d`’

‘`m32r.1d`’ Linker scripts for linking the test program on the `d10v-elf` and `m32r-elf` targets.

You can build the test program using the `d10v-elf` GCC cross-compiler like this:

```
$ d10v-elf-gcc -g -c overlays.c
$ d10v-elf-gcc -g -c ovlymgr.c
$ d10v-elf-gcc -g -c foo.c
$ d10v-elf-gcc -g -c bar.c
$ d10v-elf-gcc -g -c baz.c
$ d10v-elf-gcc -g -c grbx.c
$ d10v-elf-gcc -g overlays.o ovlymgr.o foo.o bar.o \
                           baz.o grbx.o -Wl,-Td10v.1d -o overlays
```

The build process is identical for any other architecture, except that you must substitute the appropriate compiler and linker script for the target system for `d10v-elf-gcc` and `d10v.1d`.

12 Using GDB with Different Languages

Although programming languages generally have common aspects, they are rarely expressed in the same manner. For instance, in ANSI C, dereferencing a pointer `p` is accomplished by `*p`, but in Modula-2, it is accomplished by `p^`. Values can also be represented (and displayed) differently. Hex numbers in C appear as ‘`0x1ae`’, while in Modula-2 they appear as ‘`1AEH`’.

Language-specific information is built into GDB for some languages, allowing you to express operations like the above in your program’s native language, and allowing GDB to output values in a manner consistent with the syntax of your program’s native language. The language you use to build expressions is called the *working language*.

12.1 Switching between source languages

There are two ways to control the working language—either have GDB set it automatically, or select it manually yourself. You can use the `set language` command for either purpose. On startup, GDB defaults to setting the language automatically. The working language is used to determine how expressions you type are interpreted, how values are printed, etc.

In addition to the working language, every source file that GDB knows about has its own working language. For some object file formats, the compiler might indicate which language a particular source file is in. However, most of the time GDB infers the language from the name of the file. The language of a source file controls whether C++ names are demangled—this way `backtrace` can show each frame appropriately for its own language. There is no way to set the language of a source file from within GDB, but you can set the language associated with a filename extension. See [Section 12.2 \[Displaying the language\], page 106](#).

This is most commonly a problem when you use a program, such as `cfront` or `f2c`, that generates C but is written in another language. In that case, make the program use `#line` directives in its C output; that way GDB will know the correct language of the source code of the original program, and will display that source code, not the generated C code.

12.1.1 List of filename extensions and languages

If a source file name ends in one of the following extensions, then GDB infers that its language is the one indicated.

‘.ada’	
‘.ads’	
‘.adb’	
‘.a’	Ada source file.
‘.c’	C source file
‘.C’	
‘.cc’	
‘.cp’	
‘.cpp’	
‘.cxx’	
‘.c++’	C++ source file

‘.m’	Objective-C source file
‘.f’	
‘.F’	Fortran source file
‘.mod’	Modula-2 source file
‘.s’	
‘.S’	Assembler source file. This actually behaves almost like C, but GDB does not skip over function prologues when stepping.

In addition, you may set the language associated with a filename extension. See [Section 12.2 \[Displaying the language\]](#), page 106.

12.1.2 Setting the working language

If you allow GDB to set the language automatically, expressions are interpreted the same way in your debugging session and your program.

If you wish, you may set the language manually. To do this, issue the command ‘`set language lang`’, where `lang` is the name of a language, such as `c` or `modula-2`. For a list of the supported languages, type ‘`set language`’.

Setting the language manually prevents GDB from updating the working language automatically. This can lead to confusion if you try to debug a program when the working language is not the same as the source language, when an expression is acceptable to both languages—but means different things. For instance, if the current source file were written in C, and GDB was parsing Modula-2, a command such as:

```
print a = b + c
```

might not have the effect you intended. In C, this means to add `b` and `c` and place the result in `a`. The result printed would be the value of `a`. In Modula-2, this means to compare `a` to the result of `b+c`, yielding a BOOLEAN value.

12.1.3 Having GDB infer the source language

To have GDB set the working language automatically, use ‘`set language local`’ or ‘`set language auto`’. GDB then infers the working language. That is, when your program stops in a frame (usually by encountering a breakpoint), GDB sets the working language to the language recorded for the function in that frame. If the language for a frame is unknown (that is, if the function or block corresponding to the frame was defined in a source file that does not have a recognized extension), the current working language is not changed, and GDB issues a warning.

This may not seem necessary for most programs, which are written entirely in one source language. However, program modules and libraries written in one source language can be used by a main program written in a different source language. Using ‘`set language auto`’ in this case frees you from having to set the working language manually.

12.2 Displaying the language

The following commands help you find out which language is the working language, and also what language source files were written in.

show language

Display the current working language. This is the language you can use with commands such as `print` to build and compute expressions that may involve variables in your program.

info frame

Display the source language for this frame. This language becomes the working language if you use an identifier from this frame. See [Section 6.4 \[Information about a frame\]](#), page 56, to identify the other information listed here.

info source

Display the source language of this source file. See [Chapter 13 \[Examining the Symbol Table\]](#), page 125, to identify the other information listed here.

In unusual circumstances, you may have source files with extensions not in the standard list. You can then set the extension associated with a language explicitly:

set extension-language .ext language

Set source files with extension `.ext` to be assumed to be in the source language `language`.

info extensions

List all the filename extensions and the associated languages.

12.3 Type and range checking

Warning: In this release, the GDB commands for type and range checking are included, but they do not yet have any effect. This section documents the intended facilities.

Some languages are designed to guard you against making seemingly common errors through a series of compile- and run-time checks. These include checking the type of arguments to functions and operators, and making sure mathematical overflows are caught at run time. Checks such as these help to ensure a program's correctness once it has been compiled by eliminating type mismatches, and providing active checks for range errors when your program is running.

GDB can check for conditions like the above if you wish. Although GDB does not check the statements in your program, it can check expressions entered directly into GDB for evaluation via the `print` command, for example. As with the working language, GDB can also decide whether or not to check automatically based on your program's source language. See [Section 12.4 \[Supported languages\]](#), page 109, for the default settings of supported languages.

12.3.1 An overview of type checking

Some languages, such as Modula-2, are strongly typed, meaning that the arguments to operators and functions have to be of the correct type, otherwise an error occurs. These checks prevent type mismatch errors from ever causing any run-time problems. For example,

1 + 2 \Rightarrow 3
but
error 1 + 2.3

The second example fails because the `CARDINAL 1` is not type-compatible with the `REAL 2.3`.

For the expressions you use in GDB commands, you can tell the GDB type checker to skip checking; to treat any mismatches as errors and abandon the expression; or to only issue warnings when type mismatches occur, but evaluate the expression anyway. When you choose the last of these, GDB evaluates expressions like the second example above, but also issues a warning.

Even if you turn type checking off, there may be other reasons related to type that prevent GDB from evaluating an expression. For instance, GDB does not know how to add an `int` and a `struct foo`. These particular type errors have nothing to do with the language in use, and usually arise from expressions, such as the one described above, which make little sense to evaluate anyway.

Each language defines to what degree it is strict about type. For instance, both Modula-2 and C require the arguments to arithmetical operators to be numbers. In C, enumerated types and pointers can be represented as numbers, so that they are valid arguments to mathematical operators. See [Section 12.4 \[Supported languages\]](#), page 109, for further details on specific languages.

GDB provides some additional commands for controlling the type checker:

`set check type auto`

Set type checking on or off based on the current working language. See [Section 12.4 \[Supported languages\]](#), page 109, for the default settings for each language.

`set check type on`

`set check type off`

Set type checking on or off, overriding the default setting for the current working language. Issue a warning if the setting does not match the language default. If any type mismatches occur in evaluating an expression while type checking is on, GDB prints a message and aborts evaluation of the expression.

`set check type warn`

Cause the type checker to issue warnings, but to always attempt to evaluate the expression. Evaluating the expression may still be impossible for other reasons. For example, GDB cannot add numbers and structures.

`show type` Show the current setting of the type checker, and whether or not GDB is setting it automatically.

12.3.2 An overview of range checking

In some languages (such as Modula-2), it is an error to exceed the bounds of a type; this is enforced with run-time checks. Such range checking is meant to ensure program correctness by making sure computations do not overflow, or indices on an array element access do not exceed the bounds of the array.

For expressions you use in GDB commands, you can tell GDB to treat range errors in one of three ways: ignore them, always treat them as errors and abandon the expression, or issue warnings but evaluate the expression anyway.

A range error can result from numerical overflow, from exceeding an array index bound, or when you type a constant that is not a member of any type. Some languages, however, do not treat overflows as an error. In many implementations of C, mathematical overflow causes the result to “wrap around” to lower values—for example, if m is the largest integer value, and s is the smallest, then

$$m + 1 \Rightarrow s$$

This, too, is specific to individual languages, and in some cases specific to individual compilers or machines. See [Section 12.4 \[Supported languages\], page 109](#), for further details on specific languages.

GDB provides some additional commands for controlling the range checker:

`set check range auto`

Set range checking on or off based on the current working language. See [Section 12.4 \[Supported languages\], page 109](#), for the default settings for each language.

`set check range on`

`set check range off`

Set range checking on or off, overriding the default setting for the current working language. A warning is issued if the setting does not match the language default. If a range error occurs and range checking is on, then a message is printed and evaluation of the expression is aborted.

`set check range warn`

Output messages when the GDB range checker detects a range error, but attempt to evaluate the expression anyway. Evaluating the expression may still be impossible for other reasons, such as accessing memory that the process does not own (a typical example from many Unix systems).

`show range`

Show the current setting of the range checker, and whether or not it is being set automatically by GDB.

12.4 Supported languages

GDB supports C, C++, Objective-C, Fortran, Java, assembly, Modula-2, and Ada. Some GDB features may be used in expressions regardless of the language you use: the GDB @ and :: operators, and the ‘{type}addr’ construct (see [Section 8.1 \[Expressions\], page 65](#)) can be used with the constructs of any supported language.

The following sections detail to what degree each source language is supported by GDB. These sections are not meant to be language tutorials or references, but serve only as a reference guide to what the GDB expression parser accepts, and what input and output formats should look like for different languages. There are many good books written on each of these languages; please look to these for a language reference or tutorial.

12.4.1 C and C++

Since C and C++ are so closely related, many features of GDB apply to both languages. Whenever this is the case, we discuss those languages together.

The C++ debugging facilities are jointly implemented by the C++ compiler and GDB. Therefore, to debug your C++ code effectively, you must compile your C++ programs with a supported C++ compiler, such as GNU `g++`, or the HP ANSI C++ compiler (`aCC`).

For best results when using GNU C++, use the DWARF 2 debugging format; if it doesn't work on your system, try the stabs+ debugging format. You can select those formats explicitly with the `g++` command-line options '`-gdwarf-2`' and '`-gstabs+`'. See section "Options for Debugging Your Program or GNU CC" in *Using GNU CC*.

12.4.1.1 C and C++ operators

Operators must be defined on values of specific types. For instance, `+` is defined on numbers, but not on structures. Operators are often defined on groups of types.

For the purposes of C and C++, the following definitions hold:

- *Integral types* include `int` with any of its storage-class specifiers; `char`; `enum`; and, for C++, `bool`.
- *Floating-point types* include `float`, `double`, and `long double` (if supported by the target platform).
- *Pointer types* include all types defined as `(type *)`.
- *Scalar types* include all of the above.

The following operators are supported. They are listed here in order of increasing precedence:

,	The comma or sequencing operator. Expressions in a comma-separated list are evaluated from left to right, with the result of the entire expression being the last expression evaluated.
=	Assignment. The value of an assignment expression is the value assigned. Defined on scalar types.
<code>op=</code>	Used in an expression of the form <code>a op= b</code> , and translated to <code>a = a op b</code> . <code>op=</code> and <code>=</code> have the same precedence. <code>op</code> is any one of the operators <code> </code> , <code>^</code> , <code>&</code> , <code><<</code> , <code>>></code> , <code>+</code> , <code>-</code> , <code>*</code> , <code>/</code> , <code>%</code> .
<code>? :</code>	The ternary operator. <code>a ? b : c</code> can be thought of as: if <code>a</code> then <code>b</code> else <code>c</code> . <code>a</code> should be of an integral type.
<code> </code>	Logical OR. Defined on integral types.
<code>&&</code>	Logical AND. Defined on integral types.
<code> </code>	Bitwise OR. Defined on integral types.
<code>^</code>	Bitwise exclusive-OR. Defined on integral types.
<code>&</code>	Bitwise AND. Defined on integral types.
<code>==, !=</code>	Equality and inequality. Defined on scalar types. The value of these expressions is 0 for false and non-zero for true.
<code><, >, <=, >=</code>	Less than, greater than, less than or equal, greater than or equal. Defined on scalar types. The value of these expressions is 0 for false and non-zero for true.

<code><<, >></code>	left shift, and right shift. Defined on integral types.
<code>@</code>	The GDB “artificial array” operator (see Section 8.1 [Expressions] , page 65).
<code>+, -</code>	Addition and subtraction. Defined on integral types, floating-point types and pointer types.
<code>*, /, %</code>	Multiplication, division, and modulus. Multiplication and division are defined on integral and floating-point types. Modulus is defined on integral types.
<code>++, --</code>	Increment and decrement. When appearing before a variable, the operation is performed before the variable is used in an expression; when appearing after it, the variable’s value is used before the operation takes place.
<code>*</code>	Pointer dereferencing. Defined on pointer types. Same precedence as <code>++</code> .
<code>&</code>	Address operator. Defined on variables. Same precedence as <code>++</code> . For debugging C++, GDB implements a use of ‘ <code>&</code> ’ beyond what is allowed in the C++ language itself: you can use ‘ <code>&(&ref)</code> ’ (or, if you prefer, simply ‘ <code>&&ref</code> ’) to examine the address where a C++ reference variable (declared with ‘ <code>&ref</code> ’) is stored.
<code>-</code>	Negative. Defined on integral and floating-point types. Same precedence as <code>++</code> .
<code>!</code>	Logical negation. Defined on integral types. Same precedence as <code>++</code> .
<code>~</code>	Bitwise complement operator. Defined on integral types. Same precedence as <code>++</code> .
<code>., -></code>	Structure member, and pointer-to-structure member. For convenience, GDB regards the two as equivalent, choosing whether to dereference a pointer based on the stored type information. Defined on <code>struct</code> and <code>union</code> data.
<code>.*, ->*</code>	Dereferences of pointers to members.
<code>[]</code>	Array indexing. <code>a[i]</code> is defined as <code>*(a+i)</code> . Same precedence as <code>-></code> .
<code>()</code>	Function parameter list. Same precedence as <code>-></code> .
<code>::</code>	C++ scope resolution operator. Defined on <code>struct</code> , <code>union</code> , and <code>class</code> types.
<code>::</code>	Doubled colons also represent the GDB scope operator (see Section 8.1 [Expressions] , page 65). Same precedence as <code>::</code> , above.

If an operator is redefined in the user code, GDB usually attempts to invoke the redefined version instead of using the operator’s predefined meaning.

12.4.1.2 C and C++ constants

GDB allows you to express the constants of C and C++ in the following ways:

- Integer constants are a sequence of digits. Octal constants are specified by a leading ‘0’ (i.e. zero), and hexadecimal constants by a leading ‘0x’ or ‘0X’. Constants may also end with a letter ‘l’, specifying that the constant should be treated as a `long` value.
- Floating point constants are a sequence of digits, followed by a decimal point, followed by a sequence of digits, and optionally followed by an exponent. An exponent is of the form: ‘`e[+]|-]nnn`’, where `nnn` is another sequence of digits. The ‘+’ is optional

for positive exponents. A floating-point constant may also end with a letter ‘f’ or ‘F’, specifying that the constant should be treated as being of the `float` (as opposed to the default `double`) type; or with a letter ‘l’ or ‘L’, which specifies a `long double` constant.

- Enumerated constants consist of enumerated identifiers, or their integral equivalents.
- Character constants are a single character surrounded by single quotes (‘ ’), or a number—the ordinal value of the corresponding character (usually its ASCII value). Within quotes, the single character may be represented by a letter or by escape sequences, which are of the form ‘\nnn’, where *nnn* is the octal representation of the character’s ordinal value; or of the form ‘\x’, where ‘x’ is a predefined special character—for example, ‘\n’ for newline.
- String constants are a sequence of character constants surrounded by double quotes (“ ”). Any valid character constant (as described above) may appear. Double quotes within the string must be preceded by a backslash, so for instance “a\b\c” is a string of five characters.
- Pointer constants are an integral value. You can also write pointers to constants using the C operator ‘&’.
- Array constants are comma-separated lists surrounded by braces ‘{’ and ‘}’; for example, ‘{1,2,3}’ is a three-element array of integers, ‘{{1,2}, {3,4}, {5,6}}’ is a three-by-two array, and ‘{&"hi", &"there", &"fred"}’ is a three-element array of pointers.

12.4.1.3 C++ expressions

GDB expression handling can interpret most C++ expressions.

Warning: GDB can only debug C++ code if you use the proper compiler and the proper debug format. Currently, GDB works best when debugging C++ code that is compiled with GCC 2.95.3 or with GCC 3.1 or newer, using the options ‘-gdwarf-2’ or ‘-gstabs+’. DWARF 2 is preferred over stabs+. Most configurations of GCC emit either DWARF 2 or stabs+ as their default debug format, so you usually don’t need to specify a debug format explicitly. Other compilers and/or debug formats are likely to work badly or not at all when using GDB to debug C++ code.

1. Member function calls are allowed; you can use expressions like

```
count = aml->GetOriginal(x, y)
```

2. While a member function is active (in the selected stack frame), your expressions have the same namespace available as the member function; that is, GDB allows implicit references to the class instance pointer `this` following the same rules as C++.
3. You can call overloaded functions; GDB resolves the function call to the right definition, with some restrictions. GDB does not perform overload resolution involving user-defined type conversions, calls to constructors, or instantiations of templates that do not exist in the program. It also cannot handle ellipsis argument lists or default arguments.

It does perform integral conversions and promotions, floating-point promotions, arithmetic conversions, pointer conversions, conversions of class objects to base classes, and standard conversions such as those of functions or arrays to pointers; it requires an exact match on the number of function arguments.

Overload resolution is always performed, unless you have specified `set overload-resolution off`. See [Section 12.4.1.7 \[GDB features for C++\]](#), page 114.

You must specify `set overload-resolution off` in order to use an explicit function signature to call an overloaded function, as in

```
p 'foo(char,int)'('x', 13)
```

The GDB command-completion facility can simplify this; see [Section 3.2 \[Command completion\]](#), page 17.

4. GDB understands variables declared as C++ references; you can use them in expressions just as you do in C++ source—they are automatically dereferenced.

In the parameter list shown when GDB displays a frame, the values of reference variables are not displayed (unlike other variables); this avoids clutter, since references are often used for large structures. The *address* of a reference variable is always shown, unless you have specified ‘`set print address off`’.

5. GDB supports the C++ name resolution operator `::`—your expressions can use it just as expressions in your program do. Since one scope may be defined in another, you can use `::` repeatedly if necessary, for example in an expression like ‘`scope1::scope2::name`’. GDB also allows resolving name scope by reference to source files, in both C and C++ debugging (see [Section 8.2 \[Program variables\]](#), page 66).

In addition, when used with HP’s C++ compiler, GDB supports calling virtual functions correctly, printing out virtual bases of objects, calling functions in a base subobject, casting objects, and invoking user-defined operators.

12.4.1.4 C and C++ defaults

If you allow GDB to set type and range checking automatically, they both default to `off` whenever the working language changes to C or C++. This happens regardless of whether you or GDB selects the working language.

If you allow GDB to set the language automatically, it recognizes source files whose names end with ‘.c’, ‘.C’, or ‘.cc’, etc, and when GDB enters code compiled from one of these files, it sets the working language to C or C++. See [Section 12.1.3 \[Having GDB infer the source language\]](#), page 106, for further details.

12.4.1.5 C and C++ type and range checks

By default, when GDB parses C or C++ expressions, type checking is not used. However, if you turn type checking on, GDB considers two variables type equivalent if:

- The two variables are structured and have the same structure, union, or enumerated tag.
- The two variables have the same type name, or types that have been declared equivalent through `typedef`.

Range checking, if turned on, is done on mathematical operations. Array indices are not checked, since they are often used to index a pointer that is not itself an array.

12.4.1.6 GDB and C

The `set print union` and `show print union` commands apply to the `union` type. When set to ‘on’, any `union` that is inside a `struct` or `class` is also printed. Otherwise, it appears as ‘{...}’.

The `@` operator aids in the debugging of dynamic arrays, formed with pointers and a memory allocation function. See [Section 8.1 \[Expressions\]](#), page 65.

12.4.1.7 GDB features for C++

Some GDB commands are particularly useful with C++, and some are designed specifically for use with C++. Here is a summary:

`breakpoint` menus

When you want a breakpoint in a function whose name is overloaded, GDB breakpoint menus help you specify which function definition you want. See [Section 5.1.8 \[Breakpoint menus\]](#), page 45.

`rbreak regex`

Setting breakpoints using regular expressions is helpful for setting breakpoints on overloaded functions that are not members of any special classes. See [Section 5.1.1 \[Setting breakpoints\]](#), page 34.

`catch throw`

`catch catch`

Debug C++ exception handling using these commands. See [Section 5.1.3 \[Setting catchpoints\]](#), page 39.

`ptype typename`

Print inheritance relationships as well as other information for type `typename`. See [Chapter 13 \[Examining the Symbol Table\]](#), page 125.

`set print demangle`

`show print demangle`

`set print asm-demangle`

`show print asm-demangle`

Control whether C++ symbols display in their source form, both when displaying code as C++ source and when displaying disassemblies. See [Section 8.7 \[Print settings\]](#), page 71.

`set print object`

`show print object`

Choose whether to print derived (actual) or declared types of objects. See [Section 8.7 \[Print settings\]](#), page 71.

`set print vtbl`

`show print vtbl`

Control the format for printing virtual function tables. See [Section 8.7 \[Print settings\]](#), page 71. (The `vtbl` commands do not work on programs compiled with the HP ANSI C++ compiler (aCC).)

`set overload-resolution on`

Enable overload resolution for C++ expression evaluation. The default is on. For overloaded functions, GDB evaluates the arguments and searches for a function whose signature matches the argument types, using the standard C++ conversion rules (see [Section 12.4.1.3 \[C++ expressions\]](#), page 112, for details). If it cannot find a match, it emits a message.

```
set overload-resolution off
```

Disable overload resolution for C++ expression evaluation. For overloaded functions that are not class member functions, GDB chooses the first function of the specified name that it finds in the symbol table, whether or not its arguments are of the correct type. For overloaded functions that are class member functions, GDB searches for a function whose signature *exactly* matches the argument types.

Overloaded symbol names

You can specify a particular definition of an overloaded symbol, using the same notation that is used to declare such symbols in C++: type *symbol*(*types*) rather than just *symbol*. You can also use the GDB command-line word completion facilities to list the available choices, or to finish the type list for you. See [Section 3.2 \[Command completion\]](#), page 17, for details on how to do this.

12.4.2 Objective-C

This section provides information about some commands and command options that are useful for debugging Objective-C code.

12.4.2.1 Method Names in Commands

The following commands have been extended to accept Objective-C method names as line specifications:

- clear
- break
- info line
- jump
- list

A fully qualified Objective-C method name is specified as

```
-[Class methodName]
```

where the minus sign is used to indicate an instance method and a plus sign (not shown) is used to indicate a class method. The class name *Class* and method name *methodName* are enclosed in brackets, similar to the way messages are specified in Objective-C source code. For example, to set a breakpoint at the `create` instance method of class `Fruit` in the program currently being debugged, enter:

```
break -[Fruit create]
```

To list ten program lines around the `initialize` class method, enter:

```
list +[NSText initialize]
```

In the current version of GDB, the plus or minus sign is required. In future versions of GDB, the plus or minus sign will be optional, but you can use it to narrow the search. It is also possible to specify just a method name:

```
break create
```

You must specify the complete method name, including any colons. If your program's source files contain more than one `create` method, you'll be presented with a numbered list of classes that implement that method. Indicate your choice by number, or type '0' to exit if none apply.

As another example, to clear a breakpoint established at the `makeKeyAndOrderFront:` method of the `NSWindow` class, enter:

```
clear -[NSWindow makeKeyAndOrderFront:]
```

12.4.2.2 The Print Command With Objective-C

The print command has also been extended to accept methods. For example:

```
print -[object hash]
```

will tell GDB to send the `hash` message to `object` and print the result. Also, an additional command has been added, `print-object` or `po` for short, which is meant to print the description of an object. However, this command may only work with certain Objective-C libraries that have a particular hook function, `_NSPrintForDebugger`, defined.

12.4.3 Modula-2

The extensions made to GDB to support Modula-2 only support output from the GNU Modula-2 compiler (which is currently being developed). Other Modula-2 compilers are not currently supported, and attempting to debug executables produced by them is most likely to give an error as GDB reads in the executable's symbol table.

12.4.3.1 Operators

Operators must be defined on values of specific types. For instance, `+` is defined on numbers, but not on structures. Operators are often defined on groups of types. For the purposes of Modula-2, the following definitions hold:

- *Integral types* consist of `INTEGER`, `CARDINAL`, and their subranges.
- *Character types* consist of `CHAR` and its subranges.
- *Floating-point types* consist of `REAL`.
- *Pointer types* consist of anything declared as `POINTER TO type`.
- *Scalar types* consist of all of the above.
- *Set types* consist of `SET` and `BITSET` types.
- *Boolean types* consist of `BOOLEAN`.

The following operators are supported, and appear in order of increasing precedence:

,	Function argument or array index separator.
<code>:=</code>	Assignment. The value of <code>var := value</code> is <code>value</code> .
<code><, ></code>	Less than, greater than on integral, floating-point, or enumerated types.
<code><=, >=</code>	Less than or equal to, greater than or equal to on integral, floating-point and enumerated types, or set inclusion on set types. Same precedence as <code><</code> .
<code>=, <>, #</code>	Equality and two ways of expressing inequality, valid on scalar types. Same precedence as <code><</code> . In GDB scripts, only <code><></code> is available for inequality, since <code>#</code> conflicts with the script comment character.
<code>IN</code>	Set membership. Defined on set types and the types of their members. Same precedence as <code><</code> .
<code>OR</code>	Boolean disjunction. Defined on boolean types.

AND, &	Boolean conjunction. Defined on boolean types.
@	The GDB “artificial array” operator (see Section 8.1 [Expressions], page 65).
+, -	Addition and subtraction on integral and floating-point types, or union and difference on set types.
*	Multiplication on integral and floating-point types, or set intersection on set types.
/	Division on floating-point types, or symmetric set difference on set types. Same precedence as *.
DIV, MOD	Integer division and remainder. Defined on integral types. Same precedence as *.
-	Negative. Defined on INTEGER and REAL data.
^	Pointer dereferencing. Defined on pointer types.
NOT	Boolean negation. Defined on boolean types. Same precedence as ^.
.	RECORD field selector. Defined on RECORD data. Same precedence as ^.
[]	Array indexing. Defined on ARRAY data. Same precedence as ^.
()	Procedure argument list. Defined on PROCEDURE objects. Same precedence as ^.
::, .	GDB and Modula-2 scope operators.

Warning: Sets and their operations are not yet supported, so GDB treats the use of the operator IN, or the use of operators +, -, *, /, =, , <>, #, <=, and >= on sets as an error.

12.4.3.2 Built-in functions and procedures

Modula-2 also makes available several built-in procedures and functions. In describing these, the following metavariables are used:

a	represents an ARRAY variable.
c	represents a CHAR constant or variable.
i	represents a variable or constant of integral type.
m	represents an identifier that belongs to a set. Generally used in the same function with the metavariable s. The type of s should be SET OF <i>mtype</i> (where <i>mtype</i> is the type of m).
n	represents a variable or constant of integral or floating-point type.
r	represents a variable or constant of floating-point type.
t	represents a type.
v	represents a variable.
x	represents a variable or constant of one of many types. See the explanation of the function for details.

All Modula-2 built-in procedures also return a result, described below.

ABS(<i>n</i>)	Returns the absolute value of <i>n</i> .
CAP(<i>c</i>)	If <i>c</i> is a lower case letter, it returns its upper case equivalent, otherwise it returns its argument.
CHR(<i>i</i>)	Returns the character whose ordinal value is <i>i</i> .
DEC(<i>v</i>)	Decrements the value in the variable <i>v</i> by one. Returns the new value.
DEC(<i>v, i</i>)	Decrements the value in the variable <i>v</i> by <i>i</i> . Returns the new value.
EXCL(<i>m, s</i>)	Removes the element <i>m</i> from the set <i>s</i> . Returns the new set.
FLOAT(<i>i</i>)	Returns the floating point equivalent of the integer <i>i</i> .
HIGH(<i>a</i>)	Returns the index of the last member of <i>a</i> .
INC(<i>v</i>)	Increments the value in the variable <i>v</i> by one. Returns the new value.
INC(<i>v, i</i>)	Increments the value in the variable <i>v</i> by <i>i</i> . Returns the new value.
INCL(<i>m, s</i>)	Adds the element <i>m</i> to the set <i>s</i> if it is not already there. Returns the new set.
MAX(<i>t</i>)	Returns the maximum value of the type <i>t</i> .
MIN(<i>t</i>)	Returns the minimum value of the type <i>t</i> .
ODD(<i>i</i>)	Returns boolean TRUE if <i>i</i> is an odd number.
ORD(<i>x</i>)	Returns the ordinal value of its argument. For example, the ordinal value of a character is its ASCII value (on machines supporting the ASCII character set). <i>x</i> must be of an ordered type, which include integral, character and enumerated types.
SIZE(<i>x</i>)	Returns the size of its argument. <i>x</i> can be a variable or a type.
TRUNC(<i>r</i>)	Returns the integral part of <i>r</i> .
VAL(<i>t, i</i>)	Returns the member of the type <i>t</i> whose ordinal value is <i>i</i> .

Warning: Sets and their operations are not yet supported, so GDB treats the use of procedures INCL and EXCL as an error.

12.4.3.3 Constants

GDB allows you to express the constants of Modula-2 in the following ways:

- Integer constants are simply a sequence of digits. When used in an expression, a constant is interpreted to be type-compatible with the rest of the expression. Hexadecimal integers are specified by a trailing ‘H’, and octal integers by a trailing ‘B’.
- Floating point constants appear as a sequence of digits, followed by a decimal point and another sequence of digits. An optional exponent can then be specified, in the form ‘E[+|-]nnn’, where ‘[+|-]nnn’ is the desired exponent. All of the digits of the floating point constant must be valid decimal (base 10) digits.

- Character constants consist of a single character enclosed by a pair of like quotes, either single (') or double ("). They may also be expressed by their ordinal value (their ASCII value, usually) followed by a 'C'.
- String constants consist of a sequence of characters enclosed by a pair of like quotes, either single (') or double ("). Escape sequences in the style of C are also allowed. See [Section 12.4.1.2 \[C and C++ constants\]](#), page 111, for a brief explanation of escape sequences.
- Enumerated constants consist of an enumerated identifier.
- Boolean constants consist of the identifiers TRUE and FALSE.
- Pointer constants consist of integral values only.
- Set constants are not yet supported.

12.4.3.4 Modula-2 defaults

If type and range checking are set automatically by GDB, they both default to `on` whenever the working language changes to Modula-2. This happens regardless of whether you or GDB selected the working language.

If you allow GDB to set the language automatically, then entering code compiled from a file whose name ends with ‘.mod’ sets the working language to Modula-2. See [Section 12.1.3 \[Having GDB set the language automatically\]](#), page 106, for further details.

12.4.3.5 Deviations from standard Modula-2

A few changes have been made to make Modula-2 programs easier to debug. This is done primarily via loosening its type strictness:

- Unlike in standard Modula-2, pointer constants can be formed by integers. This allows you to modify pointer variables during debugging. (In standard Modula-2, the actual address contained in a pointer variable is hidden from you; it can only be modified through direct assignment to another pointer variable or expression that returned a pointer.)
- C escape sequences can be used in strings and characters to represent non-printable characters. GDB prints out strings with these escape sequences embedded. Single non-printable characters are printed using the ‘`CHR(nn)`’ format.
- The assignment operator (`:=`) returns the value of its right-hand argument.
- All built-in procedures both modify *and* return their argument.

12.4.3.6 Modula-2 type and range checks

Warning: in this release, GDB does not yet perform type or range checking.

GDB considers two Modula-2 variables type equivalent if:

- They are of types that have been declared equivalent via a `TYPE t1 = t2` statement
- They have been declared on the same line. (Note: This is true of the GNU Modula-2 compiler, but it may not be true of other compilers.)

As long as type checking is enabled, any attempt to combine variables whose types are not equivalent is an error.

Range checking is done on all mathematical operations, assignment, array index bounds, and all built-in functions and procedures.

12.4.3.7 The scope operators :: and .

There are a few subtle differences between the Modula-2 scope operator (.) and the GDB scope operator (::). The two have similar syntax:

```
module . id
scope :: id
```

where *scope* is the name of a module or a procedure, *module* the name of a module, and *id* is any declared identifier within your program, except another module.

Using the :: operator makes GDB search the scope specified by *scope* for the identifier *id*. If it is not found in the specified scope, then GDB searches all scopes enclosing the one specified by *scope*.

Using the . operator makes GDB search the current scope for the identifier specified by *id* that was imported from the definition module specified by *module*. With this operator, it is an error if the identifier *id* was not imported from definition module *module*, or if *id* is not an identifier in *module*.

12.4.3.8 GDB and Modula-2

Some GDB commands have little use when debugging Modula-2 programs. Five subcommands of `set print` and `show print` apply specifically to C and C++: ‘vtbl’, ‘demangle’, ‘asm-demangle’, ‘object’, and ‘union’. The first four apply to C++, and the last to the C `union` type, which has no direct analogue in Modula-2.

The @ operator (see [Section 8.1 \[Expressions\], page 65](#)), while available with any language, is not useful with Modula-2. Its intent is to aid the debugging of *dynamic arrays*, which cannot be created in Modula-2 as they can in C or C++. However, because an address can be specified by an integral constant, the construct ‘{`type`}`adrexp`’ is still useful.

In GDB scripts, the Modula-2 inequality operator # is interpreted as the beginning of a comment. Use <> instead.

12.4.4 Ada

The extensions made to GDB for Ada only support output from the GNU Ada (GNAT) compiler. Other Ada compilers are not currently supported, and attempting to debug executables produced by them is most likely to be difficult.

12.4.4.1 Introduction

The Ada mode of GDB supports a fairly large subset of Ada expression syntax, with some extensions. The philosophy behind the design of this subset is

- That GDB should provide basic literals and access to operations for arithmetic, dereferencing, field selection, indexing, and subprogram calls, leaving more sophisticated computations to subprograms written into the program (which therefore may be called from GDB).
- That type safety and strict adherence to Ada language restrictions are not particularly important to the GDB user.
- That brevity is important to the GDB user.

Thus, for brevity, the debugger acts as if there were implicit `with` and `use` clauses in effect for all user-written packages, making it unnecessary to fully qualify most names with

their packages, regardless of context. Where this causes ambiguity, GDB asks the user's intent.

The debugger will start in Ada mode if it detects an Ada main program. As for other languages, it will enter Ada mode when stopped in a program that was translated from an Ada source file.

While in Ada mode, you may use '--' for comments. This is useful mostly for documenting command files. The standard GDB comment ('#') still works at the beginning of a line in Ada mode, but not in the middle (to allow based literals).

The debugger supports limited overloading. Given a subprogram call in which the function symbol has multiple definitions, it will use the number of actual parameters and some information about their types to attempt to narrow the set of definitions. It also makes very limited use of context, preferring procedures to functions in the context of the `call` command, and functions to procedures elsewhere.

12.4.4.2 Omissions from Ada

Here are the notable omissions from the subset:

- Only a subset of the attributes are supported:
 - `'First`, `'Last`, and `'Length` on array objects (not on types and subtypes).
 - `'Min` and `'Max`.
 - `'Pos` and `'Val`.
 - `'Tag`.
 - `'Range` on array objects (not subtypes), but only as the right operand of the membership (`in`) operator.
 - `'Access`, `'Unchecked_Access`, and `'Unrestricted_Access` (a GNAT extension).
 - `'Address`.
- The names in `Characters.Latin_1` are not available and concatenation is not implemented. Thus, escape characters in strings are not currently available.
- Equality tests ('=' and '/=') on arrays test for bitwise equality of representations. They will generally work correctly for strings and arrays whose elements have integer or enumeration types. They may not work correctly for arrays whose element types have user-defined equality, for arrays of real values (in particular, IEEE-conformant floating point, because of negative zeroes and NaNs), and for arrays whose elements contain unused bits with indeterminate values.
- The other component-by-component array operations (`and`, `or`, `xor`, `not`, and relational tests other than equality) are not implemented.
- There are no record or array aggregates.
- Calls to dispatching subprograms are not implemented.
- The overloading algorithm is much more limited (i.e., less selective) than that of real Ada. It makes only limited use of the context in which a subexpression appears to resolve its meaning, and it is much looser in its rules for allowing type matches. As a result, some function calls will be ambiguous, and the user will be asked to choose the proper resolution.
- The `new` operator is not implemented.

- Entry calls are not implemented.
- Aside from printing, arithmetic operations on the native VAX floating-point formats are not supported.
- It is not possible to slice a packed array.

12.4.4.3 Additions to Ada

As it does for other languages, GDB makes certain generic extensions to Ada (see [Section 8.1 \[Expressions\], page 65](#)):

- If the expression E is a variable residing in memory (typically a local variable or array element) and N is a positive integer, then $E@N$ displays the values of E and the $N-1$ adjacent variables following it in memory as an array. In Ada, this operator is generally not necessary, since its prime use is in displaying parts of an array, and slicing will usually do this in Ada. However, there are occasional uses when debugging programs in which certain debugging information has been optimized away.
- $B::var$ means “the variable named *var* that appears in function or file *B*.” When *B* is a file name, you must typically surround it in single quotes.
- The expression $\{type\} addr$ means “the variable of type *type* that appears at address *addr*.”
- A name starting with ‘\$’ is a convenience variable (see [Section 8.9 \[Convenience Vars\], page 77](#)) or a machine register (see [Section 8.10 \[Registers\], page 78](#)).

In addition, GDB provides a few other shortcuts and outright additions specific to Ada:

- The assignment statement is allowed as an expression, returning its right-hand operand as its value. Thus, you may enter

```
set x := y + 3
print A(tmp := y + 1)
```

- The semicolon is allowed as an “operator,” returning as its value the value of its right-hand operand. This allows, for example, complex conditional breaks:

```
break f
condition 1 (report(i); k += 1; A(k) > 100)
```

- Rather than use catenation and symbolic character names to introduce special characters into strings, one may instead use a special bracket notation, which is also used to print strings. A sequence of characters of the form ‘["XX"]’ within a string or character literal denotes the (single) character whose numeric encoding is XX in hexadecimal. The sequence of characters ‘[""]’ also denotes a single quotation mark in strings. For example,

```
"One line.["0a"]Next line.["0a"]"
```

contains an ASCII newline character (`Ada.Characters.Latin_1.LF`) after each period.

- The subtype used as a prefix for the attributes `'Pos`, `'Min`, and `'Max` is optional (and is ignored in any case). For example, it is valid to write

```
print 'max(x, y)
```

- When printing arrays, GDB uses positional notation when the array has a lower bound of 1, and uses a modified named notation otherwise. For example, a one-dimensional array of three integers with a lower bound of 3 might print as

```
(3 => 10, 17, 1)
```

That is, in contrast to valid Ada, only the first component has a `=>` clause.

- You may abbreviate attributes in expressions with any unique, multi-character subsequence of their names (an exact match gets preference). For example, you may use `a'len`, `a'gth`, or `a'lth` in place of `a'length`.
- Since Ada is case-insensitive, the debugger normally maps identifiers you type to lower case. The GNAT compiler uses upper-case characters for some of its internal identifiers, which are normally of no interest to users. For the rare occasions when you actually have to look at them, enclose them in angle brackets to avoid the lower-case mapping. For example,

```
gdb print <JMPBUF_SAVE>[0]
```

- Printing an object of class-wide type or dereferencing an access-to-class-wide value will display all the components of the object's specific type (as indicated by its run-time tag). Likewise, component selection on such a value will operate on the specific type of the object.

12.4.4.4 Stopping at the Very Beginning

It is sometimes necessary to debug the program during elaboration, and before reaching the main procedure. As defined in the Ada Reference Manual, the elaboration code is invoked from a procedure called `adainit`. To run your program up to the beginning of elaboration, simply use the following two commands: `tbreak adainit` and `run`.

12.4.4.5 Known Peculiarities of Ada Mode

Besides the omissions listed previously (see [Section 12.4.4.2 \[Omissions from Ada\], page 121](#)), we know of several problems with and limitations of Ada mode in GDB, some of which will be fixed with planned future releases of the debugger and the GNU Ada compiler.

- Currently, the debugger has insufficient information to determine whether certain pointers represent pointers to objects or the objects themselves. Thus, the user may have to tack an extra `.all` after an expression to get it printed properly.
- Static constants that the compiler chooses not to materialize as objects in storage are invisible to the debugger.
- Named parameter associations in function argument lists are ignored (the argument lists are treated as positional).
- Many useful library packages are currently invisible to the debugger.
- Fixed-point arithmetic, conversions, input, and output is carried out using floating-point arithmetic, and may give results that only approximate those on the host machine.
- The type of the `'Address` attribute may not be `System.Address`.
- The GNAT compiler never generates the prefix `Standard` for any of the standard symbols defined by the Ada language. GDB knows about this: it will strip the prefix from names when you use it, and will never look for a name you have so qualified among local symbols, nor match against symbols in other packages or subprograms. If you have defined entities anywhere in your program other than parameters and local variables whose simple names match names in `Standard`, GNAT's lack of qualification here can cause confusion. When this happens, you can usually resolve the confusion by qualifying the problematic names with package `Standard` explicitly.

12.5 Unsupported languages

In addition to the other fully-supported programming languages, GDB also provides a pseudo-language, called `minimal`. It does not represent a real programming language, but provides a set of capabilities close to what the C or assembly languages provide. This should allow most simple operations to be performed while debugging an application that uses a language currently not supported by GDB.

If the language is set to `auto`, GDB will automatically select this language if the current frame corresponds to an unsupported language.

13 Examining the Symbol Table

The commands described in this chapter allow you to inquire about the symbols (names of variables, functions and types) defined in your program. This information is inherent in the text of your program and does not change as your program executes. GDB finds it in your program’s symbol table, in the file indicated when you started GDB (see [Section 2.1.1 \[Choosing files\], page 12](#)), or by one of the file-management commands (see [Section 15.1 \[Commands to specify files\], page 135](#)).

Occasionally, you may need to refer to symbols that contain unusual characters, which GDB ordinarily treats as word delimiters. The most frequent case is in referring to static variables in other source files (see [Section 8.2 \[Program variables\], page 66](#)). File names are recorded in object files as debugging symbols, but GDB would ordinarily parse a typical file name, like ‘`foo.c`’, as the three words ‘`foo`’ ‘`.`’ ‘`c`’. To allow GDB to recognize ‘`foo.c`’ as a single symbol, enclose it in single quotes; for example,

```
p 'foo.c':x
```

looks up the value of `x` in the scope of the file ‘`foo.c`’.

info address symbol

Describe where the data for `symbol` is stored. For a register variable, this says which register it is kept in. For a non-register local variable, this prints the stack-frame offset at which the variable is always stored.

Note the contrast with ‘`print &symbol`’, which does not work at all for a register variable, and for a stack local variable prints the exact address of the current instantiation of the variable.

info symbol addr

Print the name of a symbol which is stored at the address `addr`. If no symbol is stored exactly at `addr`, GDB prints the nearest symbol and an offset from it:

```
(gdb) info symbol 0x54320
_initialize_vx + 396 in section .text
```

This is the opposite of the `info address` command. You can use it to find out the name of a variable or a function given its address.

whatis expr

Print the data type of expression `expr`. `expr` is not actually evaluated, and any side-effecting operations (such as assignments or function calls) inside it do not take place. See [Section 8.1 \[Expressions\], page 65](#).

whatis Print the data type of `$`, the last value in the value history.

ptype typename

Print a description of data type `typename`. `typename` may be the name of a type, or for C code it may have the form ‘`class class-name`’, ‘`struct struct-tag`’, ‘`union union-tag`’ or ‘`enum enum-tag`’.

ptype expr

ptype Print a description of the type of expression `expr`. `ptype` differs from `whatis` by printing a detailed description, instead of just the name of the type.

For example, for this variable declaration:

```
struct complex {double real; double imag;} v;
```

the two commands give this output:

```
(gdb) whatis v
type = struct complex
(gdb) ptype v
type = struct complex {
    double real;
    double imag;
}
```

As with `whatis`, using `ptype` without an argument refers to the type of \$, the last value in the value history.

`info types regexp`

`info types`

Print a brief description of all types whose names match `regexp` (or all types in your program, if you supply no argument). Each complete typename is matched as though it were a complete line; thus, ‘`i type value`’ gives information on all types in your program whose names include the string `value`, but ‘`i type ^value$`’ gives information only on types whose complete name is `value`.

This command differs from `ptype` in two ways: first, like `whatis`, it does not print a detailed description; second, it lists all source files where a type is defined.

`info scope addr`

List all the variables local to a particular scope. This command accepts a location—a function name, a source line, or an address preceded by a ‘*’, and prints all the variables local to the scope defined by that location. For example:

```
(gdb) info scope command_line_handler
Scope for command_line_handler:
Symbol rl is an argument at stack/frame offset 8, length 4.
Symbol linebuffer is in static storage at address 0x150a18, length 4.
Symbol linelength is in static storage at address 0x150a1c, length 4.
Symbol p is a local variable in register $esi, length 4.
Symbol p1 is a local variable in register $ebx, length 4.
Symbol nline is a local variable in register $edx, length 4.
Symbol repeat is a local variable at frame offset -8, length 4.
```

This command is especially useful for determining what data to collect during a *trace experiment*, see [Section 10.1.4 \[Tracepoint Actions\]](#), page 93.

`info source`

Show information about the current source file—that is, the source file for the function containing the current point of execution:

- the name of the source file, and the directory containing it,
- the directory it was compiled in,
- its length, in lines,
- which programming language it is written in,
- whether the executable includes debugging information for that file, and if so, what format the information is in (e.g., STABS, Dwarf 2, etc.), and
- whether the debugging information includes information about preprocessor macros.

info sources

Print the names of all source files in your program for which there is debugging information, organized into two lists: files whose symbols have already been read, and files whose symbols will be read when needed.

info functions

Print the names and data types of all defined functions.

info functions regexp

Print the names and data types of all defined functions whose names contain a match for regular expression *regexp*. Thus, ‘**info fun step**’ finds all functions whose names include **step**; ‘**info fun ^step**’ finds those whose names start with **step**. If a function name contains characters that conflict with the regular expression language (eg. ‘operator*()’), they may be quoted with a backslash.

info variables

Print the names and data types of all variables that are declared outside of functions (i.e. excluding local variables).

info variables regexp

Print the names and data types of all variables (except for local variables) whose names contain a match for regular expression *regexp*.

info classes**info classes regexp**

Display all Objective-C classes in your program, or (with the *regexp* argument) all those matching a particular regular expression.

info selectors**info selectors regexp**

Display all Objective-C selectors in your program, or (with the *regexp* argument) all those matching a particular regular expression.

Some systems allow individual object files that make up your program to be replaced without stopping and restarting your program. For example, in Vx-Works you can simply recompile a defective object file and keep on running. If you are running on one of these systems, you can allow GDB to reload the symbols for automatically relinked modules:

set symbol-reloading on

Replace symbol definitions for the corresponding source file when an object file with a particular name is seen again.

set symbol-reloading off

Do not replace symbol definitions when encountering object files of the same name more than once. This is the default state; if you are not running on a system that permits automatic relinking of modules, you should leave **symbol-reloading off**, since otherwise GDB may discard symbols when linking large programs, that may contain several modules (from different directories or libraries) with the same name.

```

show symbol-reloading
    Show the current on or off setting.

set opaque-type-resolution on
    Tell GDB to resolve opaque types. An opaque type is a type declared as a
    pointer to a struct, class, or union—for example, struct MyType *—that
    is used in one source file although the full declaration of struct MyType is in
    another source file. The default is on.

    A change in the setting of this subcommand will not take effect until the next
    time symbols for a file are loaded.

set opaque-type-resolution off
    Tell GDB not to resolve opaque types. In this case, the type is printed as follows:
        {<no data fields>}

show opaque-type-resolution
    Show whether opaque types are resolved or not.

maint print symbols filename
maint print psymbols filename
maint print msymbols filename
    Write a dump of debugging symbol data into the file filename. These commands
    are used to debug the GDB symbol-reading code. Only symbols with debugging
    data are included. If you use ‘maint print symbols’, GDB includes all the
    symbols for which it has already collected full details: that is, filename reflects
    symbols for only those files whose symbols GDB has read. You can use the
    command info sources to find out which files these are. If you use ‘maint
    print psymbols’ instead, the dump shows information about symbols that GDB
    only knows partially—that is, symbols defined in files that GDB has skimmed,
    but not yet read completely. Finally, ‘maint print msymbols’ dumps just the
    minimal symbol information required for each object file from which GDB has
    read some symbols. See Section 15.1 \[Commands to specify files\], page 135, for
    a discussion of how GDB reads symbols (in the description of symbol-file).

maint info symtabs [ regexp ]
maint info psyntabs [ regexp ]
    List the struct syntab or struct partial_syntab structures whose names
    match regexp. If regexp is not given, list them all. The output includes expres-
    sions which you can copy into a GDB debugging this one to examine a particular
    structure in more detail. For example:


```
(gdb) maint info psyntabs dwarf2read
{ objfile /home/gnu/build/gdb/gdb
 ((struct objfile *) 0x82e69d0)
 { psyntab /home/gnu/src/gdb/dwarf2read.c
 ((struct partial_syntab *) 0x8474b10)
 readin no
 fullname (null)
 text addresses 0x814d3c8 -- 0x8158074
 globals (* (struct partial_symbol **) 0x8507a08 @ 9)
 statics (* (struct partial_symbol **) 0x40e95b78 @ 2882)
 dependencies (none)
 }
```


```

```
}
```

```
(gdb) maint info symtabs
```

```
(gdb)
```

We see that there is one partial symbol table whose filename contains the string ‘dwarf2read’, belonging to the ‘gdb’ executable; and we see that GDB has not read in any symtabs yet at all. If we set a breakpoint on a function, that will cause GDB to read the symtab for the compilation unit containing that function:

```
(gdb) break dwarf2_psymtab_to_syntab
```

```
Breakpoint 1 at 0x814e5da: file /home/gnu/src/gdb/dwarf2read.c,
```

```
line 1574.
```

```
(gdb) maint info symtabs
```

```
{ objfile /home/gnu/build/gdb/gdb
```

```
  ((struct objfile *) 0x82e69d0)
```

```
  { symtab /home/gnu/src/gdb/dwarf2read.c
```

```
    ((struct symtab *) 0x86c1f38)
```

```
      dirname (null)
```

```
      fullname (null)
```

```
      blockvector ((struct blockvector *) 0x86c1bd0) (primary)
```

```
      debugformat DWARF 2
```

```
    }
```

```
}
```

```
(gdb)
```


14 Altering Execution

Once you think you have found an error in your program, you might want to find out for certain whether correcting the apparent error would lead to correct results in the rest of the run. You can find the answer by experiment, using the GDB features for altering execution of the program.

For example, you can store new values into variables or memory locations, give your program a signal, restart it at a different address, or even return prematurely from a function.

14.1 Assignment to variables

To alter the value of a variable, evaluate an assignment expression. See [Section 8.1 \[Expressions\]](#), page 65. For example,

```
print x=4
```

stores the value 4 into the variable `x`, and then prints the value of the assignment expression (which is 4). See [Chapter 12 \[Using GDB with Different Languages\]](#), page 105, for more information on operators in supported languages.

If you are not interested in seeing the value of the assignment, use the `set` command instead of the `print` command. `set` is really the same as `print` except that the expression's value is not printed and is not put in the value history (see [Section 8.8 \[Value history\]](#), page 76). The expression is evaluated only for its effects.

If the beginning of the argument string of the `set` command appears identical to a `set` subcommand, use the `set variable` command instead of just `set`. This command is identical to `set` except for its lack of subcommands. For example, if your program has a variable `width`, you get an error if you try to set a new value with just '`set width=13`', because GDB has the command `set width`:

```
(gdb) whatis width
type = double
(gdb) p width
$4 = 13
(gdb) set width=47
Invalid syntax in expression.
```

The invalid expression, of course, is '`=47`'. In order to actually set the program's variable `width`, use

```
(gdb) set var width=47
```

Because the `set` command has many subcommands that can conflict with the names of program variables, it is a good idea to use the `set variable` command instead of just `set`. For example, if your program has a variable `g`, you run into problems if you try to set a new value with just '`set g=4`', because GDB has the command `set gnutarget`, abbreviated `set g`:

```
(gdb) whatis g
type = double
(gdb) p g
$1 = 1
(gdb) set g=4
(gdb) p g
$2 = 1
(gdb) r
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home smith/cc_progs/a.out
"/home smith/cc_progs/a.out": can't open to read symbols:
                                Invalid bfd target.

(gdb) show g
The current BFD target is "=4".
```

The program variable `g` did not change, and you silently set the `gnutarget` to an invalid value. In order to set the variable `g`, use

```
(gdb) set var g=4
```

GDB allows more implicit conversions in assignments than C; you can freely store an integer value into a pointer variable or vice versa, and you can convert any structure to any other structure that is the same length or shorter.

To store values into arbitrary places in memory, use the ‘`{...}`’ construct to generate a value of specified type at a specified address (see [Section 8.1 \[Expressions\], page 65](#)). For example, `{int}0x83040` refers to memory location `0x83040` as an integer (which implies a certain size and representation in memory), and

```
set {int}0x83040 = 4
```

stores the value 4 into that memory location.

14.2 Continuing at a different address

Ordinarily, when you continue your program, you do so at the place where it stopped, with the `continue` command. You can instead continue at an address of your own choosing, with the following commands:

`jump linespec`

Resume execution at line `linespec`. Execution stops again immediately if there is a breakpoint there. See [Section 7.1 \[Printing source lines\], page 59](#), for a description of the different forms of `linespec`. It is common practice to use the `tbreak` command in conjunction with `jump`. See [Section 5.1.1 \[Setting breakpoints\], page 34](#).

The `jump` command does not change the current stack frame, or the stack pointer, or the contents of any memory location or any register other than the program counter. If line `linespec` is in a different function from the one currently executing, the results may be bizarre if the two functions expect different patterns of arguments or of local variables. For this reason, the `jump` command requests confirmation if the specified line is not in the function currently executing. However, even bizarre results are predictable if you are well acquainted with the machine-language code of your program.

`jump *address`

Resume execution at the instruction at address `address`.

On many systems, you can get much the same effect as the `jump` command by storing a new value into the register `$pc`. The difference is that this does not start your program running; it only changes the address of where it *will* run when you continue. For example,

```
set $pc = 0x485
```

makes the next `continue` command or stepping command execute at address `0x485`, rather than at the address where your program stopped. See [Section 5.2 \[Continuing and stepping\], page 46](#).

The most common occasion to use the `jump` command is to back up—perhaps with more breakpoints set—over a portion of a program that has already executed, in order to examine its execution in more detail.

14.3 Giving your program a signal

`signal signal`

Resume execution where your program stopped, but immediately give it the signal `signal`. `signal` can be the name or the number of a signal. For example, on many systems `signal 2` and `signal SIGINT` are both ways of sending an interrupt signal.

Alternatively, if `signal` is zero, continue execution without giving a signal. This is useful when your program stopped on account of a signal and would ordinary see the signal when resumed with the `continue` command; ‘`signal 0`’ causes it to resume without a signal.

`signal` does not repeat when you press `\RET` a second time after executing the command.

Invoking the `signal` command is not the same as invoking the `kill` utility from the shell. Sending a signal with `kill` causes GDB to decide what to do with the signal depending on the signal handling tables (see [Section 5.3 \[Signals\], page 50](#)). The `signal` command passes the signal directly to your program.

14.4 Returning from a function

`return`
`return expression`

You can cancel execution of a function call with the `return` command. If you give an `expression` argument, its value is used as the function’s return value.

When you use `return`, GDB discards the selected stack frame (and all frames within it). You can think of this as making the discarded frame return prematurely. If you wish to specify a value to be returned, give that value as the argument to `return`.

This pops the selected stack frame (see [Section 6.3 \[Selecting a frame\], page 55](#)), and any other frames inside of it, leaving its caller as the innermost remaining frame. That frame becomes selected. The specified value is stored in the registers used for returning values of functions.

The `return` command does not resume execution; it leaves the program stopped in the state that would exist if the function had just returned. In contrast, the `finish` command (see [Section 5.2 \[Continuing and stepping\], page 46](#)) resumes execution until the selected stack frame returns naturally.

14.5 Calling program functions

`call expr`

Evaluate the expression `expr` without displaying `void` returned values.

You can use this variant of the `print` command if you want to execute a function from your program, but without cluttering the output with `void` returned values. If the result is not `void`, it is printed and saved in the value history.

14.6 Patching programs

By default, GDB opens the file containing your program's executable code (or the corefile) read-only. This prevents accidental alterations to machine code; but it also prevents you from intentionally patching your program's binary.

If you'd like to be able to patch the binary, you can specify that explicitly with the `set write` command. For example, you might want to turn on internal debugging flags, or even to make emergency repairs.

`set write on`

`set write off`

If you specify '`set write on`', GDB opens executable and core files for both reading and writing; if you specify '`set write off`' (the default), GDB opens them read-only.

If you have already loaded a file, you must load it again (using the `exec-file` or `core-file` command) after changing `set write`, for your new setting to take effect.

`show write`

Display whether executable files and core files are opened for writing as well as reading.

15 GDB Files

GDB needs to know the file name of the program to be debugged, both in order to read its symbol table and in order to start your program. To debug a core dump of a previous run, you must also tell GDB the name of the core dump file.

15.1 Commands to specify files

You may want to specify executable and core dump file names. The usual way to do this is at start-up time, using the arguments to GDB's start-up commands (see [Chapter 2 \[Getting In and Out of GDB\], page 11](#)).

Occasionally it is necessary to change to a different file during a GDB session. Or you may run GDB and forget to specify a file you want to use. In these situations the GDB commands to specify new files are useful.

file *filename*

Use *filename* as the program to be debugged. It is read for its symbols and for the contents of pure memory. It is also the program executed when you use the **run** command. If you do not specify a directory and the file is not found in the GDB working directory, GDB uses the environment variable **PATH** as a list of directories to search, just as the shell does when looking for a program to run. You can change the value of this variable, for both GDB and your program, using the **path** command.

On systems with memory-mapped files, an auxiliary file named '*filename.syms*' may hold symbol table information for *filename*. If so, GDB maps in the symbol table from '*filename.syms*', starting up more quickly. See the descriptions of the file options '**-mapped**' and '**-readnow**' (available on the command line, and with the commands **file**, **symbol-file**, or **add-symbol-file**, described below), for more information.

file *file* with no argument makes GDB discard any information it has on both executable file and the symbol table.

exec-file [*filename*]

Specify that the program to be run (but not the symbol table) is found in *filename*. GDB searches the environment variable **PATH** if necessary to locate your program. Omitting *filename* means to discard information on the executable file.

symbol-file [*filename*]

Read symbol table information from file *filename*. **PATH** is searched when necessary. Use the **file** command to get both symbol table and program to run from the same file.

symbol-file with no argument clears out GDB information on your program's symbol table.

The **symbol-file** command causes GDB to forget the contents of its convenience variables, the value history, and all breakpoints and auto-display expressions. This is because they may contain pointers to the internal data recording symbols

and data types, which are part of the old symbol table data being discarded inside GDB.

`symbol-file` does not repeat if you press `<RET>` again after executing it once.

When GDB is configured for a particular environment, it understands debugging information in whatever format is the standard generated for that environment; you may use either a GNU compiler, or other compilers that adhere to the local conventions. Best results are usually obtained from GNU compilers; for example, using `gcc` you can generate debugging information for optimized code.

For most kinds of object files, with the exception of old SVR3 systems using COFF, the `symbol-file` command does not normally read the symbol table in full right away. Instead, it scans the symbol table quickly to find which source files and which symbols are present. The details are read later, one source file at a time, as they are needed.

The purpose of this two-stage reading strategy is to make GDB start up faster. For the most part, it is invisible except for occasional pauses while the symbol table details for a particular source file are being read. (The `set verbose` command can turn these pauses into messages if desired. See [Section 19.7 \[Optional warnings and messages\]](#), page 183.)

We have not implemented the two-stage strategy for COFF yet. When the symbol table is stored in COFF format, `symbol-file` reads the symbol table data in full right away. Note that “stabs-in-COFF” still does the two-stage strategy, since the debug info is actually in stabs format.

```
symbol-file filename [ -readnow ] [ -mapped ]
file filename [ -readnow ] [ -mapped ]
```

You can override the GDB two-stage strategy for reading symbol tables by using the ‘`-readnow`’ option with any of the commands that load symbol table information, if you want to be sure GDB has the entire symbol table available.

If memory-mapped files are available on your system through the `mmap` system call, you can use another option, ‘`-mapped`’, to cause GDB to write the symbols for your program into a reusable file. Future GDB debugging sessions map in symbol information from this auxiliary symbol file (if the program has not changed), rather than spending time reading the symbol table from the executable program. Using the ‘`-mapped`’ option has the same effect as starting GDB with the ‘`-mapped`’ command-line option.

You can use both options together, to make sure the auxiliary symbol file has all the symbol information for your program.

The auxiliary symbol file for a program called `myprog` is called ‘`myprog.syms`’. Once this file exists (so long as it is newer than the corresponding executable), GDB always attempts to use it when you debug `myprog`; no special options or commands are needed.

The ‘`.syms`’ file is specific to the host machine where you run GDB. It holds an exact image of the internal GDB symbol table. It cannot be shared across multiple host platforms.

```
core-file [ filename ]
```

core Specify the whereabouts of a core dump file to be used as the “contents of memory”. Traditionally, core files contain only some parts of the address space of the process that generated them; GDB can access the executable file itself for other parts.

core-file with no argument specifies that no core file is to be used.

Note that the core file is ignored when your program is actually running under GDB. So, if you have been running your program and you wish to debug a core file instead, you must kill the subprocess in which the program is running. To do this, use the **kill** command (see [Section 4.8 \[Killing the child process\], page 28](#)).

```
add-symbol-file filename address
```

```
add-symbol-file filename address [ -readnow ] [ -mapped ]
```

```
add-symbol-file filename -ssection address ...
```

The **add-symbol-file** command reads additional symbol table information from the file *filename*. You would use this command when *filename* has been dynamically loaded (by some other means) into the program that is running. *address* should be the memory address at which the file has been loaded; GDB cannot figure this out for itself. You can additionally specify an arbitrary number of ‘-s*section* *address*’ pairs, to give an explicit section name and base address for that section. You can specify any *address* as an expression.

The symbol table of the file *filename* is added to the symbol table originally read with the **symbol-file** command. You can use the **add-symbol-file** command any number of times; the new symbol data thus read keeps adding to the old. To discard all old symbol data instead, use the **symbol-file** command without any arguments.

Although *filename* is typically a shared library file, an executable file, or some other object file which has been fully relocated for loading into a process, you can also load symbolic information from relocatable ‘.o’ files, as long as:

- the file’s symbolic information refers only to linker symbols defined in that file, not to symbols defined by other object files,
- every section the file’s symbolic information refers to has actually been loaded into the inferior, as it appears in the file, and
- you can determine the address at which every section was loaded, and provide these to the **add-symbol-file** command.

Some embedded operating systems, like Sun Chorus and VxWorks, can load relocatable files into an already running program; such systems typically make the requirements above easy to meet. However, it’s important to recognize that many native systems use complex link procedures (.linkonce section factoring and C++ constructor table assembly, for example) that make the requirements difficult to meet. In general, one cannot assume that using **add-symbol-file** to read a relocatable object file’s symbolic information will have the same effect as linking the relocatable object file into the program in the normal way.

add-symbol-file does not repeat if you press **RET** after using it.

You can use the ‘**-mapped**’ and ‘**-readnow**’ options just as with the **symbol-file** command, to change how GDB manages the symbol table information for *filename*.

add-shared-symbol-file

The **add-shared-symbol-file** command can be used only under Harris’ CXUX operating system for the Motorola 88k. GDB automatically looks for shared libraries, however if GDB does not find yours, you can run **add-shared-symbol-file**. It takes no arguments.

section The **section** command changes the base address of section SECTION of the exec file to ADDR. This can be used if the exec file does not contain section addresses, (such as in the a.out format), or when the addresses specified in the file itself are wrong. Each section must be changed separately. The **info files** command, described below, lists all the sections and their addresses.

info files

info target

info files and **info target** are synonymous; both print the current target (see [Chapter 16 \[Specifying a Debugging Target\]](#), page 147), including the names of the executable and core dump files currently in use by GDB, and the files from which symbols were loaded. The command **help target** lists all possible targets rather than current ones.

maint info sections

Another command that can give you extra information about program sections is **maint info sections**. In addition to the section information displayed by **info files**, this command displays the flags and file offset of each section in the executable and core dump files. In addition, **maint info sections** provides the following command options (which may be arbitrarily combined):

ALLOBJ Display sections for all loaded object files, including shared libraries.

sections Display info only for named sections.

section-flags

Display info only for sections for which **section-flags** are true. The section flags that GDB currently knows about are:

ALLOC Section will have space allocated in the process when loaded. Set for all sections except those containing debug information.

LOAD Section will be loaded from the file into the child process memory. Set for pre-initialized code and data, clear for .bss sections.

RELOC Section needs to be relocated before loading.

READONLY Section cannot be modified by the child process.

CODE Section contains executable code only.

DATA	Section contains data only (no executable code).
ROM	Section will reside in ROM.
CONSTRUCTOR	Section contains data for constructor/destructor lists.
HAS_CONTENTS	Section is not empty.
NEVER_LOAD	An instruction to the linker to not output the section.
COFF_SHARED_LIBRARY	A notification to the linker that the section contains COFF shared library information.
IS_COMMON	Section contains common symbols.

set trust_READONLY-sections on

Tell GDB that readonly sections in your object file really are read-only (i.e. that their contents will not change). In that case, GDB can fetch values from these sections out of the object file, rather than from the target program. For some targets (notably embedded ones), this can be a significant enhancement to debugging performance.

The default is off.

set trust_READONLY-sections off

Tell GDB not to trust readonly sections. This means that the contents of the section might change while the program is running, and must therefore be fetched from the target when needed.

All file-specifying commands allow both absolute and relative file names as arguments. GDB always converts the file name to an absolute file name and remembers it that way.

GDB supports HP-UX, SunOS, SVr4, Irix 5, and IBM RS/6000 shared libraries.

GDB automatically loads symbol definitions from shared libraries when you use the **run** command, or when you examine a core file. (Before you issue the **run** command, GDB does not understand references to a function in a shared library, however—unless you are debugging a core file).

On HP-UX, if the program loads a library explicitly, GDB automatically loads the symbols at the time of the **shl_load** call.

There are times, however, when you may wish to not automatically load symbol definitions from shared libraries, such as when they are particularly large or there are many of them.

To control the automatic loading of shared library symbols, use the commands:

set auto-solib-add mode

If *mode* is **on**, symbols from all shared object libraries will be loaded automatically when the inferior begins execution, you attach to an independently started inferior, or when the dynamic linker informs GDB that a new library

has been loaded. If *mode* is **off**, symbols must be loaded manually, using the **sharedlibrary** command. The default value is **on**.

```
show auto-solib-add
Display the current autoloading mode.
```

To explicitly load shared library symbols, use the **sharedlibrary** command:

```
info share
info sharedlibrary
Print the names of the shared libraries which are currently loaded.
```

```
sharedlibrary regex
share regex
Load shared object library symbols for files matching a Unix regular expression.
As with files loaded automatically, it only loads shared libraries required by your
program for a core file or after typing run. If regex is omitted all shared libraries
required by your program are loaded.
```

On some systems, such as HP-UX systems, GDB supports autoloading shared library symbols until a limiting threshold size is reached. This provides the benefit of allowing autoloading to remain on by default, but avoids autoloading excessively large shared libraries, up to a threshold that is initially set, but which you can modify if you wish.

Beyond that threshold, symbols from shared libraries must be explicitly loaded. To load these symbols, use the command **sharedlibrary *filename***. The base address of the shared library is determined automatically by GDB and need not be specified.

To display or set the threshold, use the commands:

```
set auto-solib-limit threshold
Set the autoloading size threshold, in an integral number of megabytes. If
threshold is nonzero and shared library autoloading is enabled, symbols from
all shared object libraries will be loaded until the total size of the loaded shared
library symbols exceeds this threshold. Otherwise, symbols must be loaded
manually, using the sharedlibrary command. The default threshold is 100
(i.e. 100 Mb).
```

```
show auto-solib-limit
Display the current autoloading size threshold, in megabytes.
```

Shared libraries are also supported in many cross or remote debugging configurations. A copy of the target's libraries need to be present on the host system; they need to be the same as the target libraries, although the copies on the target can be stripped as long as the copies on the host are not.

You need to tell GDB where the target libraries are, so that it can load the correct copies—otherwise, it may try to load the host's libraries. GDB has two variables to specify the search directories for target libraries.

```
set solib-absolute-prefix path
If this variable is set, path will be used as a prefix for any absolute shared
library paths; many runtime loaders store the absolute paths to the shared
library in the target program's memory. If you use 'solib-absolute-prefix'
```

to find shared libraries, they need to be laid out in the same way that they are on the target, with e.g. a ‘`/usr/lib`’ hierarchy under `path`.

You can set the default value of ‘`solib-absolute-prefix`’ by using the configure-time ‘`--with-sysroot`’ option.

```
show solib-absolute-prefix
```

Display the current shared library prefix.

```
set solib-search-path path
```

If this variable is set, `path` is a colon-separated list of directories to search for shared libraries. ‘`solib-search-path`’ is used after ‘`solib-absolute-prefix`’ fails to locate the library, or if the path to the library is relative instead of absolute. If you want to use ‘`solib-search-path`’ instead of ‘`solib-absolute-prefix`’, be sure to set ‘`solib-absolute-prefix`’ to a nonexistent directory to prevent GDB from finding your host’s libraries.

```
show solib-search-path
```

Display the current shared library search path.

15.2 Debugging Information in Separate Files

GDB allows you to put a program’s debugging information in a file separate from the executable itself, in a way that allows GDB to find and load the debugging information automatically. Since debugging information can be very large — sometimes larger than the executable code itself — some systems distribute debugging information for their executables in separate files, which users can install only when they need to debug a problem.

If an executable’s debugging information has been extracted to a separate file, the executable should contain a *debug link* giving the name of the debugging information file (with no directory components), and a checksum of its contents. (The exact form of a debug link is described below.) If the full name of the directory containing the executable is `execdir`, and the executable has a debug link that specifies the name `debugfile`, then GDB will automatically search for the debugging information file in three places:

- the directory containing the executable file (that is, it will look for a file named ‘`execdir/debugfile`’,
- a subdirectory of that directory named ‘`.debug`’ (that is, the file ‘`execdir/.debug/debugfile`’, and
- a subdirectory of the global debug file directory that includes the executable’s full path, and the name from the link (that is, the file ‘`globaldebugdir/execdir/debugfile`’, where `globaldebugdir` is the global debug file directory, and `execdir` has been turned into a relative path).

GDB checks under each of these names for a debugging information file whose checksum matches that given in the link, and reads the debugging information from the first one it finds.

So, for example, if you ask GDB to debug ‘`/usr/bin/ls`’, which has a link containing the name ‘`ls.debug`’, and the global debug directory is ‘`/usr/lib/debug`’, then GDB will look for debug information in ‘`/usr/bin/ls.debug`’, ‘`/usr/bin/.debug/ls.debug`’, and ‘`/usr/lib/debug/usr/bin/ls.debug`’.

You can set the global debugging info directory's name, and view the name GDB is currently using.

set debug-file-directory *directory*

Set the directory which GDB searches for separate debugging information files to *directory*.

show debug-file-directory

Show the directory GDB searches for separate debugging information files.

A debug link is a special section of the executable file named `.gnu_debuglink`. The section must contain:

- A filename, with any leading directory components removed, followed by a zero byte,
- zero to three bytes of padding, as needed to reach the next four-byte boundary within the section, and
- a four-byte CRC checksum, stored in the same endianness used for the executable file itself. The checksum is computed on the debugging information file's full contents by the function given below, passing zero as the *crc* argument.

Any executable file format can carry a debug link, as long as it can contain a section named `.gnu_debuglink` with the contents described above.

The debugging information file itself should be an ordinary executable, containing a full set of linker symbols, sections, and debugging information. The sections of the debugging information file should have the same names, addresses and sizes as the original file, but they need not contain any data — much like a `.bss` section in an ordinary executable.

As of December 2002, there is no standard GNU utility to produce separated executable / debugging information file pairs. Ulrich Drepper's ‘`elfutils`’ package, starting with version 0.53, contains a version of the `strip` command such that the command `strip foo -f foo.debug` removes the debugging information from the executable file ‘`foo`’, places it in the file ‘`foo.debug`’, and leaves behind a debug link in ‘`foo`’.

Since there are many different ways to compute CRC's (different polynomials, reversals, byte ordering, etc.), the simplest way to describe the CRC used in `.gnu_debuglink` sections is to give the complete code for a function that computes it:

```
unsigned long
gnu_debuglink_crc32 (unsigned long crc,
                      unsigned char *buf, size_t len)
{
    static const unsigned long crc32_table[256] =
    {
        0x00000000, 0x77073096, 0xee0e612c, 0x990951ba, 0x076dc419,
        0x706af48f, 0xe963a535, 0x9e6495a3, 0x0edb8832, 0x79dcb8a4,
        0xe0d5e91e, 0x97d2d988, 0x09b64c2b, 0x7eb17cbd, 0xe7b82d07,
        0x90bf1d91, 0x1db71064, 0x6ab020f2, 0xf3b97148, 0x84be41de,
        0x1adad47d, 0x6ddde4eb, 0xf4d4b551, 0x83d385c7, 0x136c9856,
        0x646ba8c0, 0xfd62f97a, 0x8a65c9ec, 0x14015c4f, 0x63066cd9,
        0xfa0f3d63, 0x8d080df5, 0x3b6e20c8, 0x4c69105e, 0xd56041e4,
        0xa2677172, 0x3c03e4d1, 0x4b04d447, 0xd20d85fd, 0xa50ab56b,
        0x35b5a8fa, 0x42b2986c, 0xdbbbc9d6, 0xacfbc940, 0x32d86ce3,
        0x45df5c75, 0xcdcd60dcf, 0xabd13d59, 0x26d930ac, 0x51de003a,
        0xc8d75180, 0xbfd06116, 0x21b4f4b5, 0x56b3c423, 0xcfba9599,
        0xb8bda50f, 0x2802b89e, 0x5f058808, 0xc60cd9b2, 0xb10be924,
```

```

0x2f6f7c87, 0x58684c11, 0xc1611dab, 0xb6662d3d, 0x76dc4190,
0x01db7106, 0x98d220bc, 0xefd5102a, 0x71b18589, 0x06b6b51f,
0x9fbfe4a5, 0xe8b8d433, 0x7807c9a2, 0x0f00f934, 0x9609a88e,
0xe10e9818, 0x7f6a0dbb, 0x086d3d2d, 0x91646c97, 0xe6635c01,
0x6b6b51f4, 0x1c6c6162, 0x856530d8, 0xf262004e, 0x6c0695ed,
0x1b01a57b, 0x8208f4c1, 0xf50fc457, 0x65b0d9c6, 0x12b7e950,
0x8bbeb8ea, 0xfc9887c, 0x62dd1ddf, 0x15da2d49, 0x8cd37cf3,
0xfbcd44c65, 0x4db26158, 0x3ab551ce, 0xa3bc0074, 0xd4bb30e2,
0x4adfa541, 0x3dd895d7, 0xa4d1c46d, 0xd3d6f4fb, 0x4369e96a,
0x346ed9fc, 0xad678846, 0xda60b8d0, 0x44042d73, 0x33031de5,
0xaa0a4c5f, 0xdd0d7cc9, 0x5005713c, 0x270241aa, 0xbe0b1010,
0xc90c2086, 0x5768b525, 0x206f85b3, 0xb966d409, 0xce61e49f,
0x5edef90e, 0x29d9c998, 0xb0d09822, 0xc7d7a8b4, 0x59b33d17,
0x2eb40d81, 0xb7bd5c3b, 0xc0ba6cad, 0xedb88320, 0x9abfb3b6,
0x03b6e20c, 0x74b1d29a, 0xead54739, 0x9dd277af, 0x04db2615,
0x73dc1683, 0x3630b12, 0x94643b84, 0x0d6d6a3e, 0x7a6a5aa8,
0xe40ecf0b, 0x9309ff9d, 0xa00ae27, 0x7d079eb1, 0xf00f9344,
0x8708a3d2, 0x1e01f268, 0x6906c2fe, 0xf762575d, 0x806567cb,
0x196c3671, 0x6e6b06e7, 0xfed41b76, 0x89d32be0, 0x10da7a5a,
0x67dd4acc, 0x9b9df6f, 0x8ebeeff9, 0x17b7be43, 0x60b08ed5,
0xd6d6a3e8, 0xa1d1937e, 0x38d8c2c4, 0x4fdff252, 0xd1bb67f1,
0xa6bc5767, 0x3fb506dd, 0x48b2364b, 0xd80d2bda, 0xaf0a1b4c,
0x36034af6, 0x41047a60, 0xdf60efc3, 0xa867df55, 0x316e8eef,
0x4669be79, 0xcb61b38c, 0xbc66831a, 0x256fd2a0, 0x5268e236,
0xcc0c7795, 0xbb0b4703, 0x220216b9, 0x5505262f, 0xc5ba3bbe,
0xb2bd0b28, 0x2b45a92, 0x5cb36a04, 0xc2d7ffa7, 0xb5d0cf31,
0x2cd99e8b, 0x5bdeae1d, 0x9b64c2b0, 0xec63f226, 0x756aa39c,
0x026d930a, 0x9c0906a9, 0xeb0e363f, 0x72076785, 0x05005713,
0x95bf4a82, 0xe2b87a14, 0x7bb12bae, 0xcb61b38, 0x92d28e9b,
0xe5d5be0d, 0x7cdcef7, 0xbdbdf21, 0x86d3d2d4, 0xf1d4e242,
0x68ddb3f8, 0x1fda836e, 0x81be16cd, 0xf6b9265b, 0x6fb077e1,
0x18b74777, 0x88085ae6, 0xff0f6a70, 0x66063bca, 0x11010b5c,
0x8f659eff, 0xf862ae69, 0x616bffd3, 0x166ccf45, 0xa00ae278,
0xd70dd2ee, 0x4e048354, 0x3903b3c2, 0xa7672661, 0xd06016f7,
0x4969474d, 0x3e6e77db, 0xaed16a4a, 0xd9d65adc, 0x40df0b66,
0x37d83bf0, 0xa9bcae53, 0xdebb9ec5, 0x47b2cf7f, 0x30b5ffe9,
0xbdbdf21c, 0xcabac28a, 0x53b39330, 0x24b4a3a6, 0xbad03605,
0xcdd70693, 0x54de5729, 0x23d967bf, 0xb3667a2e, 0xc4614ab8,
0x5d681b02, 0x2a6f2b94, 0xb40bbe37, 0xc30c8ea1, 0x5a05df1b,
0x2d02ef8d
};

unsigned char *end;

crc = ~crc & 0xffffffff;
for (end = buf + len; buf < end; ++buf)
    crc = crc32_table[(crc ^ *buf) & 0xff] ^ (crc >> 8);
return ~crc & 0xffffffff;
}

```

15.3 Errors reading symbol files

While reading a symbol file, GDB occasionally encounters problems, such as symbol types it does not recognize, or known bugs in compiler output. By default, GDB does not notify you of such problems, since they are relatively common and primarily of interest to people debugging compilers. If you are interested in seeing information about ill-constructed symbol tables, you can either ask GDB to print only one message about each such type of problem, no matter how many times the problem occurs; or you can ask GDB to print more

messages, to see how many times the problems occur, with the `set complaints` command (see [Section 19.7 \[Optional warnings and messages\], page 183](#)).

The messages currently printed, and their meanings, include:

inner block not inside outer block in *symbol*

The symbol information shows where symbol scopes begin and end (such as at the start of a function or a block of statements). This error indicates that an inner scope block is not fully contained in its outer scope blocks.

GDB circumvents the problem by treating the inner block as if it had the same scope as the outer block. In the error message, *symbol* may be shown as “(don’t know)” if the outer block is not a function.

block at *address* out of order

The symbol information for symbol scope blocks should occur in order of increasing addresses. This error indicates that it does not do so.

GDB does not circumvent this problem, and has trouble locating symbols in the source file whose symbols it is reading. (You can often determine what source file is affected by specifying `set verbose on`. See [Section 19.7 \[Optional warnings and messages\], page 183](#).)

bad block start address patched

The symbol information for a symbol scope block has a start address smaller than the address of the preceding source line. This is known to occur in the SunOS 4.1.1 (and earlier) C compiler.

GDB circumvents the problem by treating the symbol scope block as starting on the previous source line.

bad string table offset in *symbol n*

Symbol number *n* contains a pointer into the string table which is larger than the size of the string table.

GDB circumvents the problem by considering the symbol to have the name `foo`, which may cause other problems if many symbols end up with this name.

unknown symbol type *0xnn*

The symbol information contains new data types that GDB does not yet know how to read. `0xnn` is the symbol type of the uncomprehended information, in hexadecimal.

GDB circumvents the error by ignoring this symbol information. This usually allows you to debug your program, though certain symbols are not accessible. If you encounter such a problem and feel like debugging it, you can debug `gdb` with itself, breakpoint on `complain`, then go up to the function `read_dbx_symtab` and examine `*bufp` to see the symbol.

stub type has NULL name

GDB could not find the full definition for a struct or class.

const/volatile indicator missing (ok if using g++ v1.x), got...

The symbol information for a C++ member function is missing some information that recent versions of the compiler should have output for it.

```
info mismatch between compiler and debugger
      GDB could not parse a type specification output by the compiler.
```


16 Specifying a Debugging Target

A **target** is the execution environment occupied by your program.

Often, GDB runs in the same host environment as your program; in that case, the debugging target is specified as a side effect when you use the `file` or `core` commands. When you need more flexibility—for example, running GDB on a physically separate host, or controlling a standalone system over a serial port or a realtime system over a TCP/IP connection—you can use the `target` command to specify one of the target types configured for GDB (see [Section 16.2 \[Commands for managing targets\], page 147](#)).

16.1 Active targets

There are three classes of targets: processes, core files, and executable files. GDB can work concurrently on up to three active targets, one in each class. This allows you to (for example) start a process and inspect its activity without abandoning your work on a core file.

For example, if you execute ‘`gdb a.out`’, then the executable file `a.out` is the only active target. If you designate a core file as well—presumably from a prior run that crashed and coredumped—then GDB has two active targets and uses them in tandem, looking first in the corefile target, then in the executable file, to satisfy requests for memory addresses. (Typically, these two classes of target are complementary, since core files contain only a program’s read-write memory—variables and so on—plus machine status, while executable files contain only the program text and initialized data.)

When you type `run`, your executable file becomes an active process target as well. When a process target is active, all GDB commands requesting memory addresses refer to that target; addresses in an active core file or executable file target are obscured while the process target is active.

Use the `core-file` and `exec-file` commands to select a new core file or executable target (see [Section 15.1 \[Commands to specify files\], page 135](#)). To specify as a target a process that is already running, use the `attach` command (see [Section 4.7 \[Debugging an already-running process\], page 28](#)).

16.2 Commands for managing targets

`target type parameters`

Connects the GDB host environment to a target machine or process. A target is typically a protocol for talking to debugging facilities. You use the argument `type` to specify the type or protocol of the target machine.

Further `parameters` are interpreted by the target protocol, but typically include things like device names or host names to connect with, process numbers, and baud rates.

The `target` command does not repeat if you press `(RET)` again after executing the command.

`help target`

Displays the names of all targets available. To display targets currently selected, use either `info target` or `info files` (see [Section 15.1 \[Commands to specify files\], page 135](#)).

help target name

Describe a particular target, including any parameters necessary to select it.

set gnutarget args

GDB uses its own library BFD to read your files. GDB knows whether it is reading an executable, a core, or a .o file; however, you can specify the file format with the **set gnutarget** command. Unlike most **target** commands, with **gnutarget** the **target** refers to a program, not a machine.

Warning: To specify a file format with **set gnutarget**, you must know the actual BFD name.

See [Section 15.1 \[Commands to specify files\]](#), page 135.

show gnutarget

Use the **show gnutarget** command to display what file format **gnutarget** is set to read. If you have not set **gnutarget**, GDB will determine the file format for each file automatically, and **show gnutarget** displays ‘The current BDF target is “auto”’.

Here are some common targets (available, or not, depending on the GDB configuration):

target exec program

An executable file. ‘**target exec program**’ is the same as ‘**exec-file program**’.

target core filename

A core dump file. ‘**target core filename**’ is the same as ‘**core-file filename**’.

target remote dev

Remote serial target in GDB-specific protocol. The argument *dev* specifies what serial device to use for the connection (e.g. ‘/dev/ttys’). See [Section 16.4 \[Remote debugging\]](#), page 149. **target remote** supports the **load** command. This is only useful if you have some other way of getting the stub to the target system, and you can put it somewhere in memory where it won’t get clobbered by the download.

target sim

Builtin CPU simulator. GDB includes simulators for most architectures. In general,

```
target sim
load
run
```

works; however, you cannot assume that a specific memory map, device drivers, or even basic I/O is available, although some simulators do provide these. For info about any processor-specific simulator details, see the appropriate section in [Section 18.3 \[Embedded Processors\]](#), page 166.

Some configurations may include these targets as well:

target nrom dev

NetROM ROM emulator. This target only supports downloading.

Different targets are available on different configurations of GDB; your configuration may have more or fewer targets.

Many remote targets require you to download the executable's code once you've successfully established a connection.

`load filename`

Depending on what remote debugging facilities are configured into GDB, the `load` command may be available. Where it exists, it is meant to make *filename* (an executable) available for debugging on the remote system—by downloading, or dynamic linking, for example. `load` also records the *filename* symbol table in GDB, like the `add-symbol-file` command.

If your GDB does not have a `load` command, attempting to execute it gets the error message “`You can't do that when your target is ...`”

The file is loaded at whatever address is specified in the executable. For some object file formats, you can specify the load address when you link the program; for other formats, like `a.out`, the object file format specifies a fixed address.

`load` does not repeat if you press `(RET)` again after using it.

16.3 Choosing target byte order

Some types of processors, such as the MIPS, PowerPC, and Renesas SH, offer the ability to run either big-endian or little-endian byte orders. Usually the executable or symbol will include a bit to designate the endian-ness, and you will not need to worry about which to use. However, you may still find it useful to adjust GDB's idea of processor endian-ness manually.

`set endian big`

Instruct GDB to assume the target is big-endian.

`set endian little`

Instruct GDB to assume the target is little-endian.

`set endian auto`

Instruct GDB to use the byte order associated with the executable.

`show endian`

Display GDB's current idea of the target byte order.

Note that these commands merely adjust interpretation of symbolic data on the host, and that they have absolutely no effect on the target system.

16.4 Remote debugging

If you are trying to debug a program running on a machine that cannot run GDB in the usual way, it is often useful to use remote debugging. For example, you might use remote debugging on an operating system kernel, or on a small system which does not have a general purpose operating system powerful enough to run a full-featured debugger.

Some configurations of GDB have special serial or TCP/IP interfaces to make this work with particular debugging targets. In addition, GDB comes with a generic serial protocol (specific to GDB, but not specific to any particular target system) which you can use if you

write the remote stubs—the code that runs on the remote system to communicate with GDB.

Other remote targets may be available in your configuration of GDB; use `help target` to list them.

16.5 Kernel Object Display

Some targets support kernel object display. Using this facility, GDB communicates specially with the underlying operating system and can display information about operating system-level objects such as mutexes and other synchronization objects. Exactly which objects can be displayed is determined on a per-OS basis.

Use the `set os` command to set the operating system. This tells GDB which kernel object display module to initialize:

```
(gdb) set os cisco
```

The associated command `show os` displays the operating system set with the `set os` command; if no operating system has been set, `show os` will display an empty string “”.

If `set os` succeeds, GDB will display some information about the operating system, and will create a new `info` command which can be used to query the target. The `info` command is named after the operating system:

```
(gdb) info cisco
List of Cisco Kernel Objects
Object      Description
any        Any and all objects
```

Further subcommands can be used to query about particular objects known by the kernel.

There is currently no way to determine whether a given operating system is supported other than to try setting it with `set os name`, where `name` is the name of the operating system you want to try.

17 Debugging remote programs

17.1 Connecting to a remote target

On the GDB host machine, you will need an unstripped copy of your program, since GDB needs symbol and debugging information. Start up GDB as usual, using the name of the local copy of your program as the first argument.

If you're using a serial line, you may want to give GDB the ‘`--baud`’ option, or use the `set remotebaud` command before the `target` command.

After that, use `target remote` to establish communications with the target machine. Its argument specifies how to communicate—either via a devicename attached to a direct serial line, or a TCP or UDP port (possibly to a terminal server which in turn has a serial line to the target). For example, to use a serial line connected to the device named ‘`/dev/ttyb`’:

```
target remote /dev/ttyb
```

To use a TCP connection, use an argument of the form `host:port` or `tcp:host:port`. For example, to connect to port 2828 on a terminal server named `manyfarms`:

```
target remote manyfarms:2828
```

If your remote target is actually running on the same machine as your debugger session (e.g. a simulator of your target running on the same host), you can omit the hostname. For example, to connect to port 1234 on your local machine:

```
target remote :1234
```

Note that the colon is still required here.

To use a UDP connection, use an argument of the form `udp:host:port`. For example, to connect to UDP port 2828 on a terminal server named `manyfarms`:

```
target remote udp:manyfarms:2828
```

When using a UDP connection for remote debugging, you should keep in mind that the ‘U’ stands for “Unreliable”. UDP can silently drop packets on busy or unreliable networks, which will cause havoc with your debugging session.

Now you can use all the usual commands to examine and change data and to step and continue the remote program.

Whenever GDB is waiting for the remote program, if you type the interrupt character (often `\C-C`), GDB attempts to stop the program. This may or may not succeed, depending in part on the hardware and the serial drivers the remote system uses. If you type the interrupt character once again, GDB displays this prompt:

```
Interrupted while waiting for the program.  
Give up (and stop debugging it)? (y or n)
```

If you type `y`, GDB abandons the remote debugging session. (If you decide you want to try again later, you can use ‘`target remote`’ again to connect once more.) If you type `n`, GDB goes back to waiting.

detach When you have finished debugging the remote program, you can use the `detach` command to release it from GDB control. Detaching from the target normally resumes its execution, but the results will depend on your particular remote stub. After the `detach` command, GDB is free to connect to another target.

disconnect

The **disconnect** command behaves like **detach**, except that the target is generally not resumed. It will wait for GDB (this instance or another one) to connect and continue debugging. After the **disconnect** command, GDB is again free to connect to another target.

17.2 Using the `gdbserver` program

`gdbserver` is a control program for Unix-like systems, which allows you to connect your program with a remote GDB via **target remote**—but without linking in the usual debugging stub.

`gdbserver` is not a complete replacement for the debugging stubs, because it requires essentially the same operating-system facilities that GDB itself does. In fact, a system that can run `gdbserver` to connect to a remote GDB could also run GDB locally! `gdbserver` is sometimes useful nevertheless, because it is a much smaller program than GDB itself. It is also easier to port than all of GDB, so you may be able to get started more quickly on a new system by using `gdbserver`. Finally, if you develop code for real-time systems, you may find that the tradeoffs involved in real-time operation make it more convenient to do as much development work as possible on another system, for example by cross-compiling. You can use `gdbserver` to make a similar choice for debugging.

GDB and `gdbserver` communicate via either a serial line or a TCP connection, using the standard GDB remote serial protocol.

On the target machine,

you need to have a copy of the program you want to debug. `gdbserver` does not need your program’s symbol table, so you can strip the program if necessary to save space. GDB on the host system does all the symbol handling.

To use the server, you must tell it how to communicate with GDB; the name of your program; and the arguments for your program. The usual syntax is:

```
target> gdbserver comm program [ args ... ]
```

`comm` is either a device name (to use a serial line) or a TCP hostname and portnumber. For example, to debug Emacs with the argument ‘`foo.txt`’ and communicate with GDB over the serial port ‘`/dev/com1`’:

```
target> gdbserver /dev/com1 emacs foo.txt
```

`gdbserver` waits passively for the host GDB to communicate with it.

To use a TCP connection instead of a serial line:

```
target> gdbserver host:2345 emacs foo.txt
```

The only difference from the previous example is the first argument, specifying that you are communicating with the host GDB via TCP. The ‘`host:2345`’ argument means that `gdbserver` is to expect a TCP connection from machine ‘`host`’ to local TCP port 2345. (Currently, the ‘`host`’ part is ignored.) You can choose any number you want for the port number as long as it does not conflict with any TCP ports already in use on the target system (for example, 23 is reserved for `telnet`).¹ You must use the same port number with the host GDB `target remote` command.

¹ If you choose a port number that conflicts with another service, `gdbserver` prints an error message and exits.

On some targets, `gdbserver` can also attach to running programs. This is accomplished via the `--attach` argument. The syntax is:

```
target> gdbserver comm --attach pid
```

pid is the process ID of a currently running process. It isn't necessary to point `gdbserver` at a binary for the running process.

You can debug processes by name instead of process ID if your target has the `pidof` utility:

```
target> gdbserver comm --attach `pidof PROGRAM`
```

In case more than one copy of *PROGRAM* is running, or *PROGRAM* has multiple threads, most versions of `pidof` support the `-s` option to only return the first process ID.

On the host machine,

connect to your target (see [Section 17.1 \[Connecting to a remote target\], page 151](#)). For TCP connections, you must start up `gdbserver` prior to using the `target remote` command. Otherwise you may get an error whose text depends on the host system, but which usually looks something like ‘Connection refused’. You don't need to use the `load` command in GDB when using `gdbserver`, since the program is already on the target.

17.3 Using the `gdbserve.nlm` program

`gdbserve.nlm` is a control program for NetWare systems, which allows you to connect your program with a remote GDB via `target remote`.

GDB and `gdbserve.nlm` communicate via a serial line, using the standard GDB remote serial protocol.

On the target machine,

you need to have a copy of the program you want to debug. `gdbserve.nlm` does not need your program's symbol table, so you can strip the program if necessary to save space. GDB on the host system does all the symbol handling.

To use the server, you must tell it how to communicate with GDB; the name of your program; and the arguments for your program. The syntax is:

```
load gdbserve [ BOARD=board ] [ PORT=port ]
               [ BAUD=baud ] program [ args ... ]
```

board and *port* specify the serial line; *baud* specifies the baud rate used by the connection. *port* and *node* default to 0, *baud* defaults to 9600 bps.

For example, to debug Emacs with the argument ‘`foo.txt`’ and communicate with GDB over serial port number 2 or board 1 using a 19200 bps connection:

```
load gdbserve BOARD=1 PORT=2 BAUD=19200 emacs foo.txt
```

On the GDB host machine, connect to your target (see [Section 17.1 \[Connecting to a remote target\], page 151](#)).

17.4 Remote configuration

The following configuration options are available when debugging remote programs:

```
set remote hardware-watchpoint-limit limit
set remote hardware-breakpoint-limit limit
```

Restrict GDB to using *limit* remote hardware breakpoint or watchpoints. A limit of -1, the default, is treated as unlimited.

17.5 Implementing a remote stub

The stub files provided with GDB implement the target side of the communication protocol, and the GDB side is implemented in the GDB source file ‘`remote.c`’. Normally, you can simply allow these subroutines to communicate, and ignore the details. (If you’re implementing your own stub file, you can still ignore the details: start with one of the existing stub files. ‘`sparc-stub.c`’ is the best organized, and therefore the easiest to read.)

To debug a program running on another machine (the debugging *target* machine), you must first arrange for all the usual prerequisites for the program to run by itself. For example, for a C program, you need:

1. A startup routine to set up the C runtime environment; these usually have a name like ‘`crt0`’. The startup routine may be supplied by your hardware supplier, or you may have to write your own.
2. A C subroutine library to support your program’s subroutine calls, notably managing input and output.
3. A way of getting your program to the other machine—for example, a download program. These are often supplied by the hardware manufacturer, but you may have to write your own from hardware documentation.

The next step is to arrange for your program to use a serial port to communicate with the machine where GDB is running (the *host* machine). In general terms, the scheme looks like this:

On the host,

GDB already understands how to use this protocol; when everything else is set up, you can simply use the ‘`target remote`’ command (see [Chapter 16 \[Specifying a Debugging Target\], page 147](#)).

On the target,

you must link with your program a few special-purpose subroutines that implement the GDB remote serial protocol. The file containing these subroutines is called a *debugging stub*.

On certain remote targets, you can use an auxiliary program `gdbserver` instead of linking a stub into your program. See [Section 17.2 \[Using the `gdbserver` program\], page 152](#), for details.

The debugging stub is specific to the architecture of the remote machine; for example, use ‘`sparc-stub.c`’ to debug programs on SPARC boards.

These working remote stubs are distributed with GDB:

`i386-stub.c`

For Intel 386 and compatible architectures.

`m68k-stub.c`

For Motorola 680x0 architectures.

sh-stub.c

For Renesas SH architectures.

sparc-stub.c

For SPARC architectures.

sparcl-stub.c

For Fujitsu SPARCLITE architectures.

The ‘README’ file in the GDB distribution may list other recently added stubs.

17.5.1 What the stub can do for you

The debugging stub for your architecture supplies these three subroutines:

set_debug_traps

This routine arranges for `handle_exception` to run when your program stops. You must call this subroutine explicitly near the beginning of your program.

handle_exception

This is the central workhorse, but your program never calls it explicitly—the setup code arranges for `handle_exception` to run when a trap is triggered.

`handle_exception` takes control when your program stops during execution (for example, on a breakpoint), and mediates communications with GDB on the host machine. This is where the communications protocol is implemented; `handle_exception` acts as the GDB representative on the target machine. It begins by sending summary information on the state of your program, then continues to execute, retrieving and transmitting any information GDB needs, until you execute a GDB command that makes your program resume; at that point, `handle_exception` returns control to your own code on the target machine.

breakpoint

Use this auxiliary subroutine to make your program contain a breakpoint. Depending on the particular situation, this may be the only way for GDB to get control. For instance, if your target machine has some sort of interrupt button, you won’t need to call this; pressing the interrupt button transfers control to `handle_exception`—in effect, to GDB. On some machines, simply receiving characters on the serial port may also trigger a trap; again, in that situation, you don’t need to call `breakpoint` from your own program—simply running ‘target remote’ from the host GDB session gets control.

Call `breakpoint` if none of these is true, or if you simply want to make certain your program stops at a predetermined point for the start of your debugging session.

17.5.2 What you must do for the stub

The debugging stubs that come with GDB are set up for a particular chip architecture, but they have no information about the rest of your debugging target machine.

First of all you need to tell the stub how to communicate with the serial port.

```
int getDebugChar()
```

Write this subroutine to read a single character from the serial port. It may be identical to `getchar` for your target system; a different name is used to allow you to distinguish the two if you wish.

```
void putDebugChar(int)
```

Write this subroutine to write a single character to the serial port. It may be identical to `putchar` for your target system; a different name is used to allow you to distinguish the two if you wish.

If you want GDB to be able to stop your program while it is running, you need to use an interrupt-driven serial driver, and arrange for it to stop when it receives a ^C ('\003', the control-C character). That is the character which GDB uses to tell the remote system to stop.

Getting the debugging target to return the proper status to GDB probably requires changes to the standard stub; one quick and dirty way is to just execute a breakpoint instruction (the “dirty” part is that GDB reports a SIGTRAP instead of a SIGINT).

Other routines you need to supply are:

```
void exceptionHandler (int exception_number, void *exception_address)
```

Write this function to install `exception_address` in the exception handling tables. You need to do this because the stub does not have any way of knowing what the exception handling tables on your target system are like (for example, the processor’s table might be in ROM, containing entries which point to a table in RAM). `exception_number` is the exception number which should be changed; its meaning is architecture-dependent (for example, different numbers might represent divide by zero, misaligned access, etc). When this exception occurs, control should be transferred directly to `exception_address`, and the processor state (stack, registers, and so on) should be just as it is when a processor exception occurs. So if you want to use a jump instruction to reach `exception_address`, it should be a simple jump, not a jump to subroutine.

For the 386, `exception_address` should be installed as an interrupt gate so that interrupts are masked while the handler runs. The gate should be at privilege level 0 (the most privileged level). The SPARC and 68k stubs are able to mask interrupts themselves without help from `exceptionHandler`.

```
void flush_i_cache()
```

On SPARC and SPARCLITE only, write this subroutine to flush the instruction cache, if any, on your target machine. If there is no instruction cache, this subroutine may be a no-op.

On target machines that have instruction caches, GDB requires this function to make certain that the state of your program is stable.

You must also make sure this library routine is available:

```
void *memset(void *, int, int)
```

This is the standard library function `memset` that sets an area of memory to a known value. If you have one of the free versions of `libc.a`, `memset` can be found there; otherwise, you must either obtain it from your hardware manufacturer, or write your own.

If you do not use the GNU C compiler, you may need other standard library subroutines as well; this varies from one stub to another, but in general the stubs are likely to use any of the common library subroutines which `gcc` generates as inline code.

17.5.3 Putting it all together

In summary, when your program is ready to debug, you must follow these steps.

1. Make sure you have defined the supporting low-level routines (see [Section 17.5.2 \[What you must do for the stub\], page 155](#)):

```
getDebugChar, putDebugChar,  
flush_i_cache, memset, exceptionHandler.
```

2. Insert these lines near the top of your program:

```
set_debug_traps();  
breakpoint();
```

3. For the 680x0 stub only, you need to provide a variable called `exceptionHook`. Normally you just use:

```
void (*exceptionHook)() = 0;
```

but if before calling `set_debug_traps`, you set it to point to a function in your program, that function is called when GDB continues after stopping on a trap (for example, bus error). The function indicated by `exceptionHook` is called with one parameter: an `int` which is the exception number.

4. Compile and link together: your program, the GDB debugging stub for your target architecture, and the supporting subroutines.
5. Make sure you have a serial connection between your target machine and the GDB host, and identify the serial port on the host.
6. Download your program to your target machine (or get it there by whatever means the manufacturer provides), and start it.
7. Start GDB on the host, and connect to the target (see [Section 17.1 \[Connecting to a remote target\], page 151](#)).

18 Configuration-Specific Information

While nearly all GDB commands are available for all native and cross versions of the debugger, there are some exceptions. This chapter describes things that are only available in certain configurations.

There are three major categories of configurations: native configurations, where the host and target are the same, embedded operating system configurations, which are usually the same for several different processor architectures, and bare embedded processors, which are quite different from each other.

18.1 Native

This section describes details specific to particular native configurations.

18.1.1 HP-UX

On HP-UX systems, if you refer to a function or variable name that begins with a dollar sign, GDB searches for a user or system name first, before it searches for a convenience variable.

18.1.2 BSD libkvm Interface

BSD-derived systems (FreeBSD/NetBSD/OpenBSD) have a kernel memory interface that provides a uniform interface for accessing kernel virtual memory images, including live systems and crash dumps. GDB uses this interface to allow you to debug live kernels and kernel crash dumps on many native BSD configurations. This is implemented as a special `kvm` debugging target. For debugging a live system, load the currently running kernel into GDB and connect to the `kvm` target:

```
(gdb) target kvm
```

For debugging crash dumps, provide the file name of the crash dump as an argument:

```
(gdb) target kvm /var/crash/bsd.0
```

Once connected to the `kvm` target, the following commands are available:

`kvm pcb` Set current context from pcb address.

`kvm proc` Set current context from proc address. This command isn't available on modern FreeBSD systems.

18.1.3 SVR4 process information

Many versions of SVR4 provide a facility called '/proc' that can be used to examine the image of a running process using file-system subroutines. If GDB is configured for an operating system with this facility, the command `info proc` is available to report on several kinds of information about the process running your program. `info proc` works only on SVR4 systems that include the `procfs` code. This includes OSF/1 (Digital Unix), Solaris, Irix, and Unixware, but not HP-UX or GNU/Linux, for example.

`info proc` Summarize available information about the process.

`info proc mappings`

Report on the address ranges accessible in the program, with information on whether your program may read, write, or execute each range.

18.1.4 Features for Debugging DJGPP Programs

DJGPP is the port of GNU development tools to MS-DOS and MS-Windows. DJGPP programs are 32-bit protected-mode programs that use the *DPMI* (DOS Protected-Mode Interface) API to run on top of real-mode DOS systems and their emulations.

GDB supports native debugging of DJGPP programs, and defines a few commands specific to the DJGPP port. This subsection describes those commands.

info dos This is a prefix of DJGPP-specific commands which print information about the target system and important OS structures.

info dos sysinfo

This command displays assorted information about the underlying platform: the CPU type and features, the OS version and flavor, the DPMI version, and the available conventional and DPMI memory.

info dos gdt

info dos ldt

info dos idt

These 3 commands display entries from, respectively, Global, Local, and Interrupt Descriptor Tables (GDT, LDT, and IDT). The descriptor tables are data structures which store a descriptor for each segment that is currently in use. The segment's selector is an index into a descriptor table; the table entry for that index holds the descriptor's base address and limit, and its attributes and access rights.

A typical DJGPP program uses 3 segments: a code segment, a data segment (used for both data and the stack), and a DOS segment (which allows access to DOS/BIOS data structures and absolute addresses in conventional memory). However, the DPMI host will usually define additional segments in order to support the DPMI environment.

These commands allow to display entries from the descriptor tables. Without an argument, all entries from the specified table are displayed. An argument, which should be an integer expression, means display a single entry whose index is given by the argument. For example, here's a convenient way to display information about the debugged program's data segment:

```
(gdb) info dos ldt $ds
0x13f: base=0x11970000 limit=0x0009ffff 32-Bit Data (Read/Write, Exp-up)
```

This comes in handy when you want to see whether a pointer is outside the data segment's limit (i.e. garbled).

info dos pde

info dos pte

These two commands display entries from, respectively, the Page Directory and the Page Tables. Page Directories and Page Tables are data structures which control how virtual memory addresses are mapped into physical addresses. A Page Table includes an entry for every page of memory that is mapped into the program's address space; there may be several Page Tables, each one holding up to 4096 entries. A Page Directory has up to 4096 entries, one each for every Page Table that is currently in use.

Without an argument, `info dos pde` displays the entire Page Directory, and `info dos pte` displays all the entries in all of the Page Tables. An argument, an integer expression, given to the `info dos pde` command means display only that entry from the Page Directory table. An argument given to the `info dos pte` command means display entries from a single Page Table, the one pointed to by the specified entry in the Page Directory.

These commands are useful when your program uses DMA (Direct Memory Access), which needs physical addresses to program the DMA controller.

These commands are supported only with some DPMI servers.

`info dos address-pte addr`

This command displays the Page Table entry for a specified linear address. The argument linear address `addr` should already have the appropriate segment's base address added to it, because this command accepts addresses which may belong to *any* segment. For example, here's how to display the Page Table entry for the page where the variable `i` is stored:

```
(gdb) info dos address-pte __djgpp_base_address + (char *)&i
Page Table entry for address 0x11a00d30:
Base=0x02698000 Dirty Acc. Not-Cached Write-Back Usr Read-Write +0xd30
```

This says that `i` is stored at offset `0xd30` from the page whose physical base address is `0x02698000`, and prints all the attributes of that page.

Note that you must cast the addresses of variables to a `char *`, since otherwise the value of `__djgpp_base_address`, the base address of all variables and functions in a DJGPP program, will be added using the rules of C pointer arithmetics: if `i` is declared an `int`, GDB will add 4 times the value of `__djgpp_base_address` to the address of `i`.

Here's another example, it displays the Page Table entry for the transfer buffer:

```
(gdb) info dos address-pte *((unsigned *)&_go32_info_block + 3)
Page Table entry for address 0x29110:
Base=0x00029000 Dirty Acc. Not-Cached Write-Back Usr Read-Write +0x110
```

(The `+ 3` offset is because the transfer buffer's address is the 3rd member of the `_go32_info_block` structure.) The output of this command clearly shows that addresses in conventional memory are mapped 1:1, i.e. the physical and linear addresses are identical.

This command is supported only with some DPMI servers.

18.1.5 Features for Debugging MS Windows PE executables

GDB supports native debugging of MS Windows programs, including DLLs with and without symbolic debugging information. There are various additional Cygwin-specific commands, described in this subsection. The subsubsection see [Section 18.1.5.1 \[Non-debug DLL symbols\]](#), [page 162](#) describes working with DLLs that have no debugging symbols.

`info w32` This is a prefix of MS Windows specific commands which print information about the target system and important OS structures.

`info w32 selector`

This command displays information returned by the Win32 API `GetThreadSelectorEntry` function. It takes an optional argument that is

evaluated to a long value to give the information about this given selector. Without argument, this command displays information about the the six segment registers.

info dll This is a Cygwin specific alias of info shared.

dll-symbols

This command loads symbols from a dll similarly to add-sym command but without the need to specify a base address.

set new-console mode

If *mode* is **on** the debugger will be started in a new console on next start. If *mode* is **off**, the debugger will be started in the same console as the debugger.

show new-console

Displays whether a new console is used when the debugger is started.

set new-group mode

This boolean value controls whether the debugger should start a new group or stay in the same group as the debugger. This affects the way the Windows OS handles Ctrl-C.

show new-group

Displays current value of new-group boolean.

set debugevents

This boolean value adds debug output concerning events seen by the debugger.

set debugexec

This boolean value adds debug output concerning execute events seen by the debugger.

set debugexceptions

This boolean value adds debug ouptut concerning exception events seen by the debugger.

set debugmemory

This boolean value adds debug ouptut concerning memory events seen by the debugger.

set shell This boolean values specifies whether the debugger is called via a shell or directly (default value is on).

show shell

Displays if the debugger will be started with a shell.

18.1.5.1 Support for DLLs without debugging symbols

Very often on windows, some of the DLLs that your program relies on do not include symbolic debugging information (for example, ‘kernel32.dll’). When GDB doesn’t recognize any debugging symbols in a DLL, it relies on the minimal amount of symbolic information contained in the DLL’s export table. This subsubsection describes working with such symbols, known internally to GDB as “minimal symbols”.

Note that before the debugged program has started execution, no DLLs will have been loaded. The easiest way around this problem is simply to start the program — either by

setting a breakpoint or letting the program run once to completion. It is also possible to force GDB to load a particular DLL before starting the executable — see the shared library information in see [Section 15.1 \[Files\], page 135](#) or the `dll-symbols` command in see [Section 18.1.5 \[Cygwin Native\], page 161](#). Currently, explicitly loading symbols from a DLL with no debugging information will cause the symbol names to be duplicated in GDB’s lookup table, which may adversely affect symbol lookup performance.

18.1.5.2 DLL name prefixes

In keeping with the naming conventions used by the Microsoft debugging tools, DLL export symbols are made available with a prefix based on the DLL name, for instance `KERNEL32!CreateFileA`. The plain name is also entered into the symbol table, so `CreateFileA` is often sufficient. In some cases there will be name clashes within a program (particularly if the executable itself includes full debugging symbols) necessitating the use of the fully qualified name when referring to the contents of the DLL. Use single-quotes around the name to avoid the exclamation mark (“!”) being interpreted as a language operator.

Note that the internal name of the DLL may be all upper-case, even though the file name of the DLL is lower-case, or vice-versa. Since symbols within GDB are *case-sensitive* this may cause some confusion. If in doubt, try the `info functions` and `info variables` commands or even `maint print msymbols` (see see [Chapter 13 \[Symbols\], page 125](#)). Here’s an example:

```
(gdb) info function CreateFileA
All functions matching regular expression "CreateFileA":

Non-debugging symbols:
0x77e885f4  CreateFileA
0x77e885f4  KERNEL32!CreateFileA

(gdb) info function !
All functions matching regular expression "!":

Non-debugging symbols:
0x6100114c  cygwin1!__assert
0x61004034  cygwin1!_dll_crt0@0
0x61004240  cygwin1!dll_crt0(per_process *)
[etc...]
```

18.1.5.3 Working with minimal symbols

Symbols extracted from a DLL’s export table do not contain very much type information. All that GDB can do is guess whether a symbol refers to a function or variable depending on the linker section that contains the symbol. Also note that the actual contents of the memory contained in a DLL are not available unless the program is running. This means that you cannot examine the contents of a variable or disassemble a function within a DLL without a running program.

Variables are generally treated as pointers and dereferenced automatically. For this reason, it is often necessary to prefix a variable name with the address-of operator (“`&`”) and provide explicit type information in the command. Here’s an example of the type of problem:

```
(gdb) print 'cygwin1!__argv'
$1 = 268572168
```

```
(gdb) x 'cygwin1!__argv'
0x10021610:      "\230y\"
```

And two possible solutions:

```
(gdb) print ((char **)'cygwin1!__argv')[0]
$2 = 0x22fd98 "/cygdrive/c/mydirectory/myprogram"
(gdb) x/2x &'cygwin1!__argv'
0x610c0aa8 <cygwin1!__argv>:    0x10021608      0x00000000
(gdb) x/x 0x10021608
0x10021608:      0x0022fd98
(gdb) x/s 0x0022fd98
0x22fd98:      "/cygdrive/c/mydirectory/myprogram"
```

Setting a break point within a DLL is possible even before the program starts execution. However, under these circumstances, GDB can't examine the initial instructions of the function in order to skip the function's frame set-up code. You can work around this by using “*&” to set the breakpoint at a raw memory address:

```
(gdb) break *&'python22!PyOS_Readline'
Breakpoint 1 at 0x1e04eff0
```

The author of these extensions is not entirely convinced that setting a break point within a shared DLL like ‘kernel32.dll’ is completely safe.

18.2 Embedded Operating Systems

This section describes configurations involving the debugging of embedded operating systems that are available for several different architectures.

GDB includes the ability to debug programs running on various real-time operating systems.

18.2.1 Using GDB with VxWorks

target vxworks *machinename*

A VxWorks system, attached via TCP/IP. The argument *machinename* is the target system’s machine name or IP address.

On VxWorks, **load** links *filename* dynamically on the current target system as well as adding its symbols in GDB.

GDB enables developers to spawn and debug tasks running on networked VxWorks targets from a Unix host. Already-running tasks spawned from the VxWorks shell can also be debugged. GDB uses code that runs on both the Unix host and on the VxWorks target. The program **gdb** is installed and executed on the Unix host. (It may be installed with the name **vxgdb**, to distinguish it from a GDB for debugging programs on the host itself.)

VxWorks-timeout *args*

All VxWorks-based targets now support the option **vxworks-timeout**. This option is set by the user, and *args* represents the number of seconds GDB waits for responses to rpc’s. You might use this if your VxWorks target is a slow software simulator or is on the far side of a thin network line.

The following information on connecting to VxWorks was current when this manual was produced; newer releases of VxWorks may use revised procedures.

To use GDB with VxWorks, you must rebuild your VxWorks kernel to include the remote debugging interface routines in the VxWorks library ‘**rdb.a**’. To do this, define **INCLUDE_RDB** in the VxWorks configuration file ‘**configAll.h**’ and rebuild your VxWorks kernel. The resulting kernel contains ‘**rdb.a**’, and spawns the source debugging task **tRdbTask** when VxWorks is booted. For more information on configuring and remaking VxWorks, see the manufacturer’s manual.

Once you have included ‘**rdb.a**’ in your VxWorks system image and set your Unix execution search path to find GDB, you are ready to run GDB. From your Unix host, run **gdb** (or **vxgdb**, depending on your installation).

GDB comes up showing the prompt:

```
(vxgdb)
```

18.2.1.1 Connecting to VxWorks

The GDB command **target** lets you connect to a VxWorks target on the network. To connect to a target whose host name is “tt”, type:

```
(vxgdb) target vxworks tt
```

GDB displays messages like these:

```
Attaching remote machine across net...
Connected to tt.
```

GDB then attempts to read the symbol tables of any object modules loaded into the VxWorks target since it was last booted. GDB locates these files by searching the directories listed in the command search path (see [Section 4.4 \[Your program’s environment\], page 26](#)); if it fails to find an object file, it displays a message such as:

```
prog.o: No such file or directory.
```

When this happens, add the appropriate directory to the search path with the GDB command **path**, and execute the **target** command again.

18.2.1.2 VxWorks download

If you have connected to the VxWorks target and you want to debug an object that has not yet been loaded, you can use the GDB **load** command to download a file from Unix to VxWorks incrementally. The object file given as an argument to the **load** command is actually opened twice: first by the VxWorks target in order to download the code, then by GDB in order to read the symbol table. This can lead to problems if the current working directories on the two systems differ. If both systems have NFS mounted the same filesystems, you can avoid these problems by using absolute paths. Otherwise, it is simplest to set the working directory on both systems to the directory in which the object file resides, and then to reference the file by its name, without any path. For instance, a program ‘**prog.o**’ may reside in ‘**vxpath/vw/demo/rdb**’ in VxWorks and in ‘**hostpath/vw/demo/rdb**’ on the host. To load this program, type this on VxWorks:

```
-> cd "vxpath/vw/demo/rdb"
```

Then, in GDB, type:

```
(vxgdb) cd hostpath/vw/demo/rdb
(vxgdb) load prog.o
```

GDB displays a response similar to this:

```
Reading symbol data from wherever/vw/demo/rdb/prog.o.... done.
```

You can also use the `load` command to reload an object module after editing and recompiling the corresponding source file. Note that this makes GDB delete all currently-defined breakpoints, auto-displays, and convenience variables, and to clear the value history. (This is necessary in order to preserve the integrity of debugger's data structures that reference the target system's symbol table.)

18.2.1.3 Running tasks

You can also attach to an existing task using the `attach` command as follows:

```
(vxgdb) attach task
```

where *task* is the VxWorks hexadecimal task ID. The task can be running or suspended when you attach to it. Running tasks are suspended at the time of attachment.

18.3 Embedded Processors

This section goes into details specific to particular embedded configurations.

18.3.1 ARM

```
target rdi dev
```

ARM Angel monitor, via RDI library interface to ADP protocol. You may use this target to communicate with both boards running the Angel monitor, or with the EmbeddedICE JTAG debug device.

```
target rdp dev
```

ARM Demon monitor.

18.3.2 Renesas H8/300

```
target hms dev
```

A Renesas SH, H8/300, or H8/500 board, attached via serial line to your host. Use special commands `device` and `speed` to control the serial line and the communications speed used.

```
target e7000 dev
```

E7000 emulator for Renesas H8 and SH.

```
target sh3 dev
```

```
target sh3e dev
```

Renesas SH-3 and SH-3E target systems.

When you select remote debugging to a Renesas SH, H8/300, or H8/500 board, the `load` command downloads your program to the Renesas board and also opens it as the current executable target for GDB on your host (like the `file` command).

GDB needs to know these things to talk to your Renesas SH, H8/300, or H8/500:

1. that you want to use ‘`target hms`’, the remote debugging interface for Renesas microprocessors, or ‘`target e7000`’, the in-circuit emulator for the Renesas SH and the Renesas 300H. (‘`target hms`’ is the default when GDB is configured specifically for the Renesas SH, H8/300, or H8/500.)

2. what serial device connects your host to your Renesas board (the first serial device available on your host is the default).
3. what speed to use over the serial device.

18.3.2.1 Connecting to Renesas boards

Use the special GDB command ‘`device port`’ if you need to explicitly set the serial device. The default *port* is the first available port on your host. This is only necessary on Unix hosts, where it is typically something like ‘`/dev/ttya`’.

GDB has another special command to set the communications speed: ‘`speed bps`’. This command also is only used from Unix hosts; on DOS hosts, set the line speed as usual from outside GDB with the DOS `mode` command (for instance, `mode com2:9600,n,8,1,p` for a 9600 bps connection).

The ‘`device`’ and ‘`speed`’ commands are available only when you use a Unix host to debug your Renesas microprocessor programs. If you use a DOS host, GDB depends on an auxiliary terminate-and-stay-resident program called `asynctsr` to communicate with the development board through a PC serial port. You must also use the DOS `mode` command to set up the serial port on the DOS side.

The following sample session illustrates the steps needed to start a program under GDB control on an H8/300. The example uses a sample H8/300 program called ‘`t.x`’. The procedure is the same for the Renesas SH and the H8/500.

First hook up your development board. In this example, we use a board attached to serial port COM2; if you use a different serial port, substitute its name in the argument of the `mode` command. When you call `asynctsr`, the auxiliary comms program used by the debugger, you give it just the numeric part of the serial port’s name; for example, ‘`asyncstr 2`’ below runs `asyncstr` on COM2.

```
C:\H8300\TEST> asynctsr 2
C:\H8300\TEST> mode com2:9600,n,8,1,p

Resident portion of MODE loaded

COM2: 9600, n, 8, 1, p
```

Warning: We have noticed a bug in PC-NFS that conflicts with `asynctsr`. If you also run PC-NFS on your DOS host, you may need to disable it, or even boot without it, to use `asynctsr` to control your development board.

Now that serial communications are set up, and the development board is connected, you can start up GDB. Call `gdb` with the name of your program as the argument. GDB prompts you, as usual, with the prompt ‘`(gdb)`’. Use two special commands to begin your debugging session: ‘`target hms`’ to specify cross-debugging to the Renesas board, and the `load` command to download your program to the board. `load` displays the names of the program’s sections, and a ‘*’ for each 2K of data downloaded. (If you want to refresh GDB data on symbols or on the executable file without downloading, use the GDB commands `file` or `symbol-file`. These commands, and `load` itself, are described in [Section 15.1 \[Commands to specify files\], page 135](#).)

```
(eg-C:\H8300\TEST) gdb t.x
GDB is free software and you are welcome to distribute copies
of it under certain conditions; type "show copying" to see
```

```

the conditions.
There is absolutely no warranty for GDB; type "show warranty"
for details.
GDB 6.2.50_20041017, Copyright 1992 Free Software Foundation, Inc...
(gdb) target hms
Connected to remote H8/300 HMS system.
(gdb) load t.x
.text    : 0x8000 .. 0xabde *****
.data    : 0xabde .. 0xad30 *
.stack   : 0xf000 .. 0xf014 *

```

At this point, you're ready to run or debug your program. From here on, you can use all the usual GDB commands. The `break` command sets breakpoints; the `run` command starts your program; `print` or `x` display data; the `continue` command resumes execution after stopping at a breakpoint. You can use the `help` command at any time to find out more about GDB commands.

Remember, however, that *operating system* facilities aren't available on your development board; for example, if your program hangs, you can't send an interrupt—but you can press the RESET switch!

Use the RESET button on the development board

- to interrupt your program (don't use `ctrl-C` on the DOS host—it has no way to pass an interrupt signal to the development board); and
- to return to the GDB command prompt after your program finishes normally. The communications protocol provides no other way for GDB to detect program completion.

In either case, GDB sees the effect of a RESET on the development board as a “normal exit” of your program.

18.3.2.2 Using the E7000 in-circuit emulator

You can use the E7000 in-circuit emulator to develop code for either the Renesas SH or the H8/300H. Use one of these forms of the ‘`target e7000`’ command to connect GDB to your E7000:

`target e7000 port speed`

Use this form if your E7000 is connected to a serial port. The `port` argument identifies what serial port to use (for example, ‘`com2`’). The third argument is the line speed in bits per second (for example, ‘`9600`’).

`target e7000 hostname`

If your E7000 is installed as a host on a TCP/IP network, you can just specify its hostname; GDB uses `telnet` to connect.

18.3.2.3 Special GDB commands for Renesas micros

Some GDB commands are available only for the H8/300:

```

set machine h8300
set machine h8300h

```

Condition GDB for one of the two variants of the H8/300 architecture with ‘`set machine`’. You can use ‘`show machine`’ to check which variant is currently in effect.

18.3.3 H8/500

```
set memory mod  
show memory
```

Specify which H8/500 memory model (*mod*) you are using with ‘`set memory`’; check which memory model is in effect with ‘`show memory`’. The accepted values for *mod* are `small`, `big`, `medium`, and `compact`.

18.3.4 Renesas M32R/D

```
target m32r dev
```

Renesas M32R/D ROM monitor.

```
target m32rsdi dev
```

Renesas M32R SDI server, connected via parallel port to the board.

18.3.5 M68k

The Motorola m68k configuration includes ColdFire support, and target command for the following ROM monitors.

```
target abug dev
```

ABug ROM monitor for M68K.

```
target cpu32bug dev
```

CPU32BUG monitor, running on a CPU32 (M68K) board.

```
target dbug dev
```

dBUG ROM monitor for Motorola ColdFire.

```
target est dev
```

EST-300 ICE monitor, running on a CPU32 (M68K) board.

```
target rom68k dev
```

ROM 68K monitor, running on an M68K IDP board.

```
target rombug dev
```

ROMBUG ROM monitor for OS/9000.

18.3.6 MIPS Embedded

GDB can use the MIPS remote debugging protocol to talk to a MIPS board attached to a serial line. This is available when you configure GDB with ‘`--target=mips-idt-ecoff`’.

Use these GDB commands to specify the connection to your target board:

```
target mips port
```

To run a program on the board, start up `gdb` with the name of your program as the argument. To connect to the board, use the command ‘`target mips port`’, where *port* is the name of the serial port connected to the board. If the program has not already been downloaded to the board, you may use the `load` command to download it. You can then use all the usual GDB commands.

For example, this sequence connects to the target board through a serial port, and loads and runs a program called *prog* through the debugger:

```
host$ gdb prog
GDB is free software and ...
(gdb) target mips /dev/ttyb
(gdb) load prog
(gdb) run
```

target mips *hostname:portnumber*

On some GDB host configurations, you can specify a TCP connection (for instance, to a serial line managed by a terminal concentrator) instead of a serial port, using the syntax ‘*hostname:portnumber*’.

target pmon *port*

PMON ROM monitor.

target ddb *port*

NEC’s DDB variant of PMON for Vr4300.

target lsi *port*

LSI variant of PMON.

target r3900 *dev*

Densan DVE-R3900 ROM monitor for Toshiba R3900 Mips.

target array *dev*

Array Tech LSI33K RAID controller board.

GDB also supports these special commands for MIPS targets:

```
set processor args
show processor
```

Use the **set processor** command to set the type of MIPS processor when you want to access processor-type-specific registers. For example, **set processor r3041** tells GDB to use the CPU registers appropriate for the 3041 chip. Use the **show processor** command to see what MIPS processor GDB is using. Use the **info reg** command to see what registers GDB is using.

```
set mipsfpu double
set mipsfpu single
set mipsfpu none
show mipsfpu
```

If your target board does not support the MIPS floating point coprocessor, you should use the command ‘**set mipsfpu none**’ (if you need this, you may wish to put the command in your GDB init file). This tells GDB how to find the return value of functions which return floating point values. It also allows GDB to avoid saving the floating point registers when calling functions on the board. If you are using a floating point coprocessor with only single precision floating point support, as on the R4650 processor, use the command ‘**set mipsfpu single**’. The default double precision floating point coprocessor may be selected using ‘**set mipsfpu double**’.

In previous versions the only choices were double precision or no floating point, so ‘**set mipsfpu on**’ will select double precision and ‘**set mipsfpu off**’ will select no floating point.

As usual, you can inquire about the **mipsfpu** variable with ‘**show mipsfpu**’.

```
set remotedebug n
show remotedebug
```

You can see some debugging information about communications with the board by setting the `remotedebug` variable. If you set it to 1 using ‘`set remotedebug 1`’, every packet is displayed. If you set it to 2, every character is displayed. You can check the current value at any time with the command ‘`show remotedebug`’.

```
set timeout seconds
set retransmit-timeout seconds
show timeout
show retransmit-timeout
```

You can control the timeout used while waiting for a packet, in the MIPS remote protocol, with the `set timeout seconds` command. The default is 5 seconds. Similarly, you can control the timeout used while waiting for an acknowledgement of a packet with the `set retransmit-timeout seconds` command. The default is 3 seconds. You can inspect both values with `show timeout` and `show retransmit-timeout`. (These commands are *only* available when GDB is configured for ‘`--target=mips-idt-ecoff`’.)

The timeout set by `set timeout` does not apply when GDB is waiting for your program to stop. In that case, GDB waits forever because it has no way of knowing how long the program is going to run before stopping.

18.3.7 OpenRISC 1000

See OR1k Architecture document (www.opencores.org) for more information about platform and commands.

`target jtag jtag://host:port`

Connects to remote JTAG server. JTAG remote server can be either an or1ksim or JTAG server, connected via parallel port to the board.

Example: `target jtag jtag://localhost:9999`

`or1ksim command`

If connected to `or1ksim` OpenRISC 1000 Architectural Simulator, proprietary commands can be executed.

`info or1k spr`

Displays spr groups.

`info or1k spr group`

`info or1k spr groupno`

Displays register names in selected group.

`info or1k spr group register`

`info or1k spr register`

`info or1k spr groupno registerno`

`info or1k spr registerno`

Shows information about specified spr register.

```
spr group register value
spr register value
spr groupno registerno value
spr registerno value
```

Writes value to specified spr register.

Some implementations of OpenRISC 1000 Architecture also have hardware trace. It is very similar to GDB trace, except it does not interfere with normal program execution and is thus much faster. Hardware breakpoints/watchpoint triggers can be set using:

\$LEA/\$LDATA

Load effective address/data

\$SEA/\$SDATA

Store effective address/data

\$AEA/\$ADATA

Access effective address (\$SEA or \$LEA) or data (\$SDATA/\$LDATA)

\$FETCH Fetch data

When triggered, it can capture low level data, like: PC, LSEA, LDATA, SDATA, READSPR, WRITESPR, INSTR.

htrace commands:

hwatch conditional

Set hardware watchpoint on combination of Load/Store Effective Address(es) or Data. For example:

```
hwatch ($LEA == my_var) && ($LDATA < 50) || ($SEA == my_var) &&
($SDATA >= 50)
```

```
hwatch ($LEA == my_var) && ($LDATA < 50) || ($SEA == my_var) &&
($SDATA >= 50)
```

htrace info

Display information about current HW trace configuration.

htrace trigger conditional

Set starting criteria for HW trace.

htrace qualifier conditional

Set acquisition qualifier for HW trace.

htrace stop conditional

Set HW trace stopping criteria.

htrace record [data]*

Selects the data to be recorded, when qualifier is met and HW trace was triggered.

htrace enable

htrace disable

Enables/disables the HW trace.

```
htrace rewind [filename]
    Clears currently recorded trace data.

    If filename is specified, new trace file is made and any newly collected data will
    be written there.

htrace print [start [len]]
    Prints trace buffer, using current record configuration.

htrace mode continuous
    Set continuous trace mode.

htrace mode suspend
    Set suspend trace mode.
```

18.3.8 PowerPC

```
target dink32 dev
    DINK32 ROM monitor.

target ppcbug dev
target ppcbug1 dev
    PPCBUG ROM monitor for PowerPC.

target sds dev
    SDS monitor, running on a PowerPC board (such as Motorola's ADS).
```

18.3.9 HP PA Embedded

```
target op50n dev
    OP50N monitor, running on an OKI HPPA board.

target w89k dev
    W89K monitor, running on a Winbond HPPA board.
```

18.3.10 Renesas SH

```
target hms dev
    A Renesas SH board attached via serial line to your host. Use special commands
    device and speed to control the serial line and the communications speed used.

target e7000 dev
    E7000 emulator for Renesas SH.

target sh3 dev
target sh3e dev
    Renesas SH-3 and SH-3E target systems.
```

18.3.11 Tsquare Sparclet

GDB enables developers to debug tasks running on Sparclet targets from a Unix host. GDB uses code that runs on both the Unix host and on the Sparclet target. The program `gdb` is installed and executed on the Unix host.

```
remotetimeout args
    GDB supports the option remotetimeout. This option is set by the user, and
    args represents the number of seconds GDB waits for responses.
```

When compiling for debugging, include the options ‘-g’ to get debug information and ‘-Ttext’ to relocate the program to where you wish to load it on the target. You may also want to add the options ‘-n’ or ‘-N’ in order to reduce the size of the sections. Example:

```
sparclet-aout-gcc prog.c -Ttext 0x12010000 -g -o prog -N
```

You can use `objdump` to verify that the addresses are what you intended:

```
sparclet-aout-objdump --headers --syms prog
```

Once you have set your Unix execution search path to find GDB, you are ready to run GDB. From your Unix host, run `gdb` (or `sparclet-aout-gdb`, depending on your installation).

GDB comes up showing the prompt:

```
(gdb$)
```

18.3.11.1 Setting file to debug

The GDB command `file` lets you choose with program to debug.

```
(gdb$) file prog
```

GDB then attempts to read the symbol table of ‘prog’. GDB locates the file by searching the directories listed in the command search path. If the file was compiled with debug information (option “-g”), source files will be searched as well. GDB locates the source files by searching the directories listed in the directory search path (see [Section 4.4 \[Your program’s environment\], page 26](#)). If it fails to find a file, it displays a message such as:

```
prog: No such file or directory.
```

When this happens, add the appropriate directories to the search paths with the GDB commands `path` and `dir`, and execute the `target` command again.

18.3.11.2 Connecting to Sparclet

The GDB command `target` lets you connect to a Sparclet target. To connect to a target on serial port “`ttya`”, type:

```
(gdb$) target sparclet /dev/ttya
Remote target sparclet connected to /dev/ttya
main () at ../prog.c:3
```

GDB displays messages like these:

```
Connected to ttya.
```

18.3.11.3 Sparclet download

Once connected to the Sparclet target, you can use the GDB `load` command to download the file from the host to the target. The file name and load offset should be given as arguments to the `load` command. Since the file format is `aout`, the program must be loaded to the starting address. You can use `objdump` to find out what this value is. The load offset is an offset which is added to the VMA (virtual memory address) of each of the file’s sections. For instance, if the program ‘`prog`’ was linked to text address `0x1201000`, with data at `0x12010160` and bss at `0x12010170`, in GDB, type:

```
(gdb$) load prog 0x12010000
Loading section .text, size 0xdb0 vma 0x12010000
```

If the code is loaded at a different address than what the program was linked to, you may need to use the `section` and `add-symbol-file` commands to tell GDB where to map the symbol table.

18.3.11.4 Running and debugging

You can now begin debugging the task using GDB's execution control commands, `b`, `step`, `run`, etc. See the GDB manual for the list of commands.

```
(gdb$) b main
Breakpoint 1 at 0x12010000: file prog.c, line 3.
(gdb$) run
Starting program: prog
Breakpoint 1, main (argc=1, argv=0xfffff21c) at prog.c:3
3      char *symarg = 0;
(gdb$) step
4      char *execarg = "hello!";
(gdb$)
```

18.3.12 Fujitsu Sparclite

`target sparclite dev`

Fujitsu sparclite boards, used only for the purpose of loading. You must use an additional command to debug the program. For example: `target remote dev` using GDB standard remote protocol.

18.3.13 Tandem ST2000

GDB may be used with a Tandem ST2000 phone switch, running Tandem's STDBUG protocol.

To connect your ST2000 to the host system, see the manufacturer's manual. Once the ST2000 is physically attached, you can run:

`target st2000 dev speed`

to establish it as your debugging environment. `dev` is normally the name of a serial device, such as '`/dev/ttya`', connected to the ST2000 via a serial line. You can instead specify `dev` as a TCP connection (for example, to a serial line attached via a terminal concentrator) using the syntax `hostname :portnumber`.

The `load` and `attach` commands are *not* defined for this target; you must load your program into the ST2000 as you normally would for standalone operation. GDB reads debugging information (such as symbols) from a separate, debugging version of the program available on your host computer.

These auxiliary GDB commands are available to help you with the ST2000 environment:

`st2000 command`

Send a `command` to the STDBUG monitor. See the manufacturer's manual for available commands.

`connect` Connect the controlling terminal to the STDBUG command monitor. When you are done interacting with STDBUG, typing either of two character sequences gets you back to the GDB command prompt: `\RET~`. (Return, followed by tilde and period) or `\RET~\C-d` (Return, followed by tilde and control-D).

18.3.14 Zilog Z8000

When configured for debugging Zilog Z8000 targets, GDB includes a Z8000 simulator.

For the Z8000 family, '`target sim`' simulates either the Z8002 (the unsegmented variant of the Z8000 architecture) or the Z8001 (the segmented variant). The simulator recognizes which architecture is appropriate by inspecting the object code.

```
target sim args
```

Debug programs on a simulated CPU. If the simulator supports setup options, specify them via `args`.

After specifying this target, you can debug programs for the simulated CPU in the same style as programs for your host computer; use the `file` command to load a new program image, the `run` command to run your program, and so on.

As well as making available all the usual machine registers (see [Section 8.10 \[Registers\], page 78](#)), the Z8000 simulator provides three additional items of information as specially named registers:

`cycles` Counts clock-ticks in the simulator.

`insts` Counts instructions run in the simulator.

`time` Execution time in 60ths of a second.

You can refer to these values in GDB expressions with the usual conventions; for example, ‘`b fputc if $cycles>5000`’ sets a conditional breakpoint that suspends only after at least 5000 simulated clock ticks.

18.4 Architectures

This section describes characteristics of architectures that affect all uses of GDB with the architecture, both native and cross.

18.4.1 A29K

```
set rstack_high_address address
```

On AMD 29000 family processors, registers are saved in a separate register stack. There is no way for GDB to determine the extent of this stack. Normally, GDB just assumes that the stack is “large enough”. This may result in GDB referencing memory locations that do not exist. If necessary, you can get around this problem by specifying the ending address of the register stack with the `set rstack_high_address` command. The argument should be an address, which you probably want to precede with ‘`0x`’ to specify in hexadecimal.

```
show rstack_high_address
```

Display the current limit of the register stack, on AMD 29000 family processors.

18.4.2 Alpha

See the following section.

18.4.3 MIPS

Alpha- and MIPS-based computers use an unusual stack frame, which sometimes requires GDB to search backward in the object code to find the beginning of a function.

To improve response time (especially for embedded applications, where GDB may be restricted to a slow serial line for this search) you may want to limit the size of this search, using one of these commands:

```
set heuristic-fence-post limit
```

Restrict GDB to examining at most *limit* bytes in its search for the beginning of a function. A value of *0* (the default) means there is no limit. However, except for *0*, the larger the limit the more bytes `heuristic-fence-post` must search and therefore the longer it takes to run.

```
show heuristic-fence-post
```

Display the current limit.

These commands are available *only* when GDB is configured for debugging programs on Alpha or MIPS processors.

19 Controlling GDB

You can alter the way GDB interacts with you by using the `set` command. For commands controlling how GDB displays data, see [Section 8.7 \[Print settings\], page 71](#). Other settings are described here.

19.1 Prompt

GDB indicates its readiness to read a command by printing a string called the *prompt*. This string is normally ‘(gdb)’. You can change the prompt string with the `set prompt` command. For instance, when debugging GDB with GDB, it is useful to change the prompt in one of the GDB sessions so that you can always tell which one you are talking to.

Note: `set prompt` does not add a space for you after the prompt you set. This allows you to set a prompt which ends in a space or a prompt that does not.

`set prompt newprompt`

Directs GDB to use *newprompt* as its prompt string henceforth.

`show prompt`

Prints a line of the form: ‘Gdb’s prompt is: *your-prompt*’

19.2 Command editing

GDB reads its input commands via the *Readline* interface. This GNU library provides consistent behavior for programs which provide a command line interface to the user. Advantages are GNU Emacs-style or vi-style inline editing of commands, csh-like history substitution, and a storage and recall of command history across debugging sessions.

You may control the behavior of command line editing in GDB with the command `set`.

`set editing`

`set editing on`

Enable command line editing (enabled by default).

`set editing off`

Disable command line editing.

`show editing`

Show whether command line editing is enabled.

See [Chapter 27 \[Command Line Editing\], page 261](#), for more details about the Readline interface. Users unfamiliar with GNU Emacs or vi are encouraged to read that chapter.

19.3 Command history

GDB can keep track of the commands you type during your debugging sessions, so that you can be certain of precisely what happened. Use these commands to manage the GDB command history facility.

GDB uses the GNU History library, a part of the Readline package, to provide the history facility. See [Chapter 28 \[Using History Interactively\], page 279](#), for the detailed description of the History library.

Here is the description of GDB commands related to command history.

set history filename *fname*

Set the name of the GDB command history file to *fname*. This is the file where GDB reads an initial command history list, and where it writes the command history from this session when it exits. You can access this list through history expansion or through the history command editing characters listed below. This file defaults to the value of the environment variable `GDBHISTFILE`, or to ‘`./.gdb_history`’ (‘`./_gdb_history`’ on MS-DOS) if this variable is not set.

set history save

set history save on

Record command history in a file, whose name may be specified with the **set history filename** command. By default, this option is disabled.

set history save off

Stop recording command history in a file.

set history size *size*

Set the number of commands which GDB keeps in its history list. This defaults to the value of the environment variable `HISTSIZE`, or to 256 if this variable is not set.

History expansion assigns special meaning to the character `!`. See [Section 28.1.1 \[Event Designators\], page 279](#), for more details.

Since `!` is also the logical not operator in C, history expansion is off by default. If you decide to enable history expansion with the **set history expansion on** command, you may sometimes need to follow `!` (when it is used as logical not, in an expression) with a space or a tab to prevent it from being expanded. The readline history facilities do not attempt substitution on the strings `!=` and `!(``, even when history expansion is enabled.

The commands to control history expansion are:

set history expansion on

set history expansion

Enable history expansion. History expansion is off by default.

set history expansion off

Disable history expansion.

show history

show history filename

show history save

show history size

show history expansion

These commands display the state of the GDB history parameters. **show history** by itself displays all four states.

show commands

Display the last ten commands in the command history.

show commands *n*

Print ten commands centered on command number *n*.

show commands +

Print ten commands just after the commands last printed.

19.4 Screen size

Certain commands to GDB may produce large amounts of information output to the screen. To help you read all of it, GDB pauses and asks you for input at the end of each page of output. Type `<RET>` when you want to continue the output, or `q` to discard the remaining output. Also, the screen width setting determines when to wrap lines of output. Depending on what is being printed, GDB tries to break the line at a readable place, rather than simply letting it overflow onto the following line.

Normally GDB knows the size of the screen from the terminal driver software. For example, on Unix GDB uses the termcap data base together with the value of the `TERM` environment variable and the `stty rows` and `stty cols` settings. If this is not correct, you can override it with the `set height` and `set width` commands:

```
set height lpp
show height
set width cpl
show width
```

These `set` commands specify a screen height of *lpp* lines and a screen width of *cpl* characters. The associated `show` commands display the current settings.

If you specify a height of zero lines, GDB does not pause during output no matter how long the output is. This is useful if output is to a file or to an editor buffer.

Likewise, you can specify ‘`set width 0`’ to prevent GDB from wrapping its output.

19.5 Numbers

You can always enter numbers in octal, decimal, or hexadecimal in GDB by the usual conventions: octal numbers begin with ‘0’, decimal numbers end with ‘.’, and hexadecimal numbers begin with ‘0x’. Numbers that begin with none of these are, by default, entered in base 10; likewise, the default display for numbers—when no particular format is specified—is base 10. You can change the default base for both input and output with the `set radix` command.

`set input-radix base`

Set the default base for numeric input. Supported choices for *base* are decimal 8, 10, or 16. *base* must itself be specified either unambiguously or using the current default radix; for example, any of

```
set radix 012
set radix 10.
set radix 0xa
```

sets the base to decimal. On the other hand, ‘`set radix 10`’ leaves the radix unchanged no matter what it was.

`set output-radix base`

Set the default base for numeric display. Supported choices for *base* are decimal 8, 10, or 16. *base* must itself be specified either unambiguously or using the current default radix.

`show input-radix`

Display the current default base for numeric input.

```
show output-radix
    Display the current default base for numeric display.
```

19.6 Configuring the current ABI

GDB can determine the *ABI* (Application Binary Interface) of your application automatically. However, sometimes you need to override its conclusions. Use these commands to manage GDB's view of the current ABI.

One GDB configuration can debug binaries for multiple operating system targets, either via remote debugging or native emulation. GDB will autodetect the *OS ABI* (Operating System ABI) in use, but you can override its conclusion using the `set osabi` command. One example where this is useful is in debugging of binaries which use an alternate C library (e.g. uClibc for GNU/Linux) which does not have the same identifying marks that the standard C library for your platform provides.

```
show osabi
    Show the OS ABI currently in use.

set osabi
    With no argument, show the list of registered available OS ABI's.

set osabi abi
    Set the current OS ABI to abi.
```

Generally, the way that an argument of type `float` is passed to a function depends on whether the function is prototyped. For a prototyped (i.e. ANSI/ISO style) function, `float` arguments are passed unchanged, according to the architecture's convention for `float`. For unprototyped (i.e. K&R style) functions, `float` arguments are first promoted to type `double` and then passed.

Unfortunately, some forms of debug information do not reliably indicate whether a function is prototyped. If GDB calls a function that is not marked as prototyped, it consults `set coerce-float-to-double`.

```
set coerce-float-to-double
set coerce-float-to-double on
    Arguments of type float will be promoted to double when passed to an unprototyped function. This is the default setting.

set coerce-float-to-double off
    Arguments of type float will be passed directly to unprototyped functions.
```

GDB needs to know the ABI used for your program's C++ objects. The correct C++ ABI depends on which C++ compiler was used to build your application. GDB only fully supports programs with a single C++ ABI; if your program contains code using multiple C++ ABI's or if GDB can not identify your program's ABI correctly, you can tell GDB which ABI to use. Currently supported ABI's include "gnu-v2", for g++ versions before 3.0, "gnu-v3", for g++ versions 3.0 and later, and "hpaCC" for the HP ANSI C++ compiler. Other C++ compilers may use the "gnu-v2" or "gnu-v3" ABI's as well. The default setting is "auto".

```
show cp-abi
    Show the C++ ABI currently in use.
```

```
set cp-abi
```

With no argument, show the list of supported C++ ABI's.

```
set cp-abi abi
```

```
set cp-abi auto
```

Set the current C++ ABI to *abi*, or return to automatic detection.

19.7 Optional warnings and messages

By default, GDB is silent about its inner workings. If you are running on a slow machine, you may want to use the `set verbose` command. This makes GDB tell you when it does a lengthy internal operation, so you will not think it has crashed.

Currently, the messages controlled by `set verbose` are those which announce that the symbol table for a source file is being read; see `symbol-file` in [Section 15.1 \[Commands to specify files\]](#), page 135.

```
set verbose on
```

Enables GDB output of certain informational messages.

```
set verbose off
```

Disables GDB output of certain informational messages.

```
show verbose
```

Displays whether `set verbose` is on or off.

By default, if GDB encounters bugs in the symbol table of an object file, it is silent; but if you are debugging a compiler, you may find this information useful (see [Section 15.3 \[Errors reading symbol files\]](#), page 143).

```
set complaints limit
```

Permits GDB to output *limit* complaints about each type of unusual symbols before becoming silent about the problem. Set *limit* to zero to suppress all complaints; set it to a large number to prevent complaints from being suppressed.

```
show complaints
```

Displays how many symbol complaints GDB is permitted to produce.

By default, GDB is cautious, and asks what sometimes seems to be a lot of stupid questions to confirm certain commands. For example, if you try to run a program which is already running:

```
(gdb) run  
The program being debugged has been started already.  
Start it from the beginning? (y or n)
```

If you are willing to unflinchingly face the consequences of your own commands, you can disable this “feature”:

```
set confirm off
```

Disables confirmation requests.

```
set confirm on
```

Enables confirmation requests (the default).

```
show confirm
```

Displays state of confirmation requests.

19.8 Optional messages about internal happenings

set debug arch
Turns on or off display of gdbarch debugging info. The default is off.

show debug arch
Displays the current state of displaying gdbarch debugging info.

set debug event
Turns on or off display of GDB event debugging info. The default is off.

show debug event
Displays the current state of displaying GDB event debugging info.

set debug expression
Turns on or off display of GDB expression debugging info. The default is off.

show debug expression
Displays the current state of displaying GDB expression debugging info.

set debug frame
Turns on or off display of GDB frame debugging info. The default is off.

show debug frame
Displays the current state of displaying GDB frame debugging info.

set debug observer
Turns on or off display of GDB observer debugging. This includes info such as the notification of observable events.

show debug observer
Displays the current state of observer debugging.

set debug overload
Turns on or off display of GDB C++ overload debugging info. This includes info such as ranking of functions, etc. The default is off.

show debug overload
Displays the current state of displaying GDB C++ overload debugging info.

set debug remote
Turns on or off display of reports on all packets sent back and forth across the serial line to the remote machine. The info is printed on the GDB standard output stream. The default is off.

show debug remote
Displays the state of display of remote packets.

set debug serial
Turns on or off display of GDB serial debugging info. The default is off.

show debug serial
Displays the current state of displaying GDB serial debugging info.

set debug target
Turns on or off display of GDB target debugging info. This info includes what is going on at the target level of GDB, as it happens. The default is 0. Set it

to 1 to track events, and to 2 to also track the value of large memory transfers. Changes to this flag do not take effect until the next time you connect to a target or use the `run` command.

show debug target

Displays the current state of displaying GDB target debugging info.

set debug varobj

Turns on or off display of GDB variable object debugging info. The default is off.

show debug varobj

Displays the current state of displaying GDB variable object debugging info.

20 Canned Sequences of Commands

Aside from breakpoint commands (see [Section 5.1.7 \[Breakpoint command lists\]](#), page 44), GDB provides two ways to store sequences of commands for execution as a unit: user-defined commands and command files.

20.1 User-defined commands

A *user-defined command* is a sequence of GDB commands to which you assign a new name as a command. This is done with the `define` command. User commands may accept up to 10 arguments separated by whitespace. Arguments are accessed within the user command via `$arg0...$arg9`. A trivial example:

```
define adder
  print $arg0 + $arg1 + $arg2
```

To execute the command use:

```
adder 1 2 3
```

This defines the command `adder`, which prints the sum of its three arguments. Note the arguments are text substitutions, so they may reference variables, use complex expressions, or even perform inferior functions calls.

`define commandname`

Define a command named `commandname`. If there is already a command by that name, you are asked to confirm that you want to redefine it.

The definition of the command is made up of other GDB command lines, which are given following the `define` command. The end of these commands is marked by a line containing `end`.

`if` Takes a single argument, which is an expression to evaluate. It is followed by a series of commands that are executed only if the expression is true (nonzero). There can then optionally be a line `else`, followed by a series of commands that are only executed if the expression was false. The end of the list is marked by a line containing `end`.

`while` The syntax is similar to `if`: the command takes a single argument, which is an expression to evaluate, and must be followed by the commands to execute, one per line, terminated by an `end`. The commands are executed repeatedly as long as the expression evaluates to true.

`document commandname`

Document the user-defined command `commandname`, so that it can be accessed by `help`. The command `commandname` must already be defined. This command reads lines of documentation just as `define` reads the lines of the command definition, ending with `end`. After the `document` command is finished, `help` on command `commandname` displays the documentation you have written.

You may use the `document` command again to change the documentation of a command. Redefining the command with `define` does not change the documentation.

```
help user-defined
    List all user-defined commands, with the first line of the documentation (if any)
    for each.

show user
show user commandname
    Display the GDB commands used to define commandname (but not its documentation).
    If no commandname is given, display the definitions for all user-defined
    commands.

show max-user-call-depth
set max-user-call-depth
    The value of max-user-call-depth controls how many recursion levels are
    allowed in user-defined commands before GDB suspects an infinite recursion
    and aborts the command.
```

When user-defined commands are executed, the commands of the definition are not printed. An error in any command stops execution of the user-defined command.

If used interactively, commands that would ask for confirmation proceed without asking when used inside a user-defined command. Many GDB commands that normally print messages to say what they are doing omit the messages when used in a user-defined command.

20.2 User-defined command hooks

You may define *hooks*, which are a special kind of user-defined command. Whenever you run the command ‘*foo*’, if the user-defined command ‘*hook-foo*’ exists, it is executed (with no arguments) before that command.

A hook may also be defined which is run after the command you executed. Whenever you run the command ‘*foo*’, if the user-defined command ‘*hookpost-foo*’ exists, it is executed (with no arguments) after that command. Post-execution hooks may exist simultaneously with pre-execution hooks, for the same command.

It is valid for a hook to call the command which it hooks. If this occurs, the hook is not re-executed, thereby avoiding infinite recursion.

In addition, a pseudo-command, ‘*stop*’ exists. Defining (‘*hook-stop*’) makes the associated commands execute every time execution stops in your program: before breakpoint commands are run, displays are printed, or the stack frame is printed.

For example, to ignore SIGALRM signals while single-stepping, but treat them normally during normal execution, you could define:

```
define hook-stop
handle SIGALRM nopass
end

define hook-run
handle SIGALRM pass
end

define hook-continue
handle SIGALRM pass
end
```

As a further example, to hook at the begining and end of the `echo` command, and to add extra text to the beginning and end of the message, you could define:

```
define hook-echo
echo <<<---
end

define hookpost-echo
echo --->>>\n
end

(gdb) echo Hello World
<<<---Hello World--->>>
(gdb)
```

You can define a hook for any single-word command in GDB, but not for command aliases; you should define a hook for the basic command name, e.g. `backtrace` rather than `bt`. If an error occurs during the execution of your hook, execution of GDB commands stops and GDB issues a prompt (before the command that you actually typed had a chance to run).

If you try to define a hook which does not match any known command, you get a warning from the `define` command.

20.3 Command files

A command file for GDB is a file of lines that are GDB commands. Comments (lines starting with `#`) may also be included. An empty line in a command file does nothing; it does not mean to repeat the last command, as it would from the terminal.

When you start GDB, it automatically executes commands from its *init files*, normally called ‘`.gdbinit`’¹. During startup, GDB does the following:

1. Reads the init file (if any) in your home directory².
2. Processes command line options and operands.
3. Reads the init file (if any) in the current working directory.
4. Reads command files specified by the ‘`-x`’ option.

The init file in your home directory can set options (such as ‘`set complaints`’) that affect subsequent processing of command line options and operands. Init files are not executed if you use the ‘`-nx`’ option (see [Section 2.1.2 \[Choosing modes\], page 13](#)).

On some configurations of GDB, the init file is known by a different name (these are typically environments where a specialized form of GDB may need to coexist with other forms, hence a different name for the specialized version’s init file). These are the environments with special init file names:

- VxWorks (Wind River Systems real-time OS): ‘`.vxgdbinit`’
- OS68K (Enea Data Systems real-time OS): ‘`.os68gdbinit`’
- ES-1800 (Ericsson Telecom AB M68000 emulator): ‘`.esgdbinit`’

¹ The DJGPP port of GDB uses the name ‘`gdb.ini`’ instead, due to the limitations of file names imposed by DOS filesystems.

² On DOS/Windows systems, the home directory is the one pointed to by the `HOME` environment variable.

You can also request the execution of a command file with the `source` command:

```
source filename
    Execute the command file filename.
```

The lines in a command file are executed sequentially. They are not printed as they are executed. An error in any command terminates execution of the command file and control is returned to the console.

Commands that would ask for confirmation if used interactively proceed without asking when used in a command file. Many GDB commands that normally print messages to say what they are doing omit the messages when called from command files.

GDB also accepts command input from standard input. In this mode, normal output goes to standard output and error output goes to standard error. Errors in a command file supplied on standard input do not terminate execution of the command file — execution continues with the next command.

```
gdb < cmds > log 2>&1
```

(The syntax above will vary depending on the shell used.) This example will execute commands from the file ‘`cmds`’. All output and errors would be directed to ‘`log`’.

20.4 Commands for controlled output

During the execution of a command file or a user-defined command, normal GDB output is suppressed; the only output that appears is what is explicitly printed by the commands in the definition. This section describes three commands useful for generating exactly the output you want.

```
echo text
```

Print *text*. Nonprinting characters can be included in *text* using C escape sequences, such as ‘\n’ to print a newline. **No newline is printed unless you specify one.** In addition to the standard C escape sequences, a backslash followed by a space stands for a space. This is useful for displaying a string with spaces at the beginning or the end, since leading and trailing spaces are otherwise trimmed from all arguments. To print ‘ and foo = ’, use the command ‘`echo \ and foo = \`’.

A backslash at the end of *text* can be used, as in C, to continue the command onto subsequent lines. For example,

```
echo This is some text\n\
which is continued\n\
onto several lines.\n
```

produces the same output as

```
echo This is some text\n
echo which is continued\n
echo onto several lines.\n
```

```
output expression
```

Print the value of *expression* and nothing but that value: no newlines, no ‘\$nn = ’. The value is not entered in the value history either. See [Section 8.1 \[Expressions\]](#), page 65, for more information on expressions.

output/*fmt* *expression*

Print the value of *expression* in format *fmt*. You can use the same formats as for `print`. See [Section 8.4 \[Output formats\]](#), page 68, for more information.

`printf` *string*, *expressions*...

Print the values of the *expressions* under the control of *string*. The *expressions* are separated by commas and may be either numbers or pointers. Their values are printed as specified by *string*, exactly as if your program were to execute the C subroutine

```
printf (string, expressions...);
```

For example, you can print two values in hex like this:

```
printf "foo, bar-foo = 0x%x, 0x%x\n", foo, bar-foo
```

The only backslash-escape sequences that you can use in the format string are the simple ones that consist of backslash followed by a letter.

21 Command Interpreters

GDB supports multiple command interpreters, and some command infrastructure to allow users or user interface writers to switch between interpreters or run commands in other interpreters.

GDB currently supports two command interpreters, the console interpreter (sometimes called the command-line interpreter or `CLI`) and the machine interface interpreter (or `GDB/MI`). This manual describes both of these interfaces in great detail.

By default, GDB will start with the console interpreter. However, the user may choose to start GDB with another interpreter by specifying the '`-i`' or '`--interpreter`' startup options. Defined interpreters include:

- console** The traditional console or command-line interpreter. This is the most often used interpreter with GDB. With no interpreter specified at runtime, GDB will use this interpreter.
- mi** The newest GDB/MI interface (currently `mi2`). Used primarily by programs wishing to use GDB as a backend for a debugger GUI or an IDE. For more information, see [Chapter 24 \[The GDB/MI Interface\], page 203](#).
- mi2** The current GDB/MI interface.
- mi1** The GDB/MI interface included in GDB 5.1, 5.2, and 5.3.

The interpreter being used by GDB may not be dynamically switched at runtime. Although possible, this could lead to a very precarious situation. Consider an IDE using GDB/MI. If a user enters the command "`interpreter-set console`" in a console view, GDB would switch to using the console interpreter, rendering the IDE inoperable!

Although you may only choose a single interpreter at startup, you may execute commands in any interpreter from the current interpreter using the appropriate command. If you are running the console interpreter, simply use the `interpreter-exec` command:

```
interpreter-exec mi "-data-list-register-names"
```

GDB/MI has a similar command, although it is only available in versions of GDB which support GDB/MI version 2 (or greater).

22 GDB Text User Interface

The GDB Text User Interface, TUI in short, is a terminal interface which uses the `curses` library to show the source file, the assembly output, the program registers and GDB commands in separate text windows.

The TUI is enabled by invoking GDB using either ‘`gdbtui`’ or ‘`gdb -tui`’.

22.1 TUI overview

The TUI has two display modes that can be switched while GDB runs:

- A curses (or TUI) mode in which it displays several text windows on the terminal.
- A standard mode which corresponds to the GDB configured without the TUI.

In the TUI mode, GDB can display several text window on the terminal:

command This window is the GDB command window with the GDB prompt and the GDB outputs. The GDB input is still managed using readline but through the TUI. The *command* window is always visible.

source The source window shows the source file of the program. The current line as well as active breakpoints are displayed in this window.

assembly The assembly window shows the disassembly output of the program.

register This window shows the processor registers. It detects when a register is changed and when this is the case, registers that have changed are highlighted.

The source and assembly windows show the current program position by highlighting the current line and marking them with the ‘>’ marker. Breakpoints are also indicated with two markers. A first one indicates the breakpoint type:

- | | |
|---|--|
| B | Breakpoint which was hit at least once. |
| b | Breakpoint which was never hit. |
| H | Hardware breakpoint which was hit at least once. |
| h | Hardware breakpoint which was never hit. |

The second marker indicates whether the breakpoint is enabled or not:

- | | |
|---|-------------------------|
| + | Breakpoint is enabled. |
| - | Breakpoint is disabled. |

The source, assembly and register windows are attached to the thread and the frame position. They are updated when the current thread changes, when the frame changes or when the program counter changes. These three windows are arranged by the TUI according to several layouts. The layout defines which of these three windows are visible. The following layouts are available:

- `source`
- `assembly`
- `source and assembly`

- source and registers
- assembly and registers

On top of the command window a status line gives various information concerning the current process being debugged. The status line is updated when the information it shows changes. The following fields are displayed:

<i>target</i>	Indicates the current gdb target (see Chapter 16 [Specifying a Debugging Target], page 147).
<i>process</i>	Gives information about the current process or thread number. When no process is being debugged, this field is set to <code>No process</code> .
<i>function</i>	Gives the current function name for the selected frame. The name is demangled if demangling is turned on (see Section 8.7 [Print Settings], page 71). When there is no symbol corresponding to the current program counter the string <code>??</code> is displayed.
<i>line</i>	Indicates the current line number for the selected frame. When the current line number is not known the string <code>??</code> is displayed.
<i>pc</i>	Indicates the current program counter address.

22.2 TUI Key Bindings

The TUI installs several key bindings in the readline keymaps (see [Chapter 27 \[Command Line Editing\], page 261](#)). They allow to leave or enter in the TUI mode or they operate directly on the TUI layout and windows. The TUI also provides a *SingleKey* keymap which binds several keys directly to GDB commands. The following key bindings are installed for both TUI mode and the GDB standard mode.

C-x C-a

C-x a

C-x A Enter or leave the TUI mode. When the TUI mode is left, the curses window management is left and GDB operates using its standard mode writing on the terminal directly. When the TUI mode is entered, the control is given back to the curses windows. The screen is then refreshed.

C-x 1 Use a TUI layout with only one window. The layout will either be ‘source’ or ‘assembly’. When the TUI mode is not active, it will switch to the TUI mode. Think of this key binding as the Emacs *C-x 1* binding.

C-x 2 Use a TUI layout with at least two windows. When the current layout shows already two windows, a next layout with two windows is used. When a new layout is chosen, one window will always be common to the previous layout and the new one.

Think of it as the Emacs *C-x 2* binding.

C-x o Change the active window. The TUI associates several key bindings (like scrolling and arrow keys) to the active window. This command gives the focus to the next TUI window.

Think of it as the Emacs *C-x o* binding.

C-x s Use the TUI *SingleKey* keymap that binds single key to gdb commands (see [Section 22.3 \[TUI Single Key Mode\], page 197](#)).

The following key bindings are handled only by the TUI mode:

PgUp	Scroll the active window one page up.
PgDn	Scroll the active window one page down.
Up	Scroll the active window one line up.
Down	Scroll the active window one line down.
Left	Scroll the active window one column left.
Right	Scroll the active window one column right.
C-L	Refresh the screen.

In the TUI mode, the arrow keys are used by the active window for scrolling. This means they are available for readline when the active window is the command window. When the command window does not have the focus, it is necessary to use other readline key bindings such as **C-p**, **C-n**, **C-b** and **C-f**.

22.3 TUI Single Key Mode

The TUI provides a *SingleKey* mode in which it installs a particular key binding in the readline keymaps to connect single keys to some gdb commands.

c	continue
d	down
f	finish
n	next
q	exit the <i>SingleKey</i> mode.
r	run
s	step
u	up
v	info locals
w	where

Other keys temporarily switch to the GDB command prompt. The key that was pressed is inserted in the editing buffer so that it is possible to type most GDB commands without interaction with the TUI *SingleKey* mode. Once the command is entered the TUI *SingleKey* mode is restored. The only way to permanently leave this mode is by hitting **Q** or ‘**C-x S**’.

22.4 TUI specific commands

The TUI has specific commands to control the text windows. These commands are always available, that is they do not depend on the current terminal mode in which GDB runs. When GDB is in the standard mode, using these commands will automatically switch in the TUI mode.

- info win** List and give the size of all displayed windows.
- layout next**
Display the next layout.
- layout prev**
Display the previous layout.
- layout src**
Display the source window only.
- layout asm**
Display the assembly window only.
- layout split**
Display the source and assembly window.
- layout regs**
Display the register window together with the source or assembly window.
- focus next | prev | src | asm | regs | split**
Set the focus to the named window. This command allows to change the active window so that scrolling keys can be affected to another window.
- refresh** Refresh the screen. This is similar to using `\C-L` key.
- tui reg float**
Show the floating point registers in the register window.
- tui reg general**
Show the general registers in the register window.
- tui reg next**
Show the next register group. The list of register groups as well as their order is target specific. The predefined register groups are the following: `general`, `float`, `system`, `vector`, `all`, `save`, `restore`.
- tui reg system**
Show the system registers in the register window.
- update** Update the source window and the current execution point.
- winheight name +count**
winheight name -count
Change the height of the window `name` by `count` lines. Positive counts increase the height, while negative counts decrease it.

22.5 TUI configuration variables

The TUI has several configuration variables that control the appearance of windows on the terminal.

`set tui border-kind kind`

Select the border appearance for the source, assembly and register windows.
The possible values are the following:

`space` Use a space character to draw the border.

`ascii` Use ascii characters + - and | to draw the border.

`acs` Use the Alternate Character Set to draw the border. The border is drawn using character line graphics if the terminal supports them.

`set tui active-border-mode mode`

Select the attributes to display the border of the active window. The possible values are `normal`, `standout`, `reverse`, `half`, `half-standout`, `bold` and `bold-standout`.

`set tui border-mode mode`

Select the attributes to display the border of other windows. The *mode* can be one of the following:

`normal` Use normal attributes to display the border.

`standout` Use standout mode.

`reverse` Use reverse video mode.

`half` Use half bright mode.

`half-standout`

Use half bright and standout mode.

`bold` Use extra bright or bold mode.

`bold-standout`

Use extra bright or bold and standout mode.

23 Using GDB under GNU Emacs

A special interface allows you to use GNU Emacs to view (and edit) the source files for the program you are debugging with GDB.

To use this interface, use the command `M-x gdb` in Emacs. Give the executable file you want to debug as an argument. This command starts GDB as a subprocess of Emacs, with input and output through a newly created Emacs buffer.

Using GDB under Emacs is just like using GDB normally except for two things:

- All “terminal” input and output goes through the Emacs buffer.

This applies both to GDB commands and their output, and to the input and output done by the program you are debugging.

This is useful because it means that you can copy the text of previous commands and input them again; you can even use parts of the output in this way.

All the facilities of Emacs’ Shell mode are available for interacting with your program. In particular, you can send signals the usual way—for example, `C-c C-c` for an interrupt, `C-c C-z` for a stop.

- GDB displays source code through Emacs.

Each time GDB displays a stack frame, Emacs automatically finds the source file for that frame and puts an arrow (‘=>’) at the left margin of the current line. Emacs uses a separate buffer for source display, and splits the screen to show both your GDB session and the source.

Explicit GDB `list` or `search` commands still produce output as usual, but you probably have no reason to use them from Emacs.

If you specify an absolute file name when prompted for the `M-x gdb` argument, then Emacs sets your current working directory to where your program resides. If you only specify the file name, then Emacs sets your current working directory to the directory associated with the previous buffer. In this case, GDB may find your program by searching your environment’s PATH variable, but on some operating systems it might not find the source. So, although the GDB input and output session proceeds normally, the auxiliary buffer does not display the current source and line of execution.

The initial working directory of GDB is printed on the top line of the GDB I/O buffer and this serves as a default for the commands that specify files for GDB to operate on. See [Section 15.1 \[Commands to specify files\]](#), page 135.

By default, `M-x gdb` calls the program called ‘`gdb`’. If you need to call GDB by a different name (for example, if you keep several configurations around, with different names) you can customize the Emacs variable `gud-gdb-command-name` to run the one you want.

In the GDB I/O buffer, you can use these special Emacs commands in addition to the standard Shell mode commands:

`C-h m` Describe the features of Emacs’ GDB Mode.

`C-c C-s` Execute to another source line, like the GDB `step` command; also update the display window to show the current file and location.

- C-c C-n* Execute to next source line in this function, skipping all function calls, like the GDB `next` command. Then update the display window to show the current file and location.
- C-c C-i* Execute one instruction, like the GDB `stepi` command; update display window accordingly.
- C-c C-f* Execute until exit from the selected stack frame, like the GDB `finish` command.
- C-c C-r* Continue execution of your program, like the GDB `continue` command.
- C-c <* Go up the number of frames indicated by the numeric argument (see section “[Numeric Arguments](#)” in *The GNU Emacs Manual*), like the GDB `up` command.
- C-c >* Go down the number of frames indicated by the numeric argument, like the GDB `down` command.

In any source file, the Emacs command *C-x SPC* (`gud-break`) tells GDB to set a breakpoint on the source line point is on.

If you type *M-x speedbar*, then Emacs displays a separate frame which shows a backtrace when the GDB I/O buffer is current. Move point to any frame in the stack and type *<RET>* to make it become the current frame and display the associated source in the source buffer. Alternatively, click *Mouse-2* to make the selected frame become the current one.

If you accidentally delete the source-display buffer, an easy way to get it back is to type the command `f` in the GDB buffer, to request a frame display; when you run under Emacs, this recreates the source buffer if necessary to show you the context of the current frame.

The source files displayed in Emacs are in ordinary Emacs buffers which are visiting the source files in the usual way. You can edit the files with these buffers if you wish; but keep in mind that GDB communicates with Emacs in terms of line numbers. If you add or delete lines from the text, the line numbers that GDB knows cease to correspond properly with the code.

The description given here is for GNU Emacs version 21.3 and a more detailed description of its interaction with GDB is given in the Emacs manual (see section “[Debuggers](#)” in *The GNU Emacs Manual*).

24 The GDB/MI Interface

Function and Purpose

GDB/MI is a line based machine oriented text interface to GDB. It is specifically intended to support the development of systems which use the debugger as just one small component of a larger system.

This chapter is a specification of the GDB/MI interface. It is written in the form of a reference manual.

Note that GDB/MI is still under construction, so some of the features described below are incomplete and subject to change.

Notation and Terminology

This chapter uses the following notation:

- | separates two alternatives.
- [*something*] indicates that *something* is optional: it may or may not be given.
- (*group*)* means that *group* inside the parentheses may repeat zero or more times.
- (*group*)+ means that *group* inside the parentheses may repeat one or more times.
- "string" means a literal string.

Acknowledgments

In alphabetic order: Andrew Cagney, Fernando Nasser, Stan Shebs and Elena Zannoni.

24.3 GDB/MI Command Syntax

24.3.1 GDB/MI Input Syntax

```

command ↪
    cli-command | mi-command

cli-command ↪
    [ token ] cli-command nl, where cli-command is any existing GDB CLI command.

mi-command ↪
    [ token ] "--" operation ( " " option )* [ " --" ] ( " " parameter )* nl

token ↪ "any sequence of digits"

option ↪
    "--" parameter [ " " parameter ]

parameter ↪
    non-blank-sequence | c-string

operation ↪
    any of the operations described in this chapter

```

```

non-blank-sequence  $\mapsto$ 
    anything, provided it doesn't contain special characters such as "-", nl, """
    and of course " "
c-string  $\mapsto$ 
    """ seven-bit-iso-c-string-content """
nl  $\mapsto$  CR | CR-LF

```

Notes:

- The CLI commands are still handled by the MI interpreter; their output is described below.
- The *token*, when present, is passed back when the command finishes.
- Some MI commands accept optional arguments as part of the parameter list. Each option is identified by a leading ‘-’ (dash) and may be followed by an optional argument parameter. Options occur first in the parameter list and can be delimited from normal parameters using ‘--’ (this is useful when some parameters begin with a dash).

Pragmatics:

- We want easy access to the existing CLI syntax (for debugging).
- We want it to be easy to spot a MI operation.

24.3.2 GDB/MI Output Syntax

The output from GDB/MI consists of zero or more out-of-band records followed, optionally, by a single result record. This result record is for the most recent command. The sequence of output records is terminated by '(gdb)'.

If an input command was prefixed with a *token* then the corresponding output for that command will also be prefixed by that same *token*.

```

output  $\mapsto$ 
    ( out-of-band-record )* [ result-record ] "(gdb)" nl
result-record  $\mapsto$ 
    [ token ] "^" result-class ( "," result )* nl
out-of-band-record  $\mapsto$ 
    async-record | stream-record
async-record  $\mapsto$ 
    exec-async-output | status-async-output | notify-async-output
exec-async-output  $\mapsto$ 
    [ token ] "*" async-output
status-async-output  $\mapsto$ 
    [ token ] "+" async-output
notify-async-output  $\mapsto$ 
    [ token ] "=" async-output
async-output  $\mapsto$ 
    async-class ( "," result )* nl

```

```

result-class ↪
    "done" | "running" | "connected" | "error" | "exit"

async-class ↪
    "stopped" | others (where others will be added depending on the needs—this
    is still in development).

result ↪
    variable "=" value

variable ↪
    string

value ↪  const | tuple | list

const ↪  c-string

tuple ↪  "{}" | "{$ result ( , " result )* }"

list ↪  "[]" | "[ value ( , " value )* ]" | "[ result ( , " result )* ]"

stream-record ↪
    console-stream-output | target-stream-output | log-stream-output

console-stream-output ↪
    "~" c-string

target-stream-output ↪
    "@" c-string

log-stream-output ↪
    "&" c-string

nl ↪  CR | CR-LF

token ↪  any sequence of digits.

```

Notes:

- All output sequences end in a single line containing a period.
- The *token* is from the corresponding request. If an execution command is interrupted by the ‘-exec-interrupt’ command, the *token* associated with the ‘*stopped’ message is the one of the original execution command, not the one of the interrupt command.
- *status-async-output* contains on-going status information about the progress of a slow operation. It can be discarded. All status output is prefixed by ‘+’.
- *exec-async-output* contains asynchronous state change on the target (stopped, started, disappeared). All async output is prefixed by ‘*’.
- *notify-async-output* contains supplementary information that the client should handle (e.g., a new breakpoint information). All notify output is prefixed by ‘=’.
- *console-stream-output* is output that should be displayed as is in the console. It is the textual response to a CLI command. All the console output is prefixed by ‘~’.
- *target-stream-output* is the output produced by the target program. All the target output is prefixed by ‘@’.

- *log-stream-output* is output text coming from GDB’s internals, for instance messages that should be displayed as part of an error log. All the log output is prefixed by ‘&’.
- New GDB/MI commands should only output *lists* containing *values*.

See [Section 24.5.2 \[GDB/MI Stream Records\]](#), page 207, for more details about the various output records.

24.3.3 Simple Examples of GDB/MI Interaction

This subsection presents several simple examples of interaction using the GDB/MI interface. In these examples, ‘->’ means that the following line is passed to GDB/MI as input, while ‘<-’ means the output received from GDB/MI.

Target Stop

Here’s an example of stopping the inferior process:

```
-> -stop
<- (gdb)
```

and later:

```
<- *stop,reason="stop",address="0x123",source="a.c:123"
<- (gdb)
```

Simple CLI Command

Here’s an example of a simple CLI command being passed through GDB/MI and on to the CLI.

```
-> print 1+2
<- &"print 1+2\n"
<- ~"$1 = 3\n"
<- ^done
<- (gdb)
```

Command With Side Effects

```
-> -symbol-file xyz.exe
<- *breakpoint,nr="3",address="0x123",source="a.c:123"
<- (gdb)
```

A Bad Command

Here’s what happens if you pass a non-existent command:

```
-> -rubbish
<- ^error,msg="Undefined MI command: rubbish"
<- (gdb)
```

24.4 GDB/MI Compatibility with CLI

To help users familiar with GDB’s existing CLI interface, GDB/MI accepts existing CLI commands. As specified by the syntax, such commands can be directly entered into the GDB/MI interface and GDB will respond.

This mechanism is provided as an aid to developers of GDB/MI clients and not as a reliable interface into the CLI. Since the command is being interpreted in an environment that assumes GDB/MI behaviour, the exact output of such commands is likely to end up being an un-supported hybrid of GDB/MI and CLI output.

24.5 GDB/MI Output Records

24.5.1 GDB/MI Result Records

In addition to a number of out-of-band notifications, the response to a GDB/MI command includes one of the following result indications:

"^done" [", " *results*]

The synchronous operation was successful, *results* are the return values.

"^running"

The asynchronous operation was successfully started. The target is running.

"^error" ", " *c-string*

The operation failed. The *c-string* contains the corresponding error message.

24.5.2 GDB/MI Stream Records

GDB internally maintains a number of output streams: the console, the target, and the log. The output intended for each of these streams is funneled through the GDB/MI interface using *stream records*.

Each stream record begins with a unique *prefix character* which identifies its stream (see [Section 24.3.2 \[GDB/MI Output Syntax\], page 204](#)). In addition to the prefix, each stream record contains a *string-output*. This is either raw text (with an implicit new line) or a quoted C string (which does not contain an implicit newline).

"~" *string-output*

The console output stream contains text that should be displayed in the CLI console window. It contains the textual responses to CLI commands.

"@" *string-output*

The target output stream contains any textual output from the running target.

"&" *string-output*

The log stream contains debugging messages being produced by GDB's internals.

24.5.3 GDB/MI Out-of-band Records

Out-of-band records are used to notify the GDB/MI client of additional changes that have occurred. Those changes can either be a consequence of GDB/MI (e.g., a breakpoint modified) or a result of target activity (e.g., target stopped).

The following is a preliminary list of possible out-of-band records.

"*" "stop"

24.6 GDB/MI Command Description Format

The remaining sections describe blocks of commands. Each block of commands is laid out in a fashion similar to this section.

Note the the line breaks shown in the examples are here only for readability. They don't appear in the real output. Also note that the commands with a non-available example (N.A.) are not yet implemented.

Motivation

The motivation for this collection of commands.

Introduction

A brief introduction to this collection of commands as a whole.

Commands

For each command in the block, the following is described:

Synopsis

```
-command args...
```

GDB Command

The corresponding GDB CLI command.

Result

Out-of-band

Notes

Example

24.7 GDB/MI Breakpoint table commands

This section documents GDB/MI commands for manipulating breakpoints.

The `-break-after` Command

Synopsis

```
-break-after number count
```

The breakpoint number *number* is not in effect until it has been hit *count* times. To see how this is reflected in the output of the ‘`-break-list`’ command, see the description of the ‘`-break-list`’ command below.

GDB Command

The corresponding GDB command is ‘`ignore`’.

Example

```
(gdb)
-break-insert main
^done,bkpt={number="1",addr="0x000100d0",file="hello.c",line="5"}
(gdb)
-break-after 1 3
^done
(gdb)
-break-list
^done,BreakpointTable={nr_rows="1",nr_cols="6",
```

```

hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},  

{width="14",alignment="-1",col_name="type",colhdr="Type"},  

{width="4",alignment="-1",col_name="disp",colhdr="Disp"},  

{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},  

{width="10",alignment="-1",col_name="addr",colhdr="Address"},  

{width="40",alignment="2",col_name="what",colhdr="What"}],  

body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",  

addr="0x000100d0",func="main",file="hello.c",line="5",times="0",  

ignore="3"}]]}  

(gdb)

```

The **-break-condition** Command

Synopsis

```
-break-condition number expr
```

Breakpoint *number* will stop the program only if the condition in *expr* is true. The condition becomes part of the ‘-break-list’ output (see the description of the ‘-break-list’ command below).

GDB Command

The corresponding GDB command is ‘condition’.

Example

```

(gdb)
-break-condition 1 1
^done
(gdb)
-break-list
^done,BreakpointTable={nr_rows="1",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},  

{width="14",alignment="-1",col_name="type",colhdr="Type"},  

{width="4",alignment="-1",col_name="disp",colhdr="Disp"},  

{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},  

{width="10",alignment="-1",col_name="addr",colhdr="Address"},  

{width="40",alignment="2",col_name="what",colhdr="What"}],  

body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",  

addr="0x000100d0",func="main",file="hello.c",line="5",cond="1",  

times="0",ignore="3"}]]}  

(gdb)

```

The **-break-delete** Command

Synopsis

```
-break-delete ( breakpoint )+
```

Delete the breakpoint(s) whose number(s) are specified in the argument list. This is obviously reflected in the breakpoint list.

GDB command

The corresponding GDB command is ‘delete’.

Example

```

(gdb)
-break-delete 1

```

```

^done
(gdb)
-break-list
^done,BreakpointTable={nr_rows="0",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},{width="14",alignment="-1",col_name="type",colhdr="Type"},{width="4",alignment="-1",col_name="disp",colhdr="Disp"},{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},{width="10",alignment="-1",col_name="addr",colhdr="Address"},{width="40",alignment="2",col_name="what",colhdr="What"}],body=[]}
(gdb)

```

The `-break-disable` Command

Synopsis

```
-break-disable ( breakpoint )+
```

Disable the named *breakpoint*(s). The field ‘enabled’ in the break list is now set to ‘n’ for the named *breakpoint*(s).

GDB Command

The corresponding GDB command is ‘`disable`’.

Example

```

(gdb)
-break-disable 2
^done
(gdb)
-break-list
^done,BreakpointTable={nr_rows="1",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},{width="14",alignment="-1",col_name="type",colhdr="Type"},{width="4",alignment="-1",col_name="disp",colhdr="Disp"},{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},{width="10",alignment="-1",col_name="addr",colhdr="Address"},{width="40",alignment="2",col_name="what",colhdr="What"}],body=[bkpt={number="2",type="breakpoint",disp="keep",enabled="n",addr="0x000100d0",func="main",file="hello.c",line="5",times="0"}]}
(gdb)

```

The `-break-enable` Command

Synopsis

```
-break-enable ( breakpoint )+
```

Enable (previously disabled) *breakpoint*(s).

GDB Command

The corresponding GDB command is ‘`enable`’.

Example

```

(gdb)
-break-enable 2
^done

```

```
(gdb)
-break-list
^done,BreakpointTable={nr_rows="1",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"}, {width="14",alignment="-1",col_name="type",colhdr="Type"}, {width="4",alignment="-1",col_name="disp",colhdr="Disp"}, {width="3",alignment="-1",col_name="enabled",colhdr="Enb"}, {width="10",alignment="-1",col_name="addr",colhdr="Address"}, {width="40",alignment="2",col_name="what",colhdr="What"}], body=[bkpt={number="2",type="breakpoint",disp="keep",enabled="y", addr="0x000100d0",func="main",file="hello.c",line="5",times="0"}]}
(gdb)
```

The -break-info Command

Synopsis

```
-break-info breakpoint
```

Get information about a single breakpoint.

GDB command

The corresponding GDB command is ‘info break *breakpoint*’.

Example

N.A.

The -break-insert Command

Synopsis

```
-break-insert [ -t ] [ -h ] [ -r ]
[ -c condition ] [ -i ignore-count ]
[ -p thread ] [ line | addr ]
```

If specified, *line*, can be one of:

- function
- filename:linenum
- filename:function
- *address

The possible optional parameters of this command are:

‘-t’ Insert a temporary breakpoint.

‘-h’ Insert a hardware breakpoint.

‘-c *condition*’

Make the breakpoint conditional on *condition*.

‘-i *ignore-count*’

Initialize the *ignore-count*.

‘-r’ Insert a regular breakpoint in all the functions whose names match the given regular expression. Other flags are not applicable to regular expression.

Result

The result is in the form:

```
^done,bkptno="number",func="funcname",
  file="filename",line="lineno"
```

where *number* is the GDB number for this breakpoint, *funcname* is the name of the function where the breakpoint was inserted, *filename* is the name of the source file which contains this function, and *lineno* is the source line number within that file.

Note: this format is open to change.

GDB Command

The corresponding GDB commands are ‘break’, ‘tbreak’, ‘hbreak’, ‘thbreak’, and ‘rbreak’.

Example

```
(gdb)
-break-insert main
^done,bkpt={number="1",addr="0x0001072c",file="recursive2.c",line="4"}
(gdb)
-break-insert -t foo
^done,bkpt={number="2",addr="0x00010774",file="recursive2.c",line="11"}
(gdb)
-break-list
^done,BreakpointTable={nr_rows="2",nr_cols="6",
  hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},{width="14",alignment="-1",col_name="type",colhdr="Type"},{width="4",alignment="-1",col_name="disp",colhdr="Disp"},{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},{width="10",alignment="-1",col_name="addr",colhdr="Address"},{width="40",alignment="2",col_name="what",colhdr="What"}],body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",addr="0x0001072c", func="main",file="recursive2.c",line="4",times="0"},bkpt={number="2",type="breakpoint",disp="del",enabled="y",addr="0x00010774",func="foo",file="recursive2.c",line="11",times="0"}]}
(gdb)
-break-insert -r foo.*
^int foo(int, int);
^done,bkpt={number="3",addr="0x00010774",file="recursive2.c",line="11"}
(gdb)
```

The -break-list Command

Synopsis

```
-break-list
```

Displays the list of inserted breakpoints, showing the following fields:

‘Number’ number of the breakpoint

‘Type’ type of the breakpoint: ‘breakpoint’ or ‘watchpoint’

‘Disposition’

should the breakpoint be deleted or disabled when it is hit: ‘keep’ or ‘nokeep’

‘Enabled’ is the breakpoint enabled or no: ‘y’ or ‘n’

‘Address’ memory location at which the breakpoint is set

‘What’ logical location of the breakpoint, expressed by function name, file name, line number

‘Times’ number of times the breakpoint has been hit

If there are no breakpoints or watchpoints, the `BreakpointTable` body field is an empty list.

GDB Command

The corresponding GDB command is ‘`info break`’.

Example

```
(gdb)
-break-list
^done,BreakpointTable={nr_rows="2",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},{width="14",alignment="-1",col_name="type",colhdr="Type"},{width="4",alignment="-1",col_name="disp",colhdr="Disp"},{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},{width="10",alignment="-1",col_name="addr",colhdr="Address"},{width="40",alignment="2",col_name="what",colhdr="What"}],body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",addr="0x000100d0",func="main",file="hello.c",line="5",times="0"},bkpt={number="2",type="breakpoint",disp="keep",enabled="y",addr="0x00010114",func="foo",file="hello.c",line="13",times="0"}]}
(gdb)
```

Here’s an example of the result when there are no breakpoints:

```
(gdb)
-break-list
^done,BreakpointTable={nr_rows="0",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},{width="14",alignment="-1",col_name="type",colhdr="Type"},{width="4",alignment="-1",col_name="disp",colhdr="Disp"},{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},{width="10",alignment="-1",col_name="addr",colhdr="Address"},{width="40",alignment="2",col_name="what",colhdr="What"}],body=[]}
(gdb)
```

The `-break-watch` Command

Synopsis

```
-break-watch [ -a | -r ]
```

Create a watchpoint. With the ‘`-a`’ option it will create an access watchpoint, i.e. a watchpoint that triggers either on a read from or on a write to the memory location. With the ‘`-r`’ option, the watchpoint created is a *read* watchpoint, i.e. it will trigger only when the memory location is accessed for reading. Without either of the options, the watchpoint created is a regular watchpoint, i.e. it will trigger when the memory location is accessed for writing. See [Section 5.1.2 \[Setting watchpoints\], page 37](#).

Note that ‘`-break-list`’ will report a single list of watchpoints and breakpoints inserted.

GDB Command

The corresponding GDB commands are ‘`watch`’, ‘`awatch`’, and ‘`rwatch`’.

Example

Setting a watchpoint on a variable in the `main` function:

```
(gdb)
-break-watch x
^done,wpt={number="2",exp="x"}
(gdb)
-exec-continue
^running
^done,reason="watchpoint-trigger",wpt={number="2",exp="x"}, 
value={old="-268439212",new="55"}, 
frame={func="main",args=[],file="recursive2.c",line="5"}
(gdb)
```

Setting a watchpoint on a variable local to a function. GDB will stop the program execution twice: first for the variable changing value, then for the watchpoint going out of scope.

```
(gdb)
-break-watch C
^done,wpt={number="5",exp="C"}
(gdb)
-exec-continue
^running
^done,reason="watchpoint-trigger",
wpt={number="5",exp="C"},value={old="-276895068",new="3"}, 
frame={func="callee4",args=[], 
file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="13"}
(gdb)
-exec-continue
^running
^done,reason="watchpoint-scope",wpnum="5",
frame={func="callee3",args=[{name="strarg",
value="0x11940 \"A string argument.\""}], 
file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="18"}
(gdb)
```

Listing breakpoints and watchpoints, at different points in the program execution. Note that once the watchpoint goes out of scope, it is deleted.

```
(gdb)
-break-watch C
^done,wpt={number="2",exp="C"}
(gdb)
-break-list
^done,BreakpointTable={nr_rows="2",nr_cols="6",
hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"}, 
{width="14",alignment="-1",col_name="type",colhdr="Type"}, 
{width="4",alignment="-1",col_name="disp",colhdr="Disp"}, 
{width="3",alignment="-1",col_name="enabled",colhdr="Enb"}, 
{width="10",alignment="-1",col_name="addr",colhdr="Address"}, 
{width="40",alignment="2",col_name="what",colhdr="What"}], 
body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",
addr="0x00010734",func="callee4",
file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8",times="1"}, 
bkpt={number="2",type="watchpoint",disp="keep",
enabled="y",addr="",what="C",times="0"}]}
(gdb)
-exec-continue
^running
```

```

^done,reason="watchpoint-trigger",wpt={number="2",exp="C"},  

value={old="-276895068",new="3"},  

frame={func="callee4",args=[]},  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="13"}  

(gdb)  

-break-list  

^done,BreakpointTable={nr_rows="2",nr_cols="6",  

hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},  

{width="14",alignment="-1",col_name="type",colhdr="Type"},  

{width="4",alignment="-1",col_name="disp",colhdr="Disp"},  

{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},  

{width="10",alignment="-1",col_name="addr",colhdr="Address"},  

{width="40",alignment="2",col_name="what",colhdr="What"}],  

body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",  

addr="0x00010734",func="callee4",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8",times="1"},  

bkpt={number="2",type="watchpoint",disp="keep",  

enabled="y",addr="",what="C",times="-5"}]]}  

(gdb)  

-exec-continue  

^running  

^done,reason="watchpoint-scope",wpnum="2",  

frame={func="callee3",args=[{name="strarg",  

value="0x11940 \"A string argument.\\""}],  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="18"}  

(gdb)  

-break-list  

^done,BreakpointTable={nr_rows="1",nr_cols="6",  

hdr=[{width="3",alignment="-1",col_name="number",colhdr="Num"},  

{width="14",alignment="-1",col_name="type",colhdr="Type"},  

{width="4",alignment="-1",col_name="disp",colhdr="Disp"},  

{width="3",alignment="-1",col_name="enabled",colhdr="Enb"},  

{width="10",alignment="-1",col_name="addr",colhdr="Address"},  

{width="40",alignment="2",col_name="what",colhdr="What"}],  

body=[bkpt={number="1",type="breakpoint",disp="keep",enabled="y",  

addr="0x00010734",func="callee4",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8",times="1"}]}  

(gdb)

```

24.8 GDB/MI Data Manipulation

This section describes the GDB/MI commands that manipulate data: examine memory and registers, evaluate expressions, etc.

The -data-disassemble Command

Synopsis

```

-data-disassemble
[ -s start-addr -e end-addr ]
| [ -f filename -l linenum [ -n lines ] ]
-- mode

```

Where:

`'start-addr'`

is the beginning address (or \$pc)

- ‘*end-addr*’ is the end address
- ‘*filename*’ is the name of the file to disassemble
- ‘*linenum*’ is the line number to disassemble around
- ‘*lines*’ is the the number of disassembly lines to be produced. If it is -1, the whole function will be disassembled, in case no *end-addr* is specified. If *end-addr* is specified as a non-zero value, and *lines* is lower than the number of disassembly lines between *start-addr* and *end-addr*, only *lines* lines are displayed; if *lines* is higher than the number of lines between *start-addr* and *end-addr*, only the lines up to *end-addr* are displayed.
- ‘*mode*’ is either 0 (meaning only disassembly) or 1 (meaning mixed source and disassembly).

Result

The output for each instruction is composed of four fields:

- Address
- Func-name
- Offset
- Instruction

Note that whatever included in the instruction field, is not manipulated directly by GDB/MI, i.e. it is not possible to adjust its format.

GDB Command

There's no direct mapping from this command to the CLI.

Example

Disassemble from the current value of \$pc to \$pc + 20:

```
(gdb)
-data-disassemble -s $pc -e "$pc + 20" -- 0
^done,
asm_insns=[
{address="0x000107c0",func-name="main",offset="4",
inst="mov 2, %o0"}, 
{address="0x000107c4",func-name="main",offset="8",
inst="sethi %hi(0x11800), %o2"}, 
{address="0x000107c8",func-name="main",offset="12",
inst="or %o2, 0x140, %o1\t! 0x11940 <_lib_version+8>"}, 
{address="0x000107cc",func-name="main",offset="16",
inst="sethi %hi(0x11800), %o2"}, 
{address="0x000107d0",func-name="main",offset="20",
inst="or %o2, 0x168, %o4\t! 0x11968 <_lib_version+48>"}
(gdb)
```

Disassemble the whole `main` function. Line 32 is part of `main`.

```
-data-disassemble -f basics.c -l 32 -- 0
^done,asm_insns=[
```

```
{address="0x000107bc",func-name="main",offset="0",
inst="save %sp, -112, %sp"},  

{address="0x000107c0",func-name="main",offset="4",
inst="mov 2, %o0"},  

{address="0x000107c4",func-name="main",offset="8",
inst="sethi %hi(0x11800), %o2"},  

[...]  

{address="0x0001081c",func-name="main",offset="96",inst="ret "},  

{address="0x00010820",func-name="main",offset="100",inst="restore "}]  

(gdb)
```

Disassemble 3 instructions from the start of `main`:

```
(gdb)
-data-disassemble -f basics.c -l 32 -n 3 -- 0
^done,asm_insns=[  

{address="0x000107bc",func-name="main",offset="0",
inst="save %sp, -112, %sp"},  

{address="0x000107c0",func-name="main",offset="4",
inst="mov 2, %o0"},  

{address="0x000107c4",func-name="main",offset="8",
inst="sethi %hi(0x11800), %o2"}]
(gdb)
```

Disassemble 3 instructions from the start of `main` in mixed mode:

```
(gdb)
-data-disassemble -f basics.c -l 32 -n 3 -- 1
^done,asm_insns=[  

src_and_asm_line={line="31",
file="/kwikemart/marge/ezannoni/flathead-dev/devo/gdb/ \
 testsuite/gdb.mi/basics.c"},line_asm_insn=[  

{address="0x000107bc",func-name="main",offset="0",
inst="save %sp, -112, %sp"}]},  

src_and_asm_line={line="32",
file="/kwikemart/marge/ezannoni/flathead-dev/devo/gdb/ \
 testsuite/gdb.mi/basics.c"},line_asm_insn=[  

{address="0x000107c0",func-name="main",offset="4",
inst="mov 2, %o0"},  

{address="0x000107c4",func-name="main",offset="8",
inst="sethi %hi(0x11800), %o2"}}]
(gdb)
```

The `-data-evaluate-expression` Command

Synopsis

```
-data-evaluate-expression expr
```

Evaluate `expr` as an expression. The expression could contain an inferior function call. The function call will execute synchronously. If the expression contains spaces, it must be enclosed in double quotes.

GDB Command

The corresponding GDB commands are ‘`print`’, ‘`output`’, and ‘`call`’. In `gdbtk` only, there’s a corresponding ‘`gdb_eval`’ command.

Example

In the following example, the numbers that precede the commands are the *tokens* described in [Section 24.3 \[GDB/MI Command Syntax\]](#), page 203. Notice how GDB/MI returns the same tokens in its output.

```
211-data-evaluate-expression A
211^done,value="1"
(gdb)
311-data-evaluate-expression &A
311^done,value="0xeffffeb7c"
(gdb)
411-data-evaluate-expression A+3
411^done,value="4"
(gdb)
511-data-evaluate-expression "A + 3"
511^done,value="4"
(gdb)
```

The -data-list-changed-registers Command

Synopsis

```
-data-list-changed-registers
```

Display a list of the registers that have changed.

GDB Command

GDB doesn't have a direct analog for this command; gdbtk has the corresponding command ‘`gdb_changed_register_list`’.

Example

On a PPC MBX board:

```
(gdb)
-exec-continue
^running

(gdb)
*stopped,reason="breakpoint-hit",bkptno="1",frame={func="main",
args=[],file="try.c",line="5"}
(gdb)
-data-list-changed-registers
^done,changed-registers=[ "0", "1", "2", "4", "5", "6", "7", "8", "9",
"10", "11", "13", "14", "15", "16", "17", "18", "19", "20", "21", "22", "23",
"24", "25", "26", "27", "28", "30", "31", "64", "65", "66", "67", "69"]
(gdb)
```

The -data-list-register-names Command

Synopsis

```
-data-list-register-names [ ( regno )+ ]
```

Show a list of register names for the current target. If no arguments are given, it shows a list of the names of all the registers. If integer numbers are given as arguments, it will print a list of the names of the registers corresponding to the arguments. To ensure consistency between a register name and its number, the output list may include empty register names.

GDB Command

GDB does not have a command which corresponds to ‘-data-list-register-names’. In gdbtk there is a corresponding command ‘gdb_regnames’.

Example

For the PPC MBX board:

```
(gdb)
-data-list-register-names
^done,register-names=[r0","r1","r2","r3","r4","r5","r6","r7",
"r8","r9","r10","r11","r12","r13","r14","r15","r16","r17","r18",
"r19","r20","r21","r22","r23","r24","r25","r26","r27","r28","r29",
"r30","r31","f0","f1","f2","f3","f4","f5","f6","f7","f8","f9",
"f10","f11","f12","f13","f14","f15","f16","f17","f18","f19","f20",
"f21","f22","f23","f24","f25","f26","f27","f28","f29","f30","f31",
"", "pc","ps","cr","lr","ctr","xer"]
(gdb)
-data-list-register-names 1 2 3
^done,register-names=[r1","r2","r3"]
(gdb)
```

The -data-list-register-values Command

Synopsis

```
-data-list-register-values fmt [ ( regno )* ]
```

Display the registers’ contents. *fmt* is the format according to which the registers’ contents are to be returned, followed by an optional list of numbers specifying the registers to display. A missing list of numbers indicates that the contents of all the registers must be returned.

Allowed formats for *fmt* are:

x	Hexadecimal
o	Octal
t	Binary
d	Decimal
r	Raw
N	Natural

GDB Command

The corresponding GDB commands are ‘info reg’, ‘info all-reg’, and (in gdbtk) ‘gdb_fetch_registers’.

Example

For a PPC MBX board (note: line breaks are for readability only, they don’t appear in the actual output):

```
(gdb)
-data-list-register-values r 64 65
^done,register-values=[{number="64",value="0xfe00a300"},{number="65",value="0x00029002"}]
```

```
(gdb)
-data-list-register-values x
^done,register-values=[{number="0",value="0xfe0043c8"}, {number="1",value="0x3fff88"}, {number="2",value="0xfffffffffe"}, {number="3",value="0x0"}, {number="4",value="0xa"}, {number="5",value="0x3fff68"}, {number="6",value="0x3fff58"}, {number="7",value="0xfe011e98"}, {number="8",value="0x2"}, {number="9",value="0xfa202820"}, {number="10",value="0xfa202808"}, {number="11",value="0x1"}, {number="12",value="0x0"}, {number="13",value="0x4544"}, {number="14",value="0xffffdfffff"}, {number="15",value="0xffffffff"}, {number="16",value="0xfffffeff"}, {number="17",value="0xfffffffed"}, {number="18",value="0xfffffffef"}, {number="19",value="0xfffffffef"}, {number="20",value="0xfffffffef"}, {number="21",value="0xfffffffef"}, {number="22",value="0xfffffff7"}, {number="23",value="0xfffffffef"}, {number="24",value="0xfffffffef"}, {number="25",value="0xfffffffef"}, {number="26",value="0xfffffffefb"}, {number="27",value="0xfffffffef"}, {number="28",value="0xf7bfffff"}, {number="29",value="0x0"}, {number="30",value="0xfe010000"}, {number="31",value="0x0"}, {number="32",value="0x0"}, {number="33",value="0x0"}, {number="34",value="0x0"}, {number="35",value="0x0"}, {number="36",value="0x0"}, {number="37",value="0x0"}, {number="38",value="0x0"}, {number="39",value="0x0"}, {number="40",value="0x0"}, {number="41",value="0x0"}, {number="42",value="0x0"}, {number="43",value="0x0"}, {number="44",value="0x0"}, {number="45",value="0x0"}, {number="46",value="0x0"}, {number="47",value="0x0"}, {number="48",value="0x0"}, {number="49",value="0x0"}, {number="50",value="0x0"}, {number="51",value="0x0"}, {number="52",value="0x0"}, {number="53",value="0x0"}, {number="54",value="0x0"}, {number="55",value="0x0"}, {number="56",value="0x0"}, {number="57",value="0x0"}, {number="58",value="0x0"}, {number="59",value="0x0"}, {number="60",value="0x0"}, {number="61",value="0x0"}, {number="62",value="0x0"}, {number="63",value="0x0"}, {number="64",value="0xfe00a300"}, {number="65",value="0x29002"}, {number="66",value="0x202f04b5"}, {number="67",value="0xfe0043b0"}, {number="68",value="0xfe00b3e4"}, {number="69",value="0x20002b03"}]
(gdb)
```

The -data-read-memory Command

Synopsis

```
-data-read-memory [ -o byte-offset ]
    address word-format word-size
    nr-rows nr-cols [ aschar ]
```

where:

'address' An expression specifying the address of the first memory word to be read. Complex expressions containing embedded white space should be quoted using the C convention.

'word-format'

The format to be used to print the memory words. The notation is the same as for GDB's print command (see [Section 8.4 \[Output formats\], page 68](#)).

'word-size'

The size of each memory word in bytes.

'nr-rows' The number of rows in the output table.**'nr-cols'** The number of columns in the output table.**'aschar'** If present, indicates that each row should include an ASCII dump. The value of aschar is used as a padding character when a byte is not a member of the printable ASCII character set (printable ASCII characters are those whose code is between 32 and 126, inclusively).**'byte-offset'**

An offset to add to the address before fetching memory.

This command displays memory contents as a table of *nr-rows* by *nr-cols* words, each word being *word-size* bytes. In total, *nr-rows* * *nr-cols* * *word-size* bytes are read (returned as ‘total-bytes’). Should less than the requested number of bytes be returned by the target, the missing words are identified using ‘N/A’. The number of bytes read from the target is returned in ‘nr-bytes’ and the starting address used to read memory in ‘addr’.

The address of the next/previous row or page is available in ‘next-row’ and ‘prev-row’, ‘next-page’ and ‘prev-page’.

GDB Command

The corresponding GDB command is ‘x’. gdbtk has ‘gdb_get_mem’ memory read command.

Example

Read six bytes of memory starting at *bytes+6* but then offset by -6 bytes. Format as three rows of two columns. One byte per word. Display each word in hex.

```
(gdb)
9-data-read-memory -o -6 -- bytes+6 x 1 3 2
9^done,addr="0x00001390",nr-bytes="6",total-bytes="6",
next-row="0x00001396",prev-row="0x0000138e",next-page="0x00001396",
prev-page="0x0000138a",memory=[
{addr="0x00001390",data=["0x00","0x01"]},
{addr="0x00001392",data=["0x02","0x03"]},
{addr="0x00001394",data=["0x04","0x05"]}]

(gdb)
```

Read two bytes of memory starting at address *shorts + 64* and display as a single word formatted in decimal.

```
(gdb)
5-data-read-memory shorts+64 d 2 1 1
5^done,addr="0x00001510",nr-bytes="2",total-bytes="2",
next-row="0x00001512",prev-row="0x0000150e",
next-page="0x00001512",prev-page="0x0000150e",memory=[
{addr="0x00001510",data=["128"]}]

(gdb)
```

Read thirty two bytes of memory starting at *bytes+16* and format as eight rows of four columns. Include a string encoding with ‘x’ used as the non-printable character.

```
(gdb)
4-data-read-memory bytes+16 x 1 8 4 x
4^done,addr="0x000013a0",nr-bytes="32",total-bytes="32",
```

```

next-row="0x000013c0",prev-row="0x0000139c",
next-page="0x000013c0",prev-page="0x00001380",memory=[
{addr="0x000013a0",data=["0x10","0x11","0x12","0x13"],ascii="xxxx"},  

{addr="0x000013a4",data=["0x14","0x15","0x16","0x17"],ascii="xxxx"},  

{addr="0x000013a8",data=["0x18","0x19","0x1a","0x1b"],ascii="xxxx"},  

{addr="0x000013ac",data=["0x1c","0x1d","0x1e","0x1f"],ascii="xxxx"},  

{addr="0x000013b0",data=["0x20","0x21","0x22","0x23"],ascii=" !\#"},  

{addr="0x000013b4",data=["0x24","0x25","0x26","0x27"],ascii="$%&'},  

{addr="0x000013b8",data=["0x28","0x29","0x2a","0x2b"],ascii="()**+"},  

{addr="0x000013bc",data=["0x2c","0x2d","0x2e","0x2f"],ascii="-./"}]
(gdb)

```

The -display-delete Command

Synopsis

`-display-delete number`

Delete the display *number*.

GDB Command

The corresponding GDB command is ‘`delete display`’.

Example

N.A.

The -display-disable Command

Synopsis

`-display-disable number`

Disable display *number*.

GDB Command

The corresponding GDB command is ‘`disable display`’.

Example

N.A.

The -display-enable Command

Synopsis

`-display-enable number`

Enable display *number*.

GDB Command

The corresponding GDB command is ‘`enable display`’.

Example

N.A.

The -display-insert Command

Synopsis

```
-display-insert expression
```

Display *expression* every time the program stops.

GDB Command

The corresponding GDB command is ‘display’.

Example

N.A.

The -display-list Command

Synopsis

```
-display-list
```

List the displays. Do not show the current values.

GDB Command

The corresponding GDB command is ‘info display’.

Example

N.A.

The -environment-cd Command

Synopsis

```
-environment-cd pathdir
```

Set GDB’s working directory.

GDB Command

The corresponding GDB command is ‘cd’.

Example

```
(gdb)
>environment-cd /kwikemart/marge/ezannoni/flathead-dev/devo/gdb
^done
(gdb)
```

The -environment-directory Command

Synopsis

```
-environment-directory [ -r ] [ pathdir ]+
```

Add directories *pathdir* to beginning of search path for source files. If the ‘-r’ option is used, the search path is reset to the default search path. If directories *pathdir* are supplied in addition to the ‘-r’ option, the search path is first reset and then addition occurs as normal. Multiple directories may be specified, separated by blanks. Specifying multiple directories in a single command results in the directories added to the beginning of the search path in the same order they were presented in the command. If blanks are needed as

part of a directory name, double-quotes should be used around the name. In the command output, the path will show up separated by the system directory-separator character. The directory-separator character must not be used in any directory name. If no directories are specified, the current search path is displayed.

GDB Command

The corresponding GDB command is ‘dir’.

Example

```
(gdb)
-environment-directory /kwikemart/marge/ezannoni/flathead-dev/devo/gdb
^done,source-path="/kwikemart/marge/ezannoni/flathead-dev/devo/gdb:$cdir:$cwd"
(gdb)
-environment-directory ""
^done,source-path="/kwikemart/marge/ezannoni/flathead-dev/devo/gdb:$cdir:$cwd"
(gdb)
-environment-directory -r /home/jjohnstn/src/gdb /usr/src
^done,source-path="/home/jjohnstn/src/gdb:/usr/src:$cdir:$cwd"
(gdb)
-environment-directory -r
^done,source-path="$cdir:$cwd"
(gdb)
```

The -environment-path Command

Synopsis

```
-environment-path [ -r ] [ pathdir ]+
```

Add directories *pathdir* to beginning of search path for object files. If the ‘-r’ option is used, the search path is reset to the original search path that existed at gdb start-up. If directories *pathdir* are supplied in addition to the ‘-r’ option, the search path is first reset and then addition occurs as normal. Multiple directories may be specified, separated by blanks. Specifying multiple directories in a single command results in the directories added to the beginning of the search path in the same order they were presented in the command. If blanks are needed as part of a directory name, double-quotes should be used around the name. In the command output, the path will show up separated by the system directory-separator character. The directory-separator character must not be used in any directory name. If no directories are specified, the current path is displayed.

GDB Command

The corresponding GDB command is ‘path’.

Example

```
(gdb)
-environment-path
^done,path="/usr/bin"
(gdb)
-environment-path /kwikemart/marge/ezannoni/flathead-dev/ppc-eabi/gdb /bin
^done,path="/kwikemart/marge/ezannoni/flathead-dev/ppc-eabi/gdb:/bin:/usr/bin"
(gdb)
-environment-path -r /usr/local/bin
^done,path="/usr/local/bin:/usr/bin"
(gdb)
```

The `-environment-pwd` Command

Synopsis

```
-environment-pwd
```

Show the current working directory.

GDB command

The corresponding GDB command is ‘`pwd`’.

Example

```
(gdb)
-environment-pwd
^done,cwd="/kwikemart/marge/ezannoni/flathead-dev/devo/gdb"
(gdb)
```

24.9 GDB/MI Program control

Program termination

As a result of execution, the inferior program can run to completion, if it doesn’t encounter any breakpoints. In this case the output will include an exit code, if the program has exited exceptionally.

Examples

Program exited normally:

```
(gdb)
-exec-run
^running
(gdb)
x = 55
*stopped,reason="exited-normally"
(gdb)
```

Program exited exceptionally:

```
(gdb)
-exec-run
^running
(gdb)
x = 55
*stopped,reason="exited",exit-code="01"
(gdb)
```

Another way the program can terminate is if it receives a signal such as `SIGINT`. In this case, GDB/MI displays this:

```
(gdb)
*stopped,reason="exited-signalled",signal-name="SIGINT",
signal-meaning="Interrupt"
```

The `-exec-abort` Command

Synopsis

```
-exec-abort
```

Kill the inferior running program.

GDB Command

The corresponding GDB command is ‘kill’.

Example

N.A.

The -exec-arguments Command**Synopsis**

```
-exec-arguments args
```

Set the inferior program arguments, to be used in the next ‘-exec-run’.

GDB Command

The corresponding GDB command is ‘set args’.

Example

Don’t have one around.

The -exec-continue Command**Synopsis**

```
-exec-continue
```

Asynchronous command. Resumes the execution of the inferior program until a breakpoint is encountered, or until the inferior exits.

GDB Command

The corresponding GDB command is ‘continue’.

Example

```
-exec-continue
^running
(gdb)
@Hello world
*stopped,reason="breakpoint-hit",bkptno="2",frame={func="foo",args=[],file="hello.c",line="13"}
(gdb)
```

The -exec-finish Command**Synopsis**

```
-exec-finish
```

Asynchronous command. Resumes the execution of the inferior program until the current function is exited. Displays the results returned by the function.

GDB Command

The corresponding GDB command is ‘finish’.

Example

Function returning `void`.

```
-exec-finish
^running
(gdb)
@Hello from foo
*stopped,reason="function-finished",frame={func="main",args=[],file="hello.c",line="7"}
(gdb)
```

Function returning other than `void`. The name of the internal GDB variable storing the result is printed, together with the value itself.

```
-exec-finish
^running
(gdb)
*stopped,reason="function-finished",frame={addr="0x000107b0",func="foo",args=[{name="a",value="1"},{name="b",value="9"}],file="recursive2.c",line="14"},gdb-result-var="$1",return-value="0"
(gdb)
```

The `-exec-interrupt` Command

Synopsis

```
-exec-interrupt
```

Asynchronous command. Interrupts the background execution of the target. Note how the token associated with the stop message is the one for the execution command that has been interrupted. The token for the interrupt itself only appears in the ‘`^done`’ output. If the user is trying to interrupt a non-running program, an error message will be printed.

GDB Command

The corresponding GDB command is ‘`interrupt`’.

Example

```
(gdb)
111-exec-continue
111^running

(gdb)
222-exec-interrupt
222^done
(gdb)
111*stopped,signal-name="SIGINT",signal-meaning="Interrupt",
frame={addr="0x00010140",func="foo",args=[],file="try.c",line="13"}
(gdb)

(gdb)
-exec-interrupt
^error,msg="mi_cmd_exec_interrupt: Inferior not executing."
(gdb)
```

The `-exec-next` Command

Synopsis

```
-exec-next
```

Asynchronous command. Resumes execution of the inferior program, stopping when the beginning of the next source line is reached.

GDB Command

The corresponding GDB command is ‘`next`’.

Example

```
-exec-next
^running
(gdb)
*stopped,reason="end-stepping-range",line="8",file="hello.c"
(gdb)
```

The `-exec-next-instruction` Command

Synopsis

```
-exec-next-instruction
```

Asynchronous command. Executes one machine instruction. If the instruction is a function call continues until the function returns. If the program stops at an instruction in the middle of a source line, the address will be printed as well.

GDB Command

The corresponding GDB command is ‘`nexti`’.

Example

```
(gdb)
-exec-next-instruction
^running

(gdb)
*stopped,reason="end-stepping-range",
addr="0x000100d4",line="5",file="hello.c"
(gdb)
```

The `-exec-return` Command

Synopsis

```
-exec-return
```

Makes current function return immediately. Doesn’t execute the inferior. Displays the new current frame.

GDB Command

The corresponding GDB command is ‘`return`’.

Example

```
(gdb)
200-break-insert callee4
200^done,bkpt={number="1",addr="0x00010734",
```

```

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8"}
(gdb)
000-exec-run
000^running
(gdb)
000*stopped,reason="breakpoint-hit",bkptno="1",
frame={func="callee4",args=[]},
file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8"}
(gdb)
205-break-delete
205^done
(gdb)
111-exec-return
111^done,frame={level="0",func="callee3",
args=[{name="strarg",
value="0x11940 \"A string argument.\""}],
file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="18"}
(gdb)

```

The -exec-run Command

Synopsis

`-exec-run`

Asynchronous command. Starts execution of the inferior from the beginning. The inferior executes until either a breakpoint is encountered or the program exits.

GDB Command

The corresponding GDB command is ‘run’.

Example

```

(gdb)
-break-insert main
^done,bkpt={number="1",addr="0x0001072c",file="recursive2.c",line="4"}
(gdb)
-exec-run
^running
(gdb)
*stopped,reason="breakpoint-hit",bkptno="1",
frame={func="main",args=[],file="recursive2.c",line="4"}
(gdb)

```

The -exec-show-arguments Command

Synopsis

`-exec-show-arguments`

Print the arguments of the program.

GDB Command

The corresponding GDB command is ‘show args’.

Example

N.A.

The `-exec-step` Command

Synopsis

```
-exec-step
```

Asynchronous command. Resumes execution of the inferior program, stopping when the beginning of the next source line is reached, if the next source line is not a function call. If it is, stop at the first instruction of the called function.

GDB Command

The corresponding GDB command is ‘`step`’.

Example

Stepping into a function:

```
-exec-step
^running
(gdb)
*stopped,reason="end-stepping-range",
frame={func="foo",args=[{name="a",value="10"},{name="b",value="0"}],file="recursive2.c",line="11"}
(gdb)
```

Regular stepping:

```
-exec-step
^running
(gdb)
*stopped,reason="end-stepping-range",line="14",file="recursive2.c"
(gdb)
```

The `-exec-step-instruction` Command

Synopsis

```
-exec-step-instruction
```

Asynchronous command. Resumes the inferior which executes one machine instruction. The output, once GDB has stopped, will vary depending on whether we have stopped in the middle of a source line or not. In the former case, the address at which the program stopped will be printed as well.

GDB Command

The corresponding GDB command is ‘`stepi`’.

Example

```
(gdb)
-exec-step-instruction
^running

(gdb)
*stopped,reason="end-stepping-range",
frame={func="foo",args=[],file="try.c",line="10"}
(gdb)
-exec-step-instruction
^running
```

```
(gdb)
*stopped,reason="end-stepping-range",
frame={addr="0x000100f4",func="foo",args=[],file="try.c",line="10"}
(gdb)
```

The `-exec-until` Command

Synopsis

```
-exec-until [ location ]
```

Asynchronous command. Executes the inferior until the *location* specified in the argument is reached. If there is no argument, the inferior executes until a source line greater than the current one is reached. The reason for stopping in this case will be ‘*location-reached*’.

GDB Command

The corresponding GDB command is ‘`until`’.

Example

```
(gdb)
-exec-until recursive2.c:6
^running
(gdb)
x = 55
*stopped,reason="location-reached",frame={func="main",args=[],file="recursive2.c",line="6"}
(gdb)
```

The `-file-exec-and-symbols` Command

Synopsis

```
-file-exec-and-symbols file
```

Specify the executable file to be debugged. This file is the one from which the symbol table is also read. If no file is specified, the command clears the executable and symbol information. If breakpoints are set when using this command with no arguments, GDB will produce error messages. Otherwise, no output is produced, except a completion notification.

GDB Command

The corresponding GDB command is ‘`file`’.

Example

```
(gdb)
-file-exec-and-symbols /kwikemart/marge/ezannoni/TRUNK/mbx/hello.mbx
^done
(gdb)
```

The `-file-exec-file` Command

Synopsis

```
-file-exec-file file
```

Specify the executable file to be debugged. Unlike ‘`-file-exec-and-symbols`’, the symbol table is *not* read from this file. If used without argument, GDB clears the information about the executable file. No output is produced, except a completion notification.

GDB Command

The corresponding GDB command is ‘exec-file’.

Example

```
(gdb)
-file-exec-file /kwikemart/marge/ezannoni/TRUNK/mbx/hello.mbx
^done
(gdb)
```

The -file-list-exec-sections Command

Synopsis

```
-file-list-exec-sections
```

List the sections of the current executable file.

GDB Command

The GDB command ‘info file’ shows, among the rest, the same information as this command. gdbtk has a corresponding command ‘gdb_load_info’.

Example

N.A.

The -file-list-exec-source-file Command

Synopsis

```
-file-list-exec-source-file
```

List the line number, the current source file, and the absolute path to the current source file for the current executable.

GDB Command

There’s no GDB command which directly corresponds to this one.

Example

```
(gdb)
123-file-list-exec-source-file
123^done,line="1",file="foo.c",fullname="/home/bar/foo.c"
(gdb)
```

The -file-list-exec-source-files Command

Synopsis

```
-file-list-exec-source-files
```

List the source files for the current executable.

It will always output the filename, but only when GDB can find the absolute file name of a source file, will it output the fullname.

GDB Command

There’s no GDB command which directly corresponds to this one. gdbtk has an analogous command ‘gdb_listfiles’.

Example

```
(gdb)
-file-list-exec-source-files
^done,files=[
{file=foo.c,fullname=/home/foo.c},
{file=/home/bar.c,fullname=/home/bar.c},
{file=gdb_could_not_find_fullpath.c}]
(gdb)
```

The -file-list-shared-libraries Command**Synopsis**

```
-file-list-shared-libraries
```

List the shared libraries in the program.

GDB Command

The corresponding GDB command is ‘info shared’.

Example

N.A.

The -file-list-symbol-files Command**Synopsis**

```
-file-list-symbol-files
```

List symbol files.

GDB Command

The corresponding GDB command is ‘info file’ (part of it).

Example

N.A.

The -file-symbol-file Command**Synopsis**

```
-file-symbol-file file
```

Read symbol table info from the specified *file* argument. When used without arguments, clears GDB’s symbol table info. No output is produced, except for a completion notification.

GDB Command

The corresponding GDB command is ‘symbol-file’.

Example

```
(gdb)
-file-symbol-file /kwikemart/marge/ezannoni/TRUNK/mbx/hello.mbx
^done
(gdb)
```

24.10 Miscellaneous GDB commands in GDB/MI

The `-gdb-exit` Command

Synopsis

```
-gdb-exit
```

Exit GDB immediately.

GDB Command

Approximately corresponds to ‘quit’.

Example

```
(gdb)
-gdb-exit
```

The `-gdb-set` Command

Synopsis

```
-gdb-set
```

Set an internal GDB variable.

GDB Command

The corresponding GDB command is ‘set’.

Example

```
(gdb)
-gdb-set $foo=3
^done
(gdb)
```

The `-gdb-show` Command

Synopsis

```
-gdb-show
```

Show the current value of a GDB variable.

GDB command

The corresponding GDB command is ‘show’.

Example

```
(gdb)
-gdb-show annotate
^done,value="0"
(gdb)
```

The `-gdb-version` Command

Synopsis

```
-gdb-version
```

Show version information for GDB. Used mostly in testing.

GDB Command

There's no equivalent GDB command. GDB by default shows this information when you start an interactive session.

Example

```
(gdb)
-gdb-version
~GNU gdb 5.2.1
~Copyright 2000 Free Software Foundation, Inc.
~GDB is free software, covered by the GNU General Public License, and
~you are welcome to change it and/or distribute copies of it under
~certain conditions.
~Type "show copying" to see the conditions.
~There is absolutely no warranty for GDB. Type "show warranty" for
~details.
~This GDB was configured as
"--host=sparc-sun-solaris2.5.1 --target=ppc-eabi".
^done
(gdb)
```

The -interpreter-exec Command

Synopsis

```
-interpreter-exec interpreter command
```

Execute the specified *command* in the given *interpreter*.

GDB Command

The corresponding GDB command is ‘`interpreter-exec`’.

Example

```
(gdb)
-interpreter-exec console "break main"
&"During symbol reading, couldn't parse type; debugger out of date?.\n"
&"During symbol reading, bad structure-type format.\n"
~"Breakpoint 1 at 0x8074fc6: file ../../src/gdb/main.c, line 743.\n"
^done
(gdb)
```

24.11 GDB/MI Stack Manipulation Commands

The -stack-info-frame Command

Synopsis

```
-stack-info-frame
```

Get info on the current frame.

GDB Command

The corresponding GDB command is ‘`info frame`’ or ‘`frame`’ (without arguments).

Example

N.A.

The `-stack-info-depth` Command

Synopsis

```
-stack-info-depth [ max-depth ]
```

Return the depth of the stack. If the integer argument *max-depth* is specified, do not count beyond *max-depth* frames.

GDB Command

There's no equivalent GDB command.

Example

For a stack with frame levels 0 through 11:

```
(gdb)
-stack-info-depth
^done,depth="12"
(gdb)
-stack-info-depth 4
^done,depth="4"
(gdb)
-stack-info-depth 12
^done,depth="12"
(gdb)
-stack-info-depth 11
^done,depth="11"
(gdb)
-stack-info-depth 13
^done,depth="12"
(gdb)
```

The `-stack-list-arguments` Command

Synopsis

```
-stack-list-arguments show-values
[ low-frame high-frame ]
```

Display a list of the arguments for the frames between *low-frame* and *high-frame* (inclusive). If *low-frame* and *high-frame* are not provided, list the arguments for the whole call stack.

The *show-values* argument must have a value of 0 or 1. A value of 0 means that only the names of the arguments are listed, a value of 1 means that both names and values of the arguments are printed.

GDB Command

GDB does not have an equivalent command. gdbtk has a ‘`gdb_get_args`’ command which partially overlaps with the functionality of ‘`-stack-list-arguments`’.

Example

```
(gdb)
-stack-list-frames
^done,
stack=[
frame={level="0",addr="0x00010734",func="callee4",
```

```

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="8"},  

frame={level="1",addr="0x0001076c",func="callee3",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="17"},  

frame={level="2",addr="0x0001078c",func="callee2",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="22"},  

frame={level="3",addr="0x000107b4",func="callee1",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="27"},  

frame={level="4",addr="0x000107e0",func="main",  

file="../../../../devo/gdb/testsuite/gdb.mi/basics.c",line="32"}]  

(gdb)  

-stack-list-arguments 0  

^done,  

stack-args=[  

frame={level="0",args=[]},  

frame={level="1",args=[name="strarg"]},  

frame={level="2",args=[name="intarg",name="strarg"]},  

frame={level="3",args=[name="intarg",name="strarg",name="fltarg"]},  

frame={level="4",args=[]}]  

(gdb)  

-stack-list-arguments 1  

^done,  

stack-args=[  

frame={level="0",args=[]},  

frame={level="1",  

args=[{name="strarg",value="0x11940 \"A string argument.\\""}]},  

frame={level="2",args=[  

{name="intarg",value="2"},  

{name="strarg",value="0x11940 \"A string argument.\\""}]},  

frame={level="3",args=[  

{name="intarg",value="2"},  

{name="strarg",value="0x11940 \"A string argument.\\""},  

{name="fltarg",value="3.5"}]},  

frame={level="4",args=[]}]  

(gdb)  

-stack-list-arguments 0 2 2  

^done,stack-args=[frame={level="2",args=[name="intarg",name="strarg"]}]]  

(gdb)  

-stack-list-arguments 1 2 2  

^done,stack-args=[frame={level="2",  

args=[{name="intarg",value="2"},  

{name="strarg",value="0x11940 \"A string argument.\\""}]}]]  

(gdb)

```

The `-stack-list-frames` Command

Synopsis

```
-stack-list-frames [ low-frame high-frame ]
```

List the frames currently on the stack. For each frame it displays the following info:

- ‘level’ The frame number, 0 being the topmost frame, i.e. the innermost function.
- ‘addr’ The \$pc value for that frame.
- ‘func’ Function name.
- ‘file’ File name of the source file where the function lives.
- ‘line’ Line number corresponding to the \$pc.

If invoked without arguments, this command prints a backtrace for the whole stack. If given two integer arguments, it shows the frames whose levels are between the two arguments (inclusive). If the two arguments are equal, it shows the single frame at the corresponding level.

GDB Command

The corresponding GDB commands are ‘`backtrace`’ and ‘`where`’.

Example

Full stack backtrace:

```
(gdb)
-stack-list-frames
^done,stack=
[frame={level="0",addr="0x0001076c",func="foo",
  file="recursive2.c",line="11"},,
frame={level="1",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="2",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="3",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="4",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="5",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="6",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="7",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="8",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="9",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="10",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="11",addr="0x00010738",func="main",
  file="recursive2.c",line="4"}]
(gdb)
```

Show frames between `low_frame` and `high_frame`:

```
(gdb)
-stack-list-frames 3 5
^done,stack=
[frame={level="3",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="4",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"},,
frame={level="5",addr="0x000107a4",func="foo",
  file="recursive2.c",line="14"}]
(gdb)
```

Show a single frame:

```
(gdb)
-stack-list-frames 3 3
^done,stack=
[frame={level="3",addr="0x000107a4",func="foo",
```

```
    file="recursive2.c",line="14"}]
(gdb)
```

The -stack-list-locals Command

Synopsis

```
-stack-list-locals print-values
```

Display the local variable names for the current frame. With an argument of 0 or `--no-values`, prints only the names of the variables. With argument of 1 or `--all-values`, prints also their values. With argument of 2 or `--simple-values`, prints the name, type and value for simple data types and the name and type for arrays, structures and unions. In this last case, the idea is that the user can see the value of simple data types immediately and he can create variable objects for other data types if he wishes to explore their values in more detail.

GDB Command

`'info locals'` in GDB, `'gdb_get_locals'` in gdbtk.

Example

```
(gdb)
-stack-list-locals 0
^done,locals=[name="A",name="B",name="C"]
(gdb)
-stack-list-locals --all-values
^done,locals=[{name="A",value="1"},{name="B",value="2"},
  {name="C",value="{1, 2, 3}"}]
-stack-list-locals --simple-values
^done,locals=[{name="A",type="int",value="1"},
  {name="B",type="int",value="2"},{name="C",type="int [3]"}]
(gdb)
```

The -stack-select-frame Command

Synopsis

```
-stack-select-frame framenum
```

Change the current frame. Select a different frame `framenum` on the stack.

GDB Command

The corresponding GDB commands are `'frame'`, `'up'`, `'down'`, `'select-frame'`, `'up-silent'`, and `'down-silent'`.

Example

```
(gdb)
-stack-select-frame 2
^done
(gdb)
```

24.12 GDB/MI Symbol Query Commands

The -symbol-info-address Command

Synopsis

```
-symbol-info-address symbol
```

Describe where *symbol* is stored.

GDB Command

The corresponding GDB command is ‘info address’.

Example

N.A.

The -symbol-info-file Command

Synopsis

```
-symbol-info-file
```

Show the file for the symbol.

GDB Command

There’s no equivalent GDB command. gdbtk has ‘gdb_find_file’.

Example

N.A.

The -symbol-info-function Command

Synopsis

```
-symbol-info-function
```

Show which function the symbol lives in.

GDB Command

‘gdb_get_function’ in gdbtk.

Example

N.A.

The -symbol-info-line Command

Synopsis

```
-symbol-info-line
```

Show the core addresses of the code for a source line.

GDB Command

The corresponding GDB command is ‘info line’. gdbtk has the ‘gdb_get_line’ and ‘gdb_get_file’ commands.

Example

N.A.

The `-symbol-info-symbol` Command

Synopsis

```
-symbol-info-symbol addr
```

Describe what symbol is at location `addr`.

GDB Command

The corresponding GDB command is ‘`info symbol`’.

Example

N.A.

The `-symbol-list-functions` Command

Synopsis

```
-symbol-list-functions
```

List the functions in the executable.

GDB Command

‘`info functions`’ in GDB, ‘`gdb_listfunc`’ and ‘`gdb_search`’ in `gdbtk`.

Example

N.A.

The `-symbol-list-lines` Command

Synopsis

```
-symbol-list-lines filename
```

Print the list of lines that contain code and their associated program addresses for the given source filename. The entries are sorted in ascending PC order.

GDB Command

There is no corresponding GDB command.

Example

```
(gdb)
-symbol-list-lines basics.c
^done,lines=[{pc="0x08048554",line="7"},{pc="0x0804855a",line="8"}]
(gdb)
```

The `-symbol-list-types` Command

Synopsis

```
-symbol-list-types
```

List all the type names.

GDB Command

The corresponding commands are ‘`info types`’ in GDB, ‘`gdb_search`’ in `gdbtk`.

Example

N.A.

The -symbol-list-variables Command

Synopsis

`-symbol-list-variables`

List all the global and static variable names.

GDB Command

‘info variables’ in GDB, ‘gdb_search’ in gdbtk.

Example

N.A.

The -symbol-locate Command

Synopsis

`-symbol-locate`

GDB Command

‘gdb_loc’ in gdbtk.

Example

N.A.

The -symbol-type Command

Synopsis

`-symbol-type variable`

Show type of *variable*.

GDB Command

The corresponding GDB command is ‘ptype’, gdbtk has ‘gdb_obj_variable’.

Example

N.A.

24.13 GDB/MI Target Manipulation Commands

The -target-attach Command

Synopsis

`-target-attach pid | file`

Attach to a process *pid* or a file *file* outside of GDB.

GDB command

The corresponding GDB command is ‘attach’.

Example

N.A.

The -target-compare-sections Command**Synopsis**

```
-target-compare-sections [ section ]
```

Compare data of section *section* on target to the exec file. Without the argument, all sections are compared.

GDB Command

The GDB equivalent is ‘compare-sections’.

Example

N.A.

The -target-detach Command**Synopsis**

```
-target-detach
```

Disconnect from the remote target. There’s no output.

GDB command

The corresponding GDB command is ‘detach’.

Example

```
(gdb)
-target-detach
^done
(gdb)
```

The -target-disconnect Command**Synopsis**

```
-target-disconnect
```

Disconnect from the remote target. There’s no output.

GDB command

The corresponding GDB command is ‘disconnect’.

Example

```
(gdb)
-target-disconnect
^done
(gdb)
```

The `-target-download` Command

Synopsis

```
-target-download
```

Loads the executable onto the remote target. It prints out an update message every half second, which includes the fields:

`'section'` The name of the section.

`'section-sent'`

The size of what has been sent so far for that section.

`'section-size'`

The size of the section.

`'total-sent'`

The total size of what was sent so far (the current and the previous sections).

`'total-size'`

The size of the overall executable to download.

Each message is sent as status record (see [Section 24.3.2 \[GDB/MI Output Syntax\]](#), page 204).

In addition, it prints the name and size of the sections, as they are downloaded. These messages include the following fields:

`'section'` The name of the section.

`'section-size'`

The size of the section.

`'total-size'`

The size of the overall executable to download.

At the end, a summary is printed.

GDB Command

The corresponding GDB command is `'load'`.

Example

Note: each status message appears on a single line. Here the messages have been broken down so that they can fit onto a page.

```
(gdb)
-target-download
+download,{section=".text",section-size="6668",total-size="9880"}
+download,{section=".text",section-sent="512",section-size="6668",
total-sent="512",total-size="9880"}
+download,{section=".text",section-sent="1024",section-size="6668",
total-sent="1024",total-size="9880"}
+download,{section=".text",section-sent="1536",section-size="6668",
total-sent="1536",total-size="9880"}
+download,{section=".text",section-sent="2048",section-size="6668",
total-sent="2048",total-size="9880"}
+download,{section=".text",section-sent="2560",section-size="6668",
total-sent="2560",total-size="9880"}
```

```
+download,{section=".text",section-sent="3072",section-size="6668",
total-sent="3072",total-size="9880"}
+download,{section=".text",section-sent="3584",section-size="6668",
total-sent="3584",total-size="9880"}
+download,{section=".text",section-sent="4096",section-size="6668",
total-sent="4096",total-size="9880"}
+download,{section=".text",section-sent="4608",section-size="6668",
total-sent="4608",total-size="9880"}
+download,{section=".text",section-sent="5120",section-size="6668",
total-sent="5120",total-size="9880"}
+download,{section=".text",section-sent="5632",section-size="6668",
total-sent="5632",total-size="9880"}
+download,{section=".text",section-sent="6144",section-size="6668",
total-sent="6144",total-size="9880"}
+download,{section=".text",section-sent="6656",section-size="6668",
total-sent="6656",total-size="9880"}
+download,{section=".init",section-size="28",total-size="9880"}
+download,{section=".fini",section-size="28",total-size="9880"}
+download,{section=".data",section-size="3156",total-size="9880"}
+download,{section=".data",section-sent="512",section-size="3156",
total-sent="7236",total-size="9880"}
+download,{section=".data",section-sent="1024",section-size="3156",
total-sent="7748",total-size="9880"}
+download,{section=".data",section-sent="1536",section-size="3156",
total-sent="8260",total-size="9880"}
+download,{section=".data",section-sent="2048",section-size="3156",
total-sent="8772",total-size="9880"}
+download,{section=".data",section-sent="2560",section-size="3156",
total-sent="9284",total-size="9880"}
+download,{section=".data",section-sent="3072",section-size="3156",
total-sent="9796",total-size="9880"}
^done,address="0x10004",load-size="9880",transfer-rate="6586",
write-rate="429"
(gdb)
```

The `-target-exec-status` Command

Synopsis

`-target-exec-status`

Provide information on the state of the target (whether it is running or not, for instance).

GDB Command

There's no equivalent GDB command.

Example

N.A.

The `-target-list-available-targets` Command

Synopsis

`-target-list-available-targets`

List the possible targets to connect to.

GDB Command

The corresponding GDB command is ‘`help target`’.

Example

N.A.

The `-target-list-current-targets` Command

Synopsis

```
-target-list-current-targets
```

Describe the current target.

GDB Command

The corresponding information is printed by ‘`info file`’ (among other things).

Example

N.A.

The `-target-list-parameters` Command

Synopsis

```
-target-list-parameters
```

GDB Command

No equivalent.

Example

N.A.

The `-target-select` Command

Synopsis

```
-target-select type parameters ...
```

Connect GDB to the remote target. This command takes two args:

‘`type`’ The type of target, for instance ‘`async`’, ‘`remote`’, etc.

‘`parameters`’

Device names, host names and the like. See [Section 16.2 \[Commands for managing targets\]](#), page 147, for more details.

The output is a connection notification, followed by the address at which the target program is, in the following form:

```
^connected,addr="address",func="function name",  
args=[arg list]
```

GDB Command

The corresponding GDB command is ‘`target`’.

Example

```
(gdb)
-target-select async /dev/ttya
^connected,addr="0xfe00a300",func="???",args=[]
(gdb)
```

24.14 GDB/MI Thread Commands

The `-thread-info` Command

Synopsis

```
-thread-info
```

GDB command

No equivalent.

Example

N.A.

The `-thread-list-all-threads` Command

Synopsis

```
-thread-list-all-threads
```

GDB Command

The equivalent GDB command is ‘`info threads`’.

Example

N.A.

The `-thread-list-ids` Command

Synopsis

```
-thread-list-ids
```

Produces a list of the currently known GDB thread ids. At the end of the list it also prints the total number of such threads.

GDB Command

Part of ‘`info threads`’ supplies the same information.

Example

No threads present, besides the main process:

```
(gdb)
-thread-list-ids
^done,thread-ids={},number-of-threads="0"
(gdb)
```

Several threads:

```
(gdb)
/thread-list-ids
^done,thread-ids={thread-id="3",thread-id="2",thread-id="1"}, 
number-of-threads="3"
(gdb)
```

The `-thread-select` Command

Synopsis

```
-thread-select threadnum
```

Make *threadnum* the current thread. It prints the number of the new current thread, and the topmost frame for that thread.

GDB Command

The corresponding GDB command is ‘`thread`’.

Example

```
(gdb)
-exec-next
^running
(gdb)
*stopped,reason="end-stepping-range",thread-id="2",line="187",
file="../../../../devo/gdb/testsuite/gdb.threads/linux-dp.c"
(gdb)
-thread-list-ids
^done,
thread-ids={thread-id="3",thread-id="2",thread-id="1"}, 
number-of-threads="3"
(gdb)
-thread-select 3
^done,new-thread-id="3",
frame={level="0",func="vprintf",
args=[{name="format",value="0x8048e9c \">%s%c %d %c\\n"}, {name="arg",value="0x2"}],file="vprintf.c",line="31"}
(gdb)
```

24.15 GDB/MI Tracepoint Commands

The tracepoint commands are not yet implemented.

24.16 GDB/MI Variable Objects

Motivation for Variable Objects in GDB/MI

For the implementation of a variable debugger window (locals, watched expressions, etc.), we are proposing the adaptation of the existing code used by `Insight`.

The two main reasons for that are:

1. It has been proven in practice (it is already on its second generation).
2. It will shorten development time (needless to say how important it is now).

The original interface was designed to be used by Tcl code, so it was slightly changed so it could be used through GDB/MI. This section describes the GDB/MI operations that will be available and gives some hints about their use.

Note: In addition to the set of operations described here, we expect the GUI implementation of a variable window to require, at least, the following operations:

- `-gdb-show output-radix`
- `-stack-list-arguments`
- `-stack-list-locals`
- `-stack-select-frame`

Introduction to Variable Objects in GDB/MI

The basic idea behind variable objects is the creation of a named object to represent a variable, an expression, a memory location or even a CPU register. For each object created, a set of operations is available for examining or changing its properties.

Furthermore, complex data types, such as C structures, are represented in a tree format. For instance, the `struct` type variable is the root and the children will represent the struct members. If a child is itself of a complex type, it will also have children of its own. Appropriate language differences are handled for C, C++ and Java.

When returning the actual values of the objects, this facility allows for the individual selection of the display format used in the result creation. It can be chosen among: binary, decimal, hexadecimal, octal and natural. Natural refers to a default format automatically chosen based on the variable type (like decimal for an `int`, hex for pointers, etc.).

The following is the complete set of GDB/MI operations defined to access this functionality:

Operation	Description
<code>-var-create</code>	create a variable object
<code>-var-delete</code>	delete the variable object and its children
<code>-var-set-format</code>	set the display format of this variable
<code>-var-show-format</code>	show the display format of this variable
<code>-var-info-num-children</code>	tells how many children this object has
<code>-var-list-children</code>	return a list of the object's children
<code>-var-info-type</code>	show the type of this variable object
<code>-var-info-expression</code>	print what this variable object represents
<code>-var-show-attributes</code>	is this variable editable? does it exist here?
<code>-var-evaluate-expression</code>	get the value of this variable
<code>-var-assign</code>	set the value of this variable
<code>-var-update</code>	update the variable and its children

In the next subsection we describe each operation in detail and suggest how it can be used.

Description And Use of Operations on Variable Objects

The `-var-create` Command

Synopsis

```
-var-create {name | "-"}  
           {frame-addr | "*"} expression
```

This operation creates a variable object, which allows the monitoring of a variable, the result of an expression, a memory cell or a CPU register.

The *name* parameter is the string by which the object can be referenced. It must be unique. If ‘-’ is specified, the varobj system will generate a string “varNNNNNN” automatically. It will be unique provided that one does not specify *name* on that format. The command fails if a duplicate name is found.

The frame under which the expression should be evaluated can be specified by *frame-addr*. A ‘*’ indicates that the current frame should be used.

expression is any expression valid on the current language set (must not begin with a ‘*’), or one of the following:

- ‘**addr*’, where *addr* is the address of a memory cell
- ‘**addr-addr*’ — a memory address range (TBD)
- ‘\$*regname*’ — a CPU register name

Result

This operation returns the name, number of children and the type of the object created. Type is returned as a string as the ones generated by the GDB CLI:

```
name="name",numchild="N",type="type"
```

The -var-delete Command

Synopsis

```
-var-delete name
```

Deletes a previously created variable object and all of its children.

Returns an error if the object *name* is not found.

The -var-set-format Command

Synopsis

```
-var-set-format name format-spec
```

Sets the output format for the value of the object *name* to be *format-spec*.

The syntax for the *format-spec* is as follows:

```
format-spec ↪
{binary | decimal | hexadecimal | octal | natural}
```

The -var-show-format Command

Synopsis

```
-var-show-format name
```

Returns the format used to display the value of the object *name*.

```
format ↪
format-spec
```

The -var-info-num-children Command

Synopsis

```
-var-info-num-children name
```

Returns the number of children of a variable object *name*:

```
numchild=n
```

The -var-list-children Command

Synopsis

```
-var-list-children [print-values] name
```

Returns a list of the children of the specified variable object. With just the variable object name as an argument or with an optional preceding argument of 0 or --no-values, prints only the names of the variables. With an optional preceding argument of 1 or --all-values, also prints their values.

Example

```
(gdb)
-var-list-children n
numchild=n,children=[{name=name,
numchild=n,type=type},(repeats N times)]
(gdb)
-var-list-children --all-values n
numchild=n,children=[{name=name,
numchild=n,value=value,type=type},(repeats N times)]
```

The -var-info-type Command

Synopsis

```
-var-info-type name
```

Returns the type of the specified variable *name*. The type is returned as a string in the same format as it is output by the GDB CLI:

```
type=typename
```

The -var-info-expression Command

Synopsis

```
-var-info-expression name
```

Returns what is represented by the variable object *name*:

```
lang=lang-spec,exp=expression
```

where *lang-spec* is {"C" | "C++" | "Java"}.

The -var-show-attributes Command

Synopsis

```
-var-show-attributes name
```

List attributes of the specified variable object *name*:

```
status=attr [ ( ,attr )* ]
```

where *attr* is { { editable | noneditable } | TBD }.

The -var-evaluate-expression Command

Synopsis

```
-var-evaluate-expression name
```

Evaluates the expression that is represented by the specified variable object and returns its value as a string in the current format specified for the object:

```
value=value
```

Note that one must invoke `-var-list-children` for a variable before the value of a child variable can be evaluated.

The `-var-assign` Command

Synopsis

```
-var-assign name expression
```

Assigns the value of *expression* to the variable object specified by *name*. The object must be ‘`editable`’. If the variable’s value is altered by the assign, the variable will show up in any subsequent `-var-update` list.

Example

```
(gdb)
-var-assign var1 3
^done,value="3"
(gdb)
-var-update *
^done,changelist=[{"name":"var1",in_scope="true",type_changed="false"}]
(gdb)
```

The `-var-update` Command

Synopsis

```
-var-update {name | "*")}
```

Update the value of the variable object *name* by evaluating its expression after fetching all the new values from memory or registers. A ‘*’ causes all existing variable objects to be updated.

25 GDB Annotations

This chapter describes annotations in GDB. Annotations were designed to interface GDB to graphical user interfaces or other similar programs which want to interact with GDB at a relatively high level.

The annotation mechanism has largely been superseeded by GDB/MI (see Chapter 24 [GDB/MI], page 203).

25.1 What is an Annotation?

Annotations start with a newline character, two ‘control-z’ characters, and the name of the annotation. If there is no additional information associated with this annotation, the name of the annotation is followed immediately by a newline. If there is additional information, the name of the annotation is followed by a space, the additional information, and a newline. The additional information cannot contain newline characters.

Any output not beginning with a newline and two ‘control-z’ characters denotes literal output from GDB. Currently there is no need for GDB to output a newline followed by two ‘control-z’ characters, but if there was such a need, the annotations could be extended with an ‘escape’ annotation which means those three characters as output.

The annotation level, which is specified using the ‘`--annotate`’ command line option (see Section 2.1.2 [Mode Options], page 13), controls how much information GDB prints together with its prompt, values of expressions, source lines, and other types of output. Level 0 is for no annotations, level 1 is for use when GDB is run as a subprocess of GNU Emacs, level 3 is the maximum annotation suitable for programs that control GDB, and level 2 annotations have been made obsolete (see section “Limitations of the Annotation Interface” in *GDB’s Obsolete Annotations*). This chapter describes level 3 annotations.

A simple example of starting up GDB with annotations is:

```
$ gdb --annotate=3
GNU gdb 6.0
Copyright 2003 Free Software Foundation, Inc.
GDB is free software, covered by the GNU General Public License,
and you are welcome to change it and/or distribute copies of it
under certain conditions.
Type "show copying" to see the conditions.
There is absolutely no warranty for GDB. Type "show warranty"
for details.
This GDB was configured as "i386-pc-linux-gnu"

^Z^Zpre-prompt
(gdb)
^Z^Zprompt
quit

^Z^Zpost-prompt
$
```

Here ‘quit’ is input to GDB; the rest is output from GDB. The three lines beginning ‘^Z^Z’ (where ‘^Z’ denotes a ‘control-z’ character) are annotations; the rest is output from GDB.

25.2 The Server Prefix

To issue a command to GDB without affecting certain aspects of the state which is seen by users, prefix it with ‘`server`’. This means that this command will not affect the command history, nor will it affect GDB’s notion of which command to repeat if `(RET)` is pressed on a line by itself.

The server prefix does not affect the recording of values into the value history; to print a value without recording it into the value history, use the `output` command instead of the `print` command.

25.3 Annotation for GDB Input

When GDB prompts for input, it annotates this fact so it is possible to know when to send output, when the output from a given command is over, etc.

Different kinds of input each have a different *input type*. Each input type has three annotations: a `pre-` annotation, which denotes the beginning of any prompt which is being output, a plain annotation, which denotes the end of the prompt, and then a `post-` annotation which denotes the end of any echo which may (or may not) be associated with the input. For example, the `prompt` input type features the following annotations:

```
^Z^Zpre-prompt
^Z^Zprompt
^Z^Zpost-prompt
```

The input types are

`prompt` When GDB is prompting for a command (the main GDB prompt).

`commands` When GDB prompts for a set of commands, like in the `commands` command.
The annotations are repeated for each command which is input.

`overload-choice`

When GDB wants the user to select between various overloaded functions.

`query` When GDB wants the user to confirm a potentially dangerous operation.

`prompt-for-continue`

When GDB is asking the user to press return to continue. Note: Don’t expect this to work well; instead use `set height 0` to disable prompting. This is because the counting of lines is buggy in the presence of annotations.

25.4 Errors

```
^Z^Zquit
```

This annotation occurs right before GDB responds to an interrupt.

```
^Z^Zerror
```

This annotation occurs right before GDB responds to an error.

Quit and error annotations indicate that any annotations which GDB was in the middle of may end abruptly. For example, if a `value-history-begin` annotation is followed by a `error`, one cannot expect to receive the matching `value-history-end`. One cannot expect not to receive it either, however; an error annotation does not necessarily mean that GDB is immediately returning all the way to the top level.

A quit or error annotation may be preceded by

`^Z^Zerror-begin`

Any output between that and the quit or error annotation is the error message.

Warning messages are not yet annotated.

25.5 Invalidation Notices

The following annotations say that certain pieces of state may have changed.

`^Z^Zframes-invalid`

The frames (for example, output from the `backtrace` command) may have changed.

`^Z^Zbreakpoints-invalid`

The breakpoints may have changed. For example, the user just added or deleted a breakpoint.

25.6 Running the Program

When the program starts executing due to a GDB command such as `step` or `continue`,

`^Z^Zstarting`

is output. When the program stops,

`^Z^Zstopped`

is output. Before the `stopped` annotation, a variety of annotations describe how the program stopped.

`^Z^Zexited exit-status`

The program exited, and `exit-status` is the exit status (zero for successful exit, otherwise nonzero).

`^Z^Zsignalled`

The program exited with a signal. After the `^Z^Zsignalled`, the annotation continues:

```
intro-text
^Z^Zsignal-name
name
^Z^Zsignal-name-end
middle-text
^Z^Zsignal-string
string
^Z^Zsignal-string-end
end-text
```

where `name` is the name of the signal, such as `SIGILL` or `SIGSEGV`, and `string` is the explanation of the signal, such as `Illegal Instruction` or `Segmentation fault`. `intro-text`, `middle-text`, and `end-text` are for the user's benefit and have no particular format.

`^Z^Zsignal`

The syntax of this annotation is just like `signalled`, but GDB is just saying that the program received the signal, not that it was terminated with it.

`^Z^Zbreakpoint number`

The program hit breakpoint number `number`.

`^Z^Zwatchpoint number`

The program hit watchpoint number *number*.

25.7 Displaying Source

The following annotation is used instead of displaying source code:

`^Z^Zsource filename:line:character:middle:addr`

where *filename* is an absolute file name indicating which source file, *line* is the line number within that file (where 1 is the first line in the file), *character* is the character position within the file (where 0 is the first character in the file) (for most debug formats this will necessarily point to the beginning of a line), *middle* is ‘*middle*’ if *addr* is in the middle of the line, or ‘*beg*’ if *addr* is at the beginning of the line, and *addr* is the address in the target program associated with the source which is being displayed. *addr* is in the form ‘*0x*’ followed by one or more lowercase hex digits (note that this does not depend on the language).

26 Reporting Bugs in GDB

Your bug reports play an essential role in making GDB reliable.

Reporting a bug may help you by bringing a solution to your problem, or it may not. But in any case the principal function of a bug report is to help the entire community by making the next version of GDB work better. Bug reports are your contribution to the maintenance of GDB.

In order for a bug report to serve its purpose, you must include the information that enables us to fix the bug.

26.1 Have you found a bug?

If you are not sure whether you have found a bug, here are some guidelines:

- If the debugger gets a fatal signal, for any input whatever, that is a GDB bug. Reliable debuggers never crash.
- If GDB produces an error message for valid input, that is a bug. (Note that if you're cross debugging, the problem may also be somewhere in the connection to the target.)
- If GDB does not produce an error message for invalid input, that is a bug. However, you should note that your idea of "invalid input" might be our idea of "an extension" or "support for traditional practice".
- If you are an experienced user of debugging tools, your suggestions for improvement of GDB are welcome in any case.

26.2 How to report bugs

A number of companies and individuals offer support for GNU products. If you obtained GDB from a support organization, we recommend you contact that organization first.

You can find contact information for many support companies and individuals in the file ‘etc/SERVICE’ in the GNU Emacs distribution.

In any event, we also recommend that you submit bug reports for GDB. The preferred method is to submit them directly using [GDB's Bugs web page](#). Alternatively, the [e-mail gateway](#) can be used.

Do not send bug reports to ‘info-gdb’, or to ‘help-gdb’, or to any newsgroups. Most users of GDB do not want to receive bug reports. Those that do have arranged to receive ‘bug-gdb’.

The mailing list ‘bug-gdb’ has a newsgroup ‘gnu.gdb.bug’ which serves as a repeater. The mailing list and the newsgroup carry exactly the same messages. Often people think of posting bug reports to the newsgroup instead of mailing them. This appears to work, but it has one problem which can be crucial: a newsgroup posting often lacks a mail path back to the sender. Thus, if we need to ask for more information, we may be unable to reach you. For this reason, it is better to send bug reports to the mailing list.

The fundamental principle of reporting bugs usefully is this: **report all the facts.** If you are not sure whether to state a fact or leave it out, state it!

Often people omit facts because they think they know what causes the problem and assume that some details do not matter. Thus, you might assume that the name of the

variable you use in an example does not matter. Well, probably it does not, but one cannot be sure. Perhaps the bug is a stray memory reference which happens to fetch from the location where that name is stored in memory; perhaps, if the name were different, the contents of that location would fool the debugger into doing the right thing despite the bug. Play it safe and give a specific, complete example. That is the easiest thing for you to do, and the most helpful.

Keep in mind that the purpose of a bug report is to enable us to fix the bug. It may be that the bug has been reported previously, but neither you nor we can know that unless your bug report is complete and self-contained.

Sometimes people give a few sketchy facts and ask, “Does this ring a bell?” Those bug reports are useless, and we urge everyone to *refuse to respond to them* except to chide the sender to report bugs properly.

To enable us to fix the bug, you should include all these things:

- The version of GDB. GDB announces it if you start with no arguments; you can also print it at any time using `show version`.
Without this, we will not know whether there is any point in looking for the bug in the current version of GDB.
- The type of machine you are using, and the operating system name and version number.
- What compiler (and its version) was used to compile GDB—e.g. “gcc-2.8.1”.
- What compiler (and its version) was used to compile the program you are debugging—e.g. “gcc-2.8.1”, or “HP92453-01 A.10.32.03 HP C Compiler”. For GCC, you can say `gcc --version` to get this information; for other compilers, see the documentation for those compilers.
- The command arguments you gave the compiler to compile your example and observe the bug. For example, did you use ‘`-O`’? To guarantee you will not omit something important, list them all. A copy of the Makefile (or the output from make) is sufficient. If we were to try to guess the arguments, we would probably guess wrong and then we might not encounter the bug.
- A complete input script, and all necessary source files, that will reproduce the bug.
- A description of what behavior you observe that you believe is incorrect. For example, “It gets a fatal signal.”

Of course, if the bug is that GDB gets a fatal signal, then we will certainly notice it. But if the bug is incorrect output, we might not notice unless it is glaringly wrong. You might as well not give us a chance to make a mistake.

Even if the problem you experience is a fatal signal, you should still say so explicitly. Suppose something strange is going on, such as, your copy of GDB is out of synch, or you have encountered a bug in the C library on your system. (This has happened!) Your copy might crash and ours would not. If you told us to expect a crash, then when ours fails to crash, we would know that the bug was not happening for us. If you had not told us to expect a crash, then we would not be able to draw any conclusion from our observations.

To collect all this information, you can use a session recording program such as `script`, which is available on many Unix systems. Just run your GDB session inside `script` and then include the ‘`typescript`’ file with your bug report.

Another way to record a GDB session is to run GDB inside Emacs and then save the entire buffer to a file.

- If you wish to suggest changes to the GDB source, send us context diffs. If you even discuss something in the GDB source, refer to it by context, not by line number.

The line numbers in our development sources will not match those in your sources. Your line numbers would convey no useful information to us.

Here are some things that are not necessary:

- A description of the envelope of the bug.

Often people who encounter a bug spend a lot of time investigating which changes to the input file will make the bug go away and which changes will not affect it.

This is often time consuming and not very useful, because the way we will find the bug is by running a single example under the debugger with breakpoints, not by pure deduction from a series of examples. We recommend that you save your time for something else.

Of course, if you can find a simpler example to report *instead* of the original one, that is a convenience for us. Errors in the output will be easier to spot, running under the debugger will take less time, and so on.

However, simplification is not vital; if you do not want to do this, report the bug anyway and send us the entire test case you used.

- A patch for the bug.

A patch for the bug does help us if it is a good one. But do not omit the necessary information, such as the test case, on the assumption that a patch is all we need. We might see problems with your patch and decide to fix the problem another way, or we might not understand it at all.

Sometimes with a program as complicated as GDB it is very hard to construct an example that will make the program follow a certain path through the code. If you do not send us the example, we will not be able to construct one, so we will not be able to verify that the bug is fixed.

And if we cannot understand what bug you are trying to fix, or why your patch should be an improvement, we will not install it. A test case will help us to understand.

- A guess about what the bug is or what it depends on.

Such guesses are usually wrong. Even we cannot guess right about such things without first using the debugger to find the facts.

27 Command Line Editing

This chapter describes the basic features of the GNU command line editing interface.

27.1 Introduction to Line Editing

The following paragraphs describe the notation used to represent keystrokes.

The text *C-k* is read as ‘Control-K’ and describes the character produced when the $\langle \text{K} \rangle$ key is pressed while the Control key is depressed.

The text *M-k* is read as ‘Meta-K’ and describes the character produced when the Meta key (if you have one) is depressed, and the $\langle \text{K} \rangle$ key is pressed. The Meta key is labeled $\langle \text{ALT} \rangle$ on many keyboards. On keyboards with two keys labeled $\langle \text{ALT} \rangle$ (usually to either side of the space bar), the $\langle \text{ALT} \rangle$ on the left side is generally set to work as a Meta key. The $\langle \text{ALT} \rangle$ key on the right may also be configured to work as a Meta key or may be configured as some other modifier, such as a Compose key for typing accented characters.

If you do not have a Meta or $\langle \text{ALT} \rangle$ key, or another key working as a Meta key, the identical keystroke can be generated by typing $\langle \text{ESC} \rangle$ first, and then typing $\langle \text{K} \rangle$. Either process is known as *metafying* the $\langle \text{K} \rangle$ key.

The text *M-C-k* is read as ‘Meta-Control-k’ and describes the character produced by *metafying* *C-k*.

In addition, several keys have their own names. Specifically, $\langle \text{DEL} \rangle$, $\langle \text{ESC} \rangle$, $\langle \text{LFD} \rangle$, $\langle \text{SPC} \rangle$, $\langle \text{RET} \rangle$, and $\langle \text{TAB} \rangle$ all stand for themselves when seen in this text, or in an init file (see [Section 27.3 \[Readline Init File\], page 264](#)). If your keyboard lacks a $\langle \text{LFD} \rangle$ key, typing $\langle \text{C-j} \rangle$ will produce the desired character. The $\langle \text{RET} \rangle$ key may be labeled $\langle \text{Return} \rangle$ or $\langle \text{Enter} \rangle$ on some keyboards.

27.2 Readline Interaction

Often during an interactive session you type in a long line of text, only to notice that the first word on the line is misspelled. The Readline library gives you a set of commands for manipulating the text as you type it in, allowing you to just fix your typo, and not forcing you to retype the majority of the line. Using these editing commands, you move the cursor to the place that needs correction, and delete or insert the text of the corrections. Then, when you are satisfied with the line, you simply press $\langle \text{RET} \rangle$. You do not have to be at the end of the line to press $\langle \text{RET} \rangle$; the entire line is accepted regardless of the location of the cursor within the line.

27.2.1 Readline Bare Essentials

In order to enter characters into the line, simply type them. The typed character appears where the cursor was, and then the cursor moves one space to the right. If you mistype a character, you can use your erase character to back up and delete the mistyped character.

Sometimes you may mistype a character, and not notice the error until you have typed several other characters. In that case, you can type *C-b* to move the cursor to the left, and then correct your mistake. Afterwards, you can move the cursor to the right with *C-f*.

When you add text in the middle of a line, you will notice that characters to the right of the cursor are ‘pushed over’ to make room for the text that you have inserted. Likewise,

when you delete text behind the cursor, characters to the right of the cursor are ‘pulled back’ to fill in the blank space created by the removal of the text. A list of the bare essentials for editing the text of an input line follows.

C-b Move back one character.

C-f Move forward one character.

⟨DEL⟩ or *⟨Backspace⟩*

Delete the character to the left of the cursor.

C-d Delete the character underneath the cursor.

Printing characters

Insert the character into the line at the cursor.

C-_ or *C-x C-u*

Undo the last editing command. You can undo all the way back to an empty line.

(Depending on your configuration, the *⟨Backspace⟩* key be set to delete the character to the left of the cursor and the *⟨DEL⟩* key set to delete the character underneath the cursor, like *C-d*, rather than the character to the left of the cursor.)

27.2.2 Readline Movement Commands

The above table describes the most basic keystrokes that you need in order to do editing of the input line. For your convenience, many other commands have been added in addition to *C-b*, *C-f*, *C-d*, and *⟨DEL⟩*. Here are some commands for moving more rapidly about the line.

C-a Move to the start of the line.

C-e Move to the end of the line.

M-f Move forward a word, where a word is composed of letters and digits.

M-b Move backward a word.

C-l Clear the screen, reprinting the current line at the top.

Notice how *C-f* moves forward a character, while *M-f* moves forward a word. It is a loose convention that control keystrokes operate on characters while meta keystrokes operate on words.

27.2.3 Readline Killing Commands

Killing text means to delete the text from the line, but to save it away for later use, usually by *yanking* (re-inserting) it back into the line. (‘Cut’ and ‘paste’ are more recent jargon for ‘kill’ and ‘yank’.)

If the description for a command says that it ‘kills’ text, then you can be sure that you can get the text back in a different (or the same) place later.

When you use a kill command, the text is saved in a *kill-ring*. Any number of consecutive kills save all of the killed text together, so that when you yank it back, you get it all. The kill ring is not line specific; the text that you killed on a previously typed line is available to be yanked back later, when you are typing another line.

Here is the list of commands for killing text.

<i>C-k</i>	Kill the text from the current cursor position to the end of the line.
<i>M-d</i>	Kill from the cursor to the end of the current word, or, if between words, to the end of the next word. Word boundaries are the same as those used by <i>M-f</i> .
<i>M-DEL</i>	Kill from the cursor the start of the current word, or, if between words, to the start of the previous word. Word boundaries are the same as those used by <i>M-b</i> .
<i>C-w</i>	Kill from the cursor to the previous whitespace. This is different than <i>M-DEL</i> because the word boundaries differ.

Here is how to *yank* the text back into the line. Yanking means to copy the most-recently-killed text from the kill buffer.

<i>C-y</i>	Yank the most recently killed text back into the buffer at the cursor.
<i>M-y</i>	Rotate the kill-ring, and yank the new top. You can only do this if the prior command is <i>C-y</i> or <i>M-y</i> .

27.2.4 Readline Arguments

You can pass numeric arguments to Readline commands. Sometimes the argument acts as a repeat count, other times it is the *sign* of the argument that is significant. If you pass a negative argument to a command which normally acts in a forward direction, that command will act in a backward direction. For example, to kill text back to the start of the line, you might type ‘*M-- C-k*’.

The general way to pass numeric arguments to a command is to type meta digits before the command. If the first ‘digit’ typed is a minus sign (‘-’), then the sign of the argument will be negative. Once you have typed one meta digit to get the argument started, you can type the remainder of the digits, and then the command. For example, to give the *C-d* command an argument of 10, you could type ‘*M-1 0 C-d*’, which will delete the next ten characters on the input line.

27.2.5 Searching for Commands in the History

Readline provides commands for searching through the command history for lines containing a specified string. There are two search modes: *incremental* and *non-incremental*.

Incremental searches begin before the user has finished typing the search string. As each character of the search string is typed, Readline displays the next entry from the history matching the string typed so far. An incremental search requires only as many characters as needed to find the desired history entry. To search backward in the history for a particular string, type *C-r*. Typing *C-s* searches forward through the history. The characters present in the value of the `isearch-terminators` variable are used to terminate an incremental search. If that variable has not been assigned a value, the `ESC` and *C-J* characters will terminate an incremental search. *C-g* will abort an incremental search and restore the original line. When the search is terminated, the history entry containing the search string becomes the current line.

To find other matching entries in the history list, type *C-r* or *C-s* as appropriate. This will search backward or forward in the history for the next entry matching the search string typed so far. Any other key sequence bound to a Readline command will terminate the

search and execute that command. For instance, a `\RET` will terminate the search and accept the line, thereby executing the command from the history list. A movement command will terminate the search, make the last line found the current line, and begin editing.

Readline remembers the last incremental search string. If two `C-rs` are typed without any intervening characters defining a new search string, any remembered search string is used.

Non-incremental searches read the entire search string before starting to search for matching history lines. The search string may be typed by the user or be part of the contents of the current line.

27.3 Readline Init File

Although the Readline library comes with a set of Emacs-like keybindings installed by default, it is possible to use a different set of keybindings. Any user can customize programs that use Readline by putting commands in an `inputrc` file, conventionally in his home directory. The name of this file is taken from the value of the environment variable `INPUTRC`. If that variable is unset, the default is `'~/.inputrc'`.

When a program which uses the Readline library starts up, the init file is read, and the key bindings are set.

In addition, the `C-x C-r` command re-reads this init file, thus incorporating any changes that you might have made to it.

27.3.1 Readline Init File Syntax

There are only a few basic constructs allowed in the Readline init file. Blank lines are ignored. Lines beginning with a '#' are comments. Lines beginning with a '\$' indicate conditional constructs (see [Section 27.3.2 \[Conditional Init Constructs\], page 269](#)). Other lines denote variable settings and key bindings.

Variable Settings

You can modify the run-time behavior of Readline by altering the values of variables in Readline using the `set` command within the init file. The syntax is simple:

```
set variable value
```

Here, for example, is how to change from the default Emacs-like key binding to use vi line editing commands:

```
set editing-mode vi
```

Variable names and values, where appropriate, are recognized without regard to case.

A great deal of run-time behavior is changeable with the following variables.

`bell-style`

Controls what happens when Readline wants to ring the terminal bell. If set to 'none', Readline never rings the bell. If set to 'visible', Readline uses a visible bell if one is available. If set to 'audible' (the default), Readline attempts to ring the terminal's bell.

comment-begin

The string to insert at the beginning of the line when the `insert-comment` command is executed. The default value is "#".

completion-ignore-case

If set to 'on', Readline performs filename matching and completion in a case-insensitive fashion. The default value is 'off'.

completion-query-items

The number of possible completions that determines when the user is asked whether he wants to see the list of possibilities. If the number of possible completions is greater than this value, Readline will ask the user whether or not he wishes to view them; otherwise, they are simply listed. This variable must be set to an integer value greater than or equal to 0. The default limit is 100.

convert-meta

If set to 'on', Readline will convert characters with the eighth bit set to an ASCII key sequence by stripping the eighth bit and prefixing an `\ESC` character, converting them to a meta-prefixed key sequence. The default value is 'on'.

disable-completion

If set to 'On', Readline will inhibit word completion. Completion characters will be inserted into the line as if they had been mapped to `self-insert`. The default is 'off'.

editing-mode

The `editing-mode` variable controls which default set of key bindings is used. By default, Readline starts up in Emacs editing mode, where the keystrokes are most similar to Emacs. This variable can be set to either '`emacs`' or '`vi`'.

enable-keypad

When set to 'on', Readline will try to enable the application keypad when it is called. Some systems need this to enable the arrow keys. The default is 'off'.

expand-tilde

If set to 'on', tilde expansion is performed when Readline attempts word completion. The default is 'off'.

If set to 'on', the history code attempts to place point at the same location on each history line retrieved with `previous-history` or `next-history`.

horizontal-scroll-mode

This variable can be set to either 'on' or 'off'. Setting it to 'on' means that the text of the lines being edited will scroll horizontally on a single screen line when they are longer than the width of the screen, instead of wrapping onto a new screen line. By default, this variable is set to 'off'.

input-meta

If set to ‘on’, Readline will enable eight-bit input (it will not clear the eighth bit in the characters it reads), regardless of what the terminal claims it can support. The default value is ‘off’. The name `meta-flag` is a synonym for this variable.

isearch-terminators

The string of characters that should terminate an incremental search without subsequently executing the character as a command (see [Section 27.2.5 \[Searching\], page 263](#)). If this variable has not been given a value, the characters `\ESC` and `C-J` will terminate an incremental search.

keymap Sets Readline’s idea of the current keymap for key binding commands. Acceptable `keymap` names are `emacs`, `emacs-standard`, `emacs-meta`, `emacs-ctlx`, `vi`, `vi-move`, `vi-command`, and `vi-insert`. `vi` is equivalent to `vi-command`; `emacs` is equivalent to `emacs-standard`. The default value is `emacs`. The value of the `editing-mode` variable also affects the default keymap.

mark-directories

If set to ‘on’, completed directory names have a slash appended. The default is ‘on’.

mark-modified-lines

This variable, when set to ‘on’, causes Readline to display an asterisk (*) at the start of history lines which have been modified. This variable is ‘off’ by default.

mark-symlinked-directories

If set to ‘on’, completed names which are symbolic links to directories have a slash appended (subject to the value of `mark-directories`). The default is ‘off’.

match-hidden-files

This variable, when set to ‘on’, causes Readline to match files whose names begin with a ‘.’ (hidden files) when performing filename completion, unless the leading ‘.’ is supplied by the user in the filename to be completed. This variable is ‘on’ by default.

output-meta

If set to ‘on’, Readline will display characters with the eighth bit set directly rather than as a meta-prefixed escape sequence. The default is ‘off’.

page-completions

If set to ‘on’, Readline uses an internal `more`-like pager to display a screenful of possible completions at a time. This variable is ‘on’ by default.

print-completions-horizontally

If set to ‘on’, Readline will display completions with matches sorted horizontally in alphabetical order, rather than down the screen. The default is ‘off’.

show-all-if-ambiguous

This alters the default behavior of the completion functions. If set to ‘on’, words which have more than one possible completion cause the matches to be listed immediately instead of ringing the bell. The default value is ‘off’.

visible-stats

If set to ‘on’, a character denoting a file’s type is appended to the filename when listing possible completions. The default is ‘off’.

Key Bindings

The syntax for controlling key bindings in the init file is simple. First you need to find the name of the command that you want to change. The following sections contain tables of the command name, the default keybinding, if any, and a short description of what the command does.

Once you know the name of the command, simply place on a line in the init file the name of the key you wish to bind the command to, a colon, and then the name of the command. The name of the key can be expressed in different ways, depending on what you find most comfortable.

In addition to command names, readline allows keys to be bound to a string that is inserted when the key is pressed (a *macro*).

keyname: function-name or macro

keyname is the name of a key spelled out in English. For example:

```
Control-u: universal-argument
Meta-Rubout: backward-kill-word
Control-o: "> output"
```

In the above example, *C-u* is bound to the function *universal-argument*, *M-DEL* is bound to the function *backward-kill-word*, and *C-o* is bound to run the macro expressed on the right hand side (that is, to insert the text ‘> output’ into the line).

A number of symbolic character names are recognized while processing this key binding syntax: *DEL*, *ESC*, *ESCAPE*, *LFD*, *NEW-LINE*, *RET*, *RETURN*, *RUBOUT*, *SPACE*, *SPC*, and *TAB*.

"keyseq": function-name or macro

keyseq differs from *keyname* above in that strings denoting an entire key sequence can be specified, by placing the key sequence in double quotes. Some GNU Emacs style key escapes can be used, as in the following example, but the special character names are not recognized.

```
"\C-u": universal-argument
"\C-x\C-r": re-read-init-file
```

```
"\e[11~": "Function Key 1"
```

In the above example, ‘*C-u*’ is again bound to the function `universal-argument` (just as it was in the first example), ‘*C-x C-r*’ is bound to the function `re-read-init-file`, and ‘`\e[11~`’ is bound to insert the text ‘Function Key 1’.

The following GNU Emacs style escape sequences are available when specifying key sequences:

<code>\C-</code>	control prefix
<code>\M-</code>	meta prefix
<code>\e</code>	an escape character
<code>\\\</code>	backslash
<code>\"</code>	<code>\\"</code> , a double quotation mark
<code>\'</code>	<code>\'</code> , a single quote or apostrophe

In addition to the GNU Emacs style escape sequences, a second set of backslash escapes is available:

<code>\a</code>	alert (bell)
<code>\b</code>	backspace
<code>\d</code>	delete
<code>\f</code>	form feed
<code>\n</code>	newline
<code>\r</code>	carriage return
<code>\t</code>	horizontal tab
<code>\v</code>	vertical tab
<code>\nnn</code>	the eight-bit character whose value is the octal value <i>nnn</i> (one to three digits)
<code>\xHH</code>	the eight-bit character whose value is the hexadecimal value <i>HH</i> (one or two hex digits)

When entering the text of a macro, single or double quotes must be used to indicate a macro definition. Unquoted text is assumed to be a function name. In the macro body, the backslash escapes described above are expanded. Backslash will quote any other character in the macro text, including “” and ‘’. For example, the following binding will make ‘*C-x *’ insert a single ‘\’ into the line:

```
"\C-x\\": "\\\"
```

27.3.2 Conditional Init Constructs

Readline implements a facility similar in spirit to the conditional compilation features of the C preprocessor which allows key bindings and variable settings to be performed as the result of tests. There are four parser directives used.

\$if The **\$if** construct allows bindings to be made based on the editing mode, the terminal being used, or the application using Readline. The text of the test extends to the end of the line; no characters are required to isolate it.

mode The **mode=** form of the **\$if** directive is used to test whether Readline is in **emacs** or **vi** mode. This may be used in conjunction with the ‘**set keymap**’ command, for instance, to set bindings in the **emacs-standard** and **emacs-ctlx** keymaps only if Readline is starting out in **emacs** mode.

term The **term=** form may be used to include terminal-specific key bindings, perhaps to bind the key sequences output by the terminal’s function keys. The word on the right side of the ‘**=**’ is tested against both the full name of the terminal and the portion of the terminal name before the first ‘**-**’. This allows **sun** to match both **sun** and **sun-cmd**, for instance.

application The **application** construct is used to include application-specific settings. Each program using the Readline library sets the *application name*, and you can test for a particular value. This could be used to bind key sequences to functions useful for a specific program. For instance, the following command adds a key sequence that quotes the current or previous word in Bash:

```
$if Bash
# Quote the current or previous word
"\C-xq": "\eb\"\'ef\""
$endif
```

\$endif This command, as seen in the previous example, terminates an **\$if** command.

\$else Commands in this branch of the **\$if** directive are executed if the test fails.

\$include This directive takes a single filename as an argument and reads commands and bindings from that file. For example, the following directive reads from ‘**/etc/inputrc**’:

```
$include /etc/inputrc
```

27.3.3 Sample Init File

Here is an example of an **inputrc** file. This illustrates key binding, variable assignment, and conditional syntax.

```
# This file controls the behaviour of line input editing for
# programs that use the GNU Readline library. Existing
# programs include FTP, Bash, and GDB.
#
# You can re-read the inputrc file with C-x C-r.
# Lines beginning with '#' are comments.
#
# First, include any systemwide bindings and variable
# assignments from /etc/Inputrc
$include /etc/Inputrc

#
# Set various bindings for emacs mode.

set editing-mode emacs

$if mode=emacs

Meta-Control-h: backward-kill-word Text after the function name is ignored

#
# Arrow keys in keypad mode
#
#"M-OD":      backward-char
#"M-OC":      forward-char
#"M-OA":      previous-history
#"M-OB":      next-history
#
# Arrow keys in ANSI mode
#
"\M-[D":      backward-char
"\M-[C":      forward-char
"\M-[A":      previous-history
"\M-[B":      next-history
#
# Arrow keys in 8 bit keypad mode
#
#"M-\C-OD":      backward-char
#"M-\C-OC":      forward-char
#"M-\C-OA":      previous-history
#"M-\C-OB":      next-history
#
# Arrow keys in 8 bit ANSI mode
#
#"M-\C-[D":      backward-char
#"M-\C-[C":      forward-char
```

```
#"M-\C-[A":      previous-history
#"M-\C-[B":      next-history

C-q: quoted-insert

$endif

# An old-style binding. This happens to be the default.
TAB: complete

# Macros that are convenient for shell interaction
$if Bash
# edit the path
"\C-xp": "PATH=${PATH}\e\C-e\C-a\ef\C-f"
# prepare to type a quoted word --
# insert open and close double quotes
# and move to just after the open quote
"\C-x\"": "\"\"\C-b"
# insert a backslash (testing backslash escapes
# in sequences and macros)
"\C-x\\": "\\"
# Quote the current or previous word
"\C-xq": "\eb"\ef\""
# Add a binding to refresh the line, which is unbound
"\C-xr": redraw-current-line
# Edit variable on current line.
"\M-\C-v": "\C-a\C-k$\C-y\M-\C-e\C-a\C-y="
$endif

# use a visible bell if one is available
set bell-style visible

# don't strip characters to 7 bits when reading
set input-meta on

# allow iso-latin1 characters to be inserted rather
# than converted to prefix-meta sequences
set convert-meta off

# display characters with the eighth bit set directly
# rather than as meta-prefixed characters
set output-meta on

# if there are more than 150 possible completions for
# a word, ask the user if he wants to see all of them
set completion-query-items 150
```

```
# For FTP
$if Ftp
"\C-xg": "get \M-?"
"\C-xt": "put \M-?"
"\M-.": yank-last-arg
$endif
```

27.4 Bindable Readline Commands

This section describes Readline commands that may be bound to key sequences. Command names without an accompanying key sequence are unbound by default.

In the following descriptions, *point* refers to the current cursor position, and *mark* refers to a cursor position saved by the `set-mark` command. The text between the point and mark is referred to as the *region*.

27.4.1 Commands For Moving

`beginning-of-line` (C-a)

Move to the start of the current line.

`end-of-line` (C-e)

Move to the end of the line.

`forward-char` (C-f)

Move forward a character.

`backward-char` (C-b)

Move back a character.

`forward-word` (M-f)

Move forward to the end of the next word. Words are composed of letters and digits.

`backward-word` (M-b)

Move back to the start of the current or previous word. Words are composed of letters and digits.

`clear-screen` (C-l)

Clear the screen and redraw the current line, leaving the current line at the top of the screen.

`redraw-current-line` ()

Refresh the current line. By default, this is unbound.

27.4.2 Commands For Manipulating The History

`accept-line` (Newline or Return)

Accept the line regardless of where the cursor is. If this line is non-empty, it may be added to the history list for future recall with `add_history()`. If this line is a modified history line, the history line is restored to its original state.

`previous-history` (C-p)

Move ‘back’ through the history list, fetching the previous command.

next-history (C-n)

Move ‘forward’ through the history list, fetching the next command.

beginning-of-history (M-<)

Move to the first line in the history.

end-of-history (M->)

Move to the end of the input history, i.e., the line currently being entered.

reverse-search-history (C-r)

Search backward starting at the current line and moving ‘up’ through the history as necessary. This is an incremental search.

forward-search-history (C-s)

Search forward starting at the current line and moving ‘down’ through the history as necessary. This is an incremental search.

non-incremental-reverse-search-history (M-p)

Search backward starting at the current line and moving ‘up’ through the history as necessary using a non-incremental search for a string supplied by the user.

non-incremental-forward-search-history (M-n)

Search forward starting at the current line and moving ‘down’ through the history as necessary using a non-incremental search for a string supplied by the user.

history-search-forward ()

Search forward through the history for the string of characters between the start of the current line and the point. This is a non-incremental search. By default, this command is unbound.

history-search-backward ()

Search backward through the history for the string of characters between the start of the current line and the point. This is a non-incremental search. By default, this command is unbound.

yank-nth-arg (M-C-y)

Insert the first argument to the previous command (usually the second word on the previous line) at point. With an argument *n*, insert the *n*th word from the previous command (the words in the previous command begin with word 0). A negative argument inserts the *n*th word from the end of the previous command.

yank-last-arg (M-. or M-_)

Insert last argument to the previous command (the last word of the previous history entry). With an argument, behave exactly like **yank-nth-arg**. Successive calls to **yank-last-arg** move back through the history list, inserting the last argument of each line in turn.

27.4.3 Commands For Changing Text

`delete-char (C-d)`

Delete the character at point. If point is at the beginning of the line, there are no characters in the line, and the last character typed was not bound to `delete-char`, then return EOF.

`backward-delete-char (Rubout)`

Delete the character behind the cursor. A numeric argument means to kill the characters instead of deleting them.

`forward-backward-delete-char ()`

Delete the character under the cursor, unless the cursor is at the end of the line, in which case the character behind the cursor is deleted. By default, this is not bound to a key.

`quoted-insert (C-q or C-v)`

Add the next character typed to the line verbatim. This is how to insert key sequences like `C-q`, for example.

`tab-insert (M-TAB)`

Insert a tab character.

`self-insert (a, b, A, 1, !, ...)`

Insert yourself.

`transpose-chars (C-t)`

Drag the character before the cursor forward over the character at the cursor, moving the cursor forward as well. If the insertion point is at the end of the line, then this transposes the last two characters of the line. Negative arguments have no effect.

`transpose-words (M-t)`

Drag the word before point past the word after point, moving point past that word as well. If the insertion point is at the end of the line, this transposes the last two words on the line.

`upcase-word (M-u)`

Uppercase the current (or following) word. With a negative argument, uppercase the previous word, but do not move the cursor.

`downcase-word (M-l)`

Lowercase the current (or following) word. With a negative argument, lowercase the previous word, but do not move the cursor.

`capitalize-word (M-c)`

Capitalize the current (or following) word. With a negative argument, capitalize the previous word, but do not move the cursor.

`overwrite-mode ()`

Toggle overwrite mode. With an explicit positive numeric argument, switches to overwrite mode. With an explicit non-positive numeric argument, switches to insert mode. This command affects only `emacs` mode; `vi` mode does overwrite differently. Each call to `readline()` starts in insert mode.

In overwrite mode, characters bound to `self-insert` replace the text at point rather than pushing the text to the right. Characters bound to `backward-delete-char` replace the character before point with a space.

By default, this command is unbound.

27.4.4 Killing And Yanking

`kill-line` (C-k)

Kill the text from point to the end of the line.

`backward-kill-line` (C-x Rubout)

Kill backward to the beginning of the line.

`unix-line-discard` (C-u)

Kill backward from the cursor to the beginning of the current line.

`kill-whole-line` ()

Kill all characters on the current line, no matter where point is. By default, this is unbound.

`kill-word` (M-d)

Kill from point to the end of the current word, or if between words, to the end of the next word. Word boundaries are the same as `forward-word`.

`backward-kill-word` (M-~~DEL~~)

Kill the word behind point. Word boundaries are the same as `backward-word`.

`unix-word-rubout` (C-w)

Kill the word behind point, using white space as a word boundary. The killed text is saved on the kill-ring.

`delete-horizontal-space` ()

Delete all spaces and tabs around point. By default, this is unbound.

`kill-region` ()

Kill the text in the current region. By default, this command is unbound.

`copy-region-as-kill` ()

Copy the text in the region to the kill buffer, so it can be yanked right away. By default, this command is unbound.

`copy-backward-word` ()

Copy the word before point to the kill buffer. The word boundaries are the same as `backward-word`. By default, this command is unbound.

`copy-forward-word` ()

Copy the word following point to the kill buffer. The word boundaries are the same as `forward-word`. By default, this command is unbound.

`yank` (C-y)

Yank the top of the kill ring into the buffer at point.

`yank-pop` (M-y)

Rotate the kill-ring, and yank the new top. You can only do this if the prior command is `yank` or `yank-pop`.

27.4.5 Specifying Numeric Arguments

`digit-argument (M-0, M-1, ... M--)`

Add this digit to the argument already accumulating, or start a new argument.
`M--` starts a negative argument.

`universal-argument ()`

This is another way to specify an argument. If this command is followed by one or more digits, optionally with a leading minus sign, those digits define the argument. If the command is followed by digits, executing `universal-argument` again ends the numeric argument, but is otherwise ignored. As a special case, if this command is immediately followed by a character that is neither a digit or minus sign, the argument count for the next command is multiplied by four. The argument count is initially one, so executing this function the first time makes the argument count four, a second time makes the argument count sixteen, and so on. By default, this is not bound to a key.

27.4.6 Letting Readline Type For You

`complete (TAB)`

Attempt to perform completion on the text before point. The actual completion performed is application-specific. The default is filename completion.

`possible-completions (M-?)`

List the possible completions of the text before point.

`insert-completions (M-*)`

Insert all completions of the text before point that would have been generated by `possible-completions`.

`menu-complete ()`

Similar to `complete`, but replaces the word to be completed with a single match from the list of possible completions. Repeated execution of `menu-complete` steps through the list of possible completions, inserting each match in turn. At the end of the list of completions, the bell is rung (subject to the setting of `bell-style`) and the original text is restored. An argument of *n* moves *n* positions forward in the list of matches; a negative argument may be used to move backward through the list. This command is intended to be bound to `(TAB)`, but is unbound by default.

`delete-char-or-list ()`

Deletes the character under the cursor if not at the beginning or end of the line (like `delete-char`). If at the end of the line, behaves identically to `possible-completions`. This command is unbound by default.

27.4.7 Keyboard Macros

`start-kbd-macro (C-x ()`

Begin saving the characters typed into the current keyboard macro.

`end-kbd-macro (C-x))`

Stop saving the characters typed into the current keyboard macro and save the definition.

call-last-kbd-macro (C-x e)

Re-execute the last keyboard macro defined, by making the characters in the macro appear as if typed at the keyboard.

27.4.8 Some Miscellaneous Commands

re-read-init-file (C-x C-r)

Read in the contents of the *inputrc* file, and incorporate any bindings or variable assignments found there.

abort (C-g)

Abort the current editing command and ring the terminal's bell (subject to the setting of **bell-style**).

do-uppercase-version (M-a, M-b, M-x, ...)

If the metafied character x is lowercase, run the command that is bound to the corresponding uppercase character.

prefix-meta (ESC)

Metafy the next character typed. This is for keyboards without a meta key. Typing '**ESC** f' is equivalent to typing **M-f**.

undo (C-_ or C-x C-u)

Incremental undo, separately remembered for each line.

revert-line (M-r)

Undo all changes made to this line. This is like executing the **undo** command enough times to get back to the beginning.

tilde-expand (M-~)

Perform tilde expansion on the current word.

set-mark (C-@)

Set the mark to the point. If a numeric argument is supplied, the mark is set to that position.

exchange-point-and-mark (C-x C-x)

Swap the point with the mark. The current cursor position is set to the saved position, and the old cursor position is saved as the mark.

character-search (C-])

A character is read and point is moved to the next occurrence of that character. A negative count searches for previous occurrences.

character-search-backward (M-C-])

A character is read and point is moved to the previous occurrence of that character. A negative count searches for subsequent occurrences.

insert-comment (M-#)

Without a numeric argument, the value of the **comment-begin** variable is inserted at the beginning of the current line. If a numeric argument is supplied, this command acts as a toggle: if the characters at the beginning of the line do not match the value of **comment-begin**, the value is inserted, otherwise the characters in **comment-begin** are deleted from the beginning of the line. In either case, the line is accepted as if a newline had been typed.

dump-functions ()

Print all of the functions and their key bindings to the Readline output stream. If a numeric argument is supplied, the output is formatted in such a way that it can be made part of an *inputrc* file. This command is unbound by default.

dump-variables ()

Print all of the settable variables and their values to the Readline output stream. If a numeric argument is supplied, the output is formatted in such a way that it can be made part of an *inputrc* file. This command is unbound by default.

dump-macros ()

Print all of the Readline key sequences bound to macros and the strings they output. If a numeric argument is supplied, the output is formatted in such a way that it can be made part of an *inputrc* file. This command is unbound by default.

emacs-editing-mode (C-e)

When in *vi* command mode, this causes a switch to *emacs* editing mode.

vi-editing-mode (M-C-j)

When in *emacs* editing mode, this causes a switch to *vi* editing mode.

27.5 Readline vi Mode

While the Readline library does not have a full set of *vi* editing functions, it does contain enough to allow simple editing of the line. The Readline *vi* mode behaves as specified in the POSIX 1003.2 standard.

In order to switch interactively between *emacs* and *vi* editing modes, use the command *M-C-j* (bound to *emacs-editing-mode* when in *vi* mode and to *vi-editing-mode* in *emacs* mode). The Readline default is *emacs* mode.

When you enter a line in *vi* mode, you are already placed in ‘insertion’ mode, as if you had typed an ‘*i*’. Pressing *ESC* switches you into ‘command’ mode, where you can edit the text of the line with the standard *vi* movement keys, move to previous history lines with ‘*k*’ and subsequent lines with ‘*j*’, and so forth.

28 Using History Interactively

This chapter describes how to use the GNU History Library interactively, from a user's standpoint. It should be considered a user's guide.

28.1 History Expansion

The History library provides a history expansion feature that is similar to the history expansion provided by `csh`. This section describes the syntax used to manipulate the history information.

History expansions introduce words from the history list into the input stream, making it easy to repeat commands, insert the arguments to a previous command into the current input line, or fix errors in previous commands quickly.

History expansion takes place in two parts. The first is to determine which line from the history list should be used during substitution. The second is to select portions of that line for inclusion into the current one. The line selected from the history is called the *event*, and the portions of that line that are acted upon are called *words*. Various *modifiers* are available to manipulate the selected words. The line is broken into words in the same fashion that Bash does, so that several words surrounded by quotes are considered one word. History expansions are introduced by the appearance of the history expansion character, which is ‘!’ by default.

28.1.1 Event Designators

An event designator is a reference to a command line entry in the history list.

- ! Start a history substitution, except when followed by a space, tab, the end of the line, ‘=’ or ‘(’.
- `!n` Refer to command line *n*.
- `!-n` Refer to the command *n* lines back.
- `!!` Refer to the previous command. This is a synonym for ‘`-1`’.
- `!string` Refer to the most recent command starting with *string*.
- `?string[?]` Refer to the most recent command containing *string*. The trailing ‘?’ may be omitted if the *string* is followed immediately by a newline.
- `^string1^string2^` Quick Substitution. Repeat the last command, replacing *string1* with *string2*. Equivalent to `!!:s/string1/string2/`.
- `#!` The entire command line typed so far.

28.1.2 Word Designators

Word designators are used to select desired words from the event. A ‘:’ separates the event specification from the word designator. It may be omitted if the word designator begins with a ‘^’, ‘\$’, ‘*’, ‘-’, or ‘%’. Words are numbered from the beginning of the line, with the

first word being denoted by 0 (zero). Words are inserted into the current line separated by single spaces.

For example,

- !!** designates the preceding command. When you type this, the preceding command is repeated in toto.
- !!:\$** designates the last argument of the preceding command. This may be shortened to **!\$**.
- !fi:2** designates the second argument of the most recent command starting with the letters **fi**.

Here are the word designators:

- 0 (zero)** The 0th word. For many applications, this is the command word.
- n** The *n*th word.
- ^** The first argument; that is, word 1.
- \$** The last argument.
- %** The word matched by the most recent ‘?string?’ search.
- x-y** A range of words; ‘-y’ abbreviates ‘0-y’.
- *** All of the words, except the 0th. This is a synonym for ‘1-\$’. It is not an error to use ‘*’ if there is just one word in the event; the empty string is returned in that case.
- x*** Abbreviates ‘x-\$’
- x-** Abbreviates ‘x-\$’ like ‘x*’, but omits the last word.

If a word designator is supplied without an event specification, the previous command is used as the event.

28.1.3 Modifiers

After the optional word designator, you can add a sequence of one or more of the following modifiers, each preceded by a ‘:’:

- h** Remove a trailing pathname component, leaving only the head.
- t** Remove all leading pathname components, leaving the tail.
- r** Remove a trailing suffix of the form ‘*.suffix*’, leaving the basename.
- e** Remove all but the trailing suffix.
- p** Print the new command but do not execute it.

s/old/new/

Substitute *new* for the first occurrence of *old* in the event line. Any delimiter may be used in place of ‘/’. The delimiter may be quoted in *old* and *new* with a single backslash. If ‘&’ appears in *new*, it is replaced by *old*. A single backslash will quote the ‘&’. The final delimiter is optional if it is the last character on the input line.

- & Repeat the previous substitution.
- g Cause changes to be applied over the entire event line. Used in conjunction with ‘s’, as in `gs/old/new/`, or with ‘&’.

Appendix A Formatting Documentation

The GDB 4 release includes an already-formatted reference card, ready for printing with PostScript or Ghostscript, in the ‘`gdb`’ subdirectory of the main source directory¹. If you can use PostScript or Ghostscript with your printer, you can print the reference card immediately with ‘`refcard.ps`’.

The release also includes the source for the reference card. You can format it, using `TEx`, by typing:

```
make refcard.dvi
```

The GDB reference card is designed to print in *landscape* mode on US “letter” size paper; that is, on a sheet 11 inches wide by 8.5 inches high. You will need to specify this form of printing as an option to your DVI output program.

All the documentation for GDB comes as part of the machine-readable distribution. The documentation is written in Texinfo format, which is a documentation system that uses a single source file to produce both on-line information and a printed manual. You can use one of the Info formatting commands to create the on-line version of the documentation and `TEx` (or `texi2roff`) to typeset the printed version.

GDB includes an already formatted copy of the on-line Info version of this manual in the ‘`gdb`’ subdirectory. The main Info file is ‘`gdb-6.2.50_20041017/gdb/gdb.info`’, and it refers to subordinate files matching ‘`gdb.info*`’ in the same directory. If necessary, you can print out these files, or read them with any editor; but they are easier to read using the `info` subsystem in GNU Emacs or the standalone `info` program, available as part of the GNU Texinfo distribution.

If you want to format these Info files yourself, you need one of the Info formatting programs, such as `texinfo-format-buffer` or `makeinfo`.

If you have `makeinfo` installed, and are in the top level GDB source directory (‘`gdb-6.2.50_20041017`’, in the case of version 6.2.50·20041017), you can make the Info file by typing:

```
cd gdb
make gdb.info
```

If you want to typeset and print copies of this manual, you need `TEx`, a program to print its DVI output files, and ‘`texinfo.tex`’, the Texinfo definitions file.

`TEx` is a typesetting program; it does not print files directly, but produces output files called DVI files. To print a typeset document, you need a program to print DVI files. If your system has `TEx` installed, chances are it has such a program. The precise command to use depends on your system; `lpr -d` is common; another (for PostScript devices) is `dvips`. The DVI print command may require a file name without any extension or a ‘`.dvi`’ extension.

`TEx` also requires a macro definitions file called ‘`texinfo.tex`’. This file tells `TEx` how to typeset a document written in Texinfo format. On its own, `TEx` cannot either read or typeset a Texinfo file. ‘`texinfo.tex`’ is distributed with GDB and is located in the ‘`gdb-version-number/texinfo`’ directory.

If you have `TEx` and a DVI printer program installed, you can typeset and print this manual. First switch to the the ‘`gdb`’ subdirectory of the main source directory (for example, to ‘`gdb-6.2.50_20041017/gdb`’) and type:

¹ In ‘`gdb-6.2.50_20041017/gdb/refcard.ps`’ of the version 6.2.50·20041017 release.

```
make gdb.dvi
```

Then give ‘`gdb.dvi`’ to your DVI printing program.

Appendix B Installing GDB

GDB comes with a `configure` script that automates the process of preparing GDB for installation; you can then use `make` to build the `gdb` program.¹

The GDB distribution includes all the source code you need for GDB in a single directory, whose name is usually composed by appending the version number to ‘`gdb`’.

For example, the GDB version 6.2.50'20041017 distribution is in the ‘`gdb-6.2.50_20041017`’ directory. That directory contains:

```

gdb-6.2.50_20041017/configure (and supporting files)
    script for configuring GDB and all its supporting libraries

gdb-6.2.50_20041017/gdb
    the source specific to GDB itself

gdb-6.2.50_20041017/bfd
    source for the Binary File Descriptor library

gdb-6.2.50_20041017/include
    GNU include files

gdb-6.2.50_20041017/libiberty
    source for the ‘-liberty’ free software library

gdb-6.2.50_20041017/opcodes
    source for the library of opcode tables and disassemblers

gdb-6.2.50_20041017/readline
    source for the GNU command-line interface

gdb-6.2.50_20041017/glob
    source for the GNU filename pattern-matching subroutine

gdb-6.2.50_20041017/mmalloc
    source for the GNU memory-mapped malloc package

```

The simplest way to configure and build GDB is to run `configure` from the ‘`gdb-version-number`’ source directory, which in this example is the ‘`gdb-6.2.50_20041017`’ directory.

First switch to the ‘`gdb-version-number`’ source directory if you are not already in it; then run `configure`. Pass the identifier for the platform on which GDB will run as an argument.

For example:

```

cd gdb-6.2.50_20041017
./configure host
make

```

where `host` is an identifier such as ‘`sun4`’ or ‘`decstation`’, that identifies the platform where GDB will run. (You can often leave off `host`; `configure` tries to guess the correct value by examining your system.)

¹ If you have a more recent version of GDB than 6.2.50'20041017, look at the ‘`README`’ file in the sources; we may have improved the installation procedures since publishing this manual.

Running ‘`configure host`’ and then running `make` builds the ‘`bfd`’, ‘`readline`’, ‘`mmalloc`’, and ‘`libiberty`’ libraries, then `gdb` itself. The configured source files, and the binaries, are left in the corresponding source directories.

`configure` is a Bourne-shell (`/bin/sh`) script; if your system does not recognize this automatically when you run a different shell, you may need to run `sh` on it explicitly:

```
sh configure host
```

If you run `configure` from a directory that contains source directories for multiple libraries or programs, such as the ‘`gdb-6.2.50_20041017`’ source directory for version 6.2.50_20041017, `configure` creates configuration files for every directory level underneath (unless you tell it not to, with the ‘`--norecursion`’ option).

You should run the `configure` script from the top directory in the source tree, the ‘`gdb-version-number`’ directory. If you run `configure` from one of the subdirectories, you will configure only that subdirectory. That is usually not what you want. In particular, if you run the first `configure` from the ‘`gdb`’ subdirectory of the ‘`gdb-version-number`’ directory, you will omit the configuration of ‘`bfd`’, ‘`readline`’, and other sibling directories of the ‘`gdb`’ subdirectory. This leads to build errors about missing include files such as ‘`bfd/bfd.h`’.

You can install `gdb` anywhere; it has no hardwired paths. However, you should make sure that the shell on your path (named by the ‘`SHELL`’ environment variable) is publicly readable. Remember that GDB uses the shell to start your program—some systems refuse to let GDB debug child processes whose programs are not readable.

B.1 Compiling GDB in another directory

If you want to run GDB versions for several host or target machines, you need a different `gdb` compiled for each combination of host and target. `configure` is designed to make this easy by allowing you to generate each configuration in a separate subdirectory, rather than in the source directory. If your `make` program handles the ‘`VPATH`’ feature (GNU `make` does), running `make` in each of these directories builds the `gdb` program specified there.

To build `gdb` in a separate directory, run `configure` with the ‘`--srcdir`’ option to specify where to find the source. (You also need to specify a path to find `configure` itself from your working directory. If the path to `configure` would be the same as the argument to ‘`--srcdir`’, you can leave out the ‘`--srcdir`’ option; it is assumed.)

For example, with version 6.2.50_20041017, you can build GDB in a separate directory for a Sun 4 like this:

```
cd gdb-6.2.50_20041017
mkdir ..../gdb-sun4
cd ..../gdb-sun4
../gdb-6.2.50_20041017/configure sun4
make
```

When `configure` builds a configuration using a remote source directory, it creates a tree for the binaries with the same structure (and using the same names) as the tree under the source directory. In the example, you’d find the Sun 4 library ‘`libiberty.a`’ in the directory ‘`gdb-sun4/libiberty`’, and GDB itself in ‘`gdb-sun4/gdb`’.

Make sure that your path to the ‘`configure`’ script has just one instance of ‘`gdb`’ in it. If your path to ‘`configure`’ looks like ‘`..../gdb-6.2.50_20041017/gdb/configure`’, you

are configuring only one subdirectory of GDB, not the whole package. This leads to build errors about missing include files such as ‘`bfd/bfd.h`’.

One popular reason to build several GDB configurations in separate directories is to configure GDB for cross-compiling (where GDB runs on one machine—the *host*—while debugging programs that run on another machine—the *target*). You specify a cross-debugging target by giving the ‘`--target=target`’ option to `configure`.

When you run `make` to build a program or library, you must run it in a configured directory—whatever directory you were in when you called `configure` (or one of its subdirectories).

The `Makefile` that `configure` generates in each source directory also runs recursively. If you type `make` in a source directory such as ‘`gdb-6.2.50_20041017`’ (or in a separate configured directory configured with ‘`--srcdir=dirname/gdb-6.2.50_20041017`’), you will build all the required libraries, and then build GDB.

When you have multiple hosts or targets configured in separate directories, you can run `make` on them in parallel (for example, if they are NFS-mounted on each of the hosts); they will not interfere with each other.

B.2 Specifying names for hosts and targets

The specifications used for hosts and targets in the `configure` script are based on a three-part naming scheme, but some short predefined aliases are also supported. The full naming scheme encodes three pieces of information in the following pattern:

`architecture-vendor-os`

For example, you can use the alias `sun4` as a *host* argument, or as the value for *target* in a `--target=target` option. The equivalent full name is ‘`sparc-sun-sunos4`’.

The `configure` script accompanying GDB does not provide any query facility to list all supported host and target names or aliases. `configure` calls the Bourne shell script `config.sub` to map abbreviations to full names; you can read the script, if you wish, or you can use it to test your guesses on abbreviations—for example:

```
% sh config.sub i386-linux
i386-pc-linux-gnu
% sh config.sub alpha-linux
alpha-unknown-linux-gnu
% sh config.sub hp9k700
hppa1.1-hp-hpux
% sh config.sub sun4
sparc-sun-sunos4.1.1
% sh config.sub sun3
m68k-sun-sunos4.1.1
% sh config.sub i986v
Invalid configuration 'i986v': machine 'i986v' not recognized
```

`config.sub` is also distributed in the GDB source directory (‘`gdb-6.2.50_20041017`’, for version 6.2.50’20041017).

B.3 `configure` options

Here is a summary of the `configure` options and arguments that are most often useful for building GDB. `configure` also has several other options not listed here. See Info file ‘`configure.info`’, node ‘What Configure Does’, for a full explanation of `configure`.

```
configure [--help]
          [--prefix=dir]
          [--exec-prefix=dir]
          [--srcdir=dirname]
          [--norecursion] [--rm]
          [--target=target]
          host
```

You may introduce options with a single ‘-’ rather than ‘--’ if you prefer; but you may abbreviate option names if you use ‘--’.

--help Display a quick summary of how to invoke `configure`.

--prefix=dir

Configure the source to install programs and files under directory ‘*dir*’.

--exec-prefix=dir

Configure the source to install programs under directory ‘*dir*’.

--srcdir=dirname

Warning: using this option requires GNU make, or another make that implements the VPATH feature.

Use this option to make configurations in directories separate from the GDB source directories. Among other things, you can use this to build (or maintain) several configurations simultaneously, in separate directories. `configure` writes configuration specific files in the current directory, but arranges for them to use the source in the directory *dirname*. `configure` creates directories under the working directory in parallel to the source directories below *dirname*.

--norecursion

Configure only the directory level where `configure` is executed; do not propagate configuration to subdirectories.

--target=target

Configure GDB for cross-debugging programs running on the specified *target*. Without this option, GDB is configured to debug programs that run on the same machine (*host*) as GDB itself.

There is no convenient way to generate a list of all available targets.

host . . . Configure GDB to run on the specified *host*.

There is no convenient way to generate a list of all available hosts.

There are many other options available as well, but they are generally needed for special purposes only.

Appendix C Maintenance Commands

In addition to commands intended for GDB users, GDB includes a number of commands intended for GDB developers. These commands are provided here for reference.

maint info breakpoints

Using the same format as ‘**info breakpoints**’, display both the breakpoints you’ve set explicitly, and those GDB is using for internal purposes. Internal breakpoints are shown with negative breakpoint numbers. The type column identifies what kind of breakpoint is shown:

breakpoint

Normal, explicitly set breakpoint.

watchpoint

Normal, explicitly set watchpoint.

longjmp Internal breakpoint, used to handle correctly stepping through **longjmp** calls.

longjmp resume

Internal breakpoint at the target of a **longjmp**.

until Temporary internal breakpoint used by the GDB **until** command.

finish Temporary internal breakpoint used by the GDB **finish** command.

shlib events

Shared library events.

maint internal-error

maint internal-warning

Cause GDB to call the internal function **internal_error** or **internal_warning** and hence behave as though an internal error or internal warning has been detected. In addition to reporting the internal problem, these functions give the user the opportunity to either quit GDB or create a core file of the current GDB session.

```
(gdb) maint internal-error testing, 1, 2
.../maint.c:121: internal-error: testing, 1, 2
A problem internal to GDB has been detected. Further
debugging may prove unreliable.
Quit this debugging session? (y or n) n
Create a core file? (y or n) n
(gdb)
```

Takes an optional parameter that is used as the text of the error or warning message.

maint print dummy-frames

Prints the contents of GDB’s internal dummy-frame stack.

```
(gdb) b add
...
(gdb) print add(2,3)
Breakpoint 2, add (a=2, b=3) at ...
58    return (a + b);
```

```
The program being debugged stopped while in a function called from GDB.
...
(gdb) maint print dummy-frames
0x1a57c80: pc=0x01014068 fp=0x0200bddc sp=0x0200bdd6
  top=0x0200bdd4 id={stack=0x200bddc,code=0x101405c}
  call_lo=0x01014000 call_hi=0x01014001
(gdb)
```

Takes an optional file parameter.

```
maint print registers
maint print raw-registers
maint print cooked-registers
maint print register-groups
```

Print GDB's internal register data structures.

The command `maint print raw-registers` includes the contents of the raw register cache; the command `maint print cooked-registers` includes the (cooked) value of all registers; and the command `maint print register-groups` includes the groups that each register is a member of. See section “Registers” in *GDB Internals*.

Takes an optional file parameter.

```
maint print reggroups
```

Print GDB's internal register group data structures.

Takes an optional file parameter.

```
(gdb) maint print reggroups
  Group      Type
  general    user
  float      user
  all         user
  vector     user
  system     user
  save       internal
  restore    internal
```

```
maint set profile
```

```
maint show profile
```

Control profiling of GDB.

Profiling will be disabled until you use the ‘`maint set profile`’ command to enable it. When you enable profiling, the system will begin collecting timing and execution count data; when you disable profiling or exit GDB, the results will be written to a log file. Remember that if you use profiling, GDB will overwrite the profiling log file (often called ‘`gmon.out`’). If you have a record of important profiling data in a ‘`gmon.out`’ file, be sure to move it to a safe location.

Configuring with ‘`--enable-profiling`’ arranges for GDB to be compiled with the ‘`-pg`’ compiler option.

```
maint set dwarf2 max-cache-age
maint show dwarf2 max-cache-age
```

Control the DWARF 2 compilation unit cache.

In object files with inter-compilation-unit references, such as those produced by the GCC option ‘`-feliminate-dwarf2-dups`’, the DWARF 2 reader needs to frequently refer to previously read compilation units. This setting controls how long a compilation unit will remain in the cache if it is not referenced. Setting it to zero disables caching, which will slow down GDB startup but reduce memory consumption.

Appendix D GDB Remote Serial Protocol

D.1 Overview

There may be occasions when you need to know something about the protocol—for example, if there is only one serial port to your target machine, you might want your program to do something special if it recognizes a packet meant for GDB.

In the examples below, ‘->’ and ‘<-’ are used to indicate transmitted and received data respectfully.

All GDB commands and responses (other than acknowledgments) are sent as a *packet*. A *packet* is introduced with the character ‘\$’, the actual *packet-data*, and the terminating character ‘#’ followed by a two-digit *checksum*:

```
$packet-data#checksum
```

The two-digit *checksum* is computed as the modulo 256 sum of all characters between the leading ‘\$’ and the trailing ‘#’ (an eight bit unsigned checksum).

Implementors should note that prior to GDB 5.0 the protocol specification also included an optional two-digit *sequence-id*:

```
$sequence-id:packet-data#checksum
```

That *sequence-id* was appended to the acknowledgment. GDB has never output *sequence-ids*. Stubs that handle packets added since GDB 5.0 must not accept *sequence-id*.

When either the host or the target machine receives a packet, the first response expected is an acknowledgment: either ‘+’ (to indicate the package was received correctly) or ‘-’ (to request retransmission):

```
-> $packet-data#checksum
<- +
```

The host (GDB) sends *commands*, and the target (the debugging stub incorporated in your program) sends a *response*. In the case of step and continue *commands*, the response is only sent when the operation has completed (the target has again stopped).

packet-data consists of a sequence of characters with the exception of ‘#’ and ‘\$’ (see ‘X’ packet for additional exceptions).

Fields within the packet should be separated using ‘,’ ‘;’ or ‘:’. Except where otherwise noted all numbers are represented in HEX with leading zeros suppressed.

Implementors should note that prior to GDB 5.0, the character ‘:’ could not appear as the third character in a packet (as it would potentially conflict with the *sequence-id*).

Response data can be run-length encoded to save space. A ‘*’ means that the next character is an ASCII encoding giving a repeat count which stands for that many repetitions of the character preceding the ‘*’. The encoding is $n+29$, yielding a printable character where $n \geq 3$ (which is where rle starts to win). The printable characters ‘\$’, ‘#’, ‘+’ and ‘-’ or with a numeric value greater than 126 should not be used.

So:

```
"0* "
```

means the same as "0000".

The error response returned for some packets includes a two character error number. That number is not well defined.

For any *command* not supported by the stub, an empty response ('\$#00') should be returned. That way it is possible to extend the protocol. A newer GDB can tell if a packet is supported based on that response.

A stub is required to support the ‘g’, ‘G’, ‘m’, ‘M’, ‘c’, and ‘s’ commands. All other commands are optional.

D.2 Packets

The following table provides a complete list of all currently defined *commands* and their corresponding response *data*.

! — extended mode

Enable extended mode. In extended mode, the remote server is made persistent. The ‘R’ packet is used to restart the program being debugged.

Reply:

‘OK’ The remote target both supports and has enabled extended mode.

? — last signal

Indicate the reason the target halted. The reply is the same as for step and continue.

Reply: See [Section D.3 \[Stop Reply Packets\]](#), page 301, for the reply specifications.

a — reserved

Reserved for future use.

Aarglen,argnum,arg,... — set program arguments (deprecated**)**

Initialized ‘`argv[]`’ array passed into program. `arglen` specifies the number of bytes in the hex encoded byte stream `arg`. See `gdbserver` for more details.

Reply:

‘OK’

‘ENN’

bbaud — set baud (deprecated**)**

Change the serial line speed to `baud`.

JTC: *When does the transport layer state change? When it's received, or after the ACK is transmitted. In either case, there are problems if the command or the acknowledgment packet is dropped.*

Stan: *If people really wanted to add something like this, and get it working for the first time, they ought to modify ser-unix.c to send some kind of out-of-band message to a specially-setup stub and have the switch happen "in between" packets, so that from remote protocol's point of view, nothing actually happened.*

Baddr,mode — set breakpoint (deprecated**)**

Set (`mode` is ‘S’) or clear (`mode` is ‘C’) a breakpoint at `addr`.

This packet has been replaced by the ‘Z’ and ‘z’ packets (see [\[insert breakpoint or watchpoint packet\]](#), page 299).

caddr — continue

addr is address to resume. If **addr** is omitted, resume at current address.

 Reply: See [Section D.3 \[Stop Reply Packets\]](#), page 301, for the reply specifications.

Csig;addr — continue with signal

 Continue with signal **sig** (hex signal number). If ;**addr** is omitted, resume at same address.

 Reply: See [Section D.3 \[Stop Reply Packets\]](#), page 301, for the reply specifications.

d — toggle debug (**deprecated**)

 Toggle debug flag.

D — detach

 Detach GDB from the remote system. Sent to the remote target before GDB disconnects via the **detach** command.

 Reply:

 ‘**no response**’

 GDB does not check for any response after sending this packet.

e — reserved

 Reserved for future use.

E — reserved

 Reserved for future use.

f — reserved

 Reserved for future use.

FRC,EE,CF;XX — Reply to target’s F packet.

 This packet is send by GDB as reply to a F request packet sent by the target.

 This is part of the File-I/O protocol extension. See [Section D.7 \[File-I/O remote protocol extension\]](#), page 306, for the specification.

g — read registers

 Read general registers.

 Reply:

 ‘**xx...**’ Each byte of register data is described by two hex digits.
 The bytes with the register are transmitted in target byte order. The size of each register and their position within the ‘g’ packet are determined by the GDB internal macros *DEPRECATED_REGISTER_RAW_SIZE* and *REGISTER_NAME* macros. The specification of several standard g packets is specified below.

 ‘**ENN**’ for an error.

GXX... — write regs

 See [\[read registers packet\]](#), page 295, for a description of the XX... data.

 Reply:

‘OK’ for success
 ‘ENN’ for an error

h — reserved

Reserved for future use.

Hct... — set thread

Set thread for subsequent operations (‘m’, ‘M’, ‘g’, ‘G’, et.al.). *c* depends on the operation to be performed: it should be ‘c’ for step and continue operations, ‘g’ for other operations. The thread designator *t...* may be -1, meaning all the threads, a thread number, or zero which means pick any thread.

Reply:

‘OK’ for success
 ‘ENN’ for an error

iaddr,nnn — cycle step (**draft**)

Step the remote target by a single clock cycle. If ,*nnn* is present, cycle step *nnn* cycles. If *addr* is present, cycle step starting at that address.

I — signal then cycle step (**reserved**)

See [step with signal packet], page 298. See [cycle step packet], page 296.

j — reserved

Reserved for future use.

J — reserved

Reserved for future use.

k — kill request

FIXME: There is no description of how to operate when a specific thread context has been selected (i.e. does ‘k’ kill only that thread?).

K — reserved

Reserved for future use.

l — reserved

Reserved for future use.

L — reserved

Reserved for future use.

maddr,length — read memory

Read *length* bytes of memory starting at address *addr*. Neither GDB nor the stub assume that sized memory transfers are assumed using word aligned accesses.

FIXME: A word aligned memory transfer mechanism is needed.

Reply:

‘XX...’ XX... is mem contents. Can be fewer bytes than requested if able to read only part of the data. Neither GDB nor the stub assume that sized memory transfers are assumed using word aligned accesses.
FIXME: A word aligned memory transfer mechanism is needed.

‘ENN’ *NN* is errno

Maddr,length:XX... — write mem

Write *length* bytes of memory starting at address *addr*. *XX...* is the data.

Reply:

‘OK’ for success

‘ENN’ for an error (this includes the case where only part of the data was written).

n — reserved

Reserved for future use.

N — reserved

Reserved for future use.

o — reserved

Reserved for future use.

0 — reserved

phex number of register — read register packet

See [read registers packet], page 295, for a description of how the returned register value is encoded.

Reply:

‘XX...’ the register’s value

‘ENN’ for an error

‘’ Indicating an unrecognized query.

Pn...=r... — write register

Write register *n...* with value *r...*, which contains two hex digits for each byte in the register (target byte order).

Reply:

‘OK’ for success

‘ENN’ for an error

qquery — general query

Request info about *query*. In general GDB queries have a leading upper case letter. Custom vendor queries should use a company prefix (in lower case) ex: ‘qfsf.var’. *query* may optionally be followed by a ‘,’ or ‘;’ separated list. Stubs must ensure that they match the full *query* name.

Reply:

‘XX...’ Hex encoded data from query. The reply can not be empty.

‘ENN’ error reply

‘’ Indicating an unrecognized query.

Qvar=val — general set

Set value of *var* to *val*.

See [general query packet], page 297, for a discussion of naming conventions.

r — reset (deprecated)

Reset the entire system.

RXX — remote restart

Restart the program being debugged. XX, while needed, is ignored. This packet is only available in extended mode.

Reply:

'no reply'

The 'R' packet has no reply.

saddr — step

addr is address to resume. If addr is omitted, resume at same address.

Reply: See [Section D.3 \[Stop Reply Packets\], page 301](#), for the reply specifications.

Ssig;addr — step with signal

Like 'C' but step not continue.

Reply: See [Section D.3 \[Stop Reply Packets\], page 301](#), for the reply specifications.

taddr:PP,MM — search

Search backwards starting at address addr for a match with pattern PP and mask MM. PP and MM are 4 bytes. addr must be at least 3 digits.

TXX — thread alive

Find out if the thread XX is alive.

Reply:

'OK' thread is still alive

'ENN' thread is dead

u — reserved

Reserved for future use.

U — reserved

Reserved for future use.

v — verbose packet prefix

Packets starting with v are identified by a multi-letter name, up to the first ; or ? (or the end of the packet).

vCont[action[:tid]]... — extended resume

Resume the inferior. Different actions may be specified for each thread. If an action is specified with no tid, then it is applied to any threads that don't have a specific action specified; if no default action is specified then other threads should remain stopped. Specifying multiple default actions is an error; specifying no actions is also an error. Thread IDs are specified in hexadecimal. Currently supported actions are:

c Continue.

Csig Continue with signal sig. sig should be two hex digits.

s Step.

Ssig Step with signal *sig*. *sig* should be two hex digits.

The optional *addr* argument normally associated with these packets is not supported in vCont.

Reply: See [Section D.3 \[Stop Reply Packets\]](#), page 301, for the reply specifications.

vCont? — extended resume query

Query support for the vCont packet.

Reply:

‘vCont[;action]...’

The vCont packet is supported. Each *action* is a supported command in the vCont packet.

‘’ The vCont packet is not supported.

V — reserved

Reserved for future use.

w — reserved

Reserved for future use.

W — reserved

Reserved for future use.

x — reserved

Reserved for future use.

Xaddr,length:XX... — write mem (binary)

addr is address, *length* is number of bytes, *XX...* is binary data. The characters \$, #, and 0x7d are escaped using 0x7d.

Reply:

‘OK’ for success

‘ENN’ for an error

y — reserved

Reserved for future use.

Y reserved

Reserved for future use.

ztype,addr,length — remove breakpoint or watchpoint (**draft**)

Ztype,addr,length — insert breakpoint or watchpoint (**draft**)

Insert (Z) or remove (z) a *type* breakpoint or watchpoint starting at address *address* and covering the next *length* bytes.

Each breakpoint and watchpoint packet *type* is documented separately.

Implementation notes: A remote target shall return an empty string for an unrecognized breakpoint or watchpoint packet type. A remote target shall support either both or neither of a given Ztype... and ztype... packet pair. To avoid

potential problems with duplicate packets, the operations should be implemented in an idempotent way.

`z0,addr,length` — remove memory breakpoint (**draft**)

`Z0,addr,length` — insert memory breakpoint (**draft**)

Insert (`Z0`) or remove (`z0`) a memory breakpoint at address `addr` of size `length`.

A memory breakpoint is implemented by replacing the instruction at `addr` with a software breakpoint or trap instruction. The `length` is used by targets that indicates the size of the breakpoint (in bytes) that should be inserted (e.g., the ARM and MIPS can insert either a 2 or 4 byte breakpoint).

Implementation note: It is possible for a target to copy or move code that contains memory breakpoints (e.g., when implementing overlays). The behavior of this packet, in the presence of such a target, is not defined.

Reply:

‘OK’ success

‘’ not supported

‘ENN’ for an error

`z1,addr,length` — remove hardware breakpoint (**draft**)

`Z1,addr,length` — insert hardware breakpoint (**draft**)

Insert (`Z1`) or remove (`z1`) a hardware breakpoint at address `addr` of size `length`.

A hardware breakpoint is implemented using a mechanism that is not dependant on being able to modify the target’s memory.

Implementation note: A hardware breakpoint is not affected by code movement.

Reply:

‘OK’ success

‘’ not supported

‘ENN’ for an error

`z2,addr,length` — remove write watchpoint (**draft**)

`Z2,addr,length` — insert write watchpoint (**draft**)

Insert (`Z2`) or remove (`z2`) a write watchpoint.

Reply:

‘OK’ success

‘’ not supported

‘ENN’ for an error

`z3,addr,length` — remove read watchpoint (**draft**)

`Z3,addr,length` — insert read watchpoint (**draft**)

Insert (`Z3`) or remove (`z3`) a read watchpoint.

Reply:

‘OK’ success

‘’ not supported

‘ENN’ for an error

z4,addr,length — remove access watchpoint (**draft**)

Z4,addr,length — insert access watchpoint (**draft**)

Insert (**Z4**) or remove (**z4**) an access watchpoint.

Reply:

‘OK’ success

‘’ not supported

‘ENN’ for an error

D.3 Stop Reply Packets

The ‘C’, ‘c’, ‘S’, ‘s’ and ‘?’ packets can receive any of the below as a reply. In the case of the ‘C’, ‘c’, ‘S’ and ‘s’ packets, that reply is only returned when the target halts. In the below the exact meaning of ‘signal number’ is poorly defined. In general one of the UNIX signal numbering conventions is used.

‘SAA’ AA is the signal number

‘TAA n... : r... ; n... : r... ; n... : r... ;’

AA = two hex digit signal number; n... = register number (hex), r... = target byte ordered register contents, size defined by DEPRECATED_REGISTER_RAW_SIZE; n... = ‘thread’, r... = thread process ID, this is a hex integer; n... = (‘watch’ | ‘rwatch’ | ‘awatch’, r... = data address, this is a hex integer; n... = other string not starting with valid hex digit. GDB should ignore this n..., r... pair and go on to the next. This way we can extend the protocol.

‘WAA’

The process exited, and AA is the exit status. This is only applicable to certain targets.

‘XAA’

The process terminated with signal AA.

‘0XX...’

XX... is hex encoding of ASCII data. This can happen at any time while the program is running and the debugger should continue to wait for ‘W’, ‘T’, etc.

‘Fcall-id,parameter...’

call-id is the identifier which says which host system call should be called. This is just the name of the function. Translation into the correct system call is only applicable as it’s defined in GDB. See [Section D.7 \[File-I/O remote protocol extension\], page 306](#), for a list of implemented system calls.

parameter... is a list of parameters as defined for this very system call.

The target replies with this packet when it expects GDB to call a host system call on behalf of the target. GDB replies with an appropriate F packet and keeps up waiting for the next reply packet from the target. The latest ‘C’, ‘c’, ‘S’ or ‘s’ action is expected to be continued. See [Section D.7 \[File-I/O remote protocol extension\], page 306](#), for more details.

D.4 General Query Packets

The following set and query packets have already been defined.

qC — current thread

Return the current thread id.

Reply:

‘QC pid ’ Where pid is a HEX encoded 16 bit process id.

‘*’ Any other reply implies the old pid.

qfThreadInfo – all thread ids

qsThreadInfo

Obtain a list of active thread ids from the target (OS). Since there may be too many active threads to fit into one reply packet, this query works iteratively: it may require more than one query/reply sequence to obtain the entire list of threads. The first query of the sequence will be the **qfThreadInfo** query; subsequent queries in the sequence will be the **qsThreadInfo** query.

NOTE: replaces the **qL** query (see below).

Reply:

‘mid’ A single thread id

‘mid, id...’
a comma-separated list of thread ids

‘1’ (lower case ‘el’) denotes end of list.

In response to each query, the target will reply with a list of one or more thread ids, in big-endian hex, separated by commas. GDB will respond to each reply with a request for more thread ids (using the **qs** form of the query), until the target responds with 1 (lower-case el, for ‘last’).

qThreadExtraInfo, id — extra thread info

Where id is a thread-id in big-endian hex. Obtain a printable string description of a thread’s attributes from the target OS. This string may contain anything that the target OS thinks is interesting for GDB to tell the user about the thread. The string is displayed in GDB’s ‘info threads’ display. Some examples of possible thread extra info strings are “Runnable”, or “Blocked on Mutex”.

Reply:

‘XX...’ Where XX... is a hex encoding of ASCII data, comprising the printable string containing the extra information about the thread’s attributes.

qLstartflagthreadcountnextthread — query *LIST* or *threadLIST* (**deprecated**)

Obtain thread information from RTOS. Where: *startflag* (one hex digit) is one to indicate the first query and zero to indicate a subsequent query; *threadcount* (two hex digits) is the maximum number of threads the response packet can contain; and *nextthread* (eight hex digits), for subsequent queries (*startflag* is zero), is returned in the response as *argthread*.

NOTE: this query is replaced by the `qfThreadInfo` query (see above).

Reply:

`'qMcount done argthread thread...'`

Where: `count` (two hex digits) is the number of threads being returned; `done` (one hex digit) is zero to indicate more threads and one indicates no further threads; `argthreadid` (eight hex digits) is `nextthread` from the request packet; `thread...` is a sequence of thread IDs from the target. `threadid` (eight hex digits). See `remote.c:parse_threadlist_response()`.

`qCRC:addr,length` — compute CRC of memory block

Reply:

`'ENN'` An error (such as memory fault)

`'CCRC32'` A 32 bit cyclic redundancy check of the specified memory region.

`qOffsets` — query sect offs

Get section offsets that the target used when re-locating the downloaded image.

Note: while a Bss offset is included in the response, GDB ignores this and instead applies the Data offset to the Bss section.

Reply:

`'Text=xxx;Data=yyy;Bss=zzz'`

`qPmodethreadid` — thread info request

Returns information on `threadid`. Where: `mode` is a hex encoded 32 bit mode; `threadid` is a hex encoded 64 bit thread ID.

Reply:

`'*'`

See `remote.c:remote_unpack_thread_info_response()`.

`qRcmd,command` — remote command

`command` (hex encoded) is passed to the local interpreter for execution. Invalid commands should be reported using the output string. Before the final result packet, the target may also respond with a number of intermediate `Output` console output packets. *Implementors should note that providing access to a stubs's interpreter may have security implications.*

Reply:

`'OK'` A command response with no output.

`'OUTPUT'` A command response with the hex encoded output string `OUTPUT`.

`'ENN'` Indicate a badly formed request.

`'',''` When '`q`'`Rcmd`' is not recognized.

`qSymbol::` — symbol lookup

Notify the target that GDB is prepared to serve symbol lookup requests. Accept requests from the target for the values of symbols.

Reply:

'OK' The target does not need to look up any (more) symbols.

'qSymbol:*sym_name*'

The target requests the value of symbol *sym_name* (hex encoded). GDB may provide the value by using the `qSymbol:sym_value:sym_name` message, described below.

`qSymbol:sym_value:sym_name` — symbol value

Set the value of *sym_name* to *sym_value*.

sym_name (hex encoded) is the name of a symbol whose value the target has previously requested.

sym_value (hex) is the value for symbol *sym_name*. If GDB cannot supply a value for *sym_name*, then this field will be empty.

Reply:

'OK' The target does not need to look up any (more) symbols.

'qSymbol:*sym_name*'

The target requests the value of a new symbol *sym_name* (hex encoded). GDB will continue to supply the values of symbols (if available), until the target ceases to request them.

`qPart:object:read:annex:offset,length` — read special data

Read uninterpreted bytes from the target's special data area identified by the keyword *object*. Request *length* bytes starting at *offset* bytes into the data. The content and encoding of *annex* is specific to the object; it can supply additional details about what data to access.

Here are the specific requests of this form defined so far. All '`qPart:object:read:...`' requests use the same reply formats, listed below.

`qPart:auxv:read::offset,length`

Access the target's auxiliary vector. See [Section 8.13 \[Auxiliary Vector\], page 80](#). Note *annex* must be empty.

Reply:

OK The *offset* in the request is at the end of the data. There is no more data to be read.

XX... Hex encoded data bytes read. This may be fewer bytes than the *length* in the request.

E00 The request was malformed, or *annex* was invalid.

Enn The offset was invalid, or there was an error encountered reading the data. *nn* is a hex-encoded *errno* value.

"" (empty)

An empty reply indicates the *object* or *annex* string was not recognized by the stub.

`qPart:object:write:annex:offset:data...`

Write uninterpreted bytes into the target's special data area identified by the keyword *object*, starting at *offset* bytes into the data. *data...* is the hex-

encoded data to be written. The content and encoding of *annex* is specific to the object; it can supply additional details about what data to access.

No requests of this form are presently in use. This specification serves as a placeholder to document the common format that new specific request specifications ought to use.

Reply:

nn *nn* (hex encoded) is the number of bytes written. This may be fewer bytes than supplied in the request.

E00 The request was malformed, or *annex* was invalid.

Enn The offset was invalid, or there was an error encountered writing the data. *nn* is a hex-encoded **errno** value.

"" (empty)

An empty reply indicates the *object* or *annex* string was not recognized by the stub, or that the object does not support writing.

qPart:*object*:*operation*: . . .

Requests of this form may be added in the future. When a stub does not recognize the *object* keyword, or its support for *object* does not recognize the *operation* keyword, the stub must respond with an empty packet.

D.5 Register Packet Format

The following ‘g’/‘G’ packets have previously been defined. In the below, some thirty-two bit registers are transferred as sixty-four bits. Those registers should be zero/sign extended (which?) to fill the space allocated. Register bytes are transferred in target byte order. The two nibbles within a register byte are transferred most-significant - least-significant.

MIPS32

All registers are transferred as thirty-two bit quantities in the order: 32 general-purpose; sr; lo; hi; bad; cause; pc; 32 floating-point registers; fsr; fir; fp.

MIPS64

All registers are transferred as sixty-four bit quantities (including thirty-two bit registers such as **sr**). The ordering is the same as MIPS32.

D.6 Examples

Example sequence of a target being re-started. Notice how the restart does not get any direct output:

```
-> R00
<- +
target restarts
-> ?
<- +
<- T001:1234123412341234
-> +
```

Example sequence of a target being stepped by a single instruction:

```

-> G1445...
<- +
-> s
<- +
time passes
<- T001:1234123412341234
-> +
-> g
<- +
<- 1455...
-> +

```

D.7 File-I/O remote protocol extension

D.7.1 File-I/O Overview

The File I/O remote protocol extension (short: File-I/O) allows the target to use the hosts file system and console I/O when calling various system calls. System calls on the target system are translated into a remote protocol packet to the host system which then performs the needed actions and returns with an adequate response packet to the target system. This simulates file system operations even on targets that lack file systems.

The protocol is defined host- and target-system independent. It uses it's own independent representation of datatypes and values. Both, GDB and the target's GDB stub are responsible for translating the system dependent values into the unified protocol values when data is transmitted.

The communication is synchronous. A system call is possible only when GDB is waiting for the ‘C’, ‘c’, ‘S’ or ‘s’ packets. While GDB handles the request for a system call, the target is stopped to allow deterministic access to the target’s memory. Therefore File-I/O is not interruptible by target signals. It is possible to interrupt File-I/O by a user interrupt (Ctrl-C), though.

The target’s request to perform a host system call does not finish the latest ‘C’, ‘c’, ‘S’ or ‘s’ action. That means, after finishing the system call, the target returns to continuing the previous activity (continue, step). No additional continue or step request from GDB is required.

```

(gdb) continue
<- target requests 'system call X'
target is stopped, GDB executes system call
-> GDB returns result
... target continues, GDB returns to wait for the target
<- target hits breakpoint and sends a Txx packet

```

The protocol is only used for files on the host file system and for I/O on the console. Character or block special devices, pipes, named pipes or sockets or any other communication method on the host system are not supported by this protocol.

D.7.2 Protocol basics

The File-I/O protocol uses the F packet, as request as well as as reply packet. Since a File-I/O system call can only occur when GDB is waiting for the continuing or stepping target, the File-I/O request is a reply that GDB has to expect as a result of a former ‘C’, ‘c’, ‘S’ or ‘s’ packet. This F packet contains all information needed to allow GDB to call the appropriate host system call:

- A unique identifier for the requested system call.
- All parameters to the system call. Pointers are given as addresses in the target memory address space. Pointers to strings are given as pointer/length pair. Numerical values are given as they are. Numerical control values are given in a protocol specific representation.

At that point GDB has to perform the following actions.

- If parameter pointer values are given, which point to data needed as input to a system call, GDB requests this data from the target with a standard **m** packet request. This additional communication has to be expected by the target implementation and is handled as any other **m** packet.
- GDB translates all value from protocol representation to host representation as needed. Datatypes are coerced into the host types.
- GDB calls the system call
- It then coerces datatypes back to protocol representation.
- If pointer parameters in the request packet point to buffer space in which a system call is expected to copy data to, the data is transmitted to the target using a **M** or **X** packet. This packet has to be expected by the target implementation and is handled as any other **M** or **X** packet.

Eventually GDB replies with another **F** packet which contains all necessary information for the target to continue. This at least contains

- Return value.
- **errno**, if has been changed by the system call.
- “Ctrl-C” flag.

After having done the needed type and value coercion, the target continues the latest continue or step action.

D.7.3 The F request packet

The F request packet has the following format:

F*call-id,parameter...*

call-id is the identifier to indicate the host system call to be called. This is just the name of the function.

parameter... are the parameters to the system call.

Parameters are hexadecimal integer values, either the real values in case of scalar datatypes, as pointers to target buffer space in case of compound datatypes and unspecified memory areas or as pointer/length pairs in case of string parameters. These are appended to the call-id, each separated from its predecessor by a comma. All values are transmitted in ASCII string representation, pointer/length pairs separated by a slash.

D.7.4 The F reply packet

The F reply packet has the following format:

F*retcode,errno,Ctrl-C flag;call specific attachment*

retcode is the return code of the system call as hexadecimal value.

`errno` is the `errno` set by the call, in protocol specific representation. This parameter can be omitted if the call was successful.

Ctrl-C flag is only send if the user requested a break. In this case, `errno` must be send as well, even if the call was successful. The *Ctrl-C flag* itself consists of the character 'C':

`F0,0,C`

or, if the call was interupted before the host call has been performed:

`F-1,4,C`

assuming 4 is the protocol specific representation of `EINTR`.

D.7.5 Memory transfer

Structured data which is transferred using a memory read or write as e.g. a `struct stat` is expected to be in a protocol specific format with all scalar multibyte datatypes being big endian. This should be done by the target before the F packet is sent resp. by GDB before it transfers memory to the target. Transferred pointers to structured data should point to the already coerced data at any time.

D.7.6 The Ctrl-C message

A special case is, if the *Ctrl-C flag* is set in the GDB reply packet. In this case the target should behave, as if it had gotten a break message. The meaning for the target is "system call interupted by `SIGINT`". Consequentially, the target should actually stop (as with a break message) and return to GDB with a T02 packet. In this case, it's important for the target to know, in which state the system call was interrupted. Since this action is by design not an atomic operation, we have to differ between two cases:

- The system call hasn't been performed on the host yet.
- The system call on the host has been finished.

These two states can be distinguished by the target by the value of the returned `errno`. If it's the protocol representation of `EINTR`, the system call hasn't been performed. This is equivalent to the `EINTR` handling on POSIX systems. In any other case, the target may presume that the system call has been finished — successful or not — and should behave as if the break message arrived right after the system call.

GDB must behave reliable. If the system call has not been called yet, GDB may send the F reply immediately, setting `EINTR` as `errno` in the packet. If the system call on the host has been finished before the user requests a break, the full action must be finshed by GDB. This requires sending M or X packets as they fit. The F packet may only be send when either nothing has happened or the full action has been completed.

D.7.7 Console I/O

By default and if not explicitely closed by the target system, the file descriptors 0, 1 and 2 are connected to the GDB console. Output on the GDB console is handled as any other file output operation (`write(1, ...)` or `write(2, ...)`). Console input is handled by GDB so that after the target read request from file descriptor 0 all following typing is buffered until either one of the following conditions is met:

- The user presses *Ctrl-C*. The behaviour is as explained above, the `read` system call is treated as finished.

- The user presses **Enter**. This is treated as end of input with a trailing line feed.
- The user presses **Ctrl-D**. This is treated as end of input. No trailing character, especially no Ctrl-D is appended to the input.

If the user has typed more characters as fit in the buffer given to the read call, the trailing characters are buffered in GDB until either another `read(0, ...)` is requested by the target or debugging is stopped on users request.

D.7.8 The `isatty(3)` call

A special case in this protocol is the library call `isatty` which is implemented as it's own call inside of this protocol. It returns 1 to the target if the file descriptor given as parameter is attached to the GDB console, 0 otherwise. Implementing through system calls would require implementing `ioctl` and would be more complex than needed.

D.7.9 The `system(3)` call

The other special case in this protocol is the `system` call which is implemented as it's own call, too. GDB is taking over the full task of calling the necessary host calls to perform the `system` call. The return value of `system` is simplified before it's returned to the target. Basically, the only signal transmitted back is `EINTR` in case the user pressed `Ctrl-C`. Otherwise the return value consists entirely of the exit status of the called command.

Due to security concerns, the `system` call is refused to be called by GDB by default. The user has to allow this call explicitly by entering

```
'set remote system-call-allowed 1'
```

Disabling the `system` call is done by

```
'set remote system-call-allowed 0'
```

The current setting is shown by typing

```
'show remote system-call-allowed'
```

D.7.10 List of supported calls

`open`

Synopsis:

```
int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
```

Request:

```
Fopen,pathptr/len,flags,mode
```

`flags` is the bitwise or of the following values:

- | | |
|-----------------------|--|
| <code>O_CREAT</code> | If the file does not exist it will be created. The host rules apply as far as file ownership and time stamps are concerned. |
| <code>O_EXCL</code> | When used with <code>O_CREAT</code> , if the file already exists it is an error and <code>open()</code> fails. |
| <code>O_TRUNC</code> | If the file already exists and the open mode allows writing (<code>O_RDWR</code> or <code>O_WRONLY</code> is given) it will be truncated to length 0. |
| <code>O_APPEND</code> | The file is opened in append mode. |

- O_RDONLY** The file is opened for reading only.
- O_WRONLY** The file is opened for writing only.
- O_RDWR** The file is opened for reading and writing.
Each other bit is silently ignored.

mode is the bitwise or of the following values:

- S_IRUSR** User has read permission.
- S_IWUSR** User has write permission.
- S_IRGRP** Group has read permission.
- S_IWGRP** Group has write permission.
- S_IROTH** Others have read permission.
- S_IWOTH** Others have write permission.
Each other bit is silently ignored.

Return value:

```
open returns the new file descriptor or -1 if an error
occured.
```

Errors:

- EEXIST** pathname already exists and O_CREAT and O_EXCL were used.
- EISDIR** pathname refers to a directory.
- EACCES** The requested access is not allowed.
- ENAMETOOLONG** pathname was too long.
- ENOENT** A directory component in pathname does not exist.
- ENODEV** pathname refers to a device, pipe, named pipe or socket.
- EROFS** pathname refers to a file on a read-only filesystem and write access was requested.
- EFAULT** pathname is an invalid pointer value.
- ENOSPC** No space on device to create the file.
- EMFILE** The process already has the maximum number of files open.
- ENFILE** The limit on the total number of files open on the system has been reached.
- EINTR** The call was interrupted by the user.

close

Synopsis:

```
int close(int fd);
```

Request:

```
Fclose,fd
```

Return value:

`close` returns zero on success, or -1 if an error occurred.

Errors:

EBADF fd isn't a valid open file descriptor.

EINTR The call was interrupted by the user.

read

Synopsis:

```
int read(int fd, void *buf, unsigned int count);
```

Request:

```
Fread,fd,bufptr,count
```

Return value:

On success, the number of bytes read is returned.
Zero indicates end of file. If count is zero, read
returns zero as well. On error, -1 is returned.

Errors:

EBADF fd is not a valid file descriptor or is not open for reading.

EFAULT buf is an invalid pointer value.

EINTR The call was interrupted by the user.

write

Synopsis:

```
int write(int fd, const void *buf, unsigned int count);
```

Request:

```
Fwrite,fd,bufptr,count
```

Return value:

On success, the number of bytes written are returned.
Zero indicates nothing was written. On error, -1
is returned.

Errors:

EBADF fd is not a valid file descriptor or is not open for writing.

EFAULT buf is an invalid pointer value.

EFBIG An attempt was made to write a file that exceeds the host specific maximum
file size allowed.

ENOSPC No space on device to write the data.

EINTR The call was interrupted by the user.

lseek

Synopsis:

```
long lseek (int fd, long offset, int flag);
```

Request:

`Ffseek,fd,offset,flag`

`flag` is one of:

`SEEK_SET` The offset is set to offset bytes.

`SEEK_CUR` The offset is set to its current location plus offset bytes.

`SEEK_END` The offset is set to the size of the file plus offset bytes.

Return value:

On success, the resulting unsigned offset in bytes from
the beginning of the file is returned. Otherwise, a
value of -1 is returned.

Errors:

`EBADF` `fd` is not a valid open file descriptor.

`ESPIPE` `fd` is associated with the GDB console.

`EINVAL` `flag` is not a proper value.

`EINTR` The call was interrupted by the user.

rename

Synopsis:

```
int rename(const char *oldpath, const char *newpath);
```

Request:

```
Frename,oldpathptr/len,newpathptr/len
```

Return value:

On success, zero is returned. On error, -1 is returned.

Errors:

`EISDIR` `newpath` is an existing directory, but `oldpath` is not a directory.

`EEXIST` `newpath` is a non-empty directory.

`EBUSY` `oldpath` or `newpath` is a directory that is in use by some process.

`EINVAL` An attempt was made to make a directory a subdirectory of itself.

`ENOTDIR` A component used as a directory in `oldpath` or new path is not a directory. Or `oldpath` is a directory and `newpath` exists but is not a directory.

`EFAULT` `oldpathptr` or `newpathptr` are invalid pointer values.

`EACCES` No access to the file or the path of the file.

ENAMETOOLONG

`oldpath` or `newpath` was too long.

`ENOENT` A directory component in `oldpath` or `newpath` does not exist.

`EROFS` The file is on a read-only filesystem.

`ENOSPC` The device containing the file has no room for the new directory entry.

`EINTR` The call was interrupted by the user.

unlink

Synopsis:

```
int unlink(const char *pathname);
```

Request:

```
Funlink, pathnameptr/len
```

Return value:

```
On success, zero is returned. On error, -1 is returned.
```

Errors:

EACCES No access to the file or the path of the file.

EPERM The system does not allow unlinking of directories.

EBUSY The file pathname cannot be unlinked because it's being used by another process.

EFAULT pathnameptr is an invalid pointer value.

ENAMETOOLONG

pathname was too long.

ENOENT A directory component in pathname does not exist.

ENOTDIR A component of the path is not a directory.

EROFS The file is on a read-only filesystem.

EINTR The call was interrupted by the user.

stat/fstat

Synopsis:

```
int stat(const char *pathname, struct stat *buf);
int fstat(int fd, struct stat *buf);
```

Request:

```
Fstat, pathnameptr/len, bufptr
Ffstat, fd, bufptr
```

Return value:

```
On success, zero is returned. On error, -1 is returned.
```

Errors:

EBADF fd is not a valid open file.

ENOENT A directory component in pathname does not exist or the path is an empty string.

ENOTDIR A component of the path is not a directory.

EFAULT pathnameptr is an invalid pointer value.

EACCES No access to the file or the path of the file.

ENAMETOOLONG

pathname was too long.

EINTR The call was interrupted by the user.

gettimeofday

Synopsis:

```
int gettimeofday(struct timeval *tv, void *tz);
```

Request:

```
Fgettimeofday,tvptr,tzptr
```

Return value:

On success, 0 is returned, -1 otherwise.

Errors:

EINVAL tz is a non-NULL pointer.

EFAULT tvptr and/or tzptr is an invalid pointer value.

isatty

Synopsis:

```
int isatty(int fd);
```

Request:

```
Fisatty,fd
```

Return value:

Returns 1 if fd refers to the GDB console, 0 otherwise.

Errors:

EINTR The call was interrupted by the user.

system

Synopsis:

```
int system(const char *command);
```

Request:

```
Fsystem,commandptr/len
```

Return value:

The value returned is -1 on error and the return status of the command otherwise. Only the exit status of the command is returned, which is extracted from the hosts system return value by calling WEXITSTATUS(retval).

In case /bin/sh could not be executed, 127 is returned.

Errors:

EINTR The call was interrupted by the user.

D.7.11 Protocol specific representation of datatypes

Integral datatypes

The integral datatypes used in the system calls are

```
int, unsigned int, long, unsigned long, mode_t and time_t
```

Int, unsigned int, mode_t and time_t are implemented as 32 bit values in this protocol.

Long and unsigned long are implemented as 64 bit types.

See [Limits], page 317, for corresponding MIN and MAX values (similar to those in ‘limits.h’) to allow range checking on host and target.

`time_t` datatypes are defined as seconds since the Epoch.

All integral datatypes transferred as part of a memory read or write of a structured datatype e.g. a `struct stat` have to be given in big endian byte order.

Pointer values

Pointers to target data are transmitted as they are. An exception is made for pointers to buffers for which the length isn’t transmitted as part of the function call, namely strings. Strings are transmitted as a pointer/length pair, both as hex values, e.g.

```
1aaf/12
```

which is a pointer to data of length 18 bytes at position 0x1aaf. The length is defined as the full string length in bytes, including the trailing null byte. Example:

```
“hello, world” at address 0x123456
```

is transmitted as

```
123456/d
```

struct stat

The buffer of type `struct stat` used by the target and GDB is defined as follows:

```
struct stat {
    unsigned int st_dev;      /* device */
    unsigned int st_ino;       /* inode */
    mode_t st_mode;          /* protection */
    unsigned int st_nlink;    /* number of hard links */
    unsigned int st_uid;      /* user ID of owner */
    unsigned int st_gid;      /* group ID of owner */
    unsigned int st_rdev;     /* device type (if inode device) */
    unsigned long st_size;    /* total size, in bytes */
    unsigned long st_blksize; /* blocksize for filesystem I/O */
    unsigned long st_blocks;  /* number of blocks allocated */
    time_t st_atime;         /* time of last access */
    time_t st_mtime;         /* time of last modification */
    time_t st_ctime;         /* time of last change */
};
```

The integral datatypes are conforming to the definitions given in the appropriate section (see [Integral datatypes], page 314, for details) so this structure is of size 64 bytes.

The values of several fields have a restricted meaning and/or range of values.

st_dev:	0	file
	1	console

st_ino: No valid meaning for the target. Transmitted unchanged.

st_mode: Valid mode bits are described in Appendix C. Any other bits have currently no meaning for the target.

st_uid: No valid meaning for the target. Transmitted unchanged.

st_gid: No valid meaning for the target. Transmitted unchanged.

st_rdev: No valid meaning for the target. Transmitted unchanged.

```

st_atime, st_mtime, st_ctime:
    These values have a host and file system dependent
    accuracy. Especially on Windows hosts the file systems
    don't support exact timing values.

```

The target gets a struct stat of the above representation and is responsible to coerce it to the target representation before continuing.

Note that due to size differences between the host and target representation of stat members, these members could eventually get truncated on the target.

struct timeval

The buffer of type struct timeval used by the target and GDB is defined as follows:

```

struct timeval {
    time_t tv_sec; /* second */
    long   tv_usec; /* microsecond */
};

```

The integral datatypes are conforming to the definitions given in the appropriate section (see [\[Integral datatypes\]](#), page 314, for details) so this structure is of size 8 bytes.

D.7.12 Constants

The following values are used for the constants inside of the protocol. GDB and target are responsible to translate these values before and after the call as needed.

Open flags

All values are given in hexadecimal representation.

O_RDONLY	0x0
O_WRONLY	0x1
O_RDWR	0x2
O_APPEND	0x8
O_CREAT	0x200
O_TRUNC	0x400
O_EXCL	0x800

mode_t values

All values are given in octal representation.

S_IFREG	0100000
S_IFDIR	040000
S_IRUSR	0400
S_IWUSR	0200
S_IXUSR	0100
S_IRGRP	040
S_IWGRP	020
S_IXGRP	010
S_IROTH	04
S_IWOTH	02
S_IXOTH	01

Errno values

All values are given in decimal representation.

EPERM	1
ENOENT	2

EINTR	4
EBADF	9
EACCES	13
EFAULT	14
EBUSY	16
EEXIST	17
ENODEV	19
ENOTDIR	20
EISDIR	21
EINVAL	22
ENFILE	23
EMFILE	24
EFBIG	27
ENOSPC	28
ESPIPE	29
EROFS	30
ENAMETOOLONG	91
EUNKNOWN	9999

EUNKNOWN is used as a fallback error value if a host system returns any error value not in the list of supported error numbers.

Lseek flags

SEEK_SET	0
SEEK_CUR	1
SEEK_END	2

Limits

All values are given in decimal representation.

INT_MIN	-2147483648
INT_MAX	2147483647
UINT_MAX	4294967295
LONG_MIN	-9223372036854775808
LONG_MAX	9223372036854775807
ULONG_MAX	18446744073709551615

D.7.13 File-I/O Examples

Example sequence of a write call, file descriptor 3, buffer is at target address 0x1234, 6 bytes should be written:

```

<- Fwrite,3,1234,6
request memory read from target
-> m1234,6
<- XXXXXX
return "6 bytes written"
-> F6

```

Example sequence of a read call, file descriptor 3, buffer is at target address 0x1234, 6 bytes should be read:

```

<- Fread,3,1234,6
request memory write to target
-> X1234,6:XXXXXX
return "6 bytes read"
-> F6

```

Example sequence of a read call, call fails on the host due to invalid file descriptor (EBADF):

```
<- Fread,3,1234,6  
-> F-1,9
```

Example sequence of a read call, user presses Ctrl-C before syscall on host is called:

```
<- Fread,3,1234,6  
-> F-1,4,C  
<- T02
```

Example sequence of a read call, user presses Ctrl-C after syscall on host is called:

```
<- Fread,3,1234,6  
-> X1234,6:XXXXXX  
<- T02
```

Appendix E The GDB Agent Expression Mechanism

In some applications, it is not feasible for the debugger to interrupt the program’s execution long enough for the developer to learn anything helpful about its behavior. If the program’s correctness depends on its real-time behavior, delays introduced by a debugger might cause the program to fail, even when the code itself is correct. It is useful to be able to observe the program’s behavior without interrupting it.

Using GDB’s `trace` and `collect` commands, the user can specify locations in the program, and arbitrary expressions to evaluate when those locations are reached. Later, using the `tfind` command, she can examine the values those expressions had when the program hit the trace points. The expressions may also denote objects in memory — structures or arrays, for example — whose values GDB should record; while visiting a particular tracepoint, the user may inspect those objects as if they were in memory at that moment. However, because GDB records these values without interacting with the user, it can do so quickly and unobtrusively, hopefully not disturbing the program’s behavior.

When GDB is debugging a remote target, the GDB agent code running on the target computes the values of the expressions itself. To avoid having a full symbolic expression evaluator on the agent, GDB translates expressions in the source language into a simpler bytecode language, and then sends the bytecode to the agent; the agent then executes the bytecode, and records the values for GDB to retrieve later.

The bytecode language is simple; there are forty-odd opcodes, the bulk of which are the usual vocabulary of C operands (addition, subtraction, shifts, and so on) and various sizes of literals and memory reference operations. The bytecode interpreter operates strictly on machine-level values — various sizes of integers and floating point numbers — and requires no information about types or symbols; thus, the interpreter’s internal data structures are simple, and each bytecode requires only a few native machine instructions to implement it. The interpreter is small, and strict limits on the memory and time required to evaluate an expression are easy to determine, making it suitable for use by the debugging agent in real-time applications.

E.1 General Bytecode Design

The agent represents bytecode expressions as an array of bytes. Each instruction is one byte long (thus the term *bytecode*). Some instructions are followed by operand bytes; for example, the `goto` instruction is followed by a destination for the jump.

The bytecode interpreter is a stack-based machine; most instructions pop their operands off the stack, perform some operation, and push the result back on the stack for the next instruction to consume. Each element of the stack may contain either a integer or a floating point value; these values are as many bits wide as the largest integer that can be directly manipulated in the source language. Stack elements carry no record of their type; bytecode could push a value as an integer, then pop it as a floating point value. However, GDB will not generate code which does this. In C, one might define the type of a stack element as follows:

```
union agent_val {
    LONGEST l;
    DOUBLEST d;
```

};

where `LONGEST` and `DOUBLEST` are `typedef` names for the largest integer and floating point types on the machine.

By the time the bytecode interpreter reaches the end of the expression, the value of the expression should be the only value left on the stack. For tracing applications, `trace` bytecodes in the expression will have recorded the necessary data, and the value on the stack may be discarded. For other applications, like conditional breakpoints, the value may be useful.

Separate from the stack, the interpreter has two registers:

- `pc` The address of the next bytecode to execute.
- `start` The address of the start of the bytecode expression, necessary for interpreting the `goto` and `if_goto` instructions.

Neither of these registers is directly visible to the bytecode language itself, but they are useful for defining the meanings of the bytecode operations.

There are no instructions to perform side effects on the running program, or call the program's functions; we assume that these expressions are only used for unobtrusive debugging, not for patching the running code.

Most bytecode instructions do not distinguish between the various sizes of values, and operate on full-width values; the upper bits of the values are simply ignored, since they do not usually make a difference to the value computed. The exceptions to this rule are:

memory reference instructions (`refn`)

There are distinct instructions to fetch different word sizes from memory. Once on the stack, however, the values are treated as full-size integers. They may need to be sign-extended; the `ext` instruction exists for this purpose.

the sign-extension instruction (`ext n`)

These clearly need to know which portion of their operand is to be extended to occupy the full length of the word.

If the interpreter is unable to evaluate an expression completely for some reason (a memory location is inaccessible, or a divisor is zero, for example), we say that interpretation "terminates with an error". This means that the problem is reported back to the interpreter's caller in some helpful way. In general, code using agent expressions should assume that they may attempt to divide by zero, fetch arbitrary memory locations, and misbehave in other ways.

Even complicated C expressions compile to a few bytecode instructions; for example, the expression `x + y * z` would typically produce code like the following, assuming that `x` and `y` live in registers, and `z` is a global variable holding a 32-bit `int`:

```
reg 1
reg 2
const32 address of z
ref32
ext 32
mul
add
```

```
end
```

In detail, these mean:

reg 1	Push the value of register 1 (presumably holding <i>x</i>) onto the stack.
reg 2	Push the value of register 2 (holding <i>y</i>).
const32 address of <i>z</i>	Push the address of <i>z</i> onto the stack.
ref32	Fetch a 32-bit word from the address at the top of the stack; replace the address on the stack with the value. Thus, we replace the address of <i>z</i> with <i>z</i> 's value.
ext 32	Sign-extend the value on the top of the stack from 32 bits to full length. This is necessary because <i>z</i> is a signed integer.
mul	Pop the top two numbers on the stack, multiply them, and push their product. Now the top of the stack contains the value of the expression <i>y * z</i> .
add	Pop the top two numbers, add them, and push the sum. Now the top of the stack contains the value of <i>x + y * z</i> .
end	Stop executing; the value left on the stack top is the value to be recorded.

E.2 Bytecode Descriptions

Each bytecode description has the following form:

add (0x02): <i>a b</i> ⇒ <i>a+b</i>	Pop the top two stack items, <i>a</i> and <i>b</i> , as integers; push their sum, as an integer.
--	--

In this example, **add** is the name of the bytecode, and (0x02) is the one-byte value used to encode the bytecode, in hexadecimal. The phrase “*a b* ⇒ *a+b*” shows the stack before and after the bytecode executes. Beforehand, the stack must contain at least two values, *a* and *b*; since the top of the stack is to the right, *b* is on the top of the stack, and *a* is underneath it. After execution, the bytecode will have popped *a* and *b* from the stack, and replaced them with a single value, *a+b*. There may be other values on the stack below those shown, but the bytecode affects only those shown.

Here is another example:

const8 (0x22) <i>n</i>: ⇒ <i>n</i>	Push the 8-bit integer constant <i>n</i> on the stack, without sign extension.
---	--

In this example, the bytecode **const8** takes an operand *n* directly from the bytecode stream; the operand follows the **const8** bytecode itself. We write any such operands immediately after the name of the bytecode, before the colon, and describe the exact encoding of the operand in the bytecode stream in the body of the bytecode description.

For the **const8** bytecode, there are no stack items given before the ⇒; this simply means that the bytecode consumes no values from the stack. If a bytecode consumes no values, or produces no values, the list on either side of the ⇒ may be empty.

If a value is written as *a*, *b*, or *n*, then the bytecode treats it as an integer. If a value is written as *addr*, then the bytecode treats it as an address.

We do not fully describe the floating point operations here; although this design can be extended in a clean way to handle floating point values, they are not of immediate interest to the customer, so we avoid describing them, to save time.

float (0x01): \Rightarrow

Prefix for floating-point bytecodes. Not implemented yet.

add (0x02): $a b \Rightarrow a+b$

Pop two integers from the stack, and push their sum, as an integer.

sub (0x03): $a b \Rightarrow a-b$

Pop two integers from the stack, subtract the top value from the next-to-top value, and push the difference.

mul (0x04): $a b \Rightarrow a*b$

Pop two integers from the stack, multiply them, and push the product on the stack. Note that, when one multiplies two n -bit numbers yielding another n -bit number, it is irrelevant whether the numbers are signed or not; the results are the same.

div_signed (0x05): $a b \Rightarrow a/b$

Pop two signed integers from the stack; divide the next-to-top value by the top value, and push the quotient. If the divisor is zero, terminate with an error.

div_unsigned (0x06): $a b \Rightarrow a/b$

Pop two unsigned integers from the stack; divide the next-to-top value by the top value, and push the quotient. If the divisor is zero, terminate with an error.

rem_signed (0x07): $a b \Rightarrow a \text{ modulo } b$

Pop two signed integers from the stack; divide the next-to-top value by the top value, and push the remainder. If the divisor is zero, terminate with an error.

rem_unsigned (0x08): $a b \Rightarrow a \text{ modulo } b$

Pop two unsigned integers from the stack; divide the next-to-top value by the top value, and push the remainder. If the divisor is zero, terminate with an error.

lsh (0x09): $a b \Rightarrow a \ll b$

Pop two integers from the stack; let a be the next-to-top value, and b be the top value. Shift a left by b bits, and push the result.

rsh_signed (0x0a): $a b \Rightarrow (\text{signed})a \gg b$

Pop two integers from the stack; let a be the next-to-top value, and b be the top value. Shift a right by b bits, inserting copies of the top bit at the high end, and push the result.

rsh_unsigned (0x0b): $a b \Rightarrow a \gg b$

Pop two integers from the stack; let a be the next-to-top value, and b be the top value. Shift a right by b bits, inserting zero bits at the high end, and push the result.

log_not (0x0e): $a \Rightarrow !a$

Pop an integer from the stack; if it is zero, push the value one; otherwise, push the value zero.

bit_and (0x0f): $a b \Rightarrow a \& b$

Pop two integers from the stack, and push their bitwise and.

bit_or (0x10): $a b \Rightarrow a | b$

Pop two integers from the stack, and push their bitwise or.

bit_xor (0x11): $a b \Rightarrow a \hat{=} b$

Pop two integers from the stack, and push their bitwise exclusive-or.

bit_not (0x12): $a \Rightarrow \sim a$

Pop an integer from the stack, and push its bitwise complement.

equal (0x13): $a b \Rightarrow a = b$

Pop two integers from the stack; if they are equal, push the value one; otherwise, push the value zero.

less_signed (0x14): $a b \Rightarrow a < b$

Pop two signed integers from the stack; if the next-to-top value is less than the top value, push the value one; otherwise, push the value zero.

less_unsigned (0x15): $a b \Rightarrow a < b$

Pop two unsigned integers from the stack; if the next-to-top value is less than the top value, push the value one; otherwise, push the value zero.

ext (0x16) $n: a \Rightarrow a$, sign-extended from n bits

Pop an unsigned value from the stack; treating it as an n -bit twos-complement value, extend it to full length. This means that all bits to the left of bit $n-1$ (where the least significant bit is bit 0) are set to the value of bit $n-1$. Note that n may be larger than or equal to the width of the stack elements of the bytecode engine; in this case, the bytecode should have no effect.

The number of source bits to preserve, n , is encoded as a single byte unsigned integer following the **ext** bytecode.

zero_ext (0x2a) $n: a \Rightarrow a$, zero-extended from n bits

Pop an unsigned value from the stack; zero all but the bottom n bits. This means that all bits to the left of bit $n-1$ (where the least significant bit is bit 0) are set to the value of bit $n-1$.

The number of source bits to preserve, n , is encoded as a single byte unsigned integer following the **zero_ext** bytecode.

ref8 (0x17): $addr \Rightarrow a$

ref16 (0x18): $addr \Rightarrow a$

ref32 (0x19): $addr \Rightarrow a$

ref64 (0x1a): $addr \Rightarrow a$

Pop an address $addr$ from the stack. For bytecode **refn**, fetch an n -bit value from $addr$, using the natural target endianness. Push the fetched value as an unsigned integer.

Note that $addr$ may not be aligned in any particular way; the **refn** bytecodes should operate correctly for any address.

If attempting to access memory at $addr$ would cause a processor exception of some sort, terminate with an error.

ref_float (0x1b): *addr* \Rightarrow *d*
ref_double (0x1c): *addr* \Rightarrow *d*
ref_long_double (0x1d): *addr* \Rightarrow *d*
l_to_d (0x1e): *a* \Rightarrow *d*
d_to_l (0x1f): *d* \Rightarrow *a*
 Not implemented yet.

dup (0x28): *a* \Rightarrow *a a*
 Push another copy of the stack's top element.

swap (0x2b): *a b* \Rightarrow *b a*
 Exchange the top two items on the stack.

pop (0x29): *a* \Rightarrow
 Discard the top value on the stack.

if_goto (0x20) *offset: a* \Rightarrow
 Pop an integer off the stack; if it is non-zero, branch to the given offset in the bytecode string. Otherwise, continue to the next instruction in the bytecode stream. In other words, if *a* is non-zero, set the **pc** register to **start + offset**. Thus, an offset of zero denotes the beginning of the expression.
 The *offset* is stored as a sixteen-bit unsigned value, stored immediately following the **if_goto** bytecode. It is always stored most significant byte first, regardless of the target's normal endianness. The offset is not guaranteed to fall at any particular alignment within the bytecode stream; thus, on machines where fetching a 16-bit on an unaligned address raises an exception, you should fetch the offset one byte at a time.

goto (0x21) *offset:* \Rightarrow
 Branch unconditionally to *offset*; in other words, set the **pc** register to **start + offset**.
 The offset is stored in the same way as for the **if_goto** bytecode.

const8 (0x22) *n:* \Rightarrow *n*
const16 (0x23) *n:* \Rightarrow *n*
const32 (0x24) *n:* \Rightarrow *n*
const64 (0x25) *n:* \Rightarrow *n*
 Push the integer constant *n* on the stack, without sign extension. To produce a small negative value, push a small twos-complement value, and then sign-extend it using the **ext** bytecode.
 The constant *n* is stored in the appropriate number of bytes following the **constb** bytecode. The constant *n* is always stored most significant byte first, regardless of the target's normal endianness. The constant is not guaranteed to fall at any particular alignment within the bytecode stream; thus, on machines where fetching a 16-bit on an unaligned address raises an exception, you should fetch *n* one byte at a time.

reg (0x26) *n:* \Rightarrow *a*
 Push the value of register number *n*, without sign extension. The registers are numbered following GDB's conventions.

The register number *n* is encoded as a 16-bit unsigned integer immediately following the `reg` bytecode. It is always stored most significant byte first, regardless of the target’s normal endianness. The register number is not guaranteed to fall at any particular alignment within the bytecode stream; thus, on machines where fetching a 16-bit on an unaligned address raises an exception, you should fetch the register number one byte at a time.

`trace (0x0c): addr size ⇒`

Record the contents of the *size* bytes at *addr* in a trace buffer, for later retrieval by GDB.

`trace_quick (0x0d) size: addr ⇒ addr`

Record the contents of the *size* bytes at *addr* in a trace buffer, for later retrieval by GDB. *size* is a single byte unsigned integer following the `trace` opcode.

This bytecode is equivalent to the sequence `dup const8 size trace`, but we provide it anyway to save space in bytecode strings.

`trace16 (0x30) size: addr ⇒ addr`

Identical to `trace_quick`, except that *size* is a 16-bit big-endian unsigned integer, not a single byte. This should probably have been named `trace_quick16`, for consistency.

`end (0x27): ⇒`

Stop executing bytecode; the result should be the top element of the stack. If the purpose of the expression was to compute an lvalue or a range of memory, then the next-to-top of the stack is the lvalue’s address, and the top of the stack is the lvalue’s size, in bytes.

E.3 Using Agent Expressions

Here is a sketch of a full non-stop debugging cycle, showing how agent expressions fit into the process.

- The user selects trace points in the program’s code at which GDB should collect data.
- The user specifies expressions to evaluate at each trace point. These expressions may denote objects in memory, in which case those objects’ contents are recorded as the program runs, or computed values, in which case the values themselves are recorded.
- GDB transmits the tracepoints and their associated expressions to the GDB agent, running on the debugging target.
- The agent arranges to be notified when a trace point is hit. Note that, on some systems, the target operating system is completely responsible for collecting the data; see [Section E.5 \[Tracing on Symmetrix\], page 326](#).
- When execution on the target reaches a trace point, the agent evaluates the expressions associated with that trace point, and records the resulting values and memory ranges.
- Later, when the user selects a given trace event and inspects the objects and expression values recorded, GDB talks to the agent to retrieve recorded data as necessary to meet the user’s requests. If the user asks to see an object whose contents have not been recorded, GDB reports an error.

E.4 Varying Target Capabilities

Some targets don't support floating-point, and some would rather not have to deal with `long long` operations. Also, different targets will have different stack sizes, and different bytecode buffer lengths.

Thus, GDB needs a way to ask the target about itself. We haven't worked out the details yet, but in general, GDB should be able to send the target a packet asking it to describe itself. The reply should be a packet whose length is explicit, so we can add new information to the packet in future revisions of the agent, without confusing old versions of GDB, and it should contain a version number. It should contain at least the following information:

- whether floating point is supported
- whether `long long` is supported
- maximum acceptable size of bytecode stack
- maximum acceptable length of bytecode expressions
- which registers are actually available for collection
- whether the target supports disabled tracepoints

E.5 Tracing on Symmetrix

This section documents the API used by the GDB agent to collect data on Symmetrix systems.

Cygnus originally implemented these tracing features to help EMC Corporation debug their Symmetrix high-availability disk drives. The Symmetrix application code already includes substantial tracing facilities; the GDB agent for the Symmetrix system uses those facilities for its own data collection, via the API described here.

```
DTC_RESPONSE adbq_find_memory_in_frame (FRAME_DEF *frame,      [Function]
                                         char *address, char **buffer, unsigned int *size)
```

Search the trace frame *frame* for memory saved from *address*. If the memory is available, provide the address of the buffer holding it; otherwise, provide the address of the next saved area.

- If the memory at *address* was saved in *frame*, set **buffer* to point to the buffer in which that memory was saved, set **size* to the number of bytes from *address* that are saved at **buffer*, and return `OK_TARGET_RESPONSE`. (Clearly, in this case, the function will always set **size* to a value greater than zero.)
- If *frame* does not record any memory at *address*, set **size* to the distance from *address* to the start of the saved region with the lowest address higher than *address*. If there is no memory saved from any higher address, set **size* to zero. Return `NOT_FOUND_TARGET_RESPONSE`.

These two possibilities allow the caller to either retrieve the data, or walk the address space to the next saved area.

This function allows the GDB agent to map the regions of memory saved in a particular frame, and retrieve their contents efficiently.

This function also provides a clean interface between the GDB agent and the Symmetrix tracing structures, making it easier to adapt the GDB agent to future versions of the

Symmetrix system, and vice versa. This function searches all data saved in *frame*, whether the data is there at the request of a bytecode expression, or because it falls in one of the format's memory ranges, or because it was saved from the top of the stack. EMC can arbitrarily change and enhance the tracing mechanism, but as long as this function works properly, all collected memory is visible to GDB.

The function itself is straightforward to implement. A single pass over the trace frame's stack area, memory ranges, and expression blocks can yield the address of the buffer (if the requested address was saved), and also note the address of the next higher range of memory, to be returned when the search fails.

As an example, suppose the trace frame *f* has saved sixteen bytes from address 0x8000 in a buffer at 0x1000, and thirty-two bytes from address 0xc000 in a buffer at 0x1010. Here are some sample calls, and the effect each would have:

```
adbg_find_memory_in_frame (f, (char*) 0x8000, &buffer, &size)
    This would set buffer to 0x1000, set size to sixteen, and return OK_TARGET_RESPONSE, since f saves sixteen bytes from 0x8000 at 0x1000.

adbg_find_memory_in_frame (f, (char *) 0x8004, &buffer, &size)
    This would set buffer to 0x1004, set size to twelve, and return OK_TARGET_RESPONSE, since 'f' saves the twelve bytes from 0x8004 starting four bytes into the buffer at 0x1000. This shows that request addresses may fall in the middle of saved areas; the function should return the address and size of the remainder of the buffer.

adbg_find_memory_in_frame (f, (char *) 0x8100, &buffer, &size)
    This would set size to 0x3f00 and return NOT_FOUND_TARGET_RESPONSE, since there is no memory saved in f from the address 0x8100, and the next memory available is at 0x8100 + 0x3f00, or 0xc000. This shows that request addresses may fall outside of all saved memory ranges; the function should indicate the next saved area, if any.

adbg_find_memory_in_frame (f, (char *) 0x7000, &buffer, &size)
    This would set size to 0x1000 and return NOT_FOUND_TARGET_RESPONSE, since the next saved memory is at 0x7000 + 0x1000, or 0x8000.

adbg_find_memory_in_frame (f, (char *) 0xf000, &buffer, &size)
    This would set size to zero, and return NOT_FOUND_TARGET_RESPONSE. This shows how the function tells the caller that no further memory ranges have been saved.
```

As another example, here is a function which will print out the addresses of all memory saved in the trace frame *frame* on the Symmetrix INLINES console:

```
void
print_frame_addresses (FRAME_DEF *frame)
{
    char *addr;
    char *buffer;
    unsigned long size;

    addr = 0;
```

```

for (;;)
{
    /* Either find out how much memory we have here, or discover
       where the next saved region is. */
    if (adbg_find_memory_in_frame (frame, addr, &buffer, &size)
        == OK_TARGET_RESPONSE)
        printp ("saved %x to %x\n", addr, addr + size);
    if (size == 0)
        break;
    addr += size;
}
}

```

Note that there is not necessarily any connection between the order in which the data is saved in the trace frame, and the order in which `adbg_find_memory_in_frame` will return those memory ranges. The code above will always print the saved memory regions in order of increasing address, while the underlying frame structure might store the data in a random order.

[[This section should cover the rest of the Symmetrix functions the stub relies upon, too.]]

E.6 Rationale

Some of the design decisions apparent above are arguable.

What about stack overflow/underflow?

GDB should be able to query the target to discover its stack size. Given that information, GDB can determine at translation time whether a given expression will overflow the stack. But this spec isn't about what kinds of error-checking GDB ought to do.

Why are you doing everything in LONGEST?

Speed isn't important, but agent code size is; using LONGEST brings in a bunch of support code to do things like division, etc. So this is a serious concern.

First, note that you don't need different bytecodes for different operand sizes. You can generate code without *knowing* how big the stack elements actually are on the target. If the target only supports 32-bit ints, and you don't send any 64-bit bytecodes, everything just works. The observation here is that the MIPS and the Alpha have only fixed-size registers, and you can still get C's semantics even though most instructions only operate on full-sized words. You just need to make sure everything is properly sign-extended at the right times. So there is no need for 32- and 64-bit variants of the bytecodes. Just implement everything using the largest size you support.

GDB should certainly check to see what sizes the target supports, so the user can get an error earlier, rather than later. But this information is not necessary for correctness.

Why don't you have > or <= operators?

I want to keep the interpreter small, and we don't need them. We can combine the `less_` opcodes with `log_not`, and swap the order of the operands, yielding

all four asymmetrical comparison operators. For example, `(x <= y)` is `! (x > y)`, which is `! (y < x)`.

Why do you have `log_not`?

Why do you have `ext`?

Why do you have `zero_ext`?

These are all easily synthesized from other instructions, but I expect them to be used frequently, and they're simple, so I include them to keep bytecode strings short.

`log_not` is equivalent to `const8 0 equal`; it's used in half the relational operators.

`ext n` is equivalent to `const8 s-n lsh const8 s-n rsh_signed`, where `s` is the size of the stack elements; it follows `refm` and `reg` bytecodes when the value should be signed. See the next bulleted item.

`zero_ext n` is equivalent to `constm mask log_and`; it's used whenever we push the value of a register, because we can't assume the upper bits of the register aren't garbage.

Why not have sign-extending variants of the `ref` operators?

Because that would double the number of `ref` operators, and we need the `ext` bytecode anyway for accessing bitfields.

Why not have constant-address variants of the `ref` operators?

Because that would double the number of `ref` operators again, and `const32 address ref32` is only one byte longer.

Why do the `refn` operators have to support unaligned fetches?

GDB will generate bytecode that fetches multi-byte values at unaligned addresses whenever the executable's debugging information tells it to. Furthermore, GDB does not know the value the pointer will have when GDB generates the bytecode, so it cannot determine whether a particular fetch will be aligned or not.

In particular, structure bitfields may be several bytes long, but follow no alignment rules; members of packed structures are not necessarily aligned either.

In general, there are many cases where unaligned references occur in correct C code, either at the programmer's explicit request, or at the compiler's discretion. Thus, it is simpler to make the GDB agent bytecodes work correctly in all circumstances than to make GDB guess in each case whether the compiler did the usual thing.

Why are there no side-effecting operators?

Because our current client doesn't want them? That's a cheap answer. I think the real answer is that I'm afraid of implementing function calls. We should re-visit this issue after the present contract is delivered.

Why aren't the `goto ops` PC-relative?

The interpreter has the base address around anyway for PC bounds checking, and it seemed simpler.

Why is there only one offset size for the goto ops?

Offsets are currently sixteen bits. I'm not happy with this situation either:

Suppose we have multiple branch ops with different offset sizes. As I generate code left-to-right, all my jumps are forward jumps (there are no loops in expressions), so I never know the target when I emit the jump opcode. Thus, I have to either always assume the largest offset size, or do jump relaxation on the code after I generate it, which seems like a big waste of time.

I can imagine a reasonable expression being longer than 256 bytes. I can't imagine one being longer than 64k. Thus, we need 16-bit offsets. This kind of reasoning is so bogus, but relaxation is pathetic.

The other approach would be to generate code right-to-left. Then I'd always know my offset size. That might be fun.

Where is the function call bytecode?

When we add side-effects, we should add this.

Why does the reg bytecode take a 16-bit register number?

Intel's IA-64 architecture has 128 general-purpose registers, and 128 floating-point registers, and I'm sure it has some random control registers.

Why do we need trace and trace_quick?

Because GDB needs to record all the memory contents and registers an expression touches. If the user wants to evaluate an expression `x->y->z`, the agent must record the values of `x` and `x->y` as well as the value of `x->y->z`.

Don't the trace bytecodes make the interpreter less general?

They do mean that the interpreter contains special-purpose code, but that doesn't mean the interpreter can only be used for that purpose. If an expression doesn't use the `trace` bytecodes, they don't get in its way.

Why doesn't trace_quick consume its arguments the way everything else does?

In general, you do want your operators to consume their arguments; it's consistent, and generally reduces the amount of stack rearrangement necessary. However, `trace_quick` is a kludge to save space; it only exists so we needn't write `dup const8 SIZE trace` before every memory reference. Therefore, it's okay for it not to consume its arguments; it's meant for a specific context in which we know exactly what it should do with the stack. If we're going to have a kludge, it should be an effective kludge.

Why does trace16 exist?

That opcode was added by the customer that contracted Cygnus for the data tracing work. I personally think it is unnecessary; objects that large will be quite rare, so it is okay to use `dup const16 size trace` in those cases.

Whatever we decide to do with `trace16`, we should at least leave opcode 0x30 reserved, to remain compatible with the customer who added it.

Appendix F GNU GENERAL PUBLIC LICENSE

Version 2, June 1991

Copyright © 1989, 1991 Free Software Foundation, Inc.
59 Temple Place - Suite 330, Boston, MA 02111-1307, USA

Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software—to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The “Program”, below, refers to any such program or work, and a “work based on the Program” means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term “modification”.) Each licensee is addressed as “you”.

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program’s source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.
You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.
2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:
 - a. You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
 - b. You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
 - c. If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:
 - a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
 - b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
 - c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.
5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.

6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.
7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.
9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software

which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.
12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

one line to give the program's name and a brief idea of what it does.

Copyright (C) year name of author

This program is free software; you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation; either version 2 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program; if not, write to the Free Software
Foundation, Inc., 59 Temple Place - Suite 330,
Boston, MA 02111-1307, USA.

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

Gnomovision version 69, Copyright (C) year name of author
Gnomovision comes with ABSOLUTELY NO WARRANTY; for details
type 'show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.

The hypothetical commands ‘show w’ and ‘show c’ should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than ‘show w’ and ‘show c’; they could even be mouse-clicks or menu items—whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a “copyright disclaimer” for the program, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the program
'Gnomovision' (which makes passes at compilers) written by James Hacker.

signature of Ty Coon, 1 April 1989
Ty Coon, President of Vice

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

Appendix G GNU Free Documentation License

Version 1.2, November 2002

Copyright © 2000,2001,2002 Free Software Foundation, Inc.
59 Temple Place, Suite 330, Boston, MA 02111-1307, USA

Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document *free* in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondarily, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of “copyleft”, which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The “Document”, below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as “you”. You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A “Modified Version” of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A “Secondary Section” is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document’s overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The “Invariant Sections” are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released

under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The “Cover Texts” are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A “Transparent” copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not “Transparent” is called “Opaque”.

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTEX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The “Title Page” means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, “Title Page” means the text near the most prominent appearance of the work’s title, preceding the beginning of the body of the text.

A section “Entitled XYZ” means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as “Acknowledgements”, “Dedications”, “Endorsements”, or “History”.) To “Preserve the Title” of such a section when you modify the Document means that it remains a section “Entitled XYZ” according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and

that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.

- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled “Endorsements”, provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled “History” in the various original documents, forming one section Entitled “History”; likewise combine any sections Entitled “Acknowledgements”, and any sections Entitled “Dedications”. You must delete all sections Entitled “Endorsements.”

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called

an “aggregate” if the copyright resulting from the compilation is not used to limit the legal rights of the compilation’s users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document’s Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled “Acknowledgements”, “Dedications”, or “History”, the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License “or any later version” applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

G.1 ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

```
Copyright (C) year your name.  
Permission is granted to copy, distribute and/or modify this document  
under the terms of the GNU Free Documentation License, Version 1.2  
or any later version published by the Free Software Foundation;  
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover  
Texts. A copy of the license is included in the section entitled "GNU  
Free Documentation License".
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with...Texts." line with this:

```
with the Invariant Sections being list their titles, with  
the Front-Cover Texts being list, and with the Back-Cover Texts  
being list.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Index

!	
! packet	294
#	
# (a comment)	17
# in Modula-2	120
\$	
\$	76
\$\$	76
\$_ and info breakpoints	36
\$_ and info line	63
\$_, \$__, and value history	70
\$_, convenience variable	78
\$___, convenience variable	78
\$_exitcode, convenience variable	78
\$bpnum, convenience variable	34
\$cdir, convenience variable	62
\$cwdr, convenience variable	62
\$tpnum	92
\$trace_file	98
\$trace_frame	98
\$trace_func	98
\$trace_line	98
\$tracepoint	98
-	
--annotate	14
--args	14
--batch	13
--baud	14
--cd	13
--command	12
--core	12
--directory	12
--epoch	14
--exec	12
--fullname	14
--interpreter	14
--mapped	12
--nowindows	13
--nx	13
--pid	12
--quiet	13
--readnow	13
--se	12
--silent	13
--statistics	15
--symbols	12
--tty	14
--tui	14
--version	15
--windows	13
--write	15
-b	14
-break-after	208
-break-condition	209
-break-delete	209
-break-disable	210
-break-enable	210
-break-info	211
-break-insert	211
-break-list	212
-break-watch	213
-c	12
-d	12
-data-disassemble	215
-data-evaluate-expression	217
-data-list-changed-registers	218
-data-list-register-names	218
-data-list-register-values	219
-data-read-memory	220
-display-delete	222
-display-disable	222
-display-enable	222
-display-insert	222
-display-list	223
-e	12
-environment-cd	223
-environment-directory	223
-environment-path	224
-environment-pwd	225
-exec-abort	225
-exec-arguments	226
-exec-continue	226
-exec-finish	226
-exec-interrupt	227
-exec-next	227
-exec-next-instruction	228
-exec-return	228
-exec-run	229
-exec-show-arguments	229
-exec-step	230
-exec-step-instruction	230
-exec-until	231
-f	14
-file-exec-and-symbols	231
-file-exec-file	231
-file-list-exec-sections	232
-file-list-exec-source-file	232
-file-list-exec-source-files	232
-file-list-shared-libraries	233
-file-list-symbol-files	233
-file-symbol-file	233
-gdb-exit	234
-gdb-set	234

-gdb-show	234
-gdb-version	234
-interpreter-exec	235
-m	12
-n	13
-nw	13
-p	12
-q	13
-r	13
-s	12
-stack-info-depth	236
-stack-info-frame	235
-stack-list-arguments	236
-stack-list-frames	237
-stack-list-locals	239
-stack-select-frame	239
-symbol-info-address	239
-symbol-info-file	240
-symbol-info-function	240
-symbol-info-line	240
-symbol-info-symbol	241
-symbol-list-functions	241
-symbol-list-lines	241
-symbol-list-types	241
-symbol-list-variables	242
-symbol-locate	242
-symbol-type	242
-t	14
-target-attach	242
-target-compare-sections	243
-target-detach	243
-target-disconnect	243
-target-download	244
-target-exec-status	245
-target-list-available-targets	245
-target-list-current-targets	246
-target-list-parameters	246
-target-select	246
-thread-info	247
-thread-list-all-threads	247
-thread-list-ids	247
-thread-select	248
-var-assign	252
-var-create	249
-var-delete	250
-var-evaluate-expression	251
-var-info-expression	251
-var-info-num-children	250
-var-info-type	251
-var-list-children	251
-var-set-format	250
-var-show-attributes	251
-var-show-format	250
-var-update	252
-w	13
-x	12
•	
., Modula-2 scope operator	120
'debug' subdirectories	141
'esgdbinit'	189
'gdbinit'	189
.gnu_debuglink sections	142
'.o' files, reading symbols from	137
'.os68gdbinit'	189
'.vxgdbinit'	189
/	
/proc	159
:	
::, context for variables/functions	66
::, in Modula-2	120
?	
? packet	294
@	
@, referencing memory as an array	67
~	
~done	207
~error	207
~running	207
_	
_NSPrintForDebugger, and printing Objective-C objects	116
'	
"No symbol "foo" in current context"	67
{	
{type}	66

A

A packet	294
abbreviation	17
abort (C-g)	277
accept-line (Newline or Return)	272
acknowledgment, for GDB remote	293
actions	93
active targets	147
Ada	120
Ada mode, general	120
Ada, deviations from	122
Ada, omissions from	121
Ada, problems	123
adbq_find_memory_in_frame	326
add-shared-symbol-file	138
add-symbol-file	137
address of a symbol	125
advance location	49
Alpha stack	176
AMD 29K register stack	176
annotations	253
annotations for errors, warnings and interrupts	254
annotations for invalidation messages	255
annotations for prompts	254
annotations for running programs	255
annotations for source display	256
append	82
append data to a file	82
apropos	20
arguments (to your program)	25
arrays	67
arrays in expressions	65
artificial array	67
ASCII character set	83
assembly instructions	63
assignment	131
async output in GDB/MI	205
AT&T disassembly flavor	63
attach	28
attach to a program by name	153
automatic display	70
automatic overlay debugging	102
automatic thread selection	31
auxiliary vector	80
awatch	38

B

b (break)	34
b packet	294
B packet	294
backtrace	54
backtrace limit	55
backtraces	54
backward-char (C-b)	272
backward-delete-char (Rubout)	274
backward-kill-line (C-x Rubout)	275

backward-kill-word (M- DEL)	275
backward-word (M-b)	272
beginning-of-history (M-<)	273
beginning-of-line (C-a)	272
bell-style	264
break	34
break ... thread <i>threadno</i>	51
break in overloaded functions	114
break on fork/exec	40
break on load/unload of shared library	40
break, and Objective-C	115
breakpoint	255
breakpoint address adjusted	46
breakpoint commands	44
breakpoint commands for GDB/MI	208
breakpoint conditions	42
breakpoint numbers	33
breakpoint on events	33
breakpoint on memory address	33
breakpoint on variable modification	33
breakpoint ranges	33
breakpoint subroutine, remote	155
breakpointing Ada elaboration code	123
breakpoints	33
breakpoints and threads	51
breakpoints in overlays	102
breakpoints-invalid	255
bt (backtrace)	54
bug criteria	257
bug reports	257
bugs in GDB	257
built-in simulator target	148

C

c (continue)	47
c (SingleKey TUI key)	197
C and C++	109
C and C++ checks	113
C and C++ constants	111
C and C++ defaults	113
C and C++ operators	110
c packet	295
C packet	295
C++	109
C++ compilers	112
C++ exception handling	114
C++ overload debugging info	184
C++ scope resolution	66
C++ symbol decoding style	75
C++ symbol display	114
C-L	197
C-o (operate-and-get-next)	17
C-x 1	196
C-x 2	196
C-x a	196
C-x A	196
C-x C-a	196

C-x o	196	configuring GDB	285
C-x s	196	configuring GDB, and source tree subdirectories	285
call	134	confirmation	183
call overloaded functions	112	connect (to STDBUG)	175
call stack	53	console i/o as part of file-i/o	308
call-last-kbd-macro (C-x e)	277	console interpreter	193
calling functions	134	console output in GDB/MI	205
calling make	15	constants, in file-i/o protocol	316
capitalize-word (M-c)	274	continue	47
casts, in expressions	65	continuing	46
casts, to view memory	66	continuing threads	52
catch	39	control C, and remote debugging	156
catch exceptions, list active handlers	57	controlling terminal	27
catchpoints	33	convenience variables	77
catchpoints, setting	39	convenience variables for tracepoints	98
cd	27	convert-meta	265
cdir	62	copy-backward-word ()	275
character sets	82	copy-forward-word ()	275
character-search (C-])	277	copy-region-as-kill ()	275
character-search-backward (M-C-])	277	core dump file	135
charset	82	core dump file target	148
checks, range	108	core-file	136
checks, type	107	crash of debugger	257
checksum, for GDB remote	293	ctrl-c message, in file-i/o protocol	308
choosing target byte order	149	current directory	62
clear	41	current stack frame	54
clear, and Objective-C	115	current thread	29
clear-screen (C-1)	272	cwd	62
clearing breakpoints, watchpoints, catchpoints	41	Cygwin-specific commands	161
close, file-i/o system call	310		
closest symbol and offset for an address	72		
code address and its source line	63		
collect (tracepoints)	93		
collected data discarded	94	d (delete)	41
colon, doubled as scope operator	120	d (SingleKey TUI key)	197
colon-colon, context for variables/functions	66	d packet	295
command editing	261	D packet	295
command files	189	data manipulation, in GDB/MI	215
command history	179	debug formats and C++	112
command hooks	188	debug links	142
command interpreters	193	debugger crash	257
command line editing	179	debugging C++ programs	112
commands	44, 254	debugging information directory, global	141
commands for C++	114	debugging information in separate files	141
commands to STDBUG (ST2000)	175	debugging multithreaded programs (on HP-UX)	30
comment	17	debugging optimized code	23
comment-begin	265	debugging stub, example	154
common targets	148	debugging target	147
compatibility, GDB/MI and CLI	206	define	187
compilation directory	62	defining macros interactively	87
compiling, on Sparclet	174	definition, showing a macro's	87
complete	20	delete	41
complete ((<u>TAB</u>))	276	delete breakpoints	41
completion	17	delete display	71
completion of quoted strings	18	delete mem	80
completion-query-items	265	delete tracepoint	92
condition	43	delete-char (C-d)	274
conditional breakpoints	42		

delete-char-or-list ()	276	edit	60
delete-horizontal-space ()	275	editing	179
deleting breakpoints, watchpoints, catchpoints	41	editing command lines	261
demangling C++ names	75	editing source files	60
derived type of an object, printing	75	editing-mode	265
descriptor tables display	160	eight-bit characters in strings	74
detach	28	else	187
detach (remote)	151	Emacs	201
device	167	emacs-editing-mode (C-e)	278
digit-argument (M-0, M-1, ..., M--)	276	enable	42
dir	62	enable display	71
direct memory access (DMA) on MS-DOS	161	enable mem	80
directories for source files	61	enable tracepoint	92
directory	62	enable-keypad	265
directory, compilation	62	enable/disable a breakpoint	41
directory, current	62	end	44
dis (disable)	42	end-kbd-macro (C-x)	276
disable	42	end-of-history (M->)	273
disable display	71	end-of-line (C-e)	272
disable mem	80	entering numbers	181
disable tracepoint	92	environment (of your program)	26
disable-completion	265	errno values, in file-i/o protocol	316
disassemble	63	error	254
disconnect	151	error on valid input	257
display	70	error-begin	254
display disabled out of scope	71	event debugging info	184
display of expressions	70	event designators	279
DJGPP debugging	160	event handling	39
dll-symbols	162	examining data	65
DLLs with no debugging symbols	162	examining memory	69
do (down)	55	exception handlers	39
do-uppercase-version (M-a, M-b, M-x, ...)	277	exception handlers, how to list	57
document	187	exceptionHandler	156
documentation	283	exchange-point-and-mark (C-x C-x)	277
down	55	exec-file	135
Down	197	executable file	135
down-silently	56	executable file target	148
downcase-word (M-l)	274	exited	255
download to H8/300 or H8/500	166	exiting GDB	15
download to Renesas SH	166	expand macro once	87
download to Sparclet	174	expand-tilde	265
download to VxWorks	165	expanding preprocessor macros	87
dump	82	expression debugging info	184
dump all data collected at tracepoint	97	expressions	65
dump data to a file	82	expressions in Ada	120
dump-functions ()	278	expressions in C or C++	109
dump-macros ()	278	expressions in C++	112
dump-variables ()	278	expressions in Modula-2	116
dump/restore files	82		
dynamic linking	137		

E

e (edit)	60
EBCDIC character set	83
echo	190

F

f (frame)	55
f (SingleKey TUI key)	197
F packet	295
F reply packet	307
F request packet	307
fatal signal	257
fatal signals	50
FDL, GNU Free Documentation License	337
fg (resume foreground execution)	47
file	135
file-i/o examples	317
file-i/o overview	306
File-I/O remote protocol extension	306
file-i/o reply packet	307
file-i/o request packet	307
find trace snapshot	95
finish	48
flinching	183
float promotion	182
floating point	79
floating point registers	78
floating point, MIPS remote	170
flush_i_cache	156
focus	198
focus of debugging	29
foo	144
fork, debugging programs which call	31
format options	71
formatted output	68
Fortran	1
forward-backward-delete-char ()	274
forward-char (C-f)	272
forward-search	61
forward-search-history (C-s)	273
forward-word (M-f)	272
frame debugging info	184
frame number	53
frame pointer	53
frame, command	54
frame, definition	53
frame, selecting	55
frameless execution	53
frames-invalid	255
free memory information (MS-DOS)	160
fstat, file-i/o system call	313
Fujitsu	155
full symbol tables, listing GDB's internal	128
function entry/exit, wrong values of variables	66
functions without line info, and stepping	48

G

g packet	295
G packet	295
g++, GNU C++ compiler	109
garbled pointers	160
GCC and C++	112

GDB bugs, reporting	257
GDB reference card	283
'gdb.ini'	189
GDB/MI, breakpoint commands	208
GDB/MI, compatibility with CLI	206
GDB/MI, data manipulation	215
GDB/MI, input syntax	203
GDB/MI, its purpose	203
GDB/MI, out-of-band records	207
GDB/MI, output syntax	204
GDB/MI, result records	207
GDB/MI, simple examples	206
GDB/MI, stream records	207
gdbarch debugging info	184
GDBHISTFILE, environment variable	180
gdbserve.nlm	153
gdbserver	152
GDT	160
getDebugChar	156
gettimeofday, file-i/o system call	314
global debugging information directory	141
GNU C++	109
GNU Emacs	201
gnu_debuglink_crc32	142

H

h (help)	19
H packet	296
H8/300 or H8/500 download	166
handle	50
handle_exception	155
handling signals	50
hardware watchpoints	37
hbreak	35
help	19
help target	147
help user-defined	187
heuristic-fence-post (Alpha, MIPS)	177
history events	279
history expansion	279
history expansion, turn on/off	180
history file	180
history number	76
history save	180
history size	180
history substitution	180
history-preserve-point	265
history-search-backward ()	273
history-search-forward ()	273
hook	188
hookpost	188
hooks, for commands	188
hooks, post-command	188
hooks, pre-command	188
horizontal-scroll-mode	265
host character set	82
htrace	172

hwatch.....	172	info threads.....	29
I			
i (info).....	20	info threads (HP-UX)	30
i packet.....	296	info tracepoints.....	94
I packet.....	296	info types.....	126
i/o.....	27	info variables.....	127
i386.....	154	info vector.....	80
'i386-stub.c'.....	154	info w32.....	161
IBM1047 character set	83	info watchpoints.....	38
IDT.....	160	info win.....	198
if.....	187	information about tracepoints	94
ignore.....	43	inheritance.....	114
ignore count (of breakpoint)	43	init file.....	189
INCLUDE_RDB.....	164	init file name.....	189
info.....	20	initial frame.....	53
info address.....	125	initialization file, readline.....	264
info all-registers.....	78	innermost frame.....	53
info args.....	57	input syntax for GDB/MI	203
info auxv.....	80	input-meta.....	266
info breakpoints.....	36	insert-comment (M-#)	277
info catch.....	57	insert-completions (M-*)	276
info cisco.....	150	inspect.....	65
info classes.....	127	installation.....	285
info display.....	71	instructions, assembly	63
info dll.....	162	integral datatypes, in file-i/o protocol	314
info dos.....	160	Intel	154
info extensions	107	Intel disassembly flavor	63
info f (info frame)	56	interaction, readline	261
info files.....	138	internal commands	289
info float.....	79	internal GDB breakpoints	37
info frame.....	56	interpreter-exec	193
info frame, show the source language	107	interrupt	15
info functions	127	interrupting remote programs	151
info line.....	63	interrupting remote targets	156
info line, and Objective-C	115	invalid input	257
info locals.....	57	invoke another interpreter	193
info macro.....	87	isatty call, file-i/o protocol	309
info mem.....	80	isatty, file-i/o system call	314
info or1k spr.....	171	isearch-terminators	266
info proc.....	159	ISO 8859-1 character set	83
info proc mappings	159	ISO Latin 1 character set	83
info program.....	33		
info registers	78		
info scope.....	126		
info selectors	127		
info set.....	20		
info share.....	140	J	
info sharedlibrary	140	jump	132
info signals.....	50	jump, and Objective-C	115
info source.....	126		
info source, show the source language	107		
info sources.....	127		
info stack.....	54	K	
info symbol.....	125	k packet	296
info target.....	138	kernel crash dump	159
info terminal.....	27	kernel memory image	159
		kernel object display	150
		keymap	266
		kill	28
		kill ring	262
		kill-line (C-k)	275
		kill-region ()	275
		kill-whole-line ()	275

kill-word (M-d)	275
killing text	262
KOD	150
kvm	159

L

l (list)	59
languages	105
last tracepoint number	92
latest breakpoint	34
layout	198
LDT	160
leaving GDB	15
Left	197
libkvm	159
limits, in file-i/o protocol	317
linespec	59
list	59
list of supported file-i/o calls	309
list output in GDB/MI	206
list, and Objective-C	115
listing GDB's internal symbol tables	128
listing machine instructions	63
listing mapped overlays	101
load address, overlay's	99
load <i>filename</i>	149
local variables	126
locate address	68
log output in GDB/MI	206
logging GDB output	15
lseek flags, in file-i/o protocol	317
lseek, file-i/o system call	311

M

m packet	296
M packet	297
m680x0	154
'm68k-stub.c'	154
machine instructions	63
macro define	87
macro definition, showing	87
macro expand	87
macro expansion, showing the results of preprocessor	87
macro undef	88
macros, example of debugging with	88
macros, user-defined	87
maint info breakpoints	289
maint info psymtabs	128
maint info sections	138
maint info symtabs	128
maint internal-error	289
maint internal-warning	289
maint print cooked-registers	290
maint print dummy-frames	289
maint print psymbols	128

maint print raw-registers	290
maint print reggroups	290
maint print register-groups	290
maint print registers	290
maint print symbols	128
maint set dwarf2 max-cache-age	290
maint set profile	290
maint show dwarf2 max-cache-age	290
maint show profile	290
maintenance commands	289
make	15
manual overlay debugging	101
map an overlay	101
mapped	136
mapped address	99
mapped overlays	99
mark-modified-lines	266
mark-symlinked-directories	266
match-hidden-files	266
maximum value for offset of closest symbol	72
mem	80
member functions	112
memory models, H8/500	169
memory region attributes	80
memory tracing	33
memory transfer, in file-i/o protocol	308
memory, viewing as typed object	66
memory-mapped symbol file	136
memset	156
menu-complete ()	276
meta-flag	266
mi interpreter	193
mil interpreter	193
mi2 interpreter	193
minimal language	124
Minimal symbols and DLLs	162
MIPS boards	169
MIPS remote floating point	170
MIPS remotedebug protocol	171
MIPS stack	176
mode_t values, in file-i/o protocol	316
Modula-2	1
Modula-2 built-ins	117
Modula-2 checks	119
Modula-2 constants	118
Modula-2 defaults	119
Modula-2 operators	116
Modula-2, deviations from	119
Modula-2, GDB support	116
Motorola 680x0	154
MS Windows debugging	161
MS-DOS system info	160
MS-DOS-specific commands	160
multiple processes	31
multiple targets	147
multiple threads	29

N

n (next)	47
n (SingleKey TUI key)	197
names of symbols	125
namespace in C++	112
native Cygwin debugging	161
native DJGPP debugging	160
negative breakpoint numbers	37
NetROM ROM emulator target	148
New systag message	29
New systag message, on HP-UX	30
next	47
next-history (C-n)	273
nexti	49
ni (nexti)	49
non-incremental-forward-search-history (M-n)	273
non-incremental-reverse-search-history (M-p)	273
non-member C++ functions, set breakpoint in ..	35
notation, readline	261
notational conventions, for GDB/MI	203
notify output in GDB/MI	205
NULL elements in arrays	73
number of array elements to print	73
number representation	181
numbers for breakpoints	33

O

object files, relocatable, reading symbols from	137
Objective-C	115
observer debugging info	184
octal escapes in strings	74
online documentation	19
open flags, in file-i/o protocol	316
open, file-i/o system call	309
OpenRISC 1000	171
OpenRISC 1000 htrace	172
operations allowed on pending breakpoints	37
optimized code, debugging	23
optimized code, wrong values of variables	66
optional debugging messages	184
or1k boards	171
or1ksim	171
OS ABI	182
out-of-band records in GDB/MI	207
outermost frame	53
output	190
output formats	68
output syntax of GDB/MI	204
output-meta	266
overlay	101
overlay area	99
overlay example program	103
overlays	99
overlays, setting breakpoints in	102

overload-choice	254
overloaded functions, calling	112
overloaded functions, overload resolution	114
overloading	45
overloading in C++	114
overwrite-mode ()	274

P

p packet	297
P packet	297
packets, reporting on stdout	184
page tables display (MS-DOS)	160
page-completions	266
partial symbol dump	128
partial symbol tables, listing GDB's internal ..	128
Pascal	1
passcount	92
patching binaries	134
path	26
pauses in output	181
pending breakpoints	36
PgDn	197
PgUp	197
physical address from linear address	161
pipes	24
po (print-object)	116
pointer values, in file-i/o protocol	315
pointer, finding referent	73
possible-completions (M-?)	276
post-commands	254
post-overload-choice	254
post-prompt	254
post-prompt-for-continue	254
post-query	254
pre-commands	254
pre-overload-choice	254
pre-prompt	254
pre-prompt-for-continue	254
pre-query	254
prefix-meta ((ESC))	277
premature return from system calls	51
preprocessor macro expansion, showing the results of	87
pretty print arrays	73
pretty print C++ virtual function tables	76
previous-history (C-p)	272
print	65
print an Objective-C object description	116
print settings	71
print-object	116
print/don't print memory addresses	72
printf	191
printing data	65
process image	159
processes, multiple	31
profiling GDB	290
prompt	179

prompt	254
prompt-for-continue	254
protocol basics, file-i/o	306
protocol specific representation of datatypes, in file-i/o protocol	314
protocol, GDB remote serial	293
ptype	125
putDebugChar	156
pwd	27

Q

q (quit)	15
q (SingleKey TUI key)	197
q packet	297
Q packet	297
query	254
quit	254
quit [expression]	15
quoted-insert (C-q or C-v)	274
quotes in commands	18
quoting Ada internal identifiers	123
quoting names	125

R

r (run)	24
r (SingleKey TUI key)	197
r packet	298
R packet	298
raise exceptions	40
range checking	108
ranges of breakpoints	33
rbreak	35
re-read-init-file (C-x C-r)	277
read, file-i/o system call	311
reading symbols from relocatable object files	137
reading symbols immediately	136
readline	179
readnow	136
recent tracepoint number	92
recording a session script	258
redirection	27
redraw-current-line ()	272
reference card	283
reference declarations	113
refresh	198
register stack, AMD29K	176
registers	78
regular expression	35
reloading symbols	127
reloading the overlay table	101
relocatable object files, reading symbols from	137
remote connection without stubs	152
remote debugging	149
remote programs, interrupting	151
remote protocol, field separator	293
remote serial debugging summary	157

remote serial debugging, overview	154
remote serial protocol	293
remote serial stub	155
remote serial stub list	154
remote serial stub, initialization	155
remote serial stub, main routine	155
remote stub, example	154
remote stub, support routines	155
remote target	148
remotedebug , MIPS protocol	171
remotetimeout	173
remove actions from a tracepoint	93
rename, file-i/o system call	312
Renesas	155
Renesas SH download	166
repeating command sequences	17
repeating commands	17
reporting bugs in GDB	257
reprint the last value	65
response time, MIPS debugging	176
restore	82
restore data from a file	82
result records in GDB/MI	207
resuming execution	46
RET (repeat last command)	17
retransmit-timeout , MIPS protocol	171
return	133
returning from a function	133
reverse-search	61
reverse-search-history (C-r)	273
revert-line (M-r)	277
Right	197
run	24
run to main procedure	25
running	24
running and debugging Sparclet programs	175
running VxWorks tasks	166
running, on Sparclet	174
rwatch	38

S

s (SingleKey TUI key)	197
s (step)	47
s packet	298
S packet	298
save tracepoints for future sessions	98
save-tracepoints	98
saving symbol table	136
scope	120
search	61
searching source files	61
section	138
segment descriptor tables	160
select trace snapshot	95
select-frame	54
selected frame	53
selecting frame silently	54

self-insert (a, b, A, 1, !, ...)	274
separate debugging information files	141
sequence-id, for GDB remote	293
serial connections, debugging	184
serial device, Renesas micros	167
serial line speed, Renesas micros	167
serial line, <code>target remote</code>	151
serial protocol, GDB remote	293
server prefix for annotations	254
set	20
set args	25
set auto-solib-add	139
set auto-solib-limit	140
set backtrace	55
set breakpoint pending	37
set breakpoints on all functions	36
set charset	83
set check range	109
set check type	108
set coerce-float-to-double	182
set complaints	183
set confirm	183
set cp-abi	182
set debug	184
set debug-file-directory	142
set debugevents	162
set debugexceptions	162
set debugexec	162
set debugmemory	162
set disassembly-flavor	63
set editing	179
set endian	149
set environment	26
set extension-language	107
set follow-fork-mode	31
set gnutarget	148
set height	181
set history	180
set history expansion	180
set history filename	180
set host-charset	83
set input-radix	181
set language	106
set listsize	59
set logging	16
set machine	168
set max-user-call-depth	188
set memory mod	169
set mipsfpu	170
set new-console	162
set new-group	162
set opaque-type-resolution	128
set os	150
set osabi	182
set output-radix	181
set overload-resolution	114
set print	72
set processor args	170
set prompt	179
set remote hardware-breakpoint-limit	153
set remote hardware-watchpoint-limit	153
set remote system-call-allowed 0	309
set remote system-call-allowed 1	309
set remotedebug, MIPS protocol	171
set retransmit-timeout	171
set rstack_high_address	176
set shell	162
set solib-absolute-prefix	140
set solib-search-path	141
set step-mode	48
set symbol-reloading	127
set target-charset	83
set timeout	171
set tracepoint	91
set trust-readonly-sections	139
set tui active-border-mode	199
set tui border-kind	199
set tui border-mode	199
set variable	131
set verbose	183
set width	181
set write	134
set-mark (C-@)	277
set_debug_traps	155
setting variables	131
setting watchpoints	37
SH	155
'sh-stub.c'	155
share	140
shared libraries	139
sharedlibrary	140
shell	15
shell escape	15
show	20
show args	25
show auto-solib-add	140
show auto-solib-limit	140
show backtrace	55
show breakpoint pending	37
show charset	83
show check range	109
show check type	108
show complaints	183
show confirm	183
show convenience	77
show copying	21
show cp-abi	182
show debug	184
show debug-file-directory	142
show directories	62
show editing	179
show environment	26
show gnutarget	148
show height	181
show history	180
show host-charset	83

show input-radix	181	spr	171
show language	106	ST2000 auxiliary commands	175
show listsizes	59	st2000 cmd	175
show logging	16	stack frame	53
show machine	168	stack on Alpha	176
show max-user-call-depth	188	stack on MIPS	176
show mipsfpu	170	stack traces	54
show new-console	162	stacking targets	147
show new-group	162	start	25
show opaque-type-resolution	128	start a new trace experiment	94
show os	150	start-kbd-macro (C-x C-o)	276
show osabi	182	starting	24
show output-radix	181	starting	255
show paths	26	stat, file-i/o system call	313
show print	72	static members of C++ objects	76
show processor	170	status of trace data collection	94
show prompt	179	status output in GDB/MI	205
show remote system-call-allowed	309	STDBUG commands (ST2000)	175
show remotedebug, MIPS protocol	171	step	47
show retransmit-timeout	171	steipi	49
show rstack_high_address	176	stepping	46
show shell	162	stepping into functions with no line info	48
show solib-absolute-prefix	141	stop a running trace experiment	94
show solib-search-path	141	stop on C++ exceptions	39
show symbol-reloading	127	stop reply packets	301
show target-charset	83	stop, a pseudo-command	188
show timeout	171	stopped threads	51
show user	188	stopping	255
show values	77	stream records in GDB/MI	207
show verbose	183	struct stat, in file-i/o protocol	315
show version	21	struct timeval, in file-i/o protocol	316
show warranty	21	stub example, remote debugging	154
show width	181	stupid questions	183
show write	134	switching threads	29
show-all-if-ambiguous	267	switching threads automatically	31
shows	180	symbol decoding style, C++	75
si (steipi)	49	symbol dump	128
signal	133, 255	symbol from address	125
signal-name	255	symbol names	125
signal-name-end	255	symbol overloading	45
signal-string	255	symbol table	135
signal-string-end	255	symbol tables, listing GDB's internal	128
signalled	255	symbol-file	135
signals	50	symbols, reading from relocatable object files ..	137
silent	44	symbols, reading immediately	136
sim	176	sysinfo	160
simulator, Z8000	175	system call, file-i/o protocol	309
size of screen	181	system calls and thread breakpoints	51
software watchpoints	37	system, file-i/o system call	314
source	190, 256		
source line and its code address	62		
source path	61		
Sparc	155		
'sparc-stub.c'	155		
'sparcl-stub.c'	155		
Sparclet	173		
SparcLite	155		
speed	167		

T

t packet	298
T packet	298
T packet reply	301
tab-insert (M- TAB)	274
target	147, 148
target abug	169
target array	170
target byte order	149
target character set	82
target cpu32bug	169
target dbug	169
target ddb port	170
target debugging info	184
target dink32	173
target e7000, with H8/300	166
target e7000, with Renesas ICE	168
target e7000, with Renesas SH	173
target est	169
target hms, and serial protocol	167
target hms, with H8/300	166
target hms, with Renesas SH	173
target jtag	171
target lsi port	170
target m32r	169
target m32rsdi	169
target mips port	169
target op50n	173
target output in GDB/MI	205
target pmon port	170
target ppcbug	173
target ppcbug1	173
target r3900	170
target rdi	166
target rdp	166
target rom68k	169
target rombug	169
target sds	173
target sh3, with H8/300	166
target sh3, with SH	173
target sh3e, with H8/300	166
target sh3e, with SH	173
target sim, with Z8000	176
target sparclite	175
target vxworks	164
target w89k	173
tbreak	35
TCP port, target remote	151
tdump	97
terminal	27
Text User Interface	195
tfind	95
thbreak	35
this, inside C++ member functions	112
thread breakpoints	51
thread breakpoints and system calls	51
thread identifier (GDB)	29
thread identifier (GDB), on HP-UX	30

thread identifier (system)	29
thread identifier (system), on HP-UX	30
thread number	29
thread <i>threadno</i>	30
threads and watchpoints	39
threads of execution	29
threads, automatic switching	31
threads, continuing	52
threads, stopped	51
tilde-expand (M-~)	277
timeout, MIPS protocol	171
trace	91
trace experiment, status of	94
tracebacks	54
tracepoint actions	93
tracepoint data, display	97
tracepoint deletion	92
tracepoint number	92
tracepoint pass count	92
tracepoint variables	98
tracepoints	91
translating between character sets	82
transpose-chars (C-t)	274
transpose-words (M-t)	274
tstart	94
tstatus	94
tstop	94
tty	27
TUI	195
TUI commands	198
TUI configuration variables	199
TUI key bindings	196
tui reg	198
TUI single key mode	197
type casting memory	66
type checking	107
type conversions in C++	112

U

u (SingleKey TUI key)	197
u (until)	48
UDP port, target remote	151
undisplay	71
undo (C-_ or C-x C-u)	277
unions in structures, printing	74
universal-argument ()	276
unix-line-discard (C-u)	275
unix-word-rubout (C-w)	275
unknown address, locating	68
unlink, file-i/o system call	313
unmap an overlay	101
unmapped overlays	99
unset environment	26
unsupported languages	124
until	48
up	55
Up	197

up-silently	56
upcase-word (M-u)	274
update	198
user-defined command	187
user-defined macros	87

V

v (SingleKey TUI key)	197
value history	76
variable name conflict	66
variable object debugging info	185
variable objects in GDB/MI	249
variable values, wrong	66
variables, readline	264
variables, setting	131
vCont packet	298
vCont? packet	299
vector unit	79
vector, auxiliary	80
version number	21
vi-editing-mode (M-C-j)	278
visible-stats	267
VxWorks	164
vxworks-timeout	164

W

w (SingleKey TUI key)	197
watch	38
watchpoint	255
watchpoints	33
watchpoints and threads	39
whatis	125
where	54
while	187
while-stepping (tracepoints)	94
wild pointer, interpreting	73
winheight	198
word completion	17

working directory	62
working directory (of your program)	27
working language	105
write, file-i/o system call	311
writing into corefiles	134
writing into executables	134
wrong values	66

X

x (examine memory)	69
x command, default address	63
X packet	299
x(examine), and info line	63

Y

yank (C-y)	275
yank-last-arg (M-. or M-_)	273
yank-nth-arg (M-C-y)	273
yank-pop (M-y)	275
yanking text	262

Z

z packet	299
Z packets	299
z0 packet	300
Z0 packet	300
z1 packet	300
Z1 packet	300
z2 packet	300
Z2 packet	300
z3 packet	300
Z3 packet	300
z4 packet	301
Z4 packet	301
Z8000	175
Zilog Z8000 simulator	175

The body of this manual is set in
cmr10 at 10.95pt,
with headings in **cmb10 at 10.95pt**
and examples in cmtt10 at 10.95pt.
cmti10 at 10.95pt,
cmb10 at 10.95pt, and
cmsl10 at 10.95pt
are used for emphasis.

